

COGNITIVE UNCERTAINTY^{*}

Benjamin Enke

Thomas Graeber[†]

November 15, 2019

Abstract

This paper introduces a formal definition and an experimental measurement of the concept of *cognitive uncertainty*: people's subjective uncertainty about what the optimal action is. This concept allows us to bring together and partially explain a set of behavioral anomalies identified across four distinct domains of decision-making: choice under risk, choice under ambiguity, belief updating, and survey expectations about economic variables. In each of these domains, behavior in experiments and surveys tends to be insensitive to variation in probabilities, as in the classical probability weighting function. Building on existing models of noisy Bayesian cognition, we formally propose that cognitive uncertainty generates these patterns by inducing people to compress probabilities towards a mental default of 50:50. We document experimentally that the responses of individuals with higher cognitive uncertainty indeed exhibit stronger compression of probabilities in choice under risk and ambiguity, belief updating, and survey expectations. Our framework makes predictions that we test using exogenous manipulations of both cognitive uncertainty and the location of the mental default. The results provide causal evidence for the role of cognitive uncertainty in belief formation and choice, which we quantify through structural estimations.

^{*}Both authors contributed equally to this project. The experiments in this paper were pre-registered in the AEA RCT registry as trial AEARCTR-0004493. For helpful comments and discussions we are grateful to Rahul Bhui, Nicola Gennaioli, Sam Gershman, Pietro Ortoleva, Matthew Rabin, Josh Schwartzstein, Frederik Schwerter, Florian Zimmermann, and in particular Xavier Gabaix and Andrei Shleifer. Seminar audiences at Hamburg, Harvard Econ, Harvard Psych, Sloan / Nomis 2019 and SITE Experimental Economics 2019 provided valuable comments. Graeber thanks the Sloan Foundation and Enke the Foundations of Human Behavior Initiative for financial support. Enke: Harvard University, Department of Economics, and NBER, enke@fas.harvard.edu; Graeber: Harvard University, Department of Economics, graeber@fas.harvard.edu.

1 Introduction

In many contexts of economic interest, behaving optimally is difficult. Consider decision-making under uncertainty: forming and updating beliefs, and making choices among risky options. In belief formation, people need to combine prior beliefs with new information to arrive at a posterior. However, people may not know Bayes rule, succumb to computational errors, or fail to retrieve relevant information from memory. In choice under risk, people need to combine probabilities, payoffs and preferences into a certainty equivalent. However, people may not know their true preferences, or fail at adequately combining probabilities and utils. All of these issues, and potentially many more, may introduce cognitive noise, which we use as a catch-all term for the noisiness that arises from cognitive imperfections in the process of optimization.

The basic premise of our paper is that people are often aware of their own cognitive noise, which induces *cognitive uncertainty*: subjective uncertainty about what the optimal action or solution to a decision problem is. For example, people may think that they do not really know their own certainty equivalent of a lottery; they may have a nagging feeling that they do not remember what their prior information is; or they may worry that they do not know how to compute rational beliefs in light of new information. This stands in contrast to the vast majority of economic models, in which people potentially make errors, but do not exhibit deliberate doubts about the optimality of their solution. Because cognitive uncertainty reflects *internal* noise rather than stochasticity in the environment, it is distinct from standard notions of confidence. For example, an individual may assign a subjective probability of 95% to the event that the stock market will go up tomorrow (and hence be confident according to standard definitions), yet still be cognitively uncertain about whether the rational posterior belief is indeed 95%, or 84%, or 97%.

The objective of this paper is to propose and document empirically that cognitive uncertainty shapes economic behavior in systematic and quantitatively meaningful ways, and that it provides a unifying lens for understanding behavioral anomalies across different decision domains. To this effect, we first provide a formal definition and a corresponding structured experimental measurement of cognitive uncertainty. We then use these tools to bring together and partially explain a set of well-known but previously-unconnected anomalies across four economic decision domains: choice under risk, choice under ambiguity, belief updating, and survey expectations about economic variables.

Figure 1 depicts the set of behavioral anomalies that we focus on. All four functions are estimated from experimental data and share in common a characteristic inverse S-shape. First, panel A depicts the well-known probability weighting function in choice under risk that goes back to Kahneman and Tversky's (1979) prospect theory. It illustrates

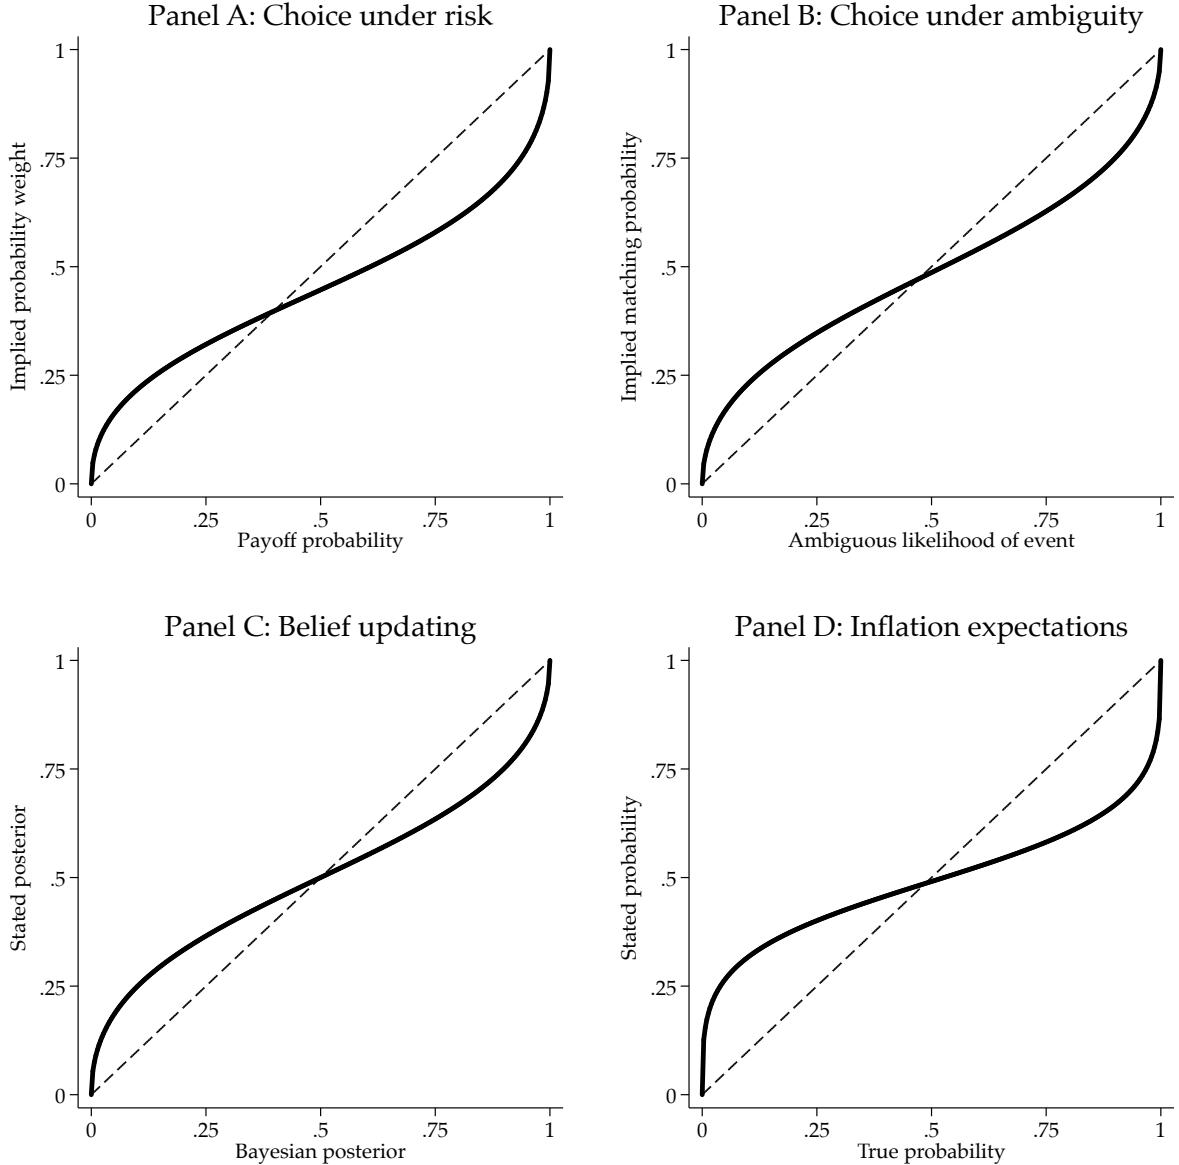


Figure 1: “Weighting functions” in choices and beliefs. Panel A depicts a probability weighting function in choice under risk, estimated from the data described in Section 3; also see Tversky and Kahneman (1992) and O’Donoghue and Somerville (2018). Panel B illustrates an “ambiguity weighting function,” where the x-axis represents the ambiguous likelihood of an event and the y-axis the matching probability (adapted from Li et al., 2019). Panel C visualizes the relationship between Bayesian posteriors and stated beliefs in binary-state balls-and-urns belief updating experiments, constructed from the data described in Section 4; also see Ambuehl and Li (2018). Finally, panel D depicts the relationship between objective probabilities and stated subjective probabilities in a survey on inflation expectations, described in Section 5; also see Fischhoff and Bruine De Bruin (1999).

how experimental subjects implicitly treat objective probabilities in choosing between different monetary gambles. Relative to an expected utility maximizer, people behave as if small probabilities were larger than they really are, and high probabilities as if they were smaller than they really are, leading to a compression effect. Second, depicted in panel B is an “ambiguity weighting function” that depicts the emerging consensus that,

in choices between monetary gambles over gains, people are ambiguity averse for likely events, yet ambiguity seeking for unlikely events. This reflects a compression effect that is labeled “a-insensitivity” in the literature (Trautmann and Van De Kuilen, 2015). Third, in panel C, we illustrate a less well-known stylized fact, which is an inverse S-shaped relationship between participants’ posterior beliefs and the Bayesian posterior in canonical “balls-and-urns” belief updating tasks of the type recently reviewed by Benjamin (2019). Finally, panel D of Figure 1 shows the relationship between objectively correct probabilities and the level of respondents’ probabilistic estimates in subjective expectations surveys about, e.g., stock market returns, inflation rates, or the shape of the income distribution. Here, again, people’s beliefs are compressed towards 50:50 (Fischhoff and Bruine De Bruin, 1999).

Each of these empirical findings represents a large and influential literature in behavioral economics. Yet, why do these four functions, drawn from different decision contexts and experimental paradigms, look so strikingly similar? Thus far, attempts to conceptualize these anomalies have focused on each decision domain in isolation.

This paper proposes cognitive uncertainty as a new lens through which these patterns can be interpreted and unified. Our experimental analysis is based on a theoretical framework that builds on the mathematical machinery of noisy Bayesian cognition models, in particular Gabaix (2019) and Khaw et al. (2017). We take a broad interpretation of these models as capturing (i) noise that primarily results from high-level reasoning in optimization rather than perceptual imperfections alone and (ii) measurable awareness of the resulting subjective uncertainty about the optimal action. In the model, people exhibit cognitive noise in translating probabilistic information into an optimal response. Similarly to standard Bayesian signal extraction models, this cognitive noise induces people to shrink objective probabilities towards a prior, or mental default. While we acknowledge that the mental default in general likely depends on a multitude of factors, we assume that in unfamiliar environments this default is influenced by an ignorance prior, which assigns equal probability mass to all states of the world *ex ante*.

Given this setup, we formally define and characterize an empirically measurable notion of cognitive uncertainty as subjective uncertainty about the optimal action. We show that this cognitive uncertainty in turn governs an individual’s degree of insensitivity to variation in probabilities, in both choice and belief formation. Building on insights from the noisy cognition literature, we then demonstrate that under the additional assumption that people perceive probabilities in log-odds space, cognitive uncertainty endogenizes the sensitivity parameter in the familiar two-parameter version of the probability weighting function proposed by Gonzalez and Wu (1999). However, our framework clarifies that we expect this weighting function to apply not only to choice under uncertainty but also to belief formation. Moreover, endogenizing the weighting function parameters

suggests that the precise shape of these functions will depend on the magnitude of cognitive noise and the location of the mental default.

This theoretical framework makes three predictions: (a) correlationally, individuals with higher cognitive uncertainty exhibit response functions that are more compressed towards 50:50 and hence more insensitive; (b) an exogenous increase in cognitive uncertainty generates more compressed response functions; and (c) an exogenous decrease in the location of the mental default shifts the entire response function downwards. All of these predictions have both a reduced-form and a structural interpretation in terms of Gonzalez and Wu's weighting function.

To test these predictions, we implement a series of pre-registered experiments with a total of $N = 2,800$ participants on Amazon Mechanical Turk (AMT). Like the motivating examples, our experiments cover the domains of choice under risk and ambiguity, balls-and-urns belief updating tasks, and survey expectations about economic variables. Our analysis rests on a new experimental paradigm to measure cognitive uncertainty. This measure is readily portable across decision domains, fast and easy to implement, and could in principle be added to a large set of experiments at little cost.

In choice under risk, each participant completes two decision screens per lottery. On the first screen, we elicit subjects' certainty equivalents for two-outcome gambles such as "Get \$20 with probability 75%; get \$0 with probability 25%" in a standard price list format. On the next screen, participants are asked how certain they are that to them the lottery is worth exactly the same as the switching interval that they stated on the previous screen. To answer this question, participants use a slider to calibrate the statement "I am certain that the lottery is worth between x and y to me." If a subject moves the slider to the very right, x and y collapse to their own switching interval in the price list. The further a subject moves the slider to the left, the wider the range of cognitive uncertainty becomes. Thus, our measure of cognitive uncertainty (i) directly reflects subjects' own assessment of uncertainty and (ii) is quantitative in nature. It is worth highlighting that existing theories – both traditional and behavioral – predict that people know their certainty equivalent. However, our data show that about 50% of the time, subjects exhibit cognitive uncertainty that is strictly wider than the switching interval of \$1.

As a test of prediction (a), we show that cognitive uncertainty is strongly correlated with the magnitude of probability weighting, in both the gain and loss domains. Specifically, as would be expected from the perspective of shrinking towards 50:50, cognitive uncertainty is positively correlated with risk taking for low probability gains and high probability losses, yet negatively correlated with risk taking for high probability gains and low probability losses. These findings refute a plausible alternative hypothesis about the effect of cognitive uncertainty, which is that people act cautiously in response to it

and hence appear universally more risk averse.

We test prediction (b) from above by introducing compound and ambiguous lotteries, which we hypothesize increase cognitive uncertainty. To illustrate, a compound lottery is a lottery that pays a non-zero amount with probability $p \sim U[0, 20]$. Similarly, an ambiguous lottery is a lottery that pays a non-zero amount with probability $p \in [0, 20]$. We verify that compound and ambiguous lotteries indeed induce substantially higher cognitive uncertainty than the corresponding reduced lotteries. Our model predicts that this increase in cognitive uncertainty translates into a more compressed weighting function. This again implies predictions about how compound or ambiguous lotteries should induce higher or lower risk aversion depending on whether one considers gains or losses, and high or low probabilities. In our experiments, we find consistent support for this hypothesis: while subjects act as if they are averse to compound lotteries or ambiguity under high probability gains and low probability losses, they are more risk seeking under compound lotteries (and “ambiguity seeking”) over small probability gains and high probability losses.

In a final step of the analysis of choice under risk, we test prediction (c) from above by exogenously manipulating the location of the mental default. To this effect, we leverage our assumption that in unfamiliar environments the default is influenced by an ignorance prior. In the two-states lotteries discussed so far, this ignorance prior is given by 50:50. To manipulate the location of the mental default, we implement a partition manipulation and translate the two-states lotteries into ten-states lotteries, without changing the objective payoff profile. We hypothesize that this shifts the mental default to an ignorance prior of 10%, which should move the entire probability weighting function closer towards zero. Our experimental results show that the probability weighting function with ten states is indeed significantly shifted towards zero, although these patterns are more pronounced for gains than for losses. These results also show that our experimental results do not just reflect a “click-in-the-middle” heuristic because we can manipulate people’s mental default in predictable ways.

In a second set of experiments, we conduct conceptually analogous exercises for belief updating. Here, we implement canonical balls-and-urns updating tasks of the type recently reviewed by Benjamin (2019). In these experiments, a computer randomly selects one of two bags according to a known base rate, yet subjects do not observe which bag got selected. The two bags both contain 100 balls, where one bag contains $q > 50$ red and $(100 - q)$ blue balls, while the other bag contains q blue and $(100 - q)$ red balls. The computer randomly draws one or more balls from the selected bag and shows these balls to the subject, who is then asked to provide a probabilistic assessment of which bag was actually drawn. Across experimental tasks, the base rate, the signal diagnosticity q and the number of random draws from the bags vary, but are always known to subjects.

After subjects state their posterior belief, we again elicit cognitive uncertainty. In a conceptually very similar fashion to choice under risk and ambiguity, we ask subjects to use a slider to calibrate the statement “I am certain that the optimal guess is between x and y .” We explain that the optimal Bayesian guess relies on the same information that is available to subjects, and combines this information in a statistically optimal way. Again, if a subject moves the slider to the very right, x and y collapse to the subject’s own previously stated belief. The further the slider is moved to the left, the wider the cognitive confidence interval becomes. In addition, we also elicit subjects’ willingness-to-pay to replace their own guess by the optimal guess as a complementary and incentivized measure of cognitive uncertainty.

Again, in contradiction to a large class of models in which agents do not exhibit doubts about the rationality of their belief updating, the vast majority of subjects indicate strictly positive cognitive uncertainty. As predicted by our model, this cognitive uncertainty is strongly correlated with compression of posterior beliefs towards 50:50. Moreover, we document that cognitive uncertainty strongly predicts the magnitude of base rate insensitivity and likelihood ratio insensitivity, two of the key underreaction anomalies highlighted in Benjamin’s (2019) meta-analysis.

To exogenously shift cognitive uncertainty, we implement compound belief updating tasks. We again hypothesize that cognitive uncertainty will be higher in compound problems, hence giving rise to more compressed belief distributions. In our experimental data, we find that cognitive uncertainty indeed increases by 33% under compound diagnosticities. Moreover, the distribution of beliefs becomes substantially more compressed towards 50:50, as predicted by our framework.

In a last step of the analysis of belief updating tasks, we exogenously vary the location of the mental default. Here, we once more employ the same partition methodology as in our risky choice experiments: we increase the number of states (bags) from two to ten, without changing the relevant Bayesian posterior. The results show that this manipulation induces a substantial and statistically significant downward shift of the entire distribution of posterior beliefs towards zero.

In the third part of the paper, we study the relationship between cognitive uncertainty and survey expectations about the performance of the stock market, inflation rates, and the structure of the national income distribution. For instance, we ask respondents to guess the probability that in a randomly selected year between 1980 and 2018 the inflation rate was less than $x\%$, where x varies across respondents. We measure cognitive uncertainty after we elicit these beliefs, using the same methodology as before. Again, we find that subjects with higher cognitive uncertainty exhibit survey expectations that are more regressive towards 50:50.

Next, we turn to estimating our model. For each decision domain, we structurally

estimate the two-parameter “weighting function” that our model endogenized, separately for above- and below-average cognitive uncertainty observations. The estimates reveal that the structural sensitivity parameter estimates of low-cognitive uncertainty observations are 60–150% higher across decision domains. Moreover, even though the underlying decision domains and experimental paradigms are very different, we always estimate fairly similar parameter values.

In the last part of the paper, we study whether variation in cognitive uncertainty is largely due to between- or within-subject variation. In a heuristic decomposition exercise, we find that – as a weak lower bound – at least 50% of the variation in our data is due to between-subjects heterogeneity. Moreover, participants’ cognitive uncertainty is highly correlated across domains, i.e., participants with high cognitive uncertainty in choice under risk (or belief updating) also exhibit high cognitive uncertainty in survey expectations. Across our different sets of experiments, this subject-level heterogeneity is correlated with observables: women, participants with low cognitive skills, and subjects with faster response times exhibit higher cognitive uncertainty.

In summary, the central contributions of our paper are (i) to introduce a formal definition and a new experimental measure of cognitive uncertainty and (ii) to provide both correlational and causal evidence for the role of cognitive uncertainty and a mental default across four domains of economic decision making, each of which has received substantial interest in the literature on its own.

Our paper builds on recent theoretical work on cognitive noise and resulting shrinkage processes, see Woodford (2012, 2019), Khaw et al. (2017), Gabaix and Laibson (2017), Gabaix (2019), Frydman and Jin (2019), Gershman and Bhui (2019) and Steiner and Stewart (2016). In contrast to some of these theories, we posit that cognitive uncertainty arises in the mental process of optimizing rather than from perceptual distortions of numeric quantities. We also do not require the neural coding of noise to be Bayes-optimal. Abstracting from these interpretive differences, we view our experiments as providing encouraging support for this emerging body of theoretical work. We show that cognitive noise is not just a low-level subconscious phenomenon but instead measurable and that it applies to, and unifies anomalies across, a broader range of economic settings than prior literature has theorized.

In the experimental literature, Butler and Loomes (2007) propose a measurement of preference imprecision in choice under risk, which is related in spirit to our measure. Agranov and Ortoleva (2017) show that experimental subjects often deliberately randomize between lottery options. However, these authors do not conceptualize their measures as cognitive uncertainty and resulting shrinking processes, and do not study the types of behavioral anomalies that we focus on. Our paper also builds on various experimental literatures on the anomalies that we attempt to bring together in this paper,

which are too voluminous to review here, such as Benjamin’s (2019) belief updating survey and his discussion of underreaction and “extreme belief aversion.” More broadly, our paper fits into the recent theoretical and experimental literature on bounded rationality that has focused on the mechanisms behind different behavioral anomalies (Bordalo et al., 2012, 2017; Enke, 2017; Enke and Zimmermann, 2019; Enke et al., 2019; Esponda and Vespa, 2016; Graeber, 2019; Martínez-Marquina et al., 2019).

The paper proceeds as follows. Section 2 lays out a theoretical framework of the role of cognitive uncertainty. Sections 3 to 5 present the experiments on choice under risk, choice under ambiguity, belief updating, and survey expectations. Section 6 provides parametric estimations of our model. Section 7 studies the correlates of cognitive uncertainty, Section 8 provides robustness checks, and Section 9 concludes.

2 A Model of Cognitive Uncertainty

2.1 Overview

Our formal framework directly builds on the cognitive imprecision models of Khaw et al. (2017) and Gabaix (2019). Following these contributions, our central assumption is the existence of cognitive noise in decision-making. In contrast to some earlier work, we interpret this noise not necessarily as reflecting low-level perceptual imperfections, but as resulting primarily from higher-level reasoning during optimization.¹ Translating a set of problem inputs (e.g., probabilities) into an optimal response (e.g., a certainty equivalent) is often difficult, which could introduce noise through various psychological mechanisms, including computational errors, retrieval from memory, or even from reading off and implementing one’s own preferences. Such cognitive noise creates cognitive uncertainty: subjective uncertainty about what the optimal action is. We focus on a notion of cognitive uncertainty that people have access to through introspection. To illustrate informally, suppose your prior belief that it rains tomorrow is 15%. Next, a weather forecast predicts that it will rain. You know from experience that the weather forecast is correct 80% of the time. What is your posterior belief that it will rain tomorrow? 45%? Really? Not 40%? Or perhaps 52%? To take another example, suppose you were asked to state your certainty equivalent of a 25% chance of getting \$15. Suppose you are small-stakes risk averse. You announce that your certainty equivalent is \$3. But is it really \$3? Or maybe \$2.50 or \$3.20?

Our second key assumption is that people represent quantities (probabilities) in log odds space, which is closely linked to the notion of a Weber’s Law that describes percep-

¹We hence follow the behavioral science tradition going back to at least Prospect Theory of recognizing that evidence about low-level processes can shed light on high-level decision-making.

tual patterns across various domains (Izard and Dehaene, 2008; Zhang and Maloney, 2012). The model’s key prediction about the effect of cognitive uncertainty on decisions (which is a form of insensitivity) does not depend on the assumption of log coding. However, as we show below, the assumption of log coding is instructive because in combination with cognitive noise it endogenously produces a canonical inverse S-shaped response function from the literature.

In this paper, we focus on noisy processing of probabilities. However, since cognitive noise may also stem from processing other problem inputs, our general notion of cognitive uncertainty could also be applied to other contexts.

2.2 Cognitive Noise, Shrinkage and their Interpretation

We focus our presentation on the case with normally distributed data and linear-quadratic utility but provide a generalization in Appendix A. Assume a decision-maker takes an action a and derives utility $u(a, x)$ that depends on a one-dimensional quantity x :

$$u(a, x) = -\frac{1}{2}(a - Bx)^2. \quad (1)$$

The quantity x may be a problem parameter explicitly presented to the decision-maker, or a value calculated by or retrieved from the agent’s memory. By “action,” we generically refer to the solution to a decision problem such as a stated posterior belief or a stated certainty equivalent. Here, x may correspond to the payout probability of a gamble in choice under risk, or to the Bayesian posterior in a belief updating task. We abstract away from both taste-based risk aversion and systematic biases in belief updating and choice – not because we think that they are unimportant but merely to keep our stylized framework as simple as possible.

The rational action is

$$a^r(x) = Bx. \quad (2)$$

We assume that the cognitive process required to identify an optimal action a is subject to cognitive noise. We model this as the agent receiving a signal $s = x + \varepsilon$ instead of having direct access to x .

It is important to emphasize that we view this “noisy perception” formalization *as if*, in that it arises *in the process of optimizing*. In choice under risk, we think of x as payout probability. Here, cognitive noise arises because combining probabilities, payoffs and preferences into a certainty equivalent is hard. In belief updating, x represents the Bayesian posterior belief that the agent attempts to compute. Here, cognitive noise arises in the complicated process of combining the available information into a rational belief. Finally, in survey expectations, x represents the true probability of an event.

Here, cognitive noise arises through the process of retrieving information from memory (Azeredo da Silveira and Woodford, 2019).

The agent is aware of the existence of cognitive noise. He *perceives* the noise term to be distributed according to $\varepsilon \sim \mathcal{N}(0, \sigma_\varepsilon^2)$. The agent's subjectively perceived cognitive noise need not be equal to the true noise that he is exposed to, $\tilde{\varepsilon} \sim \mathcal{N}(0, \sigma_{\tilde{\varepsilon}}^2)$. This assumption highlights that we do not require that an agent's cognitive uncertainty as defined below reflects the objective level of noise in his internal processing.

The agent holds a prior $x \sim \mathcal{N}(x^d, \sigma_x^2)$, where we refer to x^d as the “mental default.” The prior may be influenced by a multitude of factors and we do not model how exactly it is determined. As we discuss in greater detail below, in all of our applications x will represent a probability and we will specify x^d as influenced by a discrete ignorance prior that assigns equal probability mass to all possible states.

This assumption is plausible in our particular set of experiments because these are setups that are unfamiliar to most participants. We do not posit that the prior is always shaped by an ignorance prior.

Agents account for their cognitive noise by forming an implicit update about x . For a Bayesian agent, this creates a standard Gaussian signal extraction problem:

$$\mathbb{P}(x|s) \sim \mathcal{N}(\lambda s + (1 - \lambda)x^d, (1 - \lambda)\sigma_x^2), \quad (3)$$

with the shrinkage factor

$$\lambda = \frac{\sigma_x^2}{\sigma_x^2 + \sigma_\varepsilon^2} \in [0, 1]. \quad (4)$$

A rational agent takes an action by solving: $\max_a \mathbb{E}\left[-\frac{1}{2}(a - Bx)^2 | s\right]$. Only expectations matter in the linear first-order condition, leading to a rational action

$$a^r(s) = B(\mathbb{E}[x|s]) = B\lambda s + B(1 - \lambda)x^d \quad (5)$$

For a given x , the median action a^e across many agents with individual realizations of cognitive noise is then

$$a^e(x) = \text{Median}(a^r(s)|x) = B\lambda x + B(1 - \lambda)x^d, \quad (6)$$

which should be compared with equation (2). We see that the agent dampens his response by λ , generating shrinkage towards the default (prior). The key takeaway is that the existence of cognitive noise makes the rational action *insensitive* to variations in the problem parameter x .

2.3 Cognitive Uncertainty

Awareness of cognitive noise generates subjectively perceived uncertainty about what the optimal action is. We label this *cognitive uncertainty*. Our objective is to characterize this uncertainty *at the level of an individual action*, and derive empirical implications.

The agent's subjective uncertainty about his optimal action takes as given his individual draw of s , and reflects how the agent's rational action a^r (equation (2)) subjectively varies due to the agent's own posterior distribution of x (equation (3)), i.e., based on²

$$\mathbb{P}(a^r(x)|s) \sim \mathcal{N}(B\lambda s + B(1-\lambda)x^d, B^2(1-\lambda)\sigma_x^2). \quad (7)$$

Definition. *The agent's cognitive uncertainty is given by*

$$\sigma_{CU} = \sigma_{a^r(x)|s} = |B| \sqrt{1-\lambda} \sigma_x = |B| \frac{\sigma_\varepsilon \sigma_x}{\sqrt{\sigma_\varepsilon^2 + \sigma_x^2}}. \quad (8)$$

In our applications, cognitive uncertainty will reflect the agent's subjective uncertainty about (i) what exactly their certainty equivalent for a lottery is; (ii) what exactly the Bayesian posterior in a belief updating task is; and (iii) what (their knowledge of) the probability of some economic event is.

It is worth pointing out that cognitive uncertainty exclusively refers to a form of *internal* uncertainty about what the optimal response or behavior is. This is what distinguishes this concept from the notion of confidence, which usually involves external uncertainty. For example, an individual may assign a subjective probability of 95% to the event that the stock market will go up tomorrow (and hence be confident according to standard definitions), yet still be cognitively uncertain about whether the rational posterior belief is indeed 95%, or 84%, or 97%.

In our empirical analysis, we will measure a variant of (8). We reiterate that we do not require cognitive uncertainty to reflect the true amount of cognitive noise. Because we are concerned with empirical applications, all that matters for us is subjectively perceived cognitive uncertainty.

We make the following observation about the shrinkage factor λ :

$$\lambda = 1 - \frac{\sigma_{CU}^2}{B^2 \sigma_x^2} \quad (9)$$

That is, higher cognitive uncertainty generates more shrinking towards the default and makes people more insensitive to variation in x .³

²Note that the mean of $a^r(x)|s$ is $B(\lambda s + (1-\lambda)x^d)$ as in equation (5).

³While in literal terms our model posits shrinkage of the “input” quantity x , it also permits an equiv-

2.4 Log Coding

Following prior work in both cognitive science and economics (Gabaix, 2019; Zhang and Maloney, 2012), we assume that a probability p is transformed into a quantity q in log odds space by applying

$$q = Q(p) = \ln \frac{p}{1-p}. \quad (11)$$

In our model of probabilistic reasoning, people process probabilities with cognitive noise that occurs in log odds space. This means we now assume that the decision-relevant quantity x from (1) is a probability in log odds space q about which an agent receives a signal $s = q + \varepsilon$. This will result in shrinkage of probabilities in log odds space:

$$q(s) = \lambda s + (1 - \lambda)q^d. \quad (12)$$

In the following, we will focus on medians, which have the attractive property that for any strictly monotone function Y , $\text{Median}(Y(x)) = Y(\text{Median}(x))$. Over many draws of s – fixing x , but varying ε – the median posterior q^e about probability p after encoding in log odds space and shrinkage is:

$$q^e(q) := \text{Median}(q(s)|q) = \lambda q + (1 - \lambda)q^d. \quad (13)$$

From this we can derive the implied median posterior probability p by applying the inverse log odds function $P(q) = Q^{-1}(q) = \frac{1}{1+e^{-q}}$:

$$p^e(p) = P(q^e) = \frac{1}{1 + \exp\left(-\lambda \ln \frac{p}{1-p} - (1-\lambda) \ln \frac{p^d}{1-p^d}\right)}. \quad (14)$$

2.5 Empirical Applications and Predictions: “Weighting” Functions

Equation (14) delivers the microfoundation for our empirical analysis. Similarly to the models in Khaw et al. (2017), equation (14) can be reformulated as

$$w(p) := p^e(p) = \frac{\delta p^\lambda}{\delta p^\lambda + (1-p)^\lambda}, \quad (15)$$

where $\delta = \exp\left((1-\lambda)\ln \frac{p^d}{1-p^d}\right)$. This reformulation is instructive because it corresponds to the well-known two-parameter specification of a probability weighting function sug-

alent interpretation of shrinkage of the response a . Using $a^r(x) = Bx$ and letting $a^d = Bx^d$ we get

$$a^e(x) = B\lambda x + B(1-\lambda)x^d = \lambda a^r(x) + (1-\lambda)a^d. \quad (10)$$

gested by Gonzalez and Wu (1999). The original motivation by Gonzalez and Wu (1999) is that the log odds transformation allows a convenient characterization of the weighting function in which one parameter, λ , primarily represents the sensitivity of the weighting function to changes in probabilities, while another parameter, δ , controls the function's elevation. Our model motivates this functional form by endogenizing its parameters: λ directly corresponds to our shrinkage factor (and hence implicitly depends on cognitive uncertainty), while δ is a transformation of the constant term that jointly reflects the default and λ .

An implication of our approach, however, is that we expect this “weighting” function to adequately capture decision making not just in choice under risk, but also in choice under ambiguity, laboratory belief updating tasks, and survey expectations about economic variables. In all of these applications, we will operate under the assumption that the mental default about probabilities is influenced by an ignorance prior. We do not posit that the default is *always* affected by this ignorance prior – we just posit that this is the case in our experimental applications, with which people have no or very little prior experience.

Prediction 1. *Higher cognitive uncertainty is associated with more compressed weighting functions.*

Prediction 2. *An exogenous increase in cognitive uncertainty induces more compressed weighting functions.*

Prediction 3. *An exogenous decrease in the mental default induces the entire weighting function to move closer towards zero. That is, if the response is in the positive domain, a lower default leads to lower responses, holding fixed other problem parameters.*

Figures 13 and 14 in Appendix A illustrate these predictions graphically. We test the predictions in both reduced-form and structural analyses in which we estimate the “weighting function” from equation (15) across decision domains.

3 Choice Under Risk

3.1 Experimental Design

All experiments reported in this paper were designed with the same objective in mind: replicate standard experimental designs from the literature to elicit choices and beliefs, and supplement these tasks with a measurement of cognitive uncertainty.

3.1.1 Measuring Choice Behavior

To estimate a probability weighting function, we follow a large set of previous works and implement price lists that elicit certainty equivalents for lotteries (see, e.g. Bernheim and Sprenger, 2019; Bruhin et al., 2010; Tversky and Kahneman, 1992). Recent work suggests that these types of price lists are particularly easy for subjects to understand and give rise to more internally consistent and externally predictable choices than alternative measurement tools (Andreoni and Kuhn, 2019).

In treatment *Baseline Risk*, each subject completed a total of six price lists. On the left-hand side (Option A), a simple lottery was shown that paid y with probability p and nothing otherwise. On the right-hand side (Option B), a safe payment s was offered that increased by \$1 for each row that one proceeds down the list. As in Bruhin et al. (2010) and Bernheim and Sprenger (2019), the end points of the list were given by $s = \$0$ and $s = \$y$. Thus, each decision screen required $y + 1$ choices. A subject would typically start out by preferring Option A at the top of the list and then switch to Option B at some point as the safe payment increases.

Throughout, we enforce that subjects behave in internally consistent ways within a given choice list. That is, we do not allow for multiple switching points. This facilitates a simpler elicitation of cognitive uncertainty, as discussed below. To aid subjects' decision-making, we implemented an auto-completion mode: if a subject chose Option A in a given row, the computer implemented Option A also for all rows above this row. Likewise, if a subject chose Option B in a given row, the computer automatically and instantaneously ticked Option B in all lower rows. However, participants could always revise their decision and the auto-completion before moving on. See Figure 15 in Appendix B.1 for a screenshot of a decision screen.

The non-zero payout of the lottery y and the payout probability p were drawn randomly and independently from the sets $y \in \{15, 20, 25\}$ and $p \in \{5, 10, 25, 50, 75, 90, 95\}$. These lotteries are rather simple and well in line with previous work on probability weighting. We implemented both gain and loss gambles, where the loss amounts are the mirror images of y . In the case of loss gambles, the lowest safe payment was given by $s = -\$y$ and the highest one by $s = \$0$. In loss choice lists, subjects received a monetary endowment of $\$y$ from which potential losses were deducted. Out of the six choice lists that each subject completed, three concerned loss gambles and three gain gambles. We presented either all loss gambles or all gain gambles first, in randomized order.

Finally, with probability 1/3, a lottery choice list in treatment *Baseline Risk* was presented in a compound lottery format. We will describe, motivate and analyze these data in Section 3.3. For now we focus on the baseline (reduced) lotteries discussed above.

3.1.2 Measuring Cognitive Uncertainty

After a participant had completed a choice list, the next screen elicited their cognitive uncertainty with respect to this decision. Conceptually, we are interested in measuring the analog of σ_{CU} in the model (equation (8)). However, many people are not naturally familiar with the concept of a standard deviation. To strike a balance between conceptual clarity and quantitative interpretation on the one hand and participant comprehension on the other hand, we hence elicit an intuitive version of a subjective confidence interval.

Figure 2 provides a screenshot. Here, a participant was reminded of their valuation (switching interval) for the lottery. They were then asked to indicate how certain they are that to them the lottery is worth exactly the same as their previously indicated certainty equivalent. To answer this question, subjects used a slider to calibrate the statement “I am certain that the lottery is worth between a and b to me.” If the participant moved the slider to the very right, a and b corresponded to the previously indicated switching interval. For each of the 20 possible ticks that the slider was moved to the left, a decreased and b increased by 25 cents, in real time. In gain lotteries, a was bounded from below by zero and b bounded from above by the lottery’s upside. Analogously, for losses, a was bounded from below by the lottery’s downside and b from above by zero. The slider did not have a default value, meaning that subjects had to click somewhere on the slider in order to proceed.

This measure of cognitive uncertainty can be thought of as a subjective confidence interval.⁴ We note that we deliberately did not financially incentivize the elicitation of cognitive uncertainty. The reason is that we do not know the objective truth (subjects’ valuation for a lottery) because we do not know subjects’ true preferences.

Two remarks are in order. First, this measure of cognitive uncertainty only captures *internal* uncertainty about what the certainty equivalent is, rather than also external uncertainty that arises due to stochasticity in the environment. Second, both traditional and behavioral models that do not feature cognitive uncertainty would predict a cognitive uncertainty of zero: in these models, people may be loss averse, engage in probability weighting or be otherwise behavioral, yet they are always assumed to know their valuation for a lottery.

Throughout the paper, we normalize cognitive uncertainty to be in [0,1], where

⁴Our elicitation procedure did deliberately not specify which particular confidence interval (e.g., 95%) we are interested in. The reason is that (i) we aimed at keeping the elicitation simple and (ii) we are operating precisely under the assumption that subjects do not really know how to translate 90% or 95% confidence into an appropriate certainty equivalent. In Appendix H, we report on “calibration” experiments in which we explicitly elicit 70%, 90%, 95%, 99% and 100% confidence intervals, and compare them with our baseline measure. We find that subjects always indicate approximately identical cognitive uncertainty ranges, on average, regardless of which confidence interval we elicit. This supports our belief that providing a specific interval would not be helpful to subjects.

Decision screen 2

You will receive a **bonus of \$0.25 for a careful consideration** of the question below.

With probability 90% : **Get \$ 20**
With probability 10% : **Get \$ 0**

On the previous decision screen you indicated that this lottery is worth between getting \$17 and getting \$18 to you.

How certain are you that to you this lottery is worth **exactly** the same as getting between \$17 and \$18 for sure?



I am certain that the lottery is worth **between getting \$15.00 and getting \$20.00** to me.

Next

Figure 2: Decision screen to elicit cognitive uncertainty in choice under risk

one corresponds to the widest possible uncertainty interval. Figure 16 in Appendix B.1 shows a histogram of the distribution of cognitive uncertainty, which shows considerable variation. Average cognitive uncertainty is 0.24, with a median of 0.14 and a standard deviation of 0.21. 55% of our data indicate cognitive uncertainty that is strictly larger than the one-dollar switching interval.⁵

3.1.3 Subject Pool

All experiments reported in this paper were conducted on Amazon Mechanical Turk (AMT). AMT is becoming an increasingly used resource in experimental economics (e.g. DellaVigna and Pope, 2018; Imas et al., 2016), including in work on bounded rationality (Martínez-Marquina et al., 2019). Review papers suggest that experimental results on AMT and in the lab closely correspond to each other (Paolacci and Chandler, 2014). An important advantage of AMT that we also leverage in this study is that the pool of potential subjects is very large, which allows for high-powered analyses and a relatively large number of different treatments and tasks (Robinson et al., 2019).

We took four measures to achieve high data quality. First, our financial incentives are unusually large by AMT standards. Average realized earnings in the choice under risk experiments are \$6.10 for a median completion time of 20 minutes. This implies average hourly earnings of \$18, compared to a typical hourly wage of about \$5 on AMT. Second,

⁵As a basic validity check, in a small sample of 272 price lists, we implemented payout probabilities of $p = 0\%$ or $p = 100\%$, so that there is no external uncertainty. In these tasks, cognitive uncertainty drops considerably to an average of 0.10.

we aggressively screened out inattentive prospective subjects through comprehension questions and attention checks, described in detail below. Third, we pre-registered analyses that remove extreme outliers and speeders. Fourth, as explained above, subjects only completed six choice lists, which is considerably less than in typical choice under risk experiments.

3.1.4 Logistics and Pre-Registration

Based on the pre-registration (see below), we recruited $N = 700$ completes for treatment *Baseline Risk*. In all experiments, we restricted our sample to AMT workers that were registered in the United States, but did impose additional participation constraints such as a minimal approval rating. The common practice of selectively sampling AMT workers based on their reputation has recently been called into question as it limits the sample to the group of most experienced workers without increasing data quality (Robinson et al., 2019). After reading the instructions, participants completed a set of three comprehension questions that tested their understanding of the choice lists and the cognitive uncertainty question. Participants who answered one or more control questions incorrectly were immediately routed out of the experiment and do not count towards the number of completes. In addition, towards the end of the experiment, a screen contained a simple attention check. Subjects that answered this attention check incorrectly are excluded from the data analysis and replaced by a new complete, as specified in the pre-registration. In total, 62% of all prospective participants were screened out in the comprehension checks. Of those subjects that passed, 2% were screened out in the attention check. Thus, all of our results should be understood as being conditional on a pretty attentive participant pool – given the link between cognitive uncertainty and response times discussed in Section 7, we imagine that we would have identified even more variation in cognitive uncertainty had we not restricted the sample. Screenshots of instructions and control questions can be found in Appendix J.

In terms of timeline, subjects first completed six of the choice under risk tasks discussed above. Second, we elicited their survey expectations about various economic variables, as discussed in Section 5. Finally, participants completed a short demographic questionnaire and an eight-item Raven matrices IQ test.

Participants received a completion fee of \$1.70. In addition, each participant potentially earned a bonus. The experiment comprised three financially incentivized parts: the risky choice lists, the survey expectations questions, and the Raven IQ test. For each subject, one of these parts of the experiment was randomly selected for payment. If choice under risk was selected, a randomly selected decision from a randomly selected choice list was paid out.

All experiments reported in this paper were pre-registered in the AEA RCT registry, see <https://www.socialscienceregistry.org/trials/4493>. The pre-registration includes (i) the sample size in each treatment; (ii) data exclusion criteria such as the aforementioned attention checks or the handling of extreme outliers; and (iii) predictions about the relationship between cognitive uncertainty and our outcome measures.

3.2 Cognitive Uncertainty and the Probability Weighting Function

Because of the simple structure of our lotteries with only one non-zero payout state, an instructive non-parametric way to visualize our data is to compute *normalized certainty equivalents* as $NCE = CE/x$, where the certainty equivalent CE is defined as the midpoint of the switching interval⁶ and x is the non-zero payout. This measure hence effectively only represents the raw data, normalized by the non-zero payout. An attractive feature of NCE is that it directly corresponds to subjects' implied probability weights if one assumes that utility is linear. Because this will be instructive, these normalized certainty equivalents are negative for lotteries with losses.

For the purposes of the baseline analysis, we exclude extreme outliers as defined in the pre-registration: these are observations for which (i) the normalized certainty equivalent is strictly larger than 95% while the objective payout probability is at most 10%, or (ii) the normalized certainty equivalent is strictly less than 5% while the objective payout probability is at least 90%. For example, this excludes observations that state that the certainty equivalent for a 5% chance of receiving \$20 is strictly larger than \$19. This procedure of excluding outliers affects 3% of all data points. We report robustness checks using all data below.

Because the normalized certainty equivalents reflect implied probability weights, Figure 3 plots these normalized certainty equivalents against objective payout probabilities to visualize a heuristic version of the probability weighting function in our data. The figure distinguishes between subjects above and below average cognitive uncertainty within a given payoff probability bucket.

Focusing on the upper half of the figure (gain lotteries), first note that we replicate prior findings on the shape of the weighting function: implied probability weights are above the risk-neutral prediction for low probabilities but below the risk-neutral prediction for intermediate and high probabilities. Thus, our subject pool doesn't seem to be unusual in this regard.

More importantly, we find that subjects with higher cognitive uncertainty exhibit more pronounced probability weighting functions: still focusing on the top half, high cognitive uncertainty subjects are slightly more risk seeking for small probability gains

⁶We restrict normalized certainty equivalents to be between zero and one.

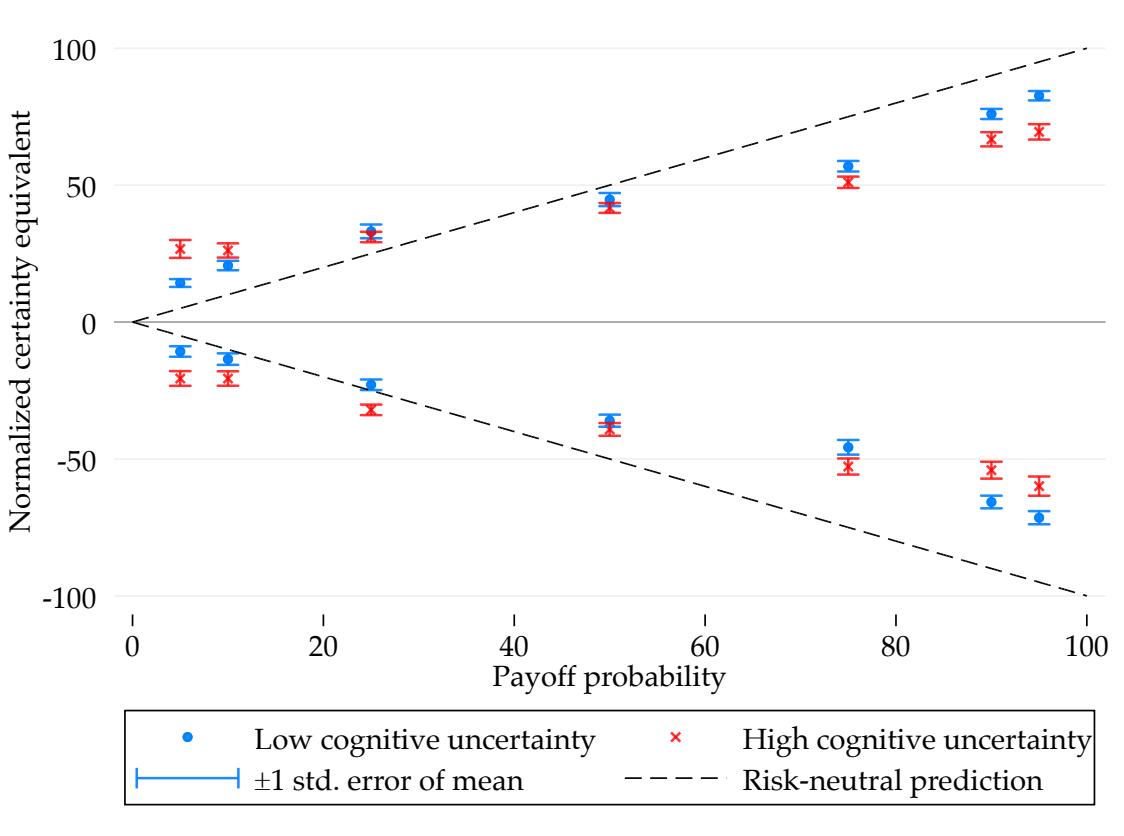


Figure 3: Probability weighting function separately for subjects above / below average cognitive uncertainty. The partition is done separately for each probability \times gains / losses bucket. The plot shows averages and corresponding standard error bars. Normalized certainty equivalents (implied probability weights) are computed as certainty equivalent divided by payout probability. The figure is based on 2,525 certainty equivalents of 700 subjects.

and more risk averse for high probability gains. Thus, overall, cognitive uncertainty is associated with more pronounced shrinking and hence a flatter relationship between implied probability weights and objective payout probabilities.⁷

Note that the heuristic probability weighting function depicted in Figure 2 crosses the 45-degree line to the left of $p = 50\%$. This pattern is well-known in the literature and in line with our hypothesis as long as subjects both (i) shrink towards 50:50 because of cognitive uncertainty and (ii) exhibit genuine preference-based risk or loss aversion, which shifts the function towards zero.

Next, we turn to the bottom panel of Figure 3, which depicts the analogous observed data for losses. By the construction of our figure, the weighting function is now given by the mirror image of the weighting function in the gain domain. Again, we see that the implied probability weights of subjects with higher cognitive uncertainty are more compressed. An attractive feature of visualizing the data as in Figure 3 is that it highlights

⁷This result resonates with the findings reported in Bruhin et al. (2010), who uncover substantial heterogeneity in individual-level probability weighting. Our notion of heterogeneity in cognitive uncertainty provides a possible micro-foundation for their results.

Table 1: Insensitivity to probability and cognitive uncertainty

	<i>Dependent variable:</i> Absolute normalized certainty equivalent					
	Gains		Losses		Pooled	
	(1)	(2)	(3)	(4)	(5)	(6)
Probability of payout	0.77*** (0.03)	0.76*** (0.03)	0.68*** (0.03)	0.68*** (0.03)	0.74*** (0.02)	0.74*** (0.02)
Probability of payout × Cognitive uncertainty	-0.67*** (0.09)	-0.67*** (0.09)	-0.56*** (0.09)	-0.55*** (0.09)	-0.64*** (0.07)	-0.64*** (0.07)
Cognitive uncertainty	25.2*** (5.88)	25.3*** (5.91)	35.7*** (5.26)	35.5*** (5.33)	31.6*** (4.12)	32.1*** (4.13)
Session FE	No	Yes	No	Yes	No	Yes
Demographic controls	No	Yes	No	Yes	No	Yes
Observations	1271	1271	1254	1254	2525	2525
R ²	0.55	0.56	0.42	0.43	0.48	0.48

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. The dependent variable is a subject's absolute normalized certainty equivalent, computed as midpoint of the switching interval divided by the non-zero payout. The sample includes choices from all baseline gambles with strictly interior payout probabilities. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

that the relationship between cognitive uncertainty and risk aversion reverses in predictable ways depending on whether the payouts are positive or negative and whether the payout probability is high or low. For instance, subjects with higher cognitive uncertainty are more risk seeking for small probability gains, but less risk seeking for small probability losses. Similarly, high cognitive uncertainty participants are more risk averse for high probability gains, yet less risk averse for high probability losses. In other words, as predicted, high cognitive uncertainty subjects exhibit a more pronounced fourfold pattern of risk attitudes, by the logic of shrinking towards 50:50.

Table 1 provides a regression analysis of these patterns. Our object of interest is the extent to which a subject's normalized certainty equivalent is (in)sensitive to variations in the probability of the non-zero payout state. Thus, we regress a participant's absolute normalized certainty equivalent on (i) the probability of receiving the non-zero gain / loss; (ii) cognitive uncertainty; and (iii) a corresponding interaction term. The regression analysis hence immediately corresponds to the setup of Figure 2.

The results show that higher cognitive uncertainty subjects respond less to variations in objective probabilities, in both the gains and the loss domain. In terms of quantitative magnitude, the regression coefficients suggest the following: if one were to increase the probability of the non-zero payout state from zero to one, then, on average, the increase in valuation for that lottery of subjects with cognitive uncertainty of zero is 55–67 percentage points higher than for subjects with cognitive uncertainty of one.

3.3 Exogenous Manipulation of Cognitive Uncertainty

3.3.1 Experimental Design

To exogenously manipulate cognitive uncertainty, we operate with compound lotteries and ambiguous lotteries. To illustrate, consider the case of compound lotteries, where an example lottery is given by: “We randomly draw an integer between 60 and 80, where each number is equally likely to be selected. Call this number n . With probability $n\%$, you receive \$20. With probability 100%- $n\%$, you receive \$0.” The analogous reduced lottery has payout probability $p = 70\%$. These two lotteries are identical under expected utility theory because EU is linear in probabilities.

Our hypothesis is that compound lotteries induce higher cognitive uncertainty. As per the model, compound lotteries should lead to more shrinking towards 50:50 – greater risk aversion under high probability gains (Halevy, 2007) and low probability losses, yet less risk aversion under low probability gains and high probability losses.

A causal interpretation of our experiments with respect to cognitive uncertainty requires the assumption that the introduction of compound lotteries affects choices *only* through cognitive uncertainty. While this is a strong assumption, we are not aware of alternative theories that would predict the nuanced pattern of how risk aversion changes as a function of reduced versus compound lotteries, depending on whether the lottery features high or low probabilities and gains or losses. As noted above, we implemented these compound lotteries as part of treatment *Baseline Risk*, where each lottery had a 1 in 3 chance of being presented in compound form. We collected 1,241 observations on compound lotteries.

To supplement this analysis of compound lotteries, and to link back to our discussion of an “ambiguity-weighting” function in the Introduction, we also ran separate experiments in which we implemented both reduced and ambiguous lotteries in a within-subjects design. Ambiguous lotteries follow the same format as compound lotteries, except that the precise distribution from which payoff probabilities are drawn is unknown. An example is: “There is a number n that lies between 60 and 80. With probability $n\%$, you receive \$20. Otherwise, you receive \$0.” We hypothesize that the introduction of ambiguity has the same effects as compound lotteries.⁸ These experiments were added to the pre-registration after the initial set of experiments was implemented. 300 subjects completed these experiments, in which each subject completed both lotteries with known payoff probabilities and ambiguous ones.⁹

⁸Consistent with our reasoning, Halevy (2007) and Gillen et al. (2019) report that the correlation between ambiguity and compound attitudes approaches one, once measurement error is accounted for.

⁹Appendix F presents an additional ambiguity experiment that we pre-registered and implemented. In these experiments, we do not elicit certainty equivalents for ambiguous lotteries but instead matching probabilities. These experiments also deliver statistically significant evidence for a correlation between

3.3.2 Results

Relative to the baseline reduced lotteries, compound and ambiguous lotteries increase stated cognitive uncertainty by 32% and 20%, on average. Figures 17 and 18 in Appendix B.1 show corresponding histograms. Thus, our experiments produce a strong “first stage.”

Figure 4 shows the results. In the top panel, we plot average normalized certainty equivalents separately for the baseline lotteries discussed above and for compound lotteries. We find that the probability weighting function is substantially more compressed under compound than under reduced lotteries, for both gains and losses. Consistent with many findings in the literature (Gillen et al., 2019; Halevy, 2007), subjects are compound lottery averse for high probability gains (and low probability losses). However, as predicted by our framework, subjects behave as if they are compound lottery loving for low probability gains and high probability losses.

The bottom panel presents analogous results for the comparison between risky and ambiguous gambles. The results are very similar to the ones for compound lotteries: subjects behave as if they are ambiguity loving for small probability gains and high probability losses, so that overall subjects’ certainty equivalents are less sensitive to variation in payout “probabilities” under ambiguity.

Table 2 provides a corresponding regression analysis. We find that subjects’ certainty equivalents are considerably less responsive to the objective payout probabilities under compound and ambiguous lotteries than under reduced lotteries, for both gains and losses. Moreover, we again find a within-treatment correlation between responsiveness to payout probabilities and cognitive uncertainty. For example, even when we restrict attention to ambiguous lotteries, the certainty equivalents of participants with higher cognitive uncertainty are significantly less responsive to variation in ambiguous likelihoods than those of subjects with low cognitive uncertainty ($p < 0.01$). This further suggests that the finding of “a-insensitivity” in the ambiguity literature (Li et al., 2019; Trautmann and Van De Kuilen, 2015) partly reflects cognitive uncertainty.

3.4 Exogenous Manipulation of the Mental Default

3.4.1 Experimental Design

In a final step of the analysis of choice under risk, we exogenously manipulate the location of the mental default. Recall that we operate under the assumption that the default is influenced by the ignorance prior. With two states of the world, the ignorance prior

cognitive uncertainty and “ambiguity-insensitivity.” We relegate these experiments to the appendix both for brevity and, as we discuss in detail in the Appendix, we now believe that they are less clean than the version that we present in the main text.

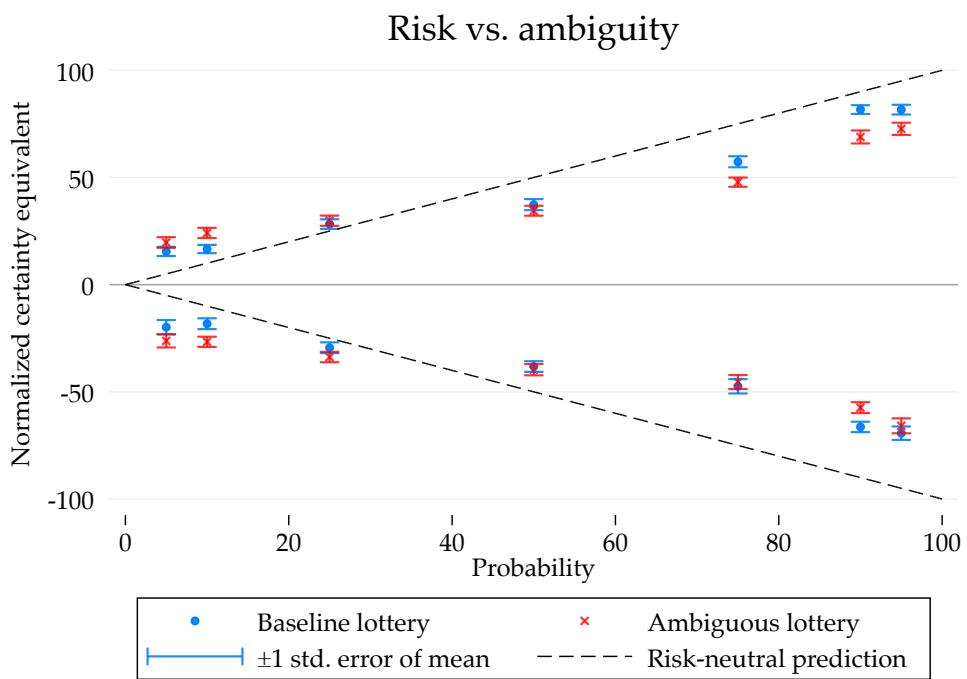
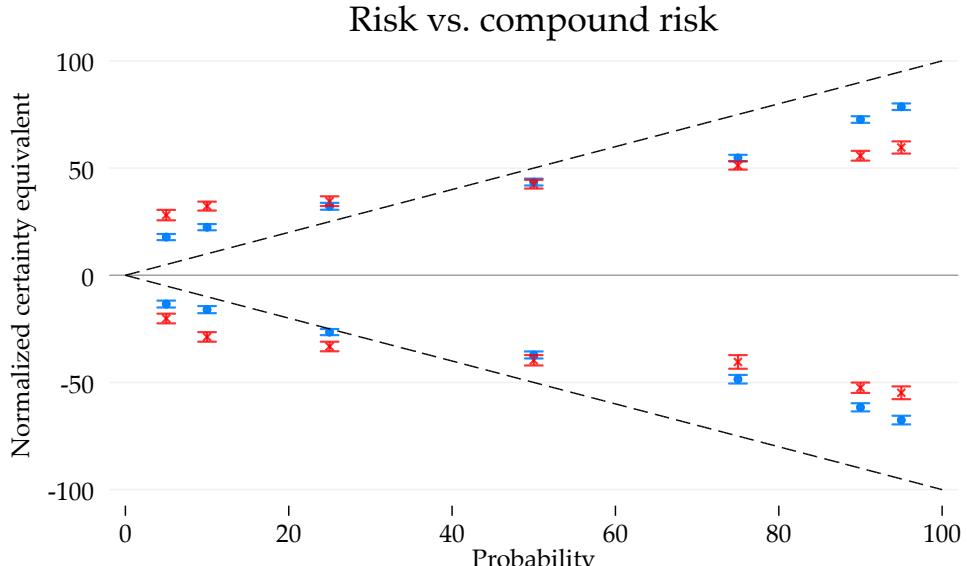


Figure 4: Top panel: Probability weighting function separately for reduced and compound lotteries. Bottom panel: “Probability” weighting function separately for reduced and ambiguous lotteries. In the bottom panel, the payout “probability” for ambiguous lotteries is denoted by the midpoint of the interval of possible payout probabilities. The plot shows averages and corresponding standard error bars. Normalized certainty equivalents (implied probability weights) are computed as certainty equivalent divided by payout probability. The top panel is based on 3,766 certainty equivalents of 700 subjects. The bottom panel is based on 1,796 certainty equivalents of 300 subjects.

Table 2: Choice under risk: Baseline versus compound / ambiguous lotteries

	Dependent variable: Absolute normalized certainty equivalent					
	Risk vs. compound risk			Risk vs. ambiguity		
	Gains (1)	Losses (2)	Pooled (3)	Gains (4)	Losses (5)	Pooled (6)
Probability of payout	0.62*** (0.02)	0.56*** (0.02)	0.72*** (0.02)	0.74*** (0.03)	0.54*** (0.04)	0.81*** (0.03)
Probability of payout \times 1 if compound lottery	-0.30*** (0.03)	-0.25*** (0.03)	-0.24*** (0.02)			
Probability of payout \times Cognitive uncertainty				-0.54*** (0.05)		-0.80*** (0.09)
1 if compound lottery	12.3*** (1.89)	12.3*** (1.84)	10.8*** (1.31)			
Cognitive uncertainty				27.6*** (3.46)		39.5*** (6.05)
Probability of payout \times 1 if ambiguous lottery					-0.20*** (0.03)	-0.16*** (0.04)
1 if ambiguous lottery					6.91*** (1.14)	8.82*** (2.31)
Session FE	No	No	Yes	No	No	Yes
Demographic controls	No	No	Yes	No	No	Yes
Observations	1918	1848	3766	889	880	1769
R ²	0.44	0.35	0.42	0.58	0.34	0.50

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. The dependent variable is a subject's absolute normalized certainty equivalent, computed as midpoint of the switching interval divided by the non-zero payout. In columns (1)–(3), the sample includes choices from the baseline and compound lotteries, where for comparability the set of baseline lotteries is restricted to lotteries with payout probabilities of 10%, 25%, 50%, 75%, and 90%, see Figure 4. In columns (4)–(6), the sample includes choices from the baseline and ambiguous lotteries. For ambiguous lotteries, we define the payout "probability" as the midpoint of the interval of possible payout probabilities. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

is 50:50. To vary the default, we implement a partition manipulation and increase the number of states to ten. Under the maintained assumption that the default is given by the ignorance prior, this means that the ignorance prior for each state is now given by 10%. We further designed this treatment variation with the objective of holding cognitive uncertainty fixed (which we verify below). Following the logic of the model in Section 2, we hence predict that the entire probability weighting function shifts towards zero as the number of states increases. This means that for the ten states lotteries we predict higher risk aversion for gains but lower risk aversion for losses.

To experimentally implement this manipulation, we replicate treatment *Baseline Risk*, but now frame probabilities in terms of number of colored balls in an urn. For example, we describe a lottery as:

Out of 100 balls, 80 are red. If a red ball gets drawn: get \$20.
20 balls are blue. If a blue ball gets drawn: get \$0.

In addition to this treatment, labeled *High Default Risk*, we also implement treatment *Low Default Risk*. Here, we implement the same lotteries as in *High Default Risk*, yet we split the zero-payout state into nine payoff-equivalent states with different probability colors. For example, the lottery above would be described as:

Out of 100 balls, 80 are red. If a red ball gets drawn: get \$20.
2 balls are blue. If a blue ball gets drawn: get \$0.
2 balls are black. If a black ball gets drawn: get \$0.
2 balls are white. If a white ball gets drawn: get \$0.
...
4 balls are yellow. If a yellow ball gets drawn: get \$0.

This lottery is identical to the one described above in terms of the objective payout profile. Still, we hypothesize that this manipulation shifts the probability weighting function towards zero.

In total, 300 subjects participated in these two treatments, which we implemented in a between-subjects design with random assignment to treatments within experimental sessions. All procedures other than the ones described here (and corresponding comprehension questions) were identical to the ones in condition *Baseline Risk*.

3.4.2 Results

First note that cognitive uncertainty does not vary across these two treatments ($p = 0.898$), see the histograms in Figure 19 in Appendix B.1. This lends credence to our implicit assumption that our experimental manipulation only affects the mental default but not cognitive uncertainty.

Figure 5 shows average normalized certainty equivalents, separately for treatments *High Default Risk* and *Low Default Risk*. We find that, in the gain domain, the probability weighting function is significantly shifted downwards towards zero with 10 states (a low default), as hypothesized. In the loss domain, our framework would predict that the weighting function is shifted upwards towards zero. We only find mixed evidence for this prediction: the weighting function appears to move up for low and intermediate probabilities but not for high probabilities.

Table 3 provides a corresponding regression analysis that confirms these visual patterns. Columns (1)–(3) analyze gain lotteries. Here, normalized certainty equivalents (observed risk tolerance) are 10 percentage points lower in the *Low Default Risk* condition. In the case of losses, the regression coefficient of the low default condition is

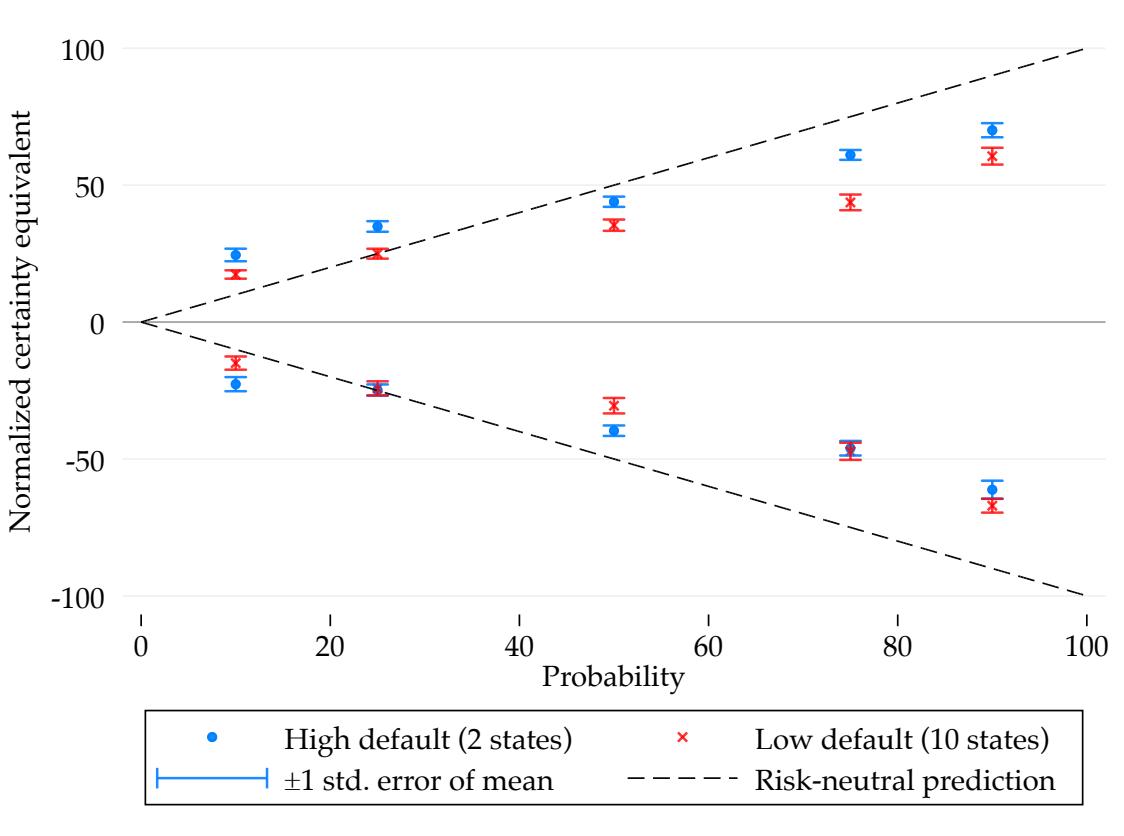


Figure 5: Probability weighting function separately for treatments *High Default Risk* and *Low Default Risk*. The plot shows averages and corresponding standard error bars. Normalized certainty equivalents (implied probability weights) are computed as certainty equivalent divided by payout probability. The figure is based on 1,757 certainty equivalents of 300 subjects.

negative – as predicted by our framework – but not statistically significant ($p = 0.15$). A potential (post-hoc) explanation for this null result is that, in all treatments, the choice data in the loss domain appear to be considerably more noisy than in the gain domain. This can be inferred from the difference in R^2 between columns (1) and (3) in Table 3 and similar patterns in all other tables above. Either way, the treatment effect of the low default is statistically significant in the pooled gains and losses sample.

4 Belief Updating

4.1 Experimental Design

Our experimental design strategy for belief updating closely mirrors the one for choice under risk: we (i) supplement an established experimental design from the literature with a measurement of cognitive uncertainty; (ii) document a correlation between cognitive uncertainty and the magnitude of compression in subjects’ beliefs; (iii) exogenously manipulate cognitive uncertainty using a compound manipulation; and (iv) vary

Table 3: Choice under risk: Treatments *Low Default Risk* and *High Default Risk*

	Dependent variable: Absolute normalized certainty equivalent					
	Gains		Losses		Pooled	
	(1)	(2)	(3)	(4)	(5)	(6)
0 if <i>High Default</i> , 1 if <i>Low Default</i>	-10.4*** (1.81)	-10.0*** (1.84)	-2.02 (2.11)	-1.87 (2.09)	-6.21*** (1.47)	-5.89*** (1.47)
Probability of payout	0.66*** (0.04)	0.66*** (0.04)	0.63*** (0.05)	0.63*** (0.05)	0.65*** (0.03)	0.65*** (0.03)
Probability of payout \times Cognitive uncertainty	-0.59*** (0.12)	-0.58*** (0.12)	-0.48*** (0.13)	-0.48*** (0.13)	-0.51*** (0.09)	-0.52*** (0.09)
Cognitive uncertainty	25.7*** (5.86)	25.2*** (5.85)	43.3*** (7.59)	43.1*** (7.79)	34.3*** (4.95)	34.1*** (5.01)
Session FE	No	Yes	No	Yes	No	Yes
Demographic controls	No	Yes	No	Yes	No	Yes
Observations	881	881	876	876	1757	1757
R ²	0.41	0.42	0.32	0.34	0.35	0.36

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. The dependent variable is a subject's absolute normalized certainty equivalent, computed as midpoint of the switching interval divided by the non-zero payout. The sample includes choices from treatments *Low Default Risk* and *High Default Risk*. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

the location of the mental default by increasing the number of states of the world.

4.1.1 Measuring Belief Updating

In designing a structured belief updating task, we follow the recent review and meta-study by Benjamin (2019) on so-called “balls-and-bags” or “bookbags-and-pokerchips” experiments. In treatment *Baseline Beliefs*, there are two bags, A and B. Both bags contain 100 balls, some of which are red and some of which are blue. The computer randomly selects one of the bags according to a pre-specified base rate. Subjects do not observe which bag was selected. Instead, the computer selects one or more of the balls from the selected bag at random (with replacement) and shows them to the subject. The subject is then asked to state a probabilistic guess that either bag was selected. We visualized this procedure for subjects using the image at the top right of Figure 6.

The three key parameters of this belief updating problem are: (i) the base rate $r \in \{10, 30, 50, 70, 90\}$ (in percent), which we operationalized as the number of cards out of 100 that had “bag A” or “bag B” written on them; (ii) the signal diagnosticity $q \in \{70, 90\}$, which is given by the number of red balls in bag A and the number of blue balls in bag B (we only implemented symmetric signal structures such that $P(red|A) = P(blue|B)$); and (iii) the number of randomly drawn balls N. These parameters were randomized across trials.

Each subject completed six belief updating tasks. In each task, they were asked to state a probabilistic belief (0-100) that bag A got selected. The computer automatically and instantaneously showed the corresponding subjective probability that bag B got selected. The decision screen contained information on the base rate, the signal diagnosticity, the number of drawn balls, and their color, see Figure 29 in Appendix C.1 for a decision screenshot.

Financial incentives were implemented through the binarized scoring rule (Hossain and Okui, 2013). Here, subjects had a chance of winning a prize of \$10. The probability of receiving the prize was given by $\pi = 100 - 0.04 * (b - t)^2$, where b is the guess (in %) and t the truth (0 or 100).

With probability 5 in 6, a belief updating task was implemented using the design discussed above, and with probability 1 in 6 in a compound design. We return to the compound data in Section 4.3 and focus on the baseline problems for now.

4.1.2 Measuring Cognitive Uncertainty

Our main measure of cognitive uncertainty in belief updating is very similar to the one for choice under risk, both conceptually and implementation-wise. The instructions explained the concept of an “optimal guess.” This guess, we explained to subjects, uses the laws of probability to compute a statistically correct statement of the probability that either bag was drawn, based on Bayes’ rule. We highlighted that this optimal guess does not rely on information that the subject does not have.

After subjects had indicated their probabilistic belief that either bag was drawn, the next decision screen elicited cognitive uncertainty. Here, we asked subjects how certain they are that their own guess equals the optimal guess for this task. Operationally, similarly to the case of choice under risk, subjects navigated a slider to calibrate the statement “I am certain that the optimal guess is between a and b ”, where a and b collapsed to the subject’s own previously indicated guess in case the slider was moved to the very right. For each of the 30 possible ticks that the slider was moved to the left, a decreased and b increased by one unit. a was bounded from below by zero and b bounded from above by 100. Again, we did not set a default: subjects had to click somewhere on the slider in order to proceed. Figure 6 shows a screenshot of the elicitation screen. For ease of interpretation, we again normalize this measure to be between zero and one.

Just like our measure of cognitive uncertainty in choice under risk, this one is not financially incentivized. However, in the case of belief updating, it is possible to devise an incentivized measure because here an objectively optimal response (the Bayesian posterior) exists. Thus, we additionally elicited a second measure of cognitive uncertainty from each participant: their willingness-to-pay (WTP) for replacing their own guess with

This decision is about the **same problem** as the one on the previous two screens:

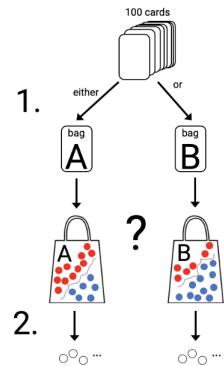
Number of "bag A" cards: **90**
 Number of "bag B" cards: **10**

Bag A contains **90 red** balls and **10 blue** balls.
 Bag B contains **10 red** balls and **90 blue** balls.

Next:

1. The computer **randomly selected one bag** by drawing a card from the deck.
2. Then, the computer randomly drew 1 ball from the secretly selected bag:

1 red ball was drawn.



Decision 3

You will receive a **bonus of \$0.25 for a careful consideration** of the question below.

On the previous screen you stated that you think it is **32% likely that bag A has been selected** and 68% likely that bag B has been selected in this task.

How **certain** are you that **the optimal guess is exactly 32%**?

Use the slider to complete the statement below.



Next

Figure 6: Decision screen to elicit cognitive uncertainty in belief updating

the optimal (Bayesian) guess. To this effect, before subjects stated their own guess, they received an endowment of \$3 for each task and then indicated how much of this budget they would at most be willing to pay to replace their guess. Subjects' WTP was elicited using a direct Becker-deGroot-Marschak elicitation mechanism. That is, we randomly drew a price $p \sim U[0, 3]$ and the guess was replaced iff $p \leq \text{WTP}$. See Figure 30 in Appendix C.1 for a screenshot.

It is worth reiterating that – just like in choice under risk – leading behavioral theories of biased belief updating predict no cognitive uncertainty: in these models, people may neglect the base rate, exhibit conservatism or be non-Bayesian in other ways, yet they never exhibit deliberate doubt about the rationality of their posterior belief. Moreover, as in choice under risk, our cognitive uncertainty measures only captures internal uncertainty about the rational solution to the decision problem, rather than also external uncertainty that is due to the stochasticity of the environment (here: which bag is selected).

To maximize statistical power, subjects' WTP and the resulting replacement of their own decision was only implemented in randomly selected 10% of all tasks. To avoid concerns about hedging, this uncertainty was resolved before subjects stated their own posterior guess. The timeline of each task was hence as follows: (i) observe game parameters; (ii) indicate WTP; (iii) find out whether own guess or Bayesian guess potentially counts for payment; (iv) state own posterior guess; and (v) indicate cognitive uncertainty range. The analysis below excludes those tasks in which a subject's guess got replaced by the optimal guess (3% of all data), though we have verified that virtually identical results hold if these (non-incentivized) guesses are included.

Figures 31 and 32 in Appendix C.1 show histograms of the cognitive uncertainty measure as well as subjects' WTP. Both measures exhibit considerable variation. Average cognitive uncertainty is 0.31, with a median of 0.33 and a standard deviation of 0.27. 85% of our data indicate strictly positive cognitive uncertainty. The average WTP is \$0.85 with a median of \$0.50 and a standard deviation of 0.93.¹⁰

The two measures exhibit a correlation of $\rho = 0.21$. While not incentivized, we view the cognitive uncertainty measure as our primary measure because (i) by its nature, and as exemplified by this paper, it is easily portable across different experimental contexts and decision situations; (ii) it is more fine-grained and exhibits more variation (26% of all WTPs are zero, perhaps due to some loss aversion vis-a-vis giving up safe money); and (iii) it is not confounded by risk aversion. Still, below we verify that all of our results are robust to using the WTP measure.

4.1.3 Logistics and Pre-Registration

Based on a pre-registration, we recruited $N = 700$ completes for treatment *Baseline Beliefs*. After reading the instructions, participants completed a set of four comprehension questions. Participants who answered one or more questions incorrectly were immediately routed out of the experiment and do not count towards the number of completes. Similarly, subjects are excluded from the analysis if they failed an attention check, as specified in the pre-registration. In total, 49% of all prospective participants were screened out in the comprehension checks. Of those subjects that passed, 6% were screened out based on the attention check.

In terms of timeline, subjects first completed the belief updating tasks discussed above. Second, we elicited their survey expectations about various economic variables, discussed in Section 5. Finally, participants completed a short demographic question-

¹⁰As a basic validity check, in a small sample of 161 updating tasks, we implemented a signal diagnosticity of $d = 100$, so that the selected bag is deterministically revealed. In these tasks, cognitive uncertainty essentially drops to zero: the distribution of both the cognitive uncertainty range and subjects' WTP has a median of zero, with means of 0.06 and 0.26.

naire and an eight-item Raven matrices IQ test. One of the three parts of the experiments (belief updating, survey expectations, or Raven test) was randomly selected for payment.

Average earnings are \$4.80 with a median completion time of 23 minutes. The experiments were pre-registered under the same AEA RCT trial as discussed above. Screenshots of the instructions and control questions can be found in Appendix J.

4.2 Cognitive Uncertainty and Belief Updating

As in the analysis of choice under risk, we begin by excluding extreme outliers to keep the analysis clean. As specified in the pre-registration, these are defined as subjective probability p_s and Bayesian posteriors p_o such that $p_s < 25 \wedge p_o > 75$ or $p_s > 75 \wedge p_o < 25$. This is the case for 5% of all data. We report robustness checks using the full sample below.

Figure 1 in the Introduction depicts the “belief weighting function” that we estimate in our data: the inverse S-shaped relationship between average stated and Bayesian posteriors that is also documented in Ambuehl and Li (2018). Figure 7 replicates this figure separately for subjects above or below average cognitive uncertainty as defined by our unincentivized cognitive uncertainty range. We see that, over the entire support of Bayesian posteriors, stated posteriors are more compressed towards 50:50 for subjects with higher cognitive uncertainty. Figure 36 in Appendix C.1 replicates this figure based on the financially incentivized WTP measure, with very similar results.

Columns (1)–(3) of Table 4 provide a corresponding econometric analysis. Here, we regress a subject’s stated posterior on (i) the Bayesian posterior; (ii) cognitive uncertainty; and (iii) their interaction term. We find that subjects with higher cognitive uncertainty respond considerably less to variation in the objectively correct answer: the quantitative magnitude of the regression coefficients suggests that the slope of the regression line between stated posterior and Bayesian posterior is 0.80 for subjects with measured cognitive uncertainty of zero, yet only 0.40 for subjects who state maximal cognitive uncertainty of one.

Grether regressions. A different way of analyzing our data is through the lens of so-called Grether regressions, see Grether (1992), El-Gamal and Grether (1995), and Benjamin (2019). This specification is derived by expressing Bayes’ rule in logarithmic form, which implies a linear relationship between the posterior odds, the prior odds, and the likelihood ratio:

$$\ln\left(\frac{b(A|s)}{b(B|s)}\right) = \beta_1 \ln\left(\frac{p(A)}{p(B)}\right) + \beta_2 \ln\left(\frac{p(s|A)}{p(s|B)}\right), \quad (16)$$

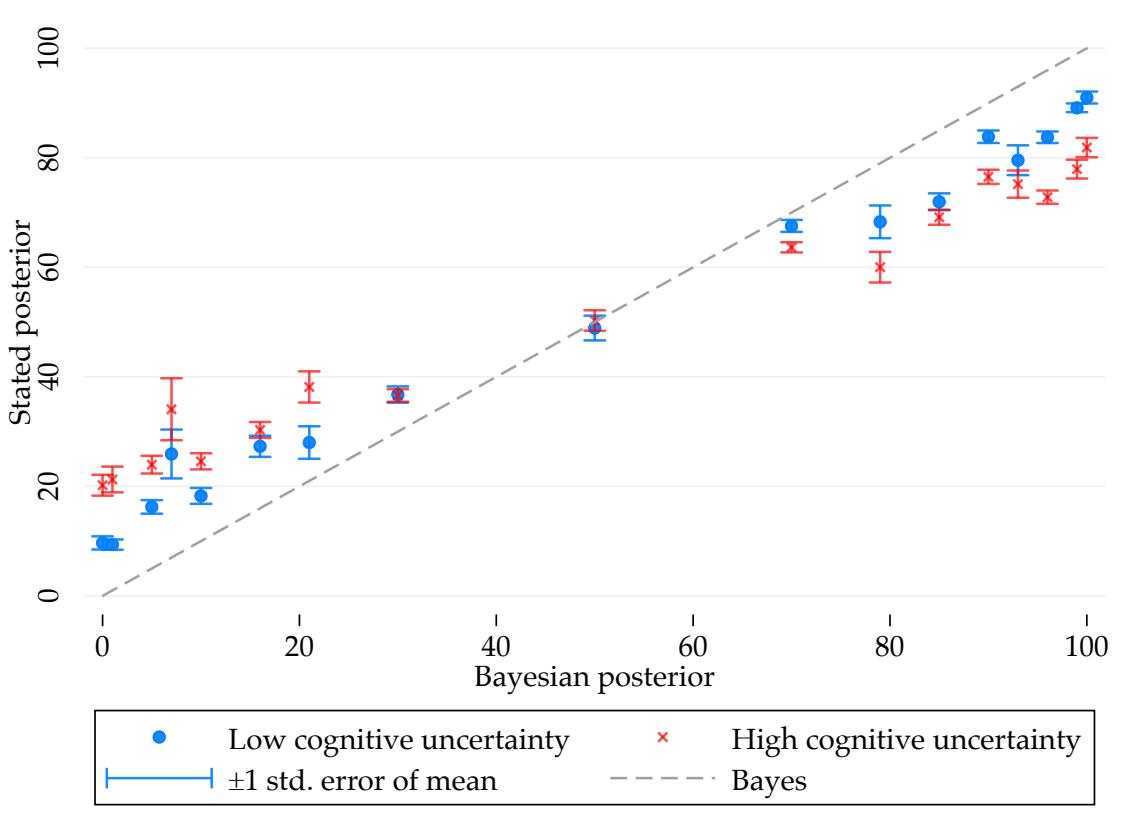


Figure 7: Relationship between average stated and Bayesian posteriors, separately for subjects above / below average cognitive uncertainty. The partition is done separately for each Bayesian posterior. Bayesian posteriors are rounded to the nearest integer. We only show buckets with more than ten observations. The figure is based on 3,187 beliefs of 700 subjects.

where $b(\cdot)$ denotes the stated posterior belief, A and B the two bags (states of the world), s a signal history, the first fraction on the right-hand side the prior odds, and the second term on the right-hand side the likelihood ratio. This formulation is attractive because it allows an assessment of the sensitivity of people's posteriors to variation in both the base rate and the likelihood ratio in a simple linear regression framework. The standard finding in the literature is that $\hat{\beta}_1 < 1$ and $\hat{\beta}_2 < 1$, even though Bayesian updating implies coefficients of one. This evidence hence points to paramount underreaction (insensitivity) to both the prior odds and the likelihood ratio (Benjamin, 2019).

Columns (4)–(7) of Table 4 implement these regressions using our data. As shown in column (4), similarly to past work, we find regression coefficients that are substantially smaller than one. In fact, our estimates are well within the range of results discussed in Benjamin's (2019) meta-study. However, as shown in columns (5)–(7), these insensitivities are significantly more pronounced for subjects with higher cognitive uncertainty: the responsiveness to the likelihood ratio (prior odds) decreases by 36% (65%) for subjects with maximal cognitive uncertainty relative to those with zero cognitive uncertainty.

These patterns suggest that (at least a part of) what this literature has identified as

Table 4: Belief updating: Baseline tasks

	Dependent variable:					
	Posterior belief			Log [Posterior odds]		
	(1)	(2)	(3)	(4)	(5)	(6)
Bayesian posterior	0.69*** (0.01)	0.83*** (0.01)	0.83*** (0.01)			
Bayesian posterior \times Cognitive uncertainty		-0.53*** (0.04)	-0.53*** (0.04)			
Cognitive uncertainty		24.8*** (2.43)	24.6*** (2.44)		-0.058 (0.07)	-0.075 (0.07)
Log [Likelihood ratio]				0.41*** (0.01)	0.46*** (0.02)	0.46*** (0.02)
Log [Prior odds]				0.42*** (0.02)	0.55*** (0.03)	0.55*** (0.03)
Log [Likelihood ratio] \times Cognitive uncertainty					-0.27*** (0.05)	-0.27*** (0.05)
Log [Prior odds] \times Cognitive uncertainty					-0.46*** (0.07)	-0.46*** (0.07)
Session FE	No	No	Yes	No	No	Yes
Demographic controls	No	No	Yes	No	No	Yes
Observations	3187	3187	3187	3104	3104	3104
R^2	0.72	0.74	0.74	0.62	0.64	0.64

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

base rate neglect or conservatism are in fact not independent psychological phenomena but instead generated by people shrinking their responses towards 50:50 due to cognitive uncertainty. Table 14 in Appendix C.2 replicates Table 4 using the WTP instead of the cognitive uncertainty measure, with very similar results.

4.3 Exogenous Manipulation of Cognitive Uncertainty

4.3.1 Experimental Design

To manipulate cognitive uncertainty, we again resort to turning “reduced” problems into compound problems. Consider belief updating problems in which the base rate is given by 50:50 and the signal diagnosticity by $d \equiv P(A|red) = P(B|blue)$. In the compound version of these problems, the base rate is again 50:50, yet the diagnosticity is the outcome of a random integer draw, $d \sim U[d - 10, d + 10]$. It is straightforward to verify that these two problems give rise to the same mean Bayesian posterior. For instance, if a red ball gets drawn, the posterior in the reduced version equals the signal diagnosticity d because the prior is 50:50. In the compound version, the posterior is equally likely to be $d - 10, d - 9, \dots, d + 10$, hence d in expectation.

As in choice under risk, we hypothesize that subjects exhibit higher cognitive uncertainty in compound than in reduced updating problems. Hence, by the logic of our framework, we expect that participants' beliefs in compound problems will be more compressed towards 50:50.¹¹

As noted above, we implemented these compound belief updating problems as part of treatment *Baseline Beliefs*, where each belief updating problem had a 1 in 6 chance of being presented in a compound form. We collected 592 observations on compound belief updating problems.

4.3.2 Results

Relative to reduced updating problems, compound signal diagnosticities increase stated cognitive uncertainty by 33% and subjects' WTP for the Bayesian guess by 43%, on average. Figures 33 and 34 in Appendix C.1 show corresponding histograms. Thus, as in choice under risk, the compound manipulation produces a strong "first stage."

Figure 8 shows the results on stated beliefs. Here, we plot average stated posteriors as a function of Bayesian posteriors, separately for baseline and compound updating problems. Because in compound problems the base rate is always 50:50, the figure only includes data from tasks with a 50:50 base rate also for the baseline updating problems. We see that subjects' posteriors are substantially more compressed towards 50:50 in compound updating problems.

Columns (1) and (2) of Table 5 provide a corresponding regression analysis. The regression coefficients suggest that the sensitivity of stated posteriors to the Bayesian posterior is 0.72 in baseline updating problem, yet only 0.21 in compound updating problems. To provide an alternative perspective on the data, we again resort to Grether regressions, see columns (4)–(6). Because in compound updating problems the base rate is fixed at 50:50, the only explanatory variable of interest here is the log likelihood ratio. The results show that subjects always underreact to variations in the likelihood ratio (the regression coefficient is smaller than one in both reduced and compound updating tasks), yet this underreaction is substantially more pronounced under compound lotteries.

4.4 Exogenous Manipulation of the Mental Default

4.4.1 Experimental Design

In a final step of the analysis of belief updating, we exogenously manipulate the location of the mental default. We again employ a partition manipulation and increase the

¹¹In contemporaneous work, Liang (2019) identifies underreaction under compound relative to reduced updating problems. This is in line with our work, but he does not measure cognitive uncertainty.

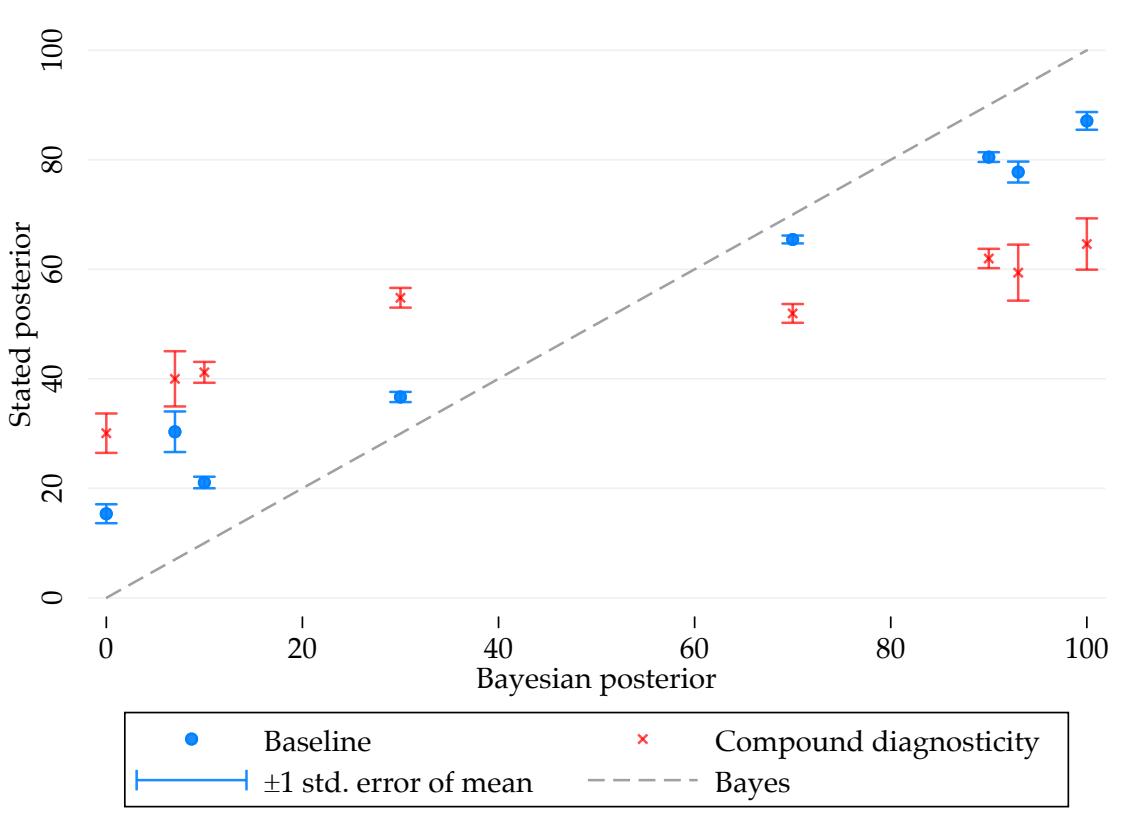


Figure 8: Stated average posteriors as a function of Bayesian posteriors, separately for reduced and compound belief updating problems. The plot shows averages and corresponding standard error bars. To allow for a valid comparison between baseline and compound updating problems, the sample is restricted to updating tasks in which the base rate is 50:50. Bayesian posteriors are rounded to the nearest integer. We only show buckets with more than ten observations. The figure is based on 1,947 beliefs of 691 subjects.

number of states to ten. Under the maintained assumption that the default is influenced by an ignorance prior, our framework predicts that the entire distribution of posterior beliefs shifts downwards.

Recall that in treatment *Baseline Beliefs*, an example updating problem is that the base rates for bags A and B are 70% and 30%, and the signal diagnosticity (number of red balls in bag A and number of blue balls in bag B) 70%. Now, in treatment *Low Default Beliefs*, we split the probability mass for bag B up into nine different bags. That is, there are now ten bags, labeled A through J. In the specific example above, the base rate for A would again be 70%, the one for B through I 3% each and the one for J 6%. Again, bag A would contain 70 red and 30 balls, and all bags B through J 30 red and 70 blue balls. That is, these bags have identical ball compositions.

Note that, regardless of what the actual draws of balls are, the Bayesian posterior for bag A having been selected is identical in the baseline version and the version with 10 bags. The reason is that under the state space {A; not A} the base rates and signal diagnosticities are identical. Thus, in treatment *Low Default Beliefs*, we asked participants

Table 5: Belief updating: Reduced versus compound signal diagnosticities

	Dependent variable:					
	Posterior belief			Log [Posterior odds]		
	(1)	(2)	(3)	(4)	(5)	(6)
Bayesian posterior	0.57*** (0.01)	0.72*** (0.02)	0.72*** (0.02)			
Bayesian posterior \times 1 if compound problem		-0.51*** (0.03)	-0.51*** (0.03)			
1 if compound problem		26.4*** (1.75)	26.5*** (1.74)		0.0058 (0.05)	0.0091 (0.05)
Log [Likelihood ratio]				0.37*** (0.01)	0.45*** (0.02)	0.45*** (0.02)
Log [Likelihood ratio] \times 1 if compound problem					-0.28*** (0.02)	-0.27*** (0.02)
Session FE	No	No	Yes	No	No	Yes
Demographic controls	No	No	Yes	No	No	Yes
Observations	1947	1947	1947	1890	1890	1890
R ²	0.51	0.60	0.60	0.47	0.53	0.53

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

to indicate their belief that bag A got selected, and the computer automatically showed the corresponding composite probability for one of the other bags having been selected.

300 subjects participated in treatment *Low Default Beliefs*, which was randomized within the same experimental sessions as treatment *Baseline Beliefs*. All procedures other than the ones described above were identical to the ones in *Baseline Beliefs*.

4.4.2 Results

Stated cognitive uncertainty is almost identical across conditions *Baseline Beliefs* and *Low Default Beliefs*, $p = 0.85$. This corroborates our implicit assumption that the experimental manipulation of increasing the number of bags only manipulates the mental default but not cognitive uncertainty.

Figure 9 shows average stated posteriors as a function of Bayesian posteriors, separately for treatments *Baseline Beliefs* and *Low Default Beliefs*. The results show that the entire distribution of subjects' beliefs is shifted down towards zero, consistent with our hypothesis that a larger state space induces shrinking towards a lower mental default. Table 6 provides a corresponding regression analysis that confirms these visual patterns.

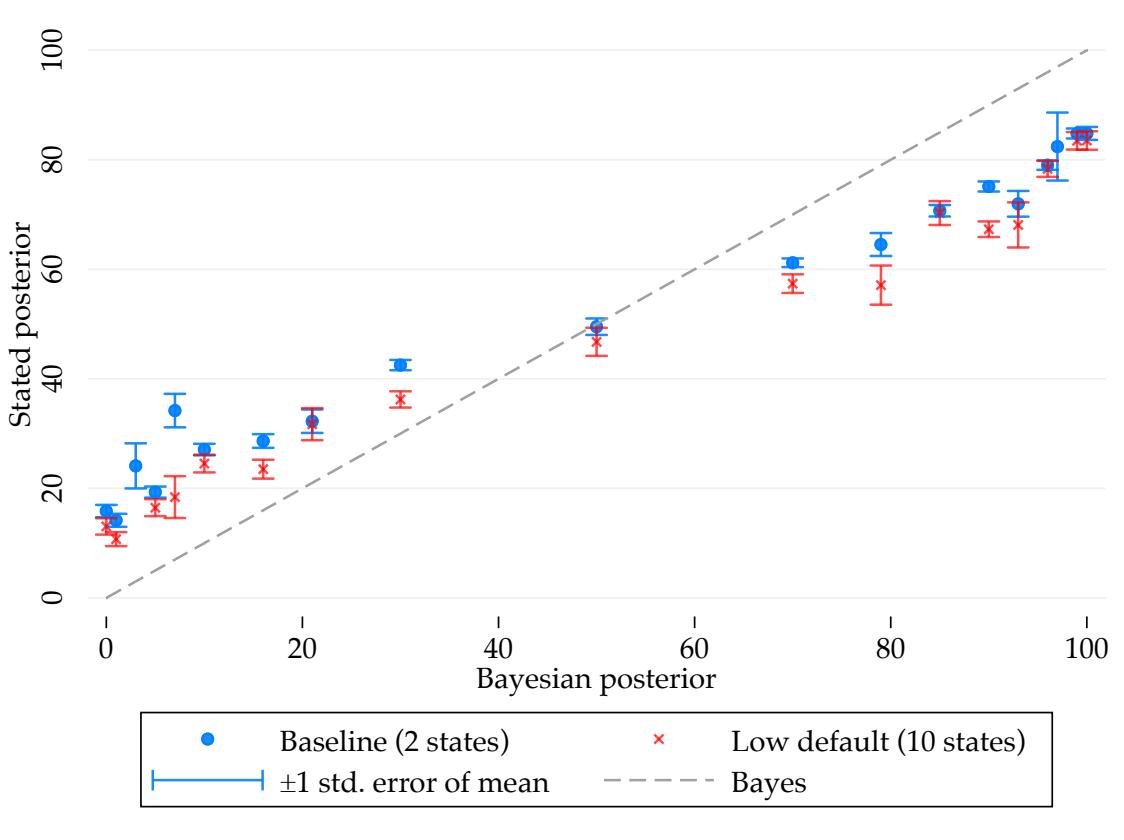


Figure 9: Stated average posteriors as a function of Bayesian posteriors, separately for treatments *Baseline Beliefs* and *Low Default Beliefs*. Bayesian posteriors are rounded to the nearest integer. We only show buckets with more than ten observations. The figure is based on 5,372 beliefs of 1,000 subjects.

5 Survey Expectations

5.1 Experimental Design

5.1.1 Measuring Beliefs

Miscalibrated survey expectations about economic variables have been documented in a wide range of contexts, including beliefs about the macroeconomy, the stock market, or personal life experiences (Hurd, 2009; Manski, 2004). A common theme is the presence of 50:50 answers (Fischhoff and Bruine De Bruin, 1999).¹²

Survey expectations are different from the contexts discussed above in that respondents are not confronted with all problem parameters that they need to provide a well-informed response. Instead, people may have limited information or imperfect memory about the value of (past or future) economic variables. However, our concept of cognitive uncertainty nevertheless applies here: people oftentimes do not know the answer to a

¹²Drerup et al. (2017) suggest that the low explanatory power of survey expectations for economic behavior might reflect that some people do not even hold meaningful belief distributions, which is reminiscent of cognitive uncertainty.

Table 6: Belief updating: Low versus high mental default

	Dependent variable:					
	Posterior belief			Log [Posterior odds]		
	(1)	(2)	(3)	(4)	(5)	(6)
0 if <i>Baseline</i> , 1 if <i>Low Default</i>	-3.75*** (0.71)	-4.12*** (0.69)	-4.26*** (0.74)	-0.25*** (0.05)	-0.27*** (0.04)	-0.28*** (0.05)
Bayesian posterior	0.64*** (0.01)	0.78*** (0.01)	0.78*** (0.01)			
Bayesian posterior × Cognitive uncertainty		-0.54*** (0.03)	-0.54*** (0.03)			
Cognitive uncertainty		24.2*** (1.94)	24.0*** (1.96)		-0.057 (0.06)	-0.062 (0.06)
Log [Likelihood ratio]				0.36*** (0.01)	0.43*** (0.01)	0.43*** (0.01)
Log [Prior odds]				0.48*** (0.02)	0.61*** (0.03)	0.61*** (0.03)
Log [Likelihood ratio] × Cognitive uncertainty					-0.27*** (0.03)	-0.27*** (0.03)
Log [Prior odds] × Cognitive uncertainty					-0.50*** (0.06)	-0.50*** (0.06)
Session FE	No	No	Yes	No	No	Yes
Demographic controls	No	No	Yes	No	No	Yes
Observations	5372	5372	5372	5226	5226	5226
R ²	0.63	0.65	0.65	0.57	0.59	0.59

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

probabilistic question, which may induce them to shrink their reported beliefs to 50:50. To illustrate the link between cognitive uncertainty and survey expectations, we elicit beliefs about three variables that have attracted attention in the literature: the structure of the national income distribution, inflation rates, and the development of the stock market.

To be able to financially incentivize participants without going through the logistical hassle of waiting for future variables to have realized, we elicited beliefs about contemporaneous or past variables. Each participant was asked three questions that elicited their beliefs about some specific aspect of the income distribution, stock returns, and the inflation rate. The question about the income distribution reads as:

Assume that in 2018, we randomly picked a household in the United States. What do you think is the probability that this household earned less than USD y in 2018, before taxes and deductions?

Beliefs about the performance of the stock market were elicited as:

The S&P 500 is an American stock market index that includes 500 of the largest companies based in the United States. We randomly picked a year X between 1980 and 2018. Imagine that someone invested \$100 into the S&P 500 at the beginning of year X . What do you think is the probability that, at the end of that same year, the value of the investment was less than \$ y ? (In other words, what do you think is the probability that the S&P 500 [lost more than $z\%$ of its value / gained less than $z\%$, or decreased in value]?)

Finally, beliefs about the inflation rate were measured as:

The inflation rate in the United States measures the percentage change in the consumer price index, which reflects the price level of a comprehensive set of consumer goods and services purchased by households. The inflation rate in a given time period captures how much more or less expensive goods and services have become on average. We randomly picked a year X between 1980 and 2018. Imagine that, at the beginning of year X , the set of products that is used to compute the inflation rate cost \$100. What do you think is the probability that, at the end of that same year, the same set of products cost less than \$ y ? (In other words, what do you think is the probability that the inflation rate in year X was lower than $z\%?$)

The order of topics was randomized across participants. Across participants, y (and hence z) varies randomly such that the true probability ranges from 0% to 100%. Subjects' beliefs were financially incentivized using the same binarized scoring rule as discussed in Section 4, except that the prize a subject could win was \$2. One of the three questions was randomly selected for payment.

5.1.2 Measuring Cognitive Uncertainty

To measure cognitive uncertainty, we again make use of the same elicitation tool as before. That is, subjects were asked how certain they are that their probabilistic guess is correct. Subjects used a slider to calibrate the statement: “I am certain that the actual probability that [...] is between a and b .”, where a and b collapsed to the subject's own previously indicated guess if the slider was moved to the very right. For each of the 30 possible ticks that the slider was moved to the left, a decreased and b increased by one probability point. a was bounded from below by zero and b bounded from above by 100. Figure 44 in Appendix D.1 shows a screenshot of the elicitation screen.

Figures 45–47 in Appendix D.1 show histograms of cognitive uncertainty for each question type. Overall, cognitive uncertainty is substantial in these contexts, in particular regarding the stock market and inflation rates.

5.1.3 Logistics and Pre-Registration

The elicitation of survey expectations took place with the same set of subjects that completed the choice under risk and belief updating tasks discussed in Sections 3 and 4. Thus, the total sample size is $N = 2,000$. As explained above, one of the three parts of the experiments (choice under risk or belief updating, survey expectations, or Raven matrices test) was randomly selected for payment.

In addition to these “backward-looking” beliefs, in separate pre-registered robustness experiments with $N = 400$ participants, we also elicit expectations about *future* realizations of inflation rates, stock market returns and the income distribution. These questions are conceptually more appropriate in that they ask about the future, but they are not financially incentivized. The results in these robustness experiments are almost identical to the ones that are reported here; we summarize these results in Appendix E.

5.2 Results

As in Section 4, and as pre-registered, we begin by excluding extreme outliers, defined as $p_s < 25 \wedge p_o > 75$ or $p_s > 75 \wedge p_o < 25$, where p_s is the subjective probability and p_o the objective one. This results in the exclusion of 5% of all data.

Figure 10 shows average beliefs as a function of objective probabilities, separately for subjects with above and below average cognitive uncertainty. Again, we see that these “survey belief weighting functions” exhibit an inverse S-shape, yet this pattern is substantially more pronounced for subjects who indicate higher cognitive uncertainty. Table 21 in Appendix D.2 provides a corresponding non-parametric econometric analysis that confirms the statistical significance of this pattern.

6 Parametric Estimations

Estimating the weighting functions. To supplement the non-parametric estimations for Sections 3–5, we now turn to estimating the model in Section 2. The most straightforward way to do so is to estimate Gonzalez and Wu’s two-parameter weighting function that we endogenized in equation (15), separately for decisions that are associated with high and low cognitive uncertainty.

As called for by our framework, we estimate this function not just for choice under uncertainty, but also for belief updating and survey expectations. In doing so, for simplicity, we again abstract away from taste-based risk or ambiguity aversion. While our data strongly suggest the presence of risk aversion, this procedure has the advantage that we estimate exactly the same function across decision domains, hence allowing for

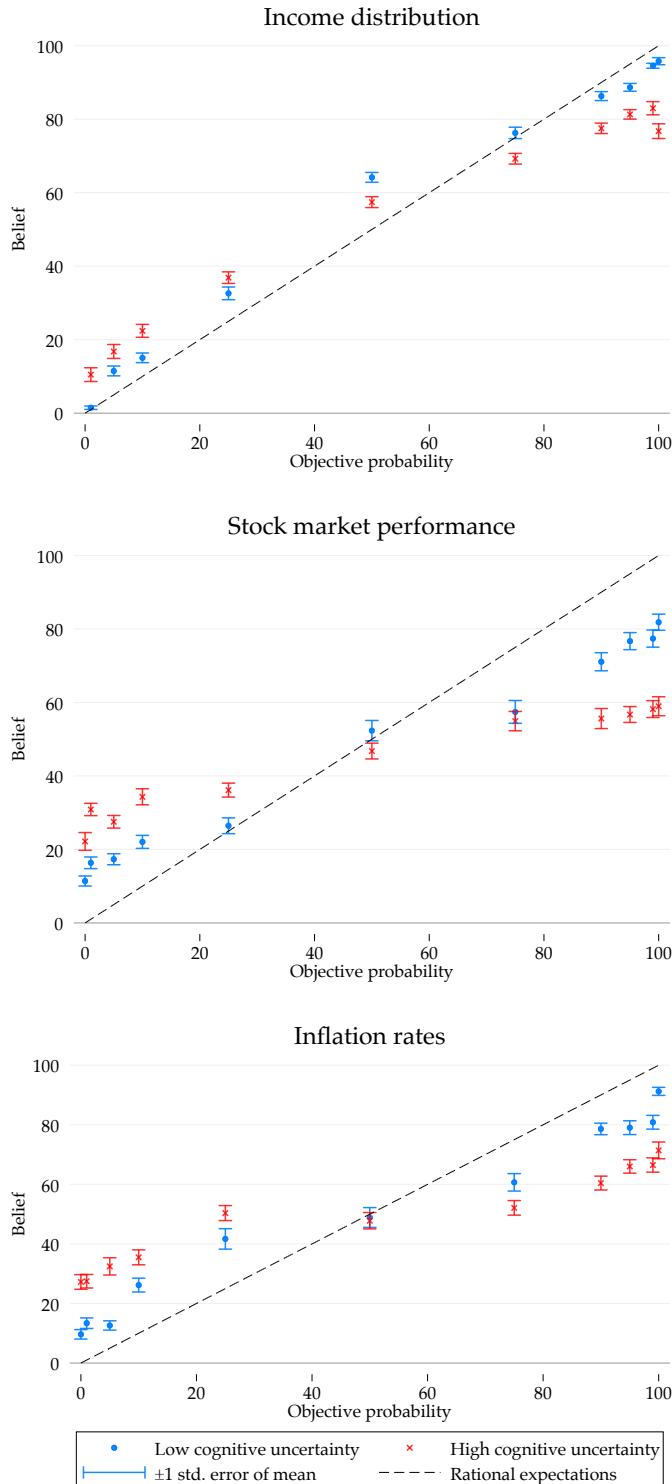


Figure 10: Survey beliefs as a function of objective probabilities, separately for subjects above / below average cognitive uncertainty. The partition is done separately for each probability bucket. In the top panel, the question asks for the probability that a randomly selected U.S. household earns less than \$x ($N = 1,974$). In the middle panel, the question asks for the probability that in a randomly selected year the S&P500 increased by less than x% ($N = 1,887$). In the bottom panel, the question asks for the probability that in a randomly selected year the inflation rate was less than x% ($N = 1,842$).

Table 7: Estimates of equation (15) across decision domains

Treatment / group	Sensitivity $\hat{\lambda}$	Elevation $\hat{\delta}$
<i>Baseline Choice under risk:</i> all observations	0.47 (0.02)	0.68 (0.02)
<i>Baseline Choice under risk:</i> high CU observations	0.33 (0.02)	0.72 (0.02)
<i>Baseline Choice under risk:</i> low CU observations	0.54 (0.02)	0.67 (0.02)
<i>Choice under ambiguity:</i> all observations	0.36 (0.02)	0.69 (0.03)
<i>Choice under ambiguity:</i> high CU observations	0.25 (0.03)	0.70 (0.03)
<i>Choice under ambiguity:</i> low CU observations	0.43 (0.03)	0.67 (0.04)
<i>Baseline Belief updating:</i> all observations	0.50 (0.01)	1.00 (0.02)
<i>Baseline Belief updating:</i> high CU observations	0.38 (0.01)	0.99 (0.02)
<i>Baseline Belief updating:</i> low CU observations	0.60 (0.02)	1.02 (0.02)
<i>Survey expectations pooled:</i> all observations	0.42 (0.01)	0.97 (0.02)
<i>Survey expectations pooled:</i> high CU observations	0.24 (0.01)	0.92 (0.02)
<i>Survey expectations pooled:</i> low CU observations	0.59 (0.01)	1.03 (0.03)

Notes. Estimates of equation (15), standard errors (clustered at subject level) reported in parentheses. CU = cognitive uncertainty (split at average). The samples include the same observations as in all baseline analyses in Sections 3–5.

a comparison of parameter estimates across domains.

Formally, for each decision domain, we estimate the parameters $\hat{\lambda}$ and $\hat{\delta}$ by minimizing the sum of squared residuals for the non-linear equation (15). In choice under risk and ambiguity, the dependent variable is a subject's normalized certainty equivalent, while in belief updating and survey expectations it is a participant's stated belief. Table 7 summarizes the estimates for the baseline conditions in each set of experiments.

In choice under risk, pooled across all subjects, we estimate a sensitivity parameter of $\hat{\lambda} = 0.47$ and an elevation parameter of $\hat{\delta} = 0.68$. These estimates are close to the estimates of $\hat{\lambda} = 0.44$ and $\hat{\delta} = 0.77$ in the classic by Gonzalez and Wu (1999), which again suggests that our sample and results are similar to previous findings. Comparing

subjects with above and below average cognitive uncertainty, we estimate that the sensitivity parameter λ is 58% higher for low than for high cognitive uncertainty subjects. We find similar patterns for choice under ambiguity, though of course a lower estimated sensitivity parameter as per the discussion in Section 3.3.

For belief updating, we again find that $\hat{\lambda}$ is significantly lower for subjects with above-average cognitive uncertainty ($\hat{\lambda} = 0.38$) than for those with below-average cognitive uncertainty ($\hat{\lambda} = 0.60$). Furthermore, the estimate of the insensitivity parameter λ is strikingly similar to that estimated in choice under risk.

Finally, for the survey expectations data, we again estimate fairly similar parameter values. In these analyses, the data are pooled across the three types of expectations. Table 22 in Appendix D.2 reports the estimates for inflation rates, stock market returns and the national income distribution separately.

In summary, the structural estimations deliver sensible parameter estimates that vary in meaningful ways with cognitive uncertainty. Moreover, even though the underlying decision domains and experimental paradigms are very different, we always estimate fairly similar parameter values, in particular regarding the sensitivity parameter λ .

Decision-level analysis. The preceding analysis focused on estimating population-level parameters. Structurally estimating the Gonzalez-Wu weighting functions at the level of individual subjects is infeasible given the small number of decisions in our experiments. However, we can make progress even at the level of individual decisions by restricting attention to the sensitivity parameter λ . In particular, recall that

$$q(a) = \lambda q(s) + (1 - \lambda)q(p^d),$$

where $q(\cdot)$ denotes the logit function. If we now explicitly impose that the mental default is given by 50% (the ignorance prior) and noting that $E[s] = x$, one can back out the sensitivity parameter λ that is implied in each decision using a (very) heuristic back-of-the-envelope calculation:

$$\hat{\lambda} \approx \frac{q(a) - q(.5)}{q(x) - q(.5)} = \frac{q(a)}{q(x)}, \quad (17)$$

where $q(a)$ is a participant's logit response (belief or normalized certainty equivalent) and $q(x)$ the logit objective probability in the respective decision domain.

Our model predicts that this sensitivity parameter λ is decreasing in cognitive uncertainty, see equation (9). Figure 11 provides binscatters of the implied $\hat{\lambda}$ against measured cognitive uncertainty. Despite the fact that our estimates of λ are very heuristic and ignore any form of decision noise or non-linear utility, we find the expected relationship: the implied λ are always negatively correlated with stated cognitive uncertainty.

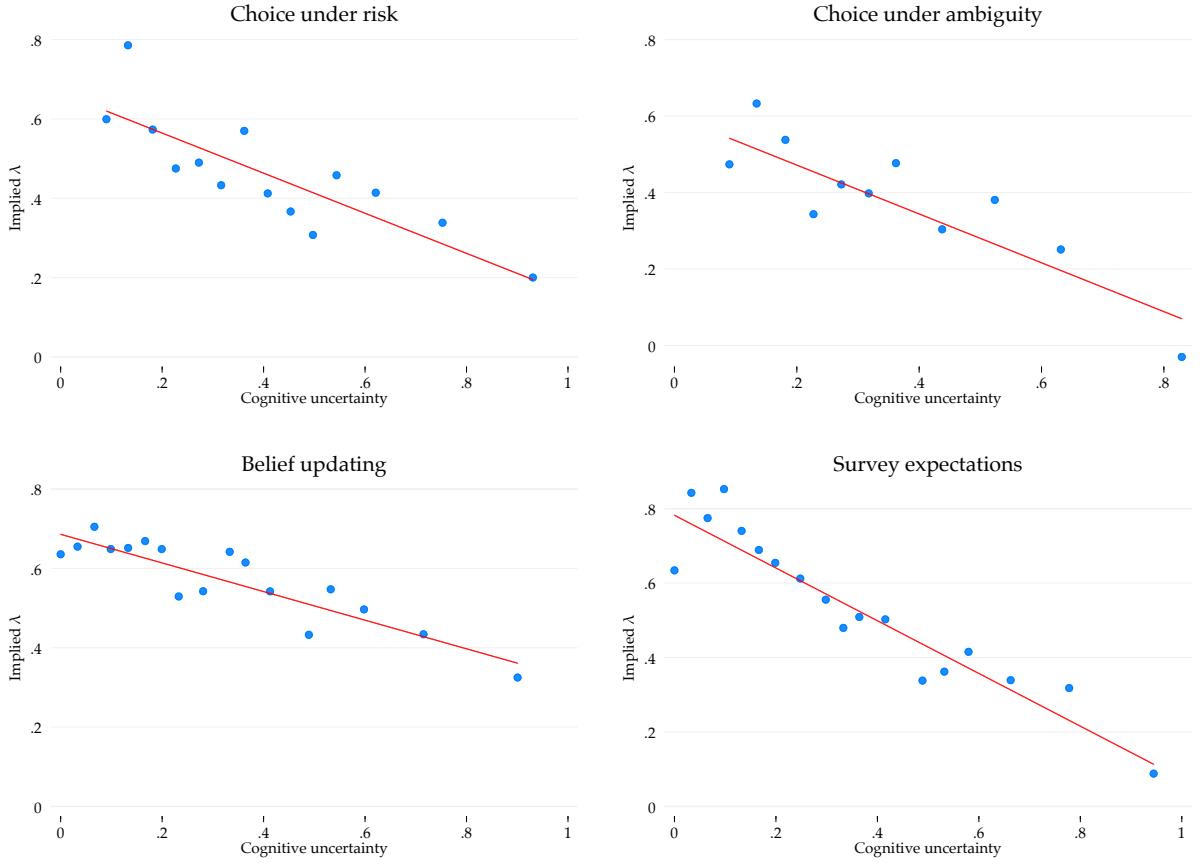


Figure 11: Correlations between implied decision-level sensitivity parameter (shrinkage factor) λ and decision-level cognitive uncertainty. The estimate of λ is computed using equation (17). The plots represent binscatters. The top left panel includes decisions from treatment *Baseline Risk* and the top right panel decisions from the treatment with ambiguous lotteries. The bottom left panel includes beliefs from treatment *Baseline Beliefs* and the bottom right panel all survey expectations. In each panel, we exclude observations with implied $\hat{\lambda} < -1$ or $\hat{\lambda} > 2$, which roughly corresponds to the 1st and 99th percentiles across decision domains.

The correlations are $\rho = -0.18$ for choice under risk, $\rho = -0.24$ for choice under ambiguity, $\rho = -0.19$ for belief updating, and $\rho = -0.33$ for survey expectations, all of which are statistically significant at the 1% level.

7 Heterogeneity in Cognitive Uncertainty

In the final part of the paper, we shed some more light on the variation in cognitive uncertainty. From an ex ante perspective, there are two plausible accounts. First, it is conceivable that individuals exhibit reasonably stable cognitive uncertainty “types,” so that a large part of the variation in cognitive uncertainty is between rather than within subjects. Second, it is conceivable that cognitive uncertainty varies dramatically across tasks, with little evidence for consistent types at the individual level.

To decompose cognitive uncertainty into between- and within-subject variation, we separately look at the choice under risk, belief updating, and survey expectations data (we exclude choice under ambiguity here because we only have 200 subjects, much less than for the other experiments). Within each dataset, we regress the collection of cognitive uncertainty statements on subject fixed effects.¹³ We find that the variance explained is 44% in choice under risk, 53% in belief updating, and 60% in survey expectations. It is worth pointing out that these numbers represent fairly weak lower bounds for the fraction of the true variation that is due to between-subject variation, as all measurement error gets soaked up by the residual and hence by “within-subject variation.”

An additional way to investigate the existence of types is to look at subjects’ consistency in cognitive uncertainty *across task domains*. Recall that each subject completed the survey expectations tasks and additionally either the risky choice or the belief updating experiments. We now compute the correlation between average subject-level cognitive uncertainty in choice under risk, belief updating and survey expectations for all subjects in the respective baseline tasks. The correlation between average cognitive uncertainty in belief updating and average cognitive uncertainty in survey expectations is $\rho = 0.57$. The correlation between cognitive uncertainty in risky choice and survey expectations is $\rho = 0.35$. Thus, subjects exhibit a reasonable degree of consistency even across different types of experiments. We conclude from these analyses that cognitive uncertainty varies in meaningful and reasonable stable ways across participants.

Next, we investigate correlates of this variation in cognitive uncertainty across individuals. For this purpose, we relate subjects’ cognitive uncertainty in each of the three decision domains to a vector of individual characteristics: the score on an eight-item Raven matrices IQ test as a proxy for cognitive skills, educational attainment, response times as a proxy for cognitive effort, gender, and age. All of these correlational analyses except for the one with response times were pre-registered. Figure 12 summarizes the results, separately for each of the types of experiments reported above. The figure shows the results of different regressions, each of which relates average (subject-level) cognitive uncertainty to a different subject-level variable, controlling for treatment fixed effects. The figure reports standardized beta coefficients, so that the y-axis shows the percent change in cognitive uncertainty that is associated with a 1% increase in the explanatory variable of interest. While the results are mixed overall, the strongest and most consistent correlations reflect that women, people who take less time to complete the experiment, and people with lower IQ test scores report higher cognitive uncertainty.

This correlational evidence suggests that the availability of cognitive resources – cognitive skills and cognitive effort – may reduce cognitive uncertainty. This raises the question whether an increase in financial incentives would increase effort and hence reduce

¹³Treatment fixed effects always explain less than 1% of the variation in the data.

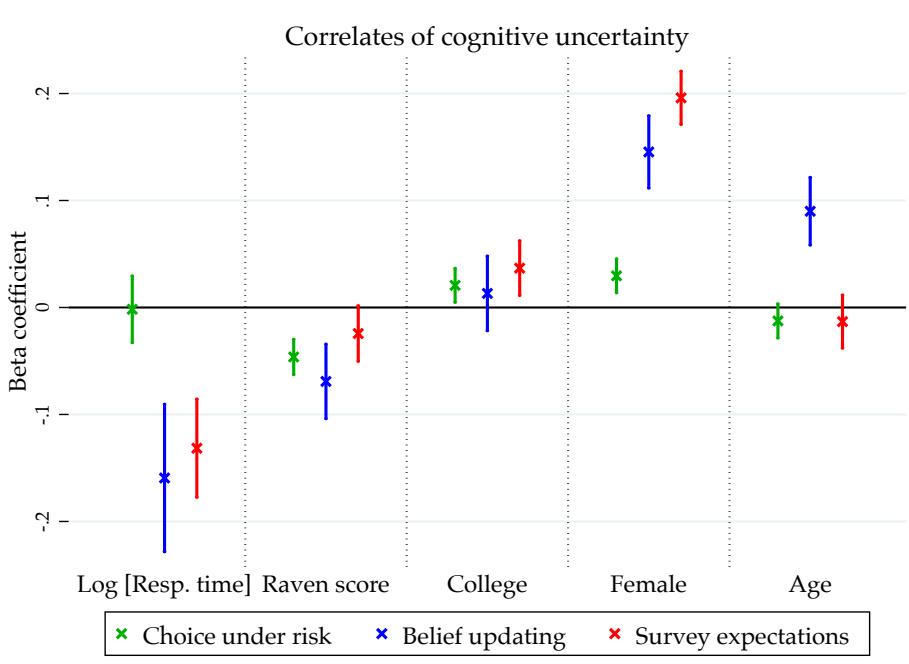


Figure 12: Correlates of average cognitive uncertainty. The figure shows the standardized beta coefficients of regressions of a subject's average cognitive uncertainty on different variables, controlling for treatment fixed effects. The values on the y-axis show the percent change in cognitive uncertainty that is associated with a 1% increase in the explanatory variable of interest. The beta coefficients are estimated conditional on treatment fixed effects. Response times are computed as total completion time within the relevant part of the experiment. $N = 1,000$ observations for choice under risk and belief updating and $N = 2,000$ observations for survey expectations.

cognitive uncertainty. To test this, we ran a final set of experiments on choice under risk ($N = 150$) and belief updating ($N = 150$). These experiments followed the same procedure as outlined in Sections 3 and 4, except that across the six tasks the probability of being payout-relevant varied within subjects. For five tasks, the probability that a task would determine payment was 1% and for one task it was 95%. In both sets of experiments, we find that subjects' response times are about 20–30% higher under higher incentives. Cognitive uncertainty, on the other hand, either remains unaffected (choice under risk) or decreases only marginally (belief updating). See Appendix G for details.

8 Robustness Checks

Additional pre-registered analyses. The pre-registration specified that we will conduct our analyses on three different samples: (i) excluding extreme outliers, as done thus far; (ii) using all data; and (iii) excluding “speeders,” defined as subjects in the bottom decile of the response time distribution. Appendices B.2 and B.3 reproduce the analysis of choice under risk on the full sample and excluding speeders. Appendices C.3 and C.4 provide analogous analyses for belief updating and Appendices D.3 and D.4 for

survey expectations. The results are always very similar. Two minor exceptions are that (i) the treatment difference between *Low Default Risk* and *High Default Risk* is statistically significant also for losses when we exclude speeders (the p-value was $p = 0.15$ in the baseline analysis above) and (ii) the interaction between cognitive uncertainty and payout probability in treatment *Baseline Risk* is marginally not significant for losses.

Censoring. In all our experiments, the use of a bounded response scale can lead to censoring of both the choice or belief that a subject states and the range of cognitive uncertainty indicated using the slider. This may affect the observed relationship between actions and cognitive uncertainty in two ways. First, choices and beliefs may be influenced by boundary effects. Assume, for example, that a subject in a belief updating task wants to state a true posterior belief of 95%. However, some form of decision noise such as trembling when submitting a response leads her to instead indicate a posterior belief that is uniformly drawn from within $+/- 10\%$ of her true posterior, i.e., she would end up with any posterior between 85% and 105% with equal probability. Since it is not possible to state a posterior greater than 100%, she will state 100% whenever she would like to state something greater than 100%, leading to an observed posterior that is *lower* than 95% in expectation. Importantly, this distortion in observed beliefs away from the boundary will be greater for someone with greater decision noise. If subjects' cognitive uncertainty statements then accurately reflect the amount of trembling, i.e., the length of the trembling interval in this case, this form of censoring will mechanically generate a positive relationship between the extent of cognitive uncertainty and shrinkage. We find that the actual amount of bunching at the upper and lower bounds of the response scales, however, is small: it is 4.28% of observations in choice under risk, 2.6% of observations in belief updating, and 6.61% in survey expectations. In Appendix I.1, we show that the observed relationship between cognitive uncertainty and choices or beliefs is virtually unaffected when excluding these observations.

Second, censoring might occur when choosing an interval on the response scale to indicate cognitive uncertainty. While the interval increases symmetrically when moving the slider to the left, it increases asymmetrically once it hits one of the response scale boundaries. One may think that subjects stop moving the slider to the left once they hit a boundary. This implies that measured cognitive uncertainty tends to be smaller for responses that are closer to a boundary, again leading to a mechanic relationship between observed cognitive uncertainty and the amount of shrinkage. In our data, we find that 23.25% of cognitive uncertainty intervals in choice under risk, 25.93% in belief updating and 16.66% in survey expectations are censored at one of the boundaries. However, as we exclude those observations, the relationship between cognitive uncertainty and choices or beliefs persists in the same way as before (see Appendix I.2), showing that

our findings are not an artifact of censoring due to bounded response scales.

9 Conclusion

This paper has formally defined and introduced an experimental measurement of *cognitive uncertainty*: people’s subjective uncertainty about the rational solution to a decision problem. As we have documented using belief updating and survey expectations data, such cognitive uncertainty does not just reflect preference uncertainty but captures a more general uncertainty about how to behave optimally.

Based on a simple formal framework that draws from existing theories, we have argued that cognitive uncertainty induces people to shrink probabilities towards a simple ignorance prior. This idea both reconciles existing evidence and makes new predictions. To argue our case, the paper has brought together decision tasks on choice under risk and ambiguity, belief updating, and survey expectations, all of which generate a stylized pattern of inverse S-shaped response functions. Across all of these perhaps seemingly-unrelated decision domains, participants with higher cognitive uncertainty exhibit more strongly compressed response functions. Moreover, in an attempt to provide causal evidence for our framework, we have exogenously manipulated both the magnitude of cognitive uncertainty and the location of the ignorance prior, and have identified predictable changes in subjects’ beliefs and behaviors in response to these treatments.

We believe that the concept of cognitive uncertainty is likely to be important also outside of the domains that we study in this paper. By providing a simple and portable experimental tool that allows to measure cognitive uncertainty in a quantitative fashion, our paper opens up the possibility for future experimental work on the relationship between cognitive uncertainty and economic decision-making.

References

- Agranov, Marina and Pietro Ortoleva**, “Stochastic choice and preferences for randomization,” *Journal of Political Economy*, 2017, 125 (1), 40–68.
- Ambuehl, Sandro and Shengwu Li**, “Belief updating and the demand for information,” *Games and Economic Behavior*, 2018, 109, 21–39.
- Andreoni, James and Michael A Kuhn**, “Is It Safe to Measure Risk Preferences? A Comparison of Four Techniques,” *Working Paper*, 2019.
- Benjamin, Daniel J**, “Errors in Probabilistic Reasoning and Judgmental Biases,” in “Handbook of Behavioral Economics” 2019.
- Bernheim, B Douglas and Charles Sprenger**, “Direct tests of cumulative prospect theory,” Available at SSRN 3350196, 2019.
- Bordalo, Pedro, Nicola Gennaioli, and Andrei Shleifer**, “Salience theory of choice under risk,” *Quarterly Journal of Economics*, 2012, 127 (3), 1243–1285.
- , — , and — , “Memory, Attention, and Choice,” *Working Paper*, 2017.
- Bruhin, Adrian, Helga Fehr-Duda, and Thomas Epper**, “Risk and rationality: Uncovering heterogeneity in probability distortion,” *Econometrica*, 2010, 78 (4), 1375–1412.
- Butler, David J and Graham C Loomes**, “Imprecision as an account of the preference reversal phenomenon,” *American Economic Review*, 2007, 97 (1), 277–297.
- da Silveira, Rava Azeredo and Michael Woodford**, “Noisy Memory and Over-Reaction to News,” in “AEA Papers and Proceedings,” Vol. 109 2019, pp. 557–61.
- DellaVigna, Stefano and Devin Pope**, “Predicting experimental results: who knows what?,” *Journal of Political Economy*, 2018, 126 (6), 2410–2456.
- Dimmock, Stephen G, Roy Kouwenberg, and Peter P Wakker**, “Ambiguity attitudes in a large representative sample,” *Management Science*, 2015, 62 (5), 1363–1380.
- Drerup, Tilman, Benjamin Enke, and Hans-Martin Von Gaudecker**, “The precision of subjective data and the explanatory power of economic models,” *Journal of Econometrics*, 2017, 200 (2), 378–389.
- El-Gamal, Mahmoud A. and David M. Grether**, “Are People Bayesian? Uncovering Behavioral Strategies,” *Journal of the American statistical Association*, 1995, 90 (432), 1137–1145.

- Enke, Benjamin**, “What You See Is All There Is,” *Working Paper*, 2017.
- and **Florian Zimmermann**, “Correlation Neglect in Belief Formation,” *Review of Economic Studies*, 2019, 86 (1), 313–332.
 - , **Frederik Schwerter, and Florian Zimmermann**, “Associative Memory and Belief Formation,” *Working Paper*, 2019.
- Esponda, Ignacio and Emanuel Vespa**, “Hypothetical Thinking: Revisiting Classic Anomalies in the Laboratory,” *Working Paper*, 2016.
- Fischhoff, Baruch and Wändi Bruine De Bruin**, “Fifty-fifty= 50%?,” *Journal of Behavioral Decision Making*, 1999, 12 (2), 149–163.
- Frydman, Cary and Lawrence Jin**, “Efficient Coding and Risky Choice,” *Working Paper*, 2019.
- Gabaix, Xavier**, “Behavioral inattention,” in Douglas Bernheim, Stefano DellaVigna, and David Laibson, eds., *Handbook of Behavioral Economics-Foundations and Applications* 2, Elsevier, 2019, p. 261.
- and **David Laibson**, “Myopia and discounting,” Technical Report, National bureau of economic research 2017.
- Gershman, Samuel J and Rahul Bhui**, “Rationally inattentive intertemporal choice,” *bioRxiv*, 2019, p. 680652.
- Gillen, Ben, Erik Snowberg, and Leeat Yariv**, “Experimenting with Measurement Error: Techniques with Applications to the Caltech Cohort Study,” *Journal of Political Economy*, 2019.
- Gonzalez, Richard and George Wu**, “On the shape of the probability weighting function,” *Cognitive psychology*, 1999, 38 (1), 129–166.
- Graeber, Thomas**, “Inattentive Inference,” *Working Paper*, 2019.
- Grether, David M**, “Testing Bayes rule and the representativeness heuristic: Some experimental evidence,” *Journal of Economic Behavior & Organization*, 1992, 17 (1), 31–57.
- Halevy, Yoram**, “Ellsberg revisited: An experimental study,” *Econometrica*, 2007, 75 (2), 503–536.
- Hossain, Tanjim and Ryo Okui**, “The binarized scoring rule,” *Review of Economic Studies*, 2013, 80 (3), 984–1001.

- Hurd, Michael D**, “Subjective probabilities in household surveys,” *Annu. Rev. Econ.*, 2009, 1 (1), 543–562.
- Imas, Alex, Michael Kuhn, and Vera Mironova**, “Waiting to choose,” *Working Paper*, 2016.
- Izard, Véronique and Stanislas Dehaene**, “Calibrating the mental number line,” *Cognition*, 2008, 106 (3), 1221–1247.
- Kahneman, Daniel and Amos Tversky**, “Prospect Theory: An Analysis of Decision under Risk,” *Econometrica*, 1979, 47 (2), 263–292.
- Khaw, Mel Win, Ziang Li, and Michael Woodford**, “Risk aversion as a perceptual bias,” Technical Report, National Bureau of Economic Research 2017.
- Li, Chen, Uyanga Turmunkh, and Peter P Wakker**, “Social and Strategic Ambiguity versus Betrayal Aversion,” *Working Paper*, 2019.
- Liang, Yucheng**, “Learning from Unknown Information Sources,” *Working Paper*, 2019.
- Manski, Charles F**, “Measuring expectations,” *Econometrica*, 2004, 72 (5), 1329–1376.
- Martínez-Marquina, Alejandro, Muriel Niederle, and Emanuel Vespa**, “Failures in Contingent Reasoning: The Role of Uncertainty,” *American Economic Review*, 2019.
- O'Donoghue, Ted and Jason Somerville**, “Modeling risk aversion in economics,” *Journal of Economic Perspectives*, 2018, 32 (2), 91–114.
- Paolacci, Gabriele and Jesse Chandler**, “Inside the Turk: Understanding Mechanical Turk as a participant pool,” *Current Directions in Psychological Science*, 2014, 23 (3), 184–188.
- Robinson, Jonathan, Cheskie Rosenzweig, Aaron J Moss, and Leib Litman**, “Tapped Out or Barely Tapped? Recommendations for How to Harness the Vast and Largely Unused Potential of the Mechanical Turk Participant Pool,” 2019.
- Steiner, Jakub and Colin Stewart**, “Perceiving prospects properly,” *American Economic Review*, 2016, 106 (7), 1601–31.
- Trautmann, Stefan T and Gijs Van De Kuilen**, “Ambiguity attitudes,” *The Wiley Blackwell handbook of judgment and decision making*, 2015, 1, 89–116.
- Tversky, Amos and Daniel Kahneman**, “Advances in prospect theory: Cumulative representation of uncertainty,” *Journal of Risk and uncertainty*, 1992, 5 (4), 297–323.

Woodford, Michael, “Prospect theory as efficient perceptual distortion,” *American Economic Review*, 2012, 102 (3), 41–46.

— , “Modeling Imprecision in Perception, Valuation and Choice,” *Working Paper*, 2019.

Zhang, Hang and Laurence T Maloney, “Ubiquitous log odds: a common representation of probability and frequency distortion in perception, action, and cognition,” *Frontiers in neuroscience*, 2012, 6, 1.

ONLINE APPENDIX

A Model Extensions and Illustrations

A.1 Nonlinear Version

We now allow the rational action to be a nonlinear function of x , so that

$$a^r = A(x). \quad (18)$$

We make the simplifying assumptions that, first, the agent still chooses an action based on the posterior expectation about x , as has been done in prior literature (Gabaix, 2019),

$$a(s) = A(\mathbb{E}[x|s]), \quad (19)$$

second, that the function A is strictly monotone, such that it can again be identified from the median action a^e ,

$$a^e(x) = \text{Median}(a(s)|x) = A(\lambda x + (1 - \lambda)x^d), \quad (20)$$

and third, that $x^d = 0$, which is merely a notational simplification. In our empirical applications we will slightly deviate from this and elicit a different type of interval that is wider than the interquartile range, but we here stick to the notation of cognitive uncertainty as denoting one perceived standard deviation around the action for simplicity.

We define cognitive uncertainty analogously to (8) as the agent's perceived uncertainty about his rational action,

$$\sigma_{CU}(x) = \left| A\left(\lambda x + \frac{1}{2}\sqrt{1-\lambda}\sigma_x\right) - A\left(\lambda x - \frac{1}{2}\sqrt{1-\lambda}\sigma_x\right) \right|. \quad (21)$$

At the median, using $a^e(x) = A(\lambda x)$ yields

$$\sigma_{CU}(x) = \left| a^e\left(x + \frac{1}{2}\frac{\sqrt{1-\lambda}}{\lambda}\sigma_x\right) - a^e\left(x - \frac{1}{2}\frac{\sqrt{1-\lambda}}{\lambda}\sigma_x\right) \right|. \quad (22)$$

A Taylor expansion of (22) gives

$$\sigma_{CU} = |a^{e'}(x)| \frac{\sqrt{1-\lambda}}{\lambda} \sigma_x. \quad (23)$$

which is the nonlinear equivalent of equation (8):

$$\frac{\lambda}{\sqrt{1-\lambda}} = \frac{|a^{e'}(x)| \sigma_x}{\sigma_{CU}}. \quad (24)$$

A.2 Illustration of Model Predictions

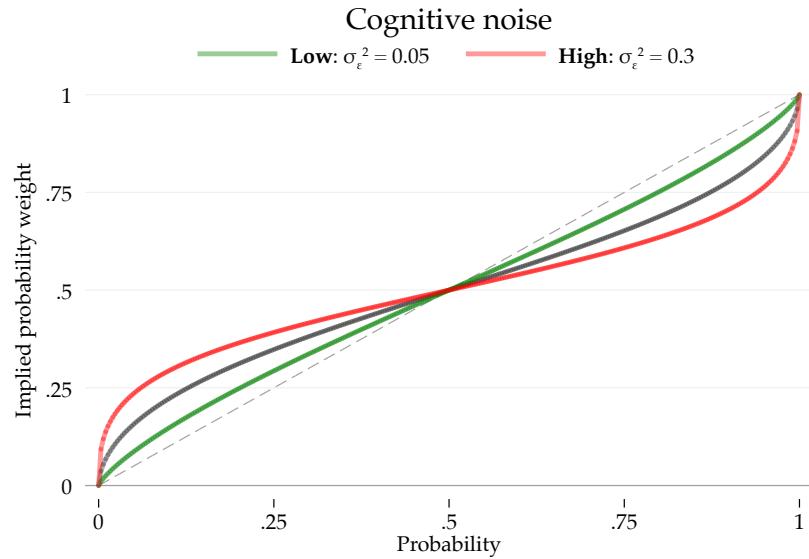


Figure 13: Illustration of model predictions 1 and 2

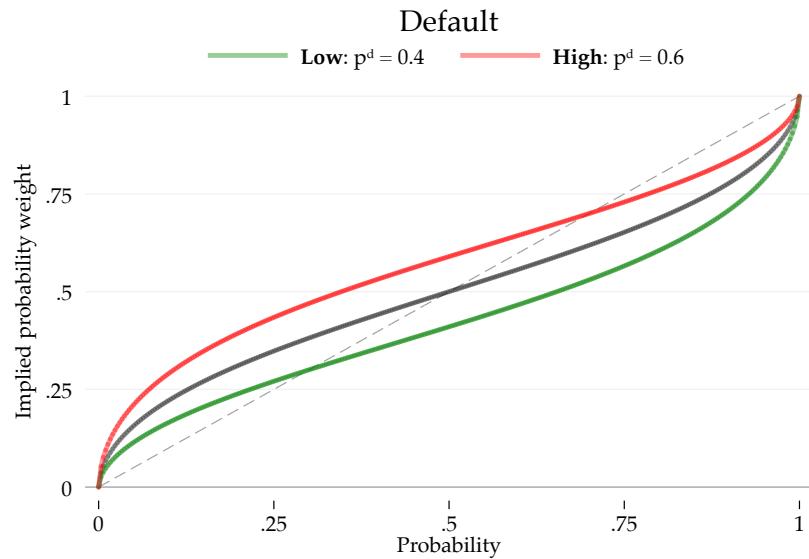


Figure 14: Illustration of model prediction 3

B Additional Details and Analyses for Choice under Risk Experiments

B.1 Additional Figures

Decision screen 1

Option A		Option B
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 0
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 1
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 2
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 3
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 4
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 5
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 6
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 7
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 8
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 9
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 10
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 11
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 12
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 13
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 14
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 15
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 16
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 17
	<input type="radio"/> <input checked="" type="radio"/>	With certainty: Get \$ 18
	<input type="radio"/> <input checked="" type="radio"/>	With certainty: Get \$ 19
	<input type="radio"/> <input checked="" type="radio"/>	With certainty: Get \$ 20

[Next](#)

Figure 15: Decision screen to elicit certainty equivalents for lotteries

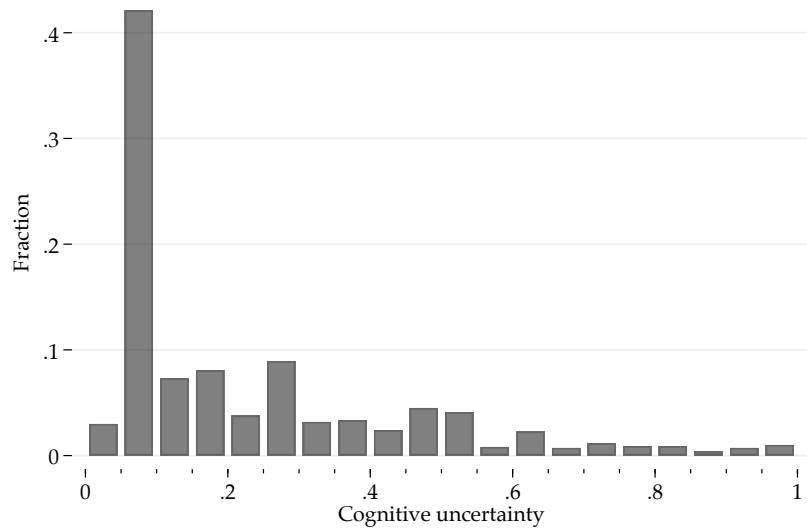


Figure 16: Histogram of cognitive uncertainty in baseline choice under risk tasks

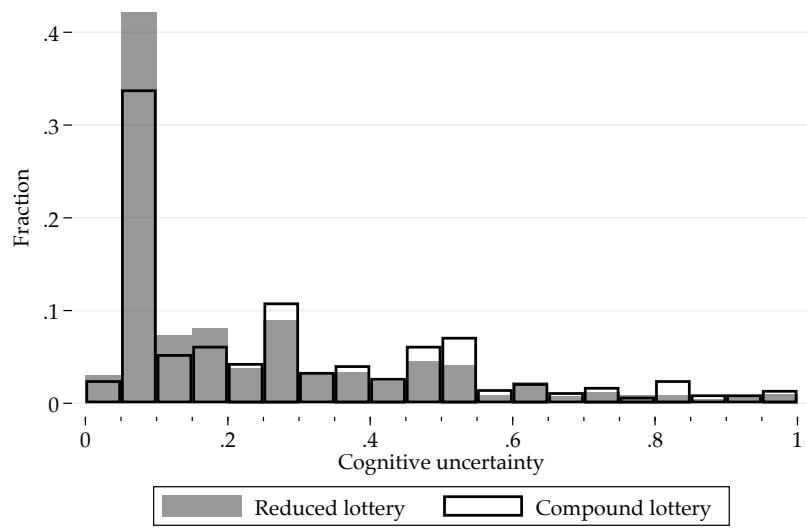


Figure 17: Histograms of cognitive uncertainty in choice under risk tasks, separately for reduced and compound lotteries

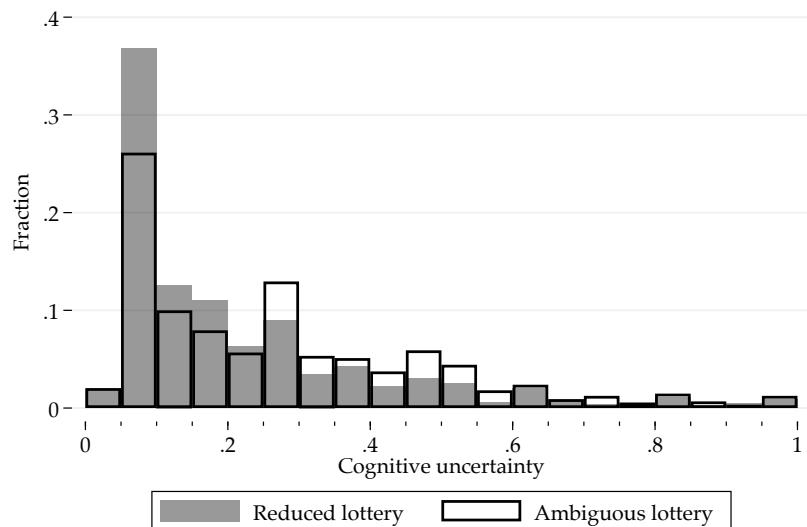


Figure 18: Histograms of cognitive uncertainty in choice under risk tasks, separately for reduced and ambiguous lotteries

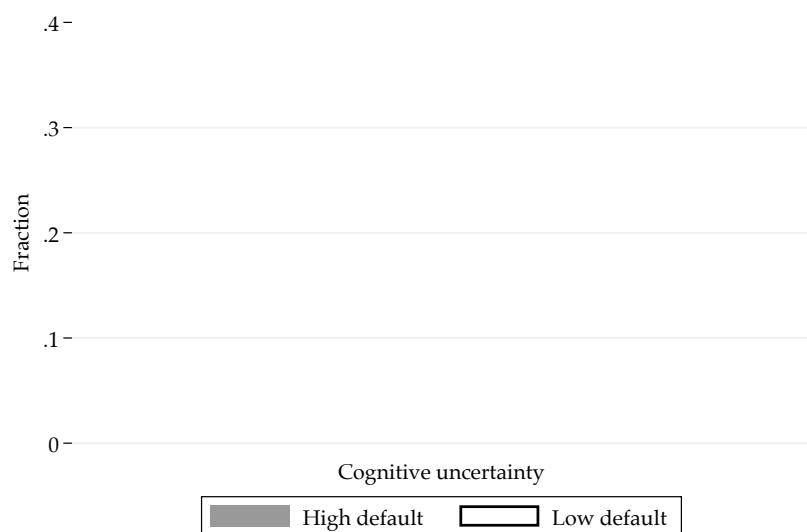


Figure 19: Histograms of cognitive uncertainty in choice under risk tasks, separately for treatments *High Default Risk* and *Low Default Risk*.

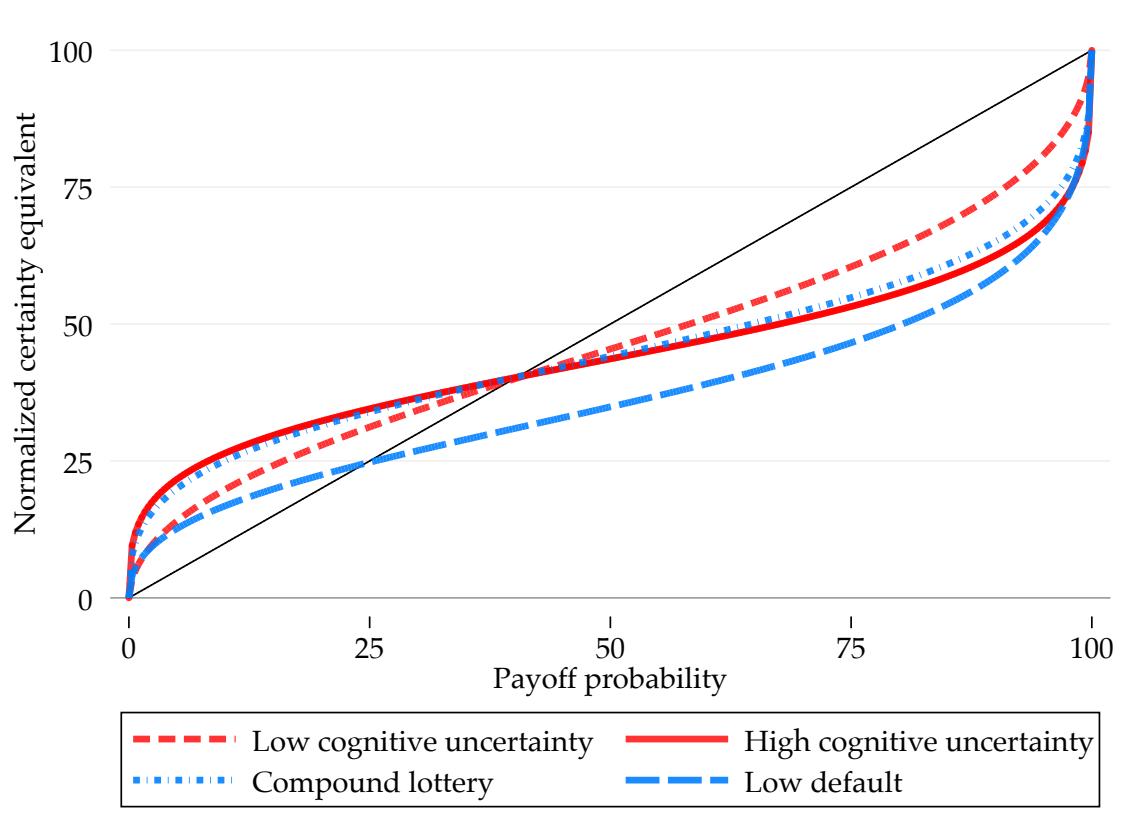


Figure 20: Estimated probability weighting functions across treatments and groups of subjects.

B.2 Results with Full Sample

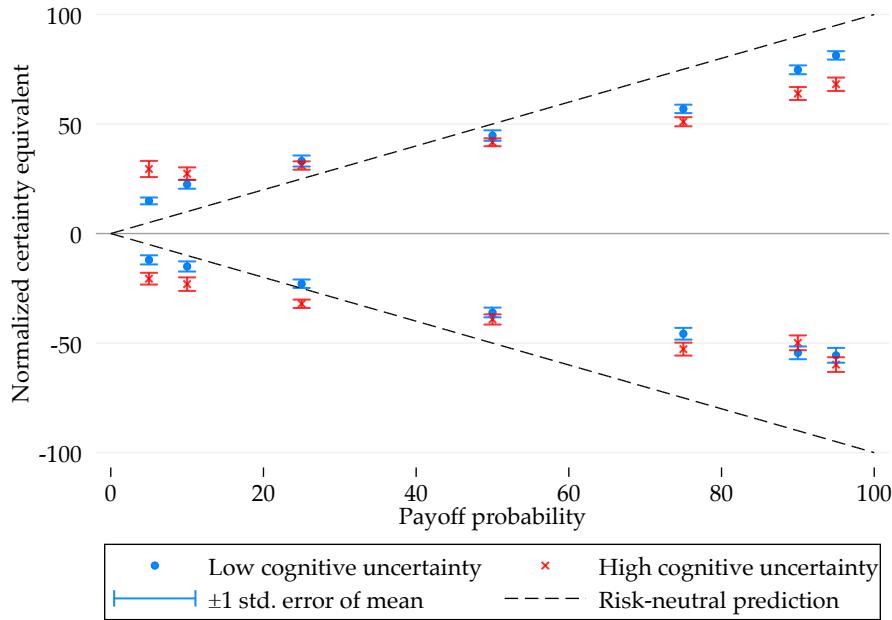


Figure 21: Probability weighting function separately for subjects above / below average cognitive uncertainty (full sample). The partition is done separately for each probability \times gains / losses bucket. The plot shows averages and corresponding standard error bars. Normalized equivalents (implied probability weights) are computed as certainty equivalent divided by payout probability. The figure is based on 2,601 certainty equivalents of 700 subjects.

Table 8: Insensitivity to probability and cognitive uncertainty (full sample)

	Dependent variable: Absolute normalized certainty equivalent					
	Gains		Losses		Pooled	
	(1)	(2)	(3)	(4)	(5)	(6)
Probability of payout	0.74*** (0.03)	0.73*** (0.03)	0.51*** (0.04)	0.51*** (0.04)	0.63*** (0.03)	0.63*** (0.03)
Probability of payout \times Cognitive uncertainty	-0.68*** (0.10)	-0.68*** (0.10)	-0.29*** (0.10)	-0.28*** (0.10)	-0.49*** (0.07)	-0.49*** (0.07)
Cognitive uncertainty	25.1*** (6.17)	25.3*** (6.17)	28.5*** (5.68)	27.7*** (5.75)	27.6*** (4.35)	28.0*** (4.37)
Session FE	No	Yes	No	Yes	No	Yes
Demographic controls	No	Yes	No	Yes	No	Yes
Observations	1286	1286	1315	1315	2601	2601
R ²	0.49	0.50	0.28	0.29	0.36	0.36

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. The dependent variable is a subject's normalized certainty equivalent, computed as midpoint of the switching interval divided by the non-zero payout. The sample includes choices from all baseline gambles with strictly interior probabilities. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

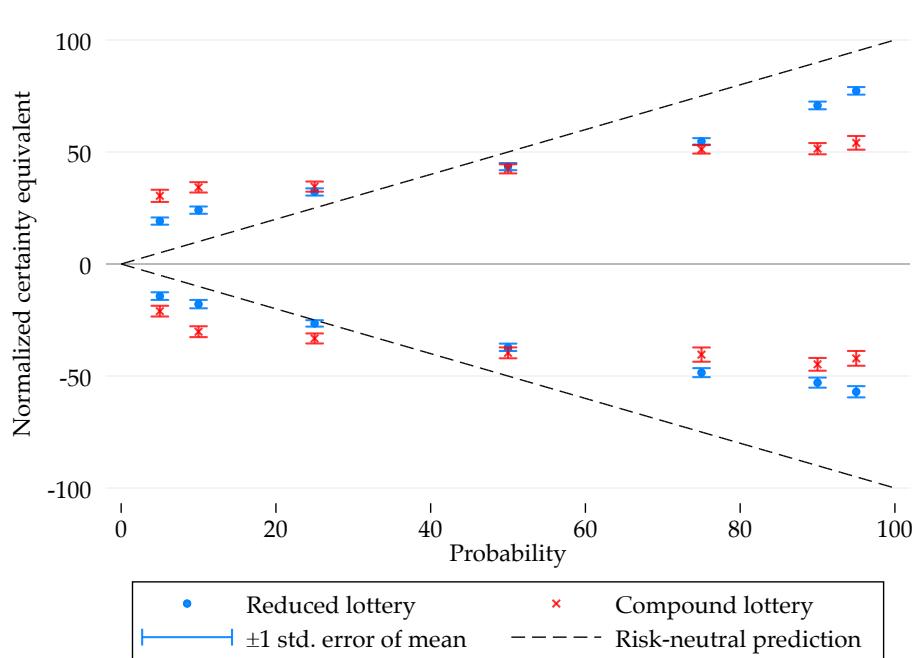


Figure 22: Probability weighting function separately for reduced and compound lotteries (full sample). The plot shows averages and corresponding standard error bars. Normalized certainty equivalents (implied probability weights) are computed as certainty equivalent divided by payout probability. The figure is based on 3,905 certainty equivalents of 700 subjects.

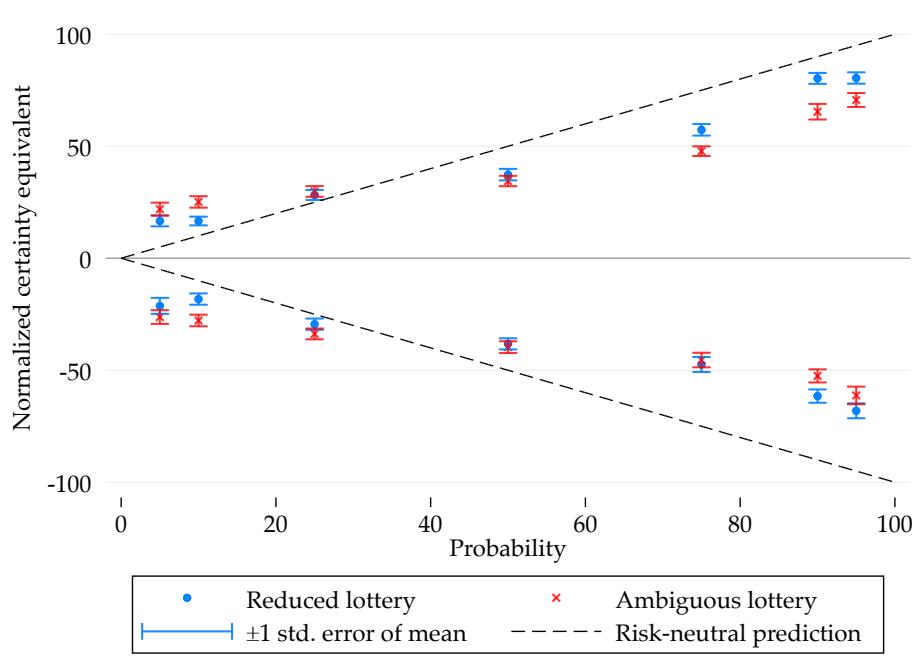


Figure 23: “Probability” weighting function separately for reduced and ambiguous lotteries (full sample). The plot shows averages and corresponding standard error bars. Normalized certainty equivalents (implied probability weights) are computed as certainty equivalent divided by payout probability. The figure is based on 1,800 certainty equivalents of 300 subjects.

Table 9: Choice under risk: Baseline versus compound lotteries (full sample)

	Dependent variable: Absolute normalized certainty equivalent					
	Gains		Losses		Pooled	
	(1)	(2)	(3)	(4)	(5)	(6)
Probability of payout	0.59*** (0.02)	0.70*** (0.03)	0.45*** (0.02)	0.50*** (0.03)	0.52*** (0.02)	0.61*** (0.02)
Probability of payout \times 1 if compound lottery	-0.34*** (0.04)	-0.31*** (0.04)	-0.25*** (0.04)	-0.23*** (0.04)	-0.29*** (0.03)	-0.27*** (0.03)
Probability of payout \times Cognitive uncertainty		-0.49*** (0.08)		-0.24*** (0.08)		-0.38*** (0.06)
1 if compound lottery	13.6*** (2.09)	12.6*** (2.08)	12.3*** (1.98)	10.5*** (1.99)	12.9*** (1.46)	11.7*** (1.44)
Cognitive uncertainty		19.3*** (4.87)		24.8*** (4.67)		22.9*** (3.70)
Session FE	No	Yes	No	Yes	No	Yes
Demographic controls	No	Yes	No	Yes	No	Yes
Observations	1958	1958	1947	1947	3905	3905
R ²	0.37	0.40	0.21	0.24	0.28	0.30

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. The dependent variable is a subject's normalized certainty equivalent, computed as midpoint of the switching interval divided by the non-zero payout. The sample includes choices from the baseline and compound gambles, where for comparability the set of baseline gambles is restricted to gambles with payout probabilities of 10%, 25%, 50%, 75%, and 90%, see Figure 4. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 10: Choice under risk: Treatments *Low Default* and *High Default* (full sample)

	Dependent variable: Absolute normalized certainty equivalent					
	Gains		Losses		Pooled	
	(1)	(2)	(3)	(4)	(5)	(6)
0 if <i>High Default</i> , 1 if <i>Low Default</i>	-12.1*** (1.93)	-11.5*** (1.98)	-2.96 (2.24)	-2.63 (2.22)	-7.53*** (1.58)	-7.05*** (1.58)
Probability of payout	0.59*** (0.05)	0.60*** (0.05)	0.54*** (0.05)	0.54*** (0.05)	0.57*** (0.04)	0.57*** (0.04)
Probability of payout \times Cognitive uncertainty	-0.53*** (0.12)	-0.53*** (0.12)	-0.38*** (0.13)	-0.39*** (0.14)	-0.43*** (0.09)	-0.45*** (0.10)
Cognitive uncertainty	23.2*** (6.00)	23.1*** (6.01)	39.7*** (7.64)	39.6*** (7.90)	31.2*** (5.14)	31.1*** (5.26)
Session FE	No	Yes	No	Yes	No	Yes
Demographic controls	No	Yes	No	Yes	No	Yes
Observations	900	900	900	900	1800	1800
R ²	0.34	0.35	0.25	0.27	0.27	0.29

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. The dependent variable is a subject's normalized certainty equivalent, computed as midpoint of the switching interval divided by the non-zero payout. The sample includes choices from treatments *Low Default* and *High Default*. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

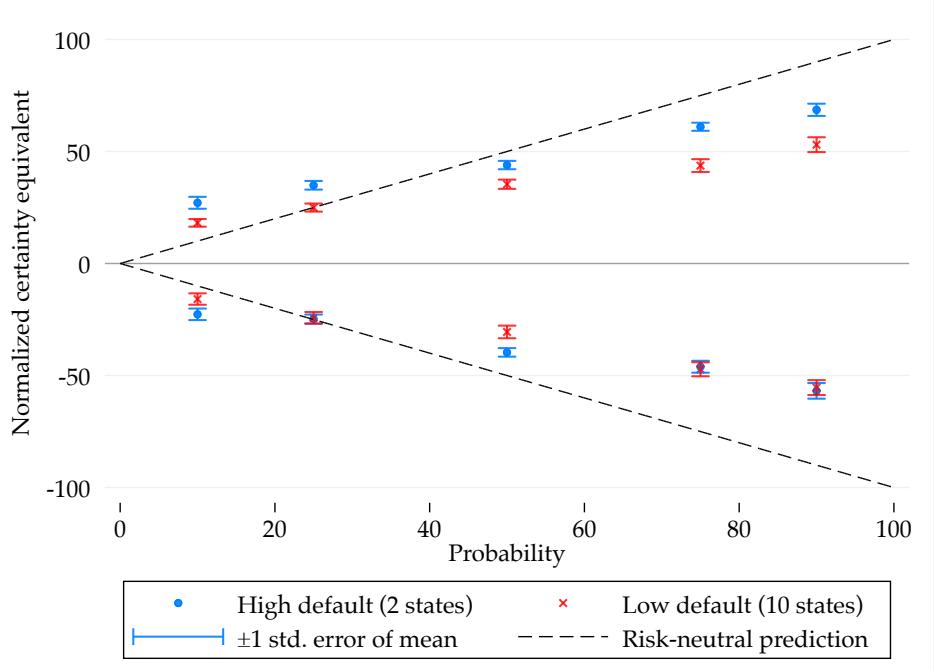


Figure 24: Probability weighting function separately for treatments *High Default Risk* and *Low Default Risk* (full sample). The plot shows averages and corresponding standard error bars. Normalized certainty equivalents (implied probability weights) are computed as certainty equivalent divided by payout probability. The figure is based on 1,800 certainty equivalents of 700 subjects.

B.3 Results excluding Speeders

Table 11: Insensitivity to probability and cognitive uncertainty (excl. speeders)

	Dependent variable: Absolute normalized certainty equivalent					
	Gains		Losses		Pooled	
	(1)	(2)	(3)	(4)	(5)	(6)
Probability of payout	0.74*** (0.03)	0.74*** (0.03)	0.52*** (0.04)	0.51*** (0.04)	0.64*** (0.03)	0.64*** (0.03)
Probability of payout × Cognitive uncertainty	-0.70*** (0.10)	-0.70*** (0.10)	-0.32*** (0.10)	-0.30*** (0.10)	-0.52*** (0.08)	-0.51*** (0.08)
Cognitive uncertainty	26.6*** (6.61)	27.1*** (6.55)	29.3*** (5.84)	28.3*** (5.90)	28.8*** (4.57)	29.1*** (4.58)
Session FE	No	Yes	No	Yes	No	Yes
Demographic controls	No	Yes	No	Yes	No	Yes
Observations	1162	1162	1187	1187	2349	2349
R ²	0.49	0.50	0.28	0.29	0.36	0.37

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. The dependent variable is a subject's normalized certainty equivalent, computed as midpoint of the switching interval divided by the non-zero payout. The sample includes choices from all baseline gambles with strictly interior probabilities. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

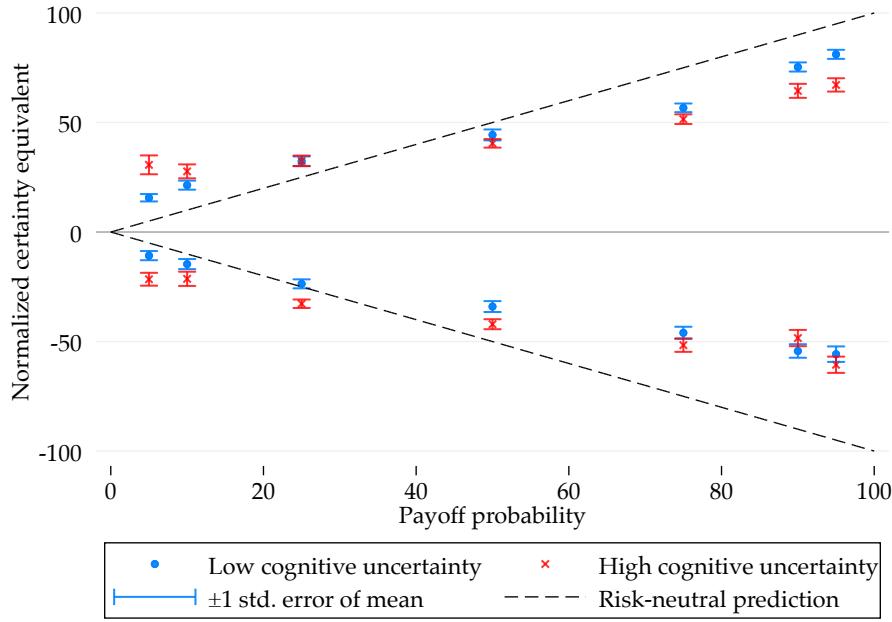


Figure 25: Probability weighting function separately for subjects above / below average cognitive uncertainty (excl. speeders). The partition is done separately for each probability \times gains / losses bucket. The plot shows averages and corresponding standard error bars. Normalized certainty equivalents (implied probability weights) are computed as certainty equivalent divided by payout probability. The figure is based on 2,349 certainty equivalents of 630 subjects.

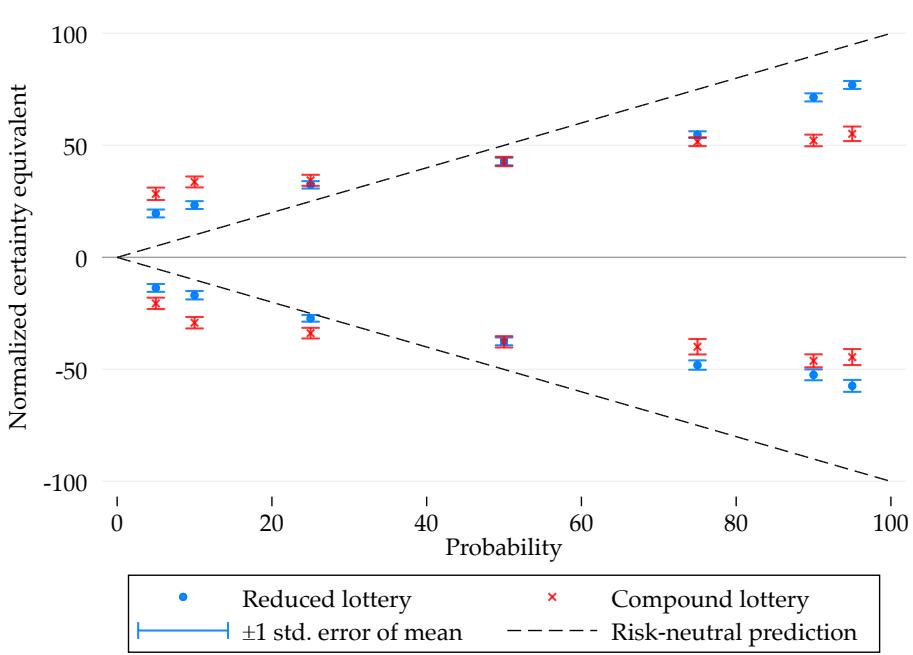


Figure 26: Probability weighting function separately for reduced and compound lotteries (excl. speeders). The plot shows averages and corresponding standard error bars. Normalized certainty equivalents (implied probability weights) are computed as certainty equivalent divided by payout probability. The figure is based on 3,519 certainty equivalents of 700 subjects.

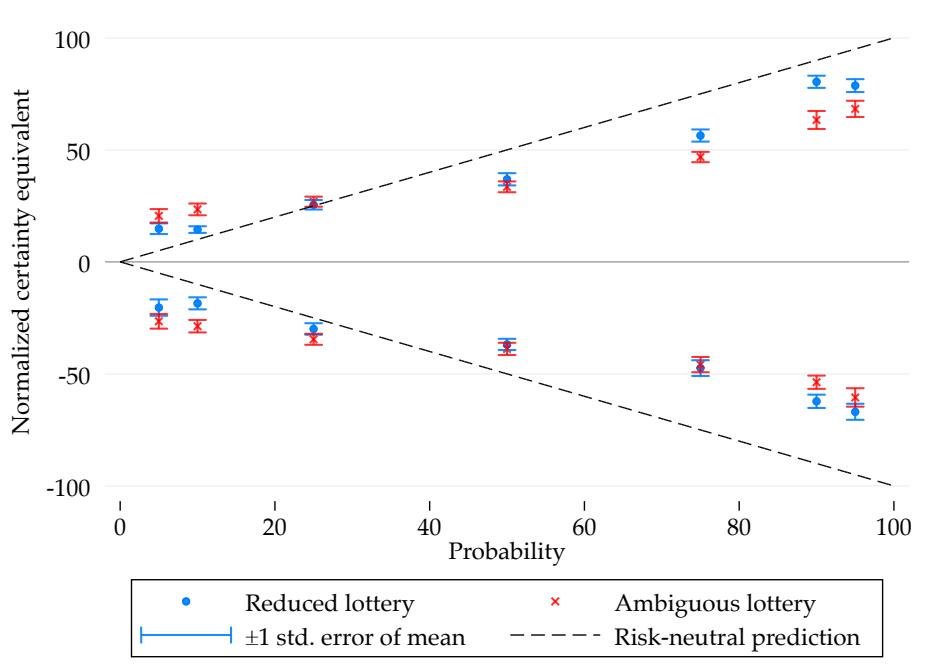


Figure 27: “Probability” weighting function separately for reduced and ambiguous lotteries (excl. speeders). The plot shows averages and corresponding standard error bars. Normalized certainty equivalents (implied probability weights) are computed as certainty equivalent divided by payout probability. The figure is based on 1,608 certainty equivalents of 268 subjects.

Table 12: Choice under risk: Baseline versus compound lotteries (excl. speeders)

	Dependent variable: Absolute normalized certainty equivalent					
	Gains		Losses		Pooled	
	(1)	(2)	(3)	(4)	(5)	(6)
Probability of payout	0.59*** (0.02)	0.70*** (0.03)	0.46*** (0.03)	0.51*** (0.03)	0.53*** (0.02)	0.61*** (0.02)
Probability of payout × 1 if compound lottery	-0.32*** (0.04)	-0.29*** (0.04)	-0.23*** (0.04)	-0.21*** (0.04)	-0.27*** (0.03)	-0.25*** (0.03)
Probability of payout × Cognitive uncertainty		-0.49*** (0.08)		-0.26*** (0.08)		-0.39*** (0.06)
1 if compound lottery	12.5*** (2.18)	11.5*** (2.18)	11.6*** (2.05)	9.55*** (2.06)	12.0*** (1.52)	10.7*** (1.50)
Cognitive uncertainty		19.9*** (5.23)		25.7*** (4.92)		23.6*** (3.95)
Session FE	No	Yes	No	Yes	No	Yes
Demographic controls	No	Yes	No	Yes	No	Yes
Observations	1766	1766	1753	1753	3519	3519
R ²	0.38	0.40	0.22	0.25	0.29	0.30

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. The dependent variable is a subject’s normalized certainty equivalent, computed as midpoint of the switching interval divided by the non-zero payout. The sample includes choices from the baseline and compound gambles, where for comparability the set of baseline gambles is restricted to gambles with payout probabilities of 10%, 25%, 50%, 75%, and 90%, see Figure 4. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

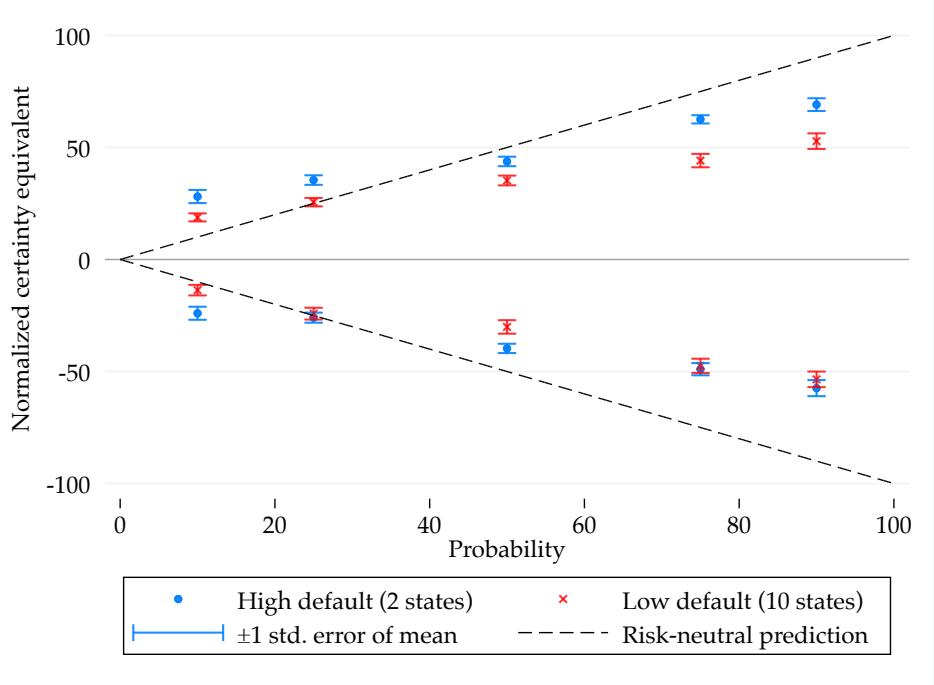


Figure 28: Probability weighting function separately for treatments *High Default Risk* and *Low Default Risk* (excl. speeders). The plot shows averages and corresponding standard error bars. Normalized certainty equivalents (implied probability weights) are computed as certainty equivalent divided by payout probability. The figure is based on 1,620 certainty equivalents of 270 subjects.

Table 13: Choice under risk: Treatments *Low Default* and *High Default* (excl. speeders)

	Dependent variable: Absolute normalized certainty equivalent					
	Gains		Losses		Pooled	
	(1)	(2)	(3)	(4)	(5)	(6)
0 if <i>High Default</i> , 1 if <i>Low Default</i>	-12.5*** (2.06)	-12.1*** (2.12)	-4.86** (2.35)	-4.53* (2.35)	-8.67*** (1.68)	-8.24*** (1.69)
Probability of payout	0.59*** (0.05)	0.59*** (0.05)	0.54*** (0.05)	0.55*** (0.05)	0.57*** (0.04)	0.57*** (0.04)
Probability of payout × Cognitive uncertainty	-0.54*** (0.13)	-0.53*** (0.13)	-0.36*** (0.14)	-0.38*** (0.14)	-0.44*** (0.10)	-0.45*** (0.10)
Cognitive uncertainty	22.6*** (6.60)	22.1*** (6.57)	39.3*** (7.82)	39.0*** (8.11)	31.2*** (5.38)	31.2*** (5.47)
Session FE	No	Yes	No	Yes	No	Yes
Demographic controls	No	Yes	No	Yes	No	Yes
Observations	810	810	810	810	1620	1620
R ²	0.33	0.35	0.26	0.28	0.28	0.29

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. The dependent variable is a subject's normalized certainty equivalent, computed as midpoint of the switching interval divided by the non-zero payout. The sample includes choices from treatments *Low Default* and *High Default*. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

C Additional Details and Analyses for Belief Updating Experiments

C.1 Additional Figures

This decision is about the **same problem** as the one on the previous screen:

Number of "bag A" cards: **90**

Number of "bag B" cards: **10**

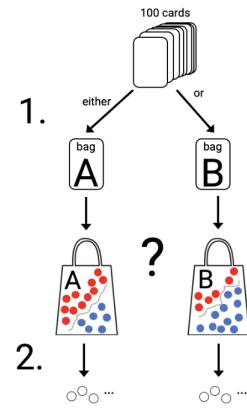
Bag A contains **90 red** balls and **10 blue** balls.

Bag B contains **10 red** balls and **90 blue** balls.

Next:

1. The computer **randomly selected one bag** by drawing a card from the deck.
2. Then, the computer randomly drew 1 ball from the secretly selected bag:

1 red ball was drawn.



Decision 2

Your task is to guess which bag was selected in this case.

Your guess:

Select a probability (between 0 and 100) that expresses **how likely** you think it is that **bag A as opposed to bag B has been selected**:

Probability of **bag A**:

32 %

Probability of **bag B**:

68 %

Submit your guess

Figure 29: Decision screen to elicit posterior belief in belief updating tasks

In this task:

Number of "bag A" cards: 90

Number of "bag B" cards: 10

Bag A contains 90 red balls and 10 blue balls.

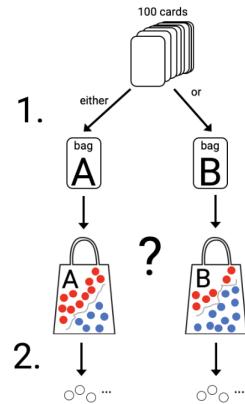
Bag B contains 10 red balls and 90 blue balls.

Next:

1. The computer randomly selected one bag by drawing a card from the deck.

2. Then, the computer randomly drew 1 ball from the secretly selected bag:

1 red ball was drawn.



Decision 1

By replacing your guess with the optimal guess you may **increase your chances of winning \$10.00**. You have a budget of \$3.00 to purchase the optimal guess in this task.

How much of your \$3.00 budget are you **willing to pay to replace your guess with the optimal guess in this task?**

Your willingness to pay for the optimal guess: 1.54\$



Figure 30: Decision screen to elicit willingness-to-pay for optimal guess in belief updating

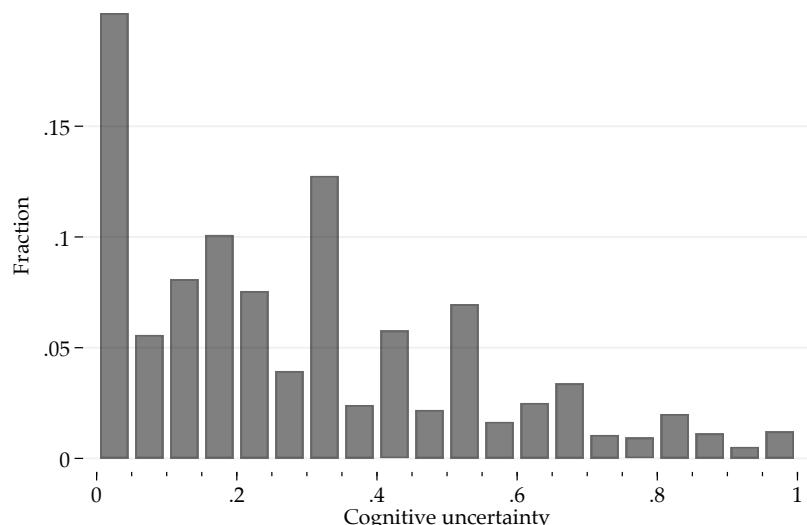


Figure 31: Histogram of cognitive uncertainty in baseline belief updating tasks

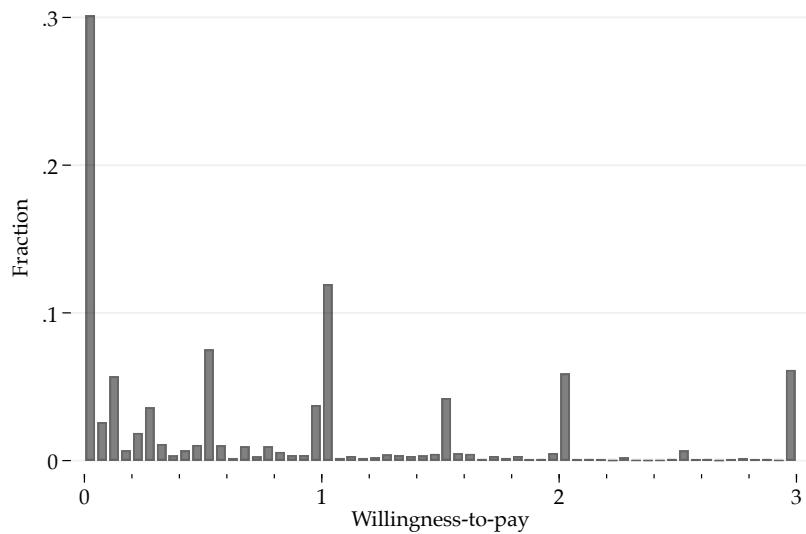


Figure 32: Histogram of willingness-to-pay to replace own guess by Bayesian posterior in baseline belief updating tasks

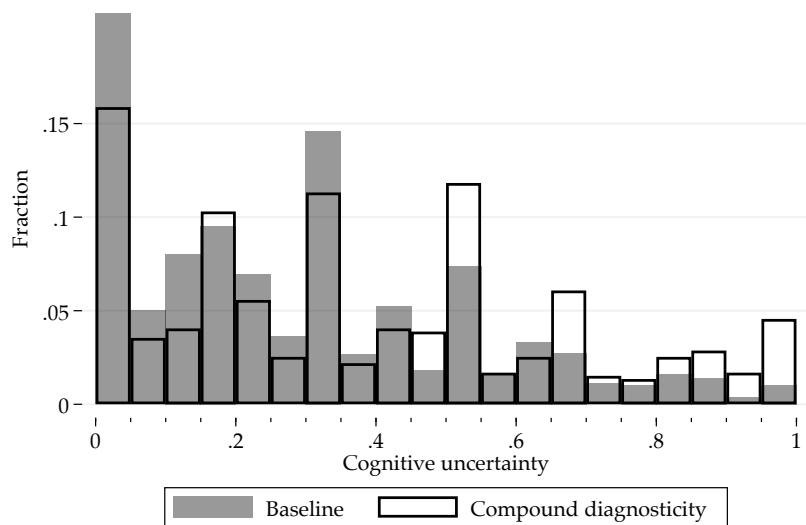


Figure 33: Histograms of cognitive uncertainty in belief updating tasks, separately for baseline and compound diagnosticities

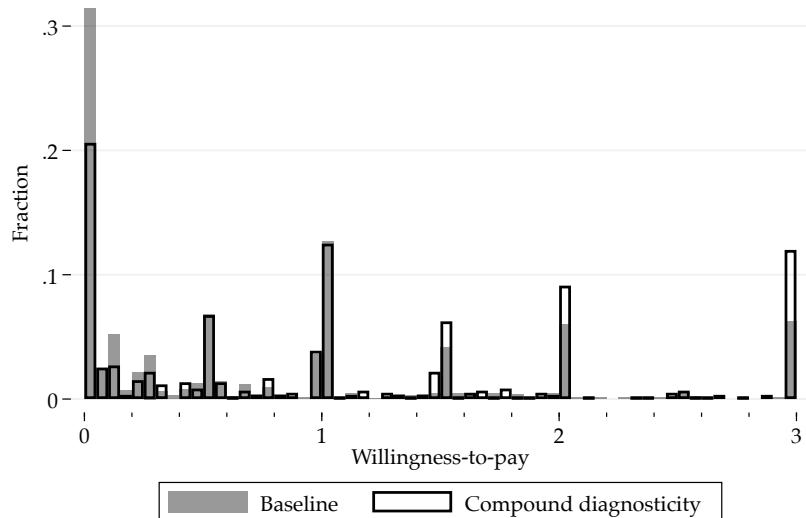


Figure 34: Histograms of willingness-to-pay to replace own guess by Bayesian posterior in belief updating tasks, separately for baseline and compound diagnosticities

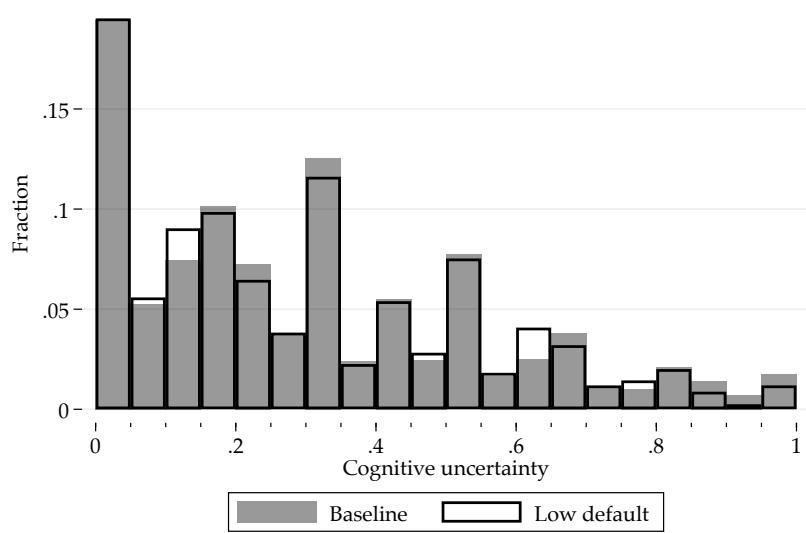


Figure 35: Histograms of cognitive uncertainty in belief updating tasks, separately for treatments *Baseline* and *Low Default Beliefs*.

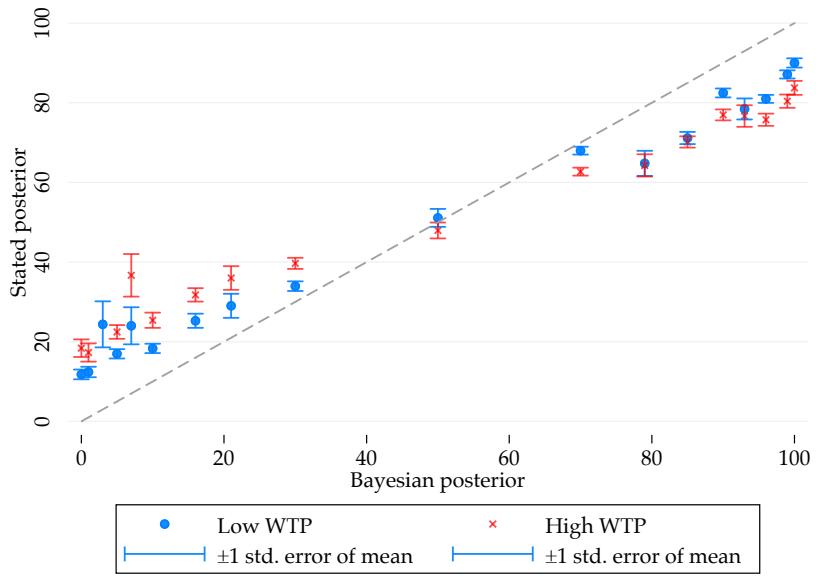


Figure 36: Relationship between stated and Bayesian posteriors, separately for subjects above / below median WTP for the Bayesian guess. The partition is done separately for each Bayesian posterior. The plot shows averages and corresponding standard error bars.

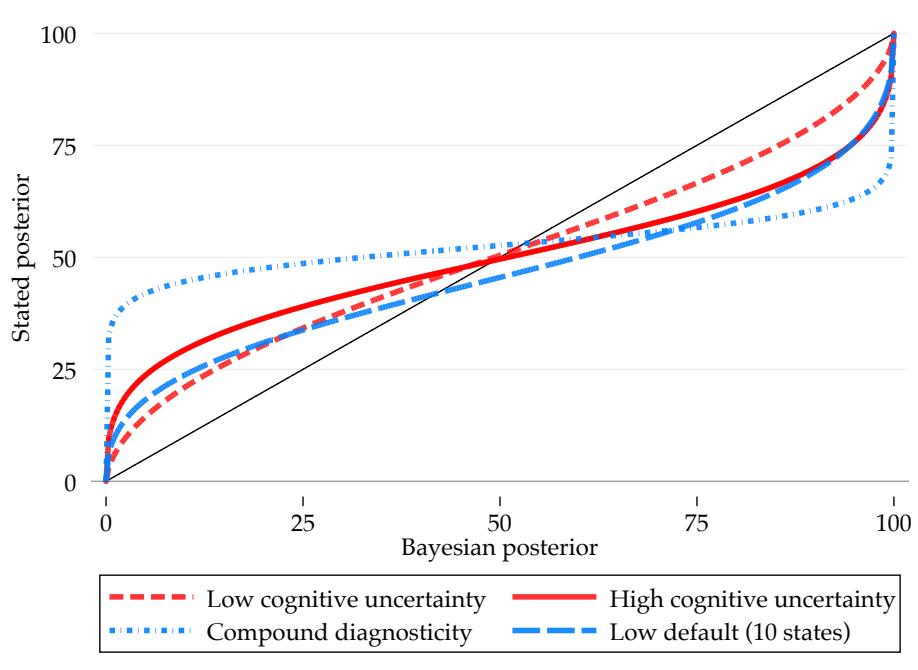


Figure 37: Estimated belief weighting functions across treatments and groups of subjects.

C.2 Additional Tables

Table 14: Belief updating: Baseline tasks: WTP measure

	Dependent variable:					
	Posterior belief			Log [Posterior odds]		
	(1)	(2)	(3)	(4)	(5)	(6)
Bayesian posterior	0.69*** (0.01)	0.76*** (0.01)	0.76*** (0.01)			
Bayesian posterior \times WTP for Bayes		-0.096*** (0.01)	-0.096*** (0.01)			
WTP for Bayesian posterior		5.49*** (0.76)	5.47*** (0.76)	0.027 (0.02)	0.024 (0.02)	
Log [Likelihood ratio]				0.41*** (0.01)	0.43*** (0.01)	0.43*** (0.01)
Log [Prior odds]				0.42*** (0.02)	0.44*** (0.03)	0.44*** (0.03)
Log [Likelihood ratio] \times WTP for Bayes					-0.042*** (0.01)	-0.043*** (0.01)
Log [Prior odds] \times WTP for Bayes					-0.028 (0.02)	-0.027 (0.02)
Session FE	No	No	Yes	No	No	Yes
Demographic controls	No	No	Yes	No	No	Yes
Observations	3187	3187	3187	3104	3104	3104
R ²	0.72	0.73	0.73	0.62	0.63	0.63

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

C.3 Results with Full Sample

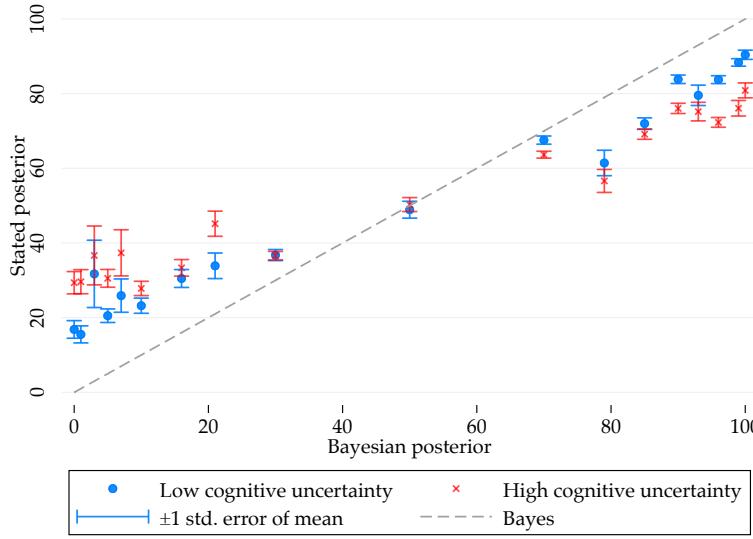


Figure 38: Relationship between average stated and Bayesian posteriors, separately for subjects above / below average cognitive uncertainty (full sample). The partition is done separately for each Bayesian posterior. Bayesian posteriors are rounded to the nearest integer. We only show buckets with more than ten observations. The figure is based on 3,310 beliefs of 700 subjects.

Table 15: Belief updating: Baseline tasks (full sample)

	Dependent variable:					
	Posterior belief			Log [Posterior odds]		
	(1)	(2)	(3)	(4)	(5)	(6)
Bayesian posterior	0.62*** (0.01)	0.77*** (0.02)	0.77*** (0.02)			
Bayesian posterior × Cognitive uncertainty		-0.56*** (0.05)	-0.56*** (0.05)			
Cognitive uncertainty		25.3*** (3.15)	25.3*** (3.18)		-0.12 (0.08)	-0.12 (0.08)
Log [Likelihood ratio]				0.36*** (0.01)	0.42*** (0.02)	0.42*** (0.02)
Log [Prior odds]				0.36*** (0.02)	0.51*** (0.03)	0.51*** (0.03)
Log [Likelihood ratio] × Cognitive uncertainty					-0.28*** (0.05)	-0.28*** (0.05)
Log [Prior odds] × Cognitive uncertainty					-0.55*** (0.07)	-0.54*** (0.07)
Session FE	No	No	Yes	No	No	Yes
Demographic controls	No	No	Yes	No	No	Yes
Observations	3310	3310	3310	3222	3222	3222
R ²	0.57	0.60	0.60	0.48	0.51	0.51

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

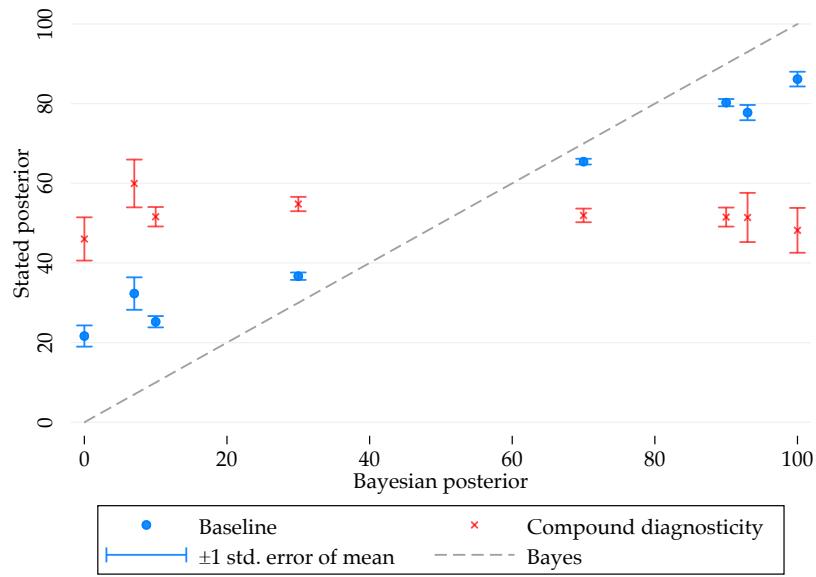


Figure 39: Stated average posteriors as a function of Bayesian posteriors, separately for reduced and compound belief updating problems (full sample). The plot shows averages and corresponding standard error bars. To allow for a valid comparison between baseline and compound updating problems, the sample is restricted to updating tasks in which the base rate is 50:50. Bayesian posteriors are rounded to the nearest integer. We only show buckets with more than ten observations. The figure is based on 2,056 beliefs of 697 subjects.

Table 16: Belief updating: Reduced versus compound signal diagnosticities (full sample)

	Dependent variable:					
	Posterior belief			Log [Posterior odds]		
	(1)	(2)	(3)	(4)	(5)	(6)
Bayesian posterior	0.44*** (0.02)	0.67*** (0.02)	0.67*** (0.02)			
Bayesian posterior × 1 if compound problem		-0.69*** (0.04)	-0.69*** (0.04)			
1 if compound problem		34.5*** (2.17)	34.7*** (2.15)		-0.046 (0.06)	-0.043 (0.06)
Log [Likelihood ratio]				0.31*** (0.01)	0.40*** (0.02)	0.40*** (0.02)
Log [Likelihood ratio] × 1 if compound problem					-0.32*** (0.03)	-0.32*** (0.03)
Session FE	No	No	Yes	No	No	Yes
Demographic controls	No	No	Yes	No	No	Yes
Observations	2056	2056	2056	1954	1954	1954
R ²	0.29	0.45	0.46	0.33	0.40	0.41

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

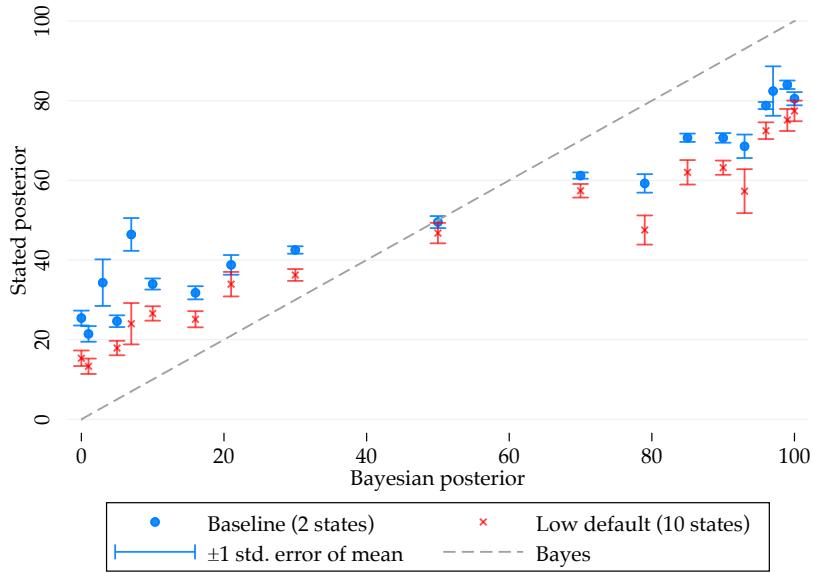


Figure 40: Stated average posteriors as a function of Bayesian posteriors, separately for treatments *Baseline Beliefs* and *Low Default Beliefs* (full sample). Bayesian posteriors are rounded to the nearest integer. We only show buckets with more than ten observations. The figure is based on 5,668 beliefs of 1,000 subjects.

Table 17: Belief updating: Low versus high mental default (full sample)

	Dependent variable:					
	Posterior belief			Log [Posterior odds]		
	(1)	(2)	(3)	(4)	(5)	(6)
0 if <i>Baseline</i> , 1 if <i>Low Default</i>	-6.94*** (0.97)	-7.22*** (0.94)	-7.67*** (1.00)	-0.41*** (0.06)	-0.43*** (0.06)	-0.46*** (0.06)
Bayesian posterior	0.54*** (0.01)	0.66*** (0.02)	0.66*** (0.02)			
Bayesian posterior × Cognitive uncertainty		-0.47*** (0.04)	-0.47*** (0.04)			
Cognitive uncertainty		19.5*** (2.46)	19.4*** (2.49)		-0.12* (0.07)	-0.12 (0.07)
Log [Likelihood ratio]				0.31*** (0.01)	0.38*** (0.02)	0.38*** (0.02)
Log [Prior odds]				0.41*** (0.02)	0.56*** (0.03)	0.56*** (0.03)
Log [Likelihood ratio] × Cognitive uncertainty					-0.29*** (0.04)	-0.29*** (0.04)
Log [Prior odds] × Cognitive uncertainty					-0.56*** (0.07)	-0.56*** (0.07)
Session FE	No	No	Yes	No	No	Yes
Demographic controls	No	No	Yes	No	No	Yes
Observations	5668	5668	5668	5473	5473	5473
R ²	0.44	0.46	0.46	0.42	0.45	0.45

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

C.4 Results excluding Speeders

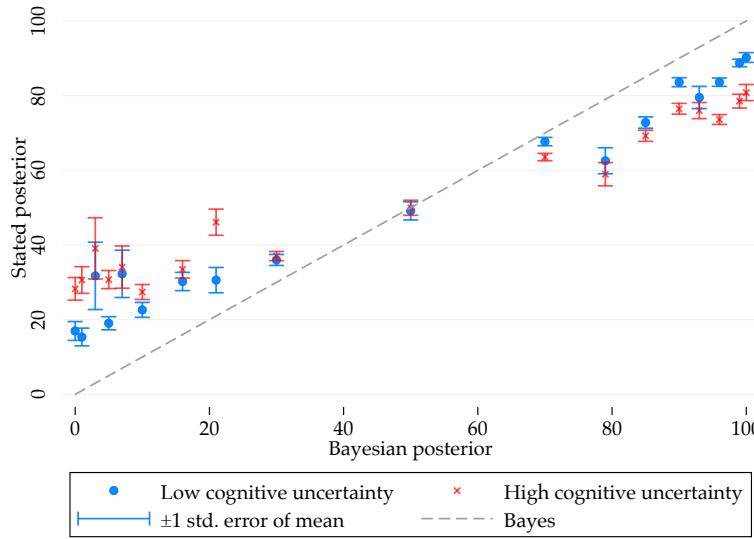


Figure 41: Relationship between average stated and Bayesian posteriors, separately for subjects above / below average cognitive uncertainty (excl. speeders). The partition is done separately for each Bayesian posterior. Bayesian posteriors are rounded to the nearest integer. We only show buckets with more than ten observations. The figure is based on 3,006 beliefs of 635 subjects.

Table 18: Belief updating: Baseline tasks (excl. speeders)

	Dependent variable:					
	Posterior belief			Log [Posterior odds]		
	(1)	(2)	(3)	(4)	(5)	(6)
Bayesian posterior	0.63*** (0.01)	0.78*** (0.02)	0.78*** (0.02)			
Bayesian posterior × Cognitive uncertainty		-0.57*** (0.05)	-0.57*** (0.05)			
Cognitive uncertainty		27.1*** (3.20)	27.3*** (3.21)		-0.066 (0.09)	-0.066 (0.09)
Log [Likelihood ratio]				0.36*** (0.01)	0.42*** (0.02)	0.42*** (0.02)
Log [Prior odds]				0.38*** (0.02)	0.53*** (0.04)	0.52*** (0.04)
Log [Likelihood ratio] × Cognitive uncertainty					-0.29*** (0.05)	-0.29*** (0.05)
Log [Prior odds] × Cognitive uncertainty					-0.55*** (0.08)	-0.55*** (0.08)
Session FE	No	No	Yes	No	No	Yes
Demographic controls	No	No	Yes	No	No	Yes
Observations	3006	3006	3006	2925	2925	2925
R ²	0.59	0.62	0.62	0.49	0.51	0.51

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

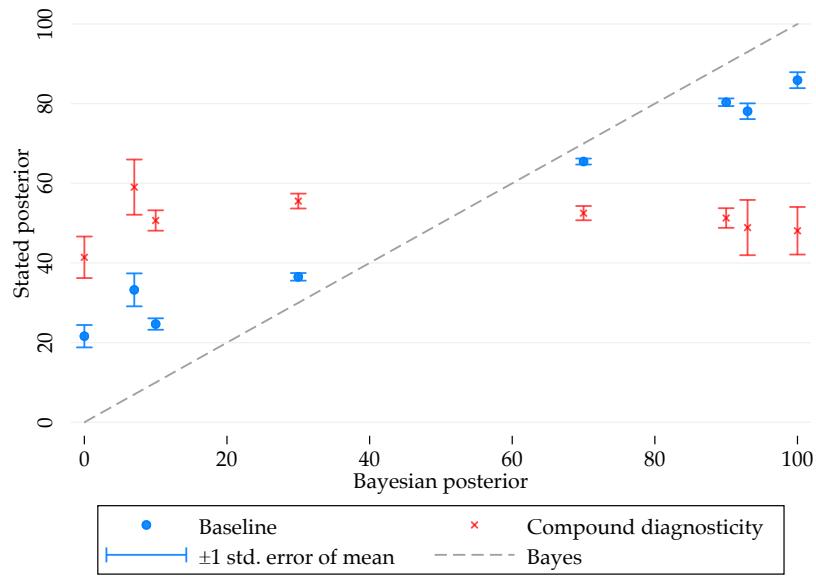


Figure 42: Stated average posteriors as a function of Bayesian posteriors, separately for reduced and compound belief updating problems (excl. speeders). The plot shows averages and corresponding standard error bars. To allow for a valid comparison between baseline and compound updating problems, the sample is restricted to updating tasks in which the base rate is 50:50. Bayesian posteriors are rounded to the nearest integer. We only show buckets with more than ten observations. The figure is based on 1,874 beliefs of 632 subjects.

Table 19: Belief updating: Reduced versus compound signal diagnosticities (excl. speeders)

	Dependent variable:					
	Posterior belief			Log [Posterior odds]		
	(1)	(2)	(3)	(4)	(5)	(6)
Bayesian posterior	0.45*** (0.02)	0.68*** (0.02)	0.68*** (0.02)			
Bayesian posterior × 1 if compound problem		-0.68*** (0.04)	-0.68*** (0.04)			
1 if compound problem		33.9*** (2.25)	34.1*** (2.24)		-0.071 (0.06)	-0.069 (0.06)
Log [Likelihood ratio]				0.31*** (0.02)	0.41*** (0.02)	0.41*** (0.02)
Log [Likelihood ratio] × 1 if compound problem					-0.31*** (0.03)	-0.31*** (0.03)
Session FE	No	No	Yes	No	No	Yes
Demographic controls	No	No	Yes	No	No	Yes
Observations	1874	1874	1874	1779	1779	1779
R ²	0.30	0.46	0.46	0.34	0.40	0.41

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

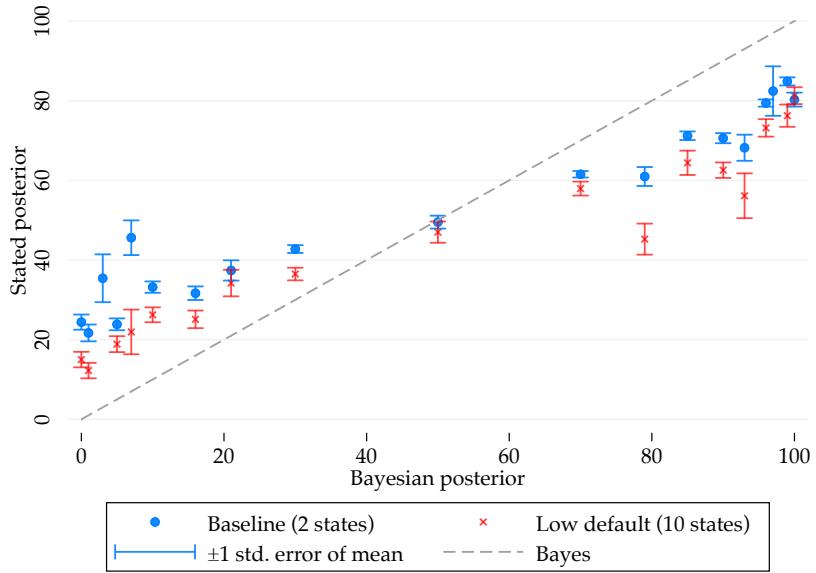


Figure 43: Stated average posteriors as a function of Bayesian posteriors, separately for treatments *Baseline Beliefs* and *Low Default Beliefs* (full sample). Bayesian posteriors are rounded to the nearest integer. We only show buckets with more than ten observations. The figure is based on 5,107 beliefs of 899 subjects.

Table 20: Belief updating: Low versus high mental default (excl. speeders)

	Dependent variable:					
	Posterior belief			Log [Posterior odds]		
	(1)	(2)	(3)	(4)	(5)	(6)
0 if <i>Baseline</i> , 1 if <i>Low Default</i>	-6.64*** (0.98)	-6.88*** (0.96)	-7.16*** (1.02)	-0.40*** (0.06)	-0.42*** (0.06)	-0.43*** (0.06)
Bayesian posterior	0.55*** (0.01)	0.67*** (0.02)	0.67*** (0.02)			
Bayesian posterior × Cognitive uncertainty		-0.48*** (0.04)	-0.48*** (0.04)			
Cognitive uncertainty		21.0*** (2.56)	20.9*** (2.59)		-0.065 (0.07)	-0.057 (0.08)
Log [Likelihood ratio]				0.32*** (0.01)	0.38*** (0.02)	0.38*** (0.02)
Log [Prior odds]				0.43*** (0.02)	0.58*** (0.03)	0.58*** (0.03)
Log [Likelihood ratio] × Cognitive uncertainty					-0.28*** (0.04)	-0.28*** (0.04)
Log [Prior odds] × Cognitive uncertainty					-0.57*** (0.07)	-0.57*** (0.07)
Session FE	No	No	Yes	No	No	Yes
Demographic controls	No	No	Yes	No	No	Yes
Observations	5107	5107	5107	4930	4930	4930
R ²	0.45	0.47	0.48	0.44	0.46	0.46

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

D Additional Details and Analyses for Survey Expectations

D.1 Additional Figures

Your certainty about your estimate

On the previous screen, you indicated that you think that in 2018, a randomly selected household in the United States **earned less than \$ 236,000 with a probability of 32 %.**

How certain are you that this probability is **exactly 32 %?**

Use the slider to complete the statement below.

I
very uncertain completely certain

I am certain that the **actual probability** that a household earned less than \$ 236,000 is **between 15 % and 49 %.**

Next

Figure 44: Decision screen to elicit cognitive uncertainty in survey expectations

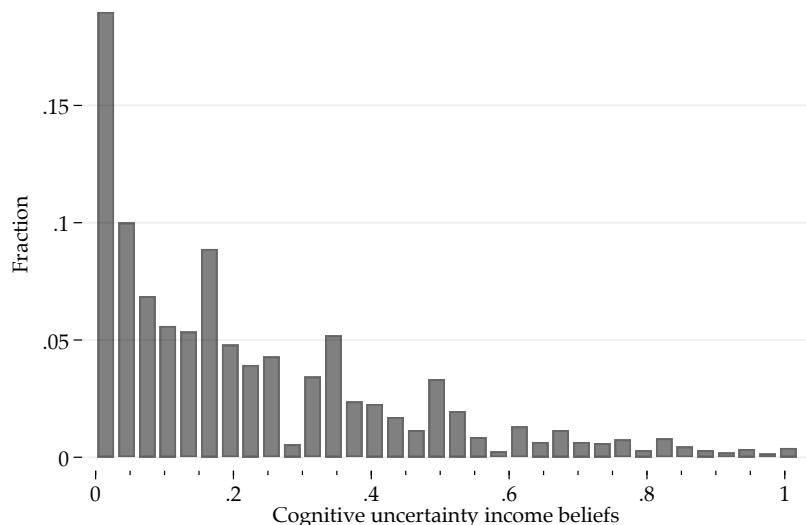


Figure 45: Histogram of cognitive uncertainty in survey expectations about income distribution

.15 -

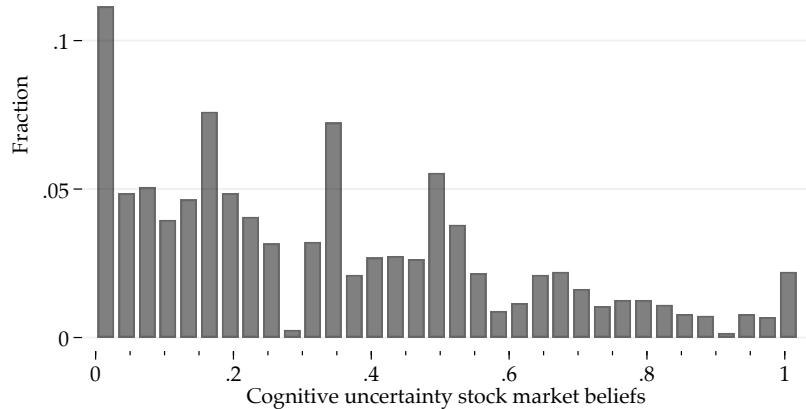


Figure 46: Histogram of cognitive uncertainty in survey expectations about the stock market

.15 -

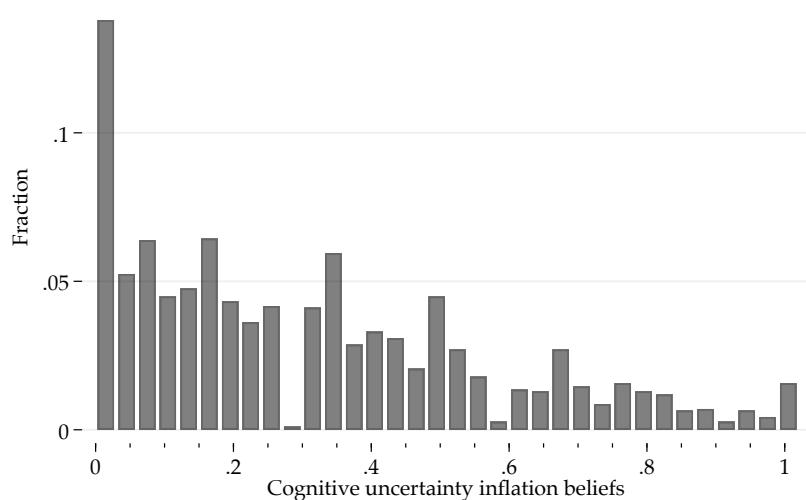


Figure 47: Histogram of cognitive uncertainty in survey expectations about inflation rates

D.2 Additional Tables

Table 21: Survey expectations and cognitive uncertainty

	Dependent variable: Probability estimate about:					
	Income distr.		Stock market		Inflation rate	
	(1)	(2)	(3)	(4)	(5)	(6)
Objective probability	0.92*** (0.01)	0.93*** (0.01)	0.74*** (0.02)	0.74*** (0.02)	0.80*** (0.02)	0.80*** (0.02)
Objective probability \times Cognitive uncertainty	-0.60*** (0.04)	-0.60*** (0.04)	-0.70*** (0.04)	-0.70*** (0.04)	-0.81*** (0.04)	-0.80*** (0.04)
Cognitive uncertainty	29.4*** (2.44)	29.0*** (2.50)	34.2*** (2.09)	34.6*** (2.13)	39.1*** (2.67)	38.5*** (2.74)
Session FE	No	Yes	No	Yes	No	Yes
Demographic controls	No	Yes	No	Yes	No	Yes
Observations	1980	1980	1892	1892	1848	1848
R ²	0.84	0.84	0.54	0.55	0.56	0.56

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. In columns (1)–(2), the question about income distribution asks participants for the probability that a randomly selected U.S. household earns less than x . In columns (3)–(4), the question about the stock market asks participants for the probability that in a randomly selected year the S&P500 increased by less than $x\%$. In columns (5)–(6), the question about inflation rates asks participants for the probability that in a randomly selected year the inflation rate was less than $x\%$. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 22: Estimates of “survey expectations weighting function”

Task / group	Sensitivity $\hat{\lambda}$	Elevation $\hat{\delta}$
<i>Income distribution: high CU</i>	0.51 (0.02)	1.08 (0.04)
<i>Income distribution: low CU</i>	0.75 (0.02)	1.27 (0.04)
<i>Stock market performance: high CU</i>	0.19 (0.01)	0.82 (0.03)
<i>Stock market performance: low CU</i>	0.45 (0.02)	0.84 (0.04)
<i>Inflation rates: high CU</i>	0.22 (0.01)	0.97 (0.03)
<i>Inflation rates: low CU</i>	0.47 (0.02)	0.98 (0.05)

Notes. Estimates of equation (15) for survey expectations, standard errors (clustered at subject level) reported in parentheses. CU = cognitive uncertainty (split at average).

D.3 Results with Full Sample

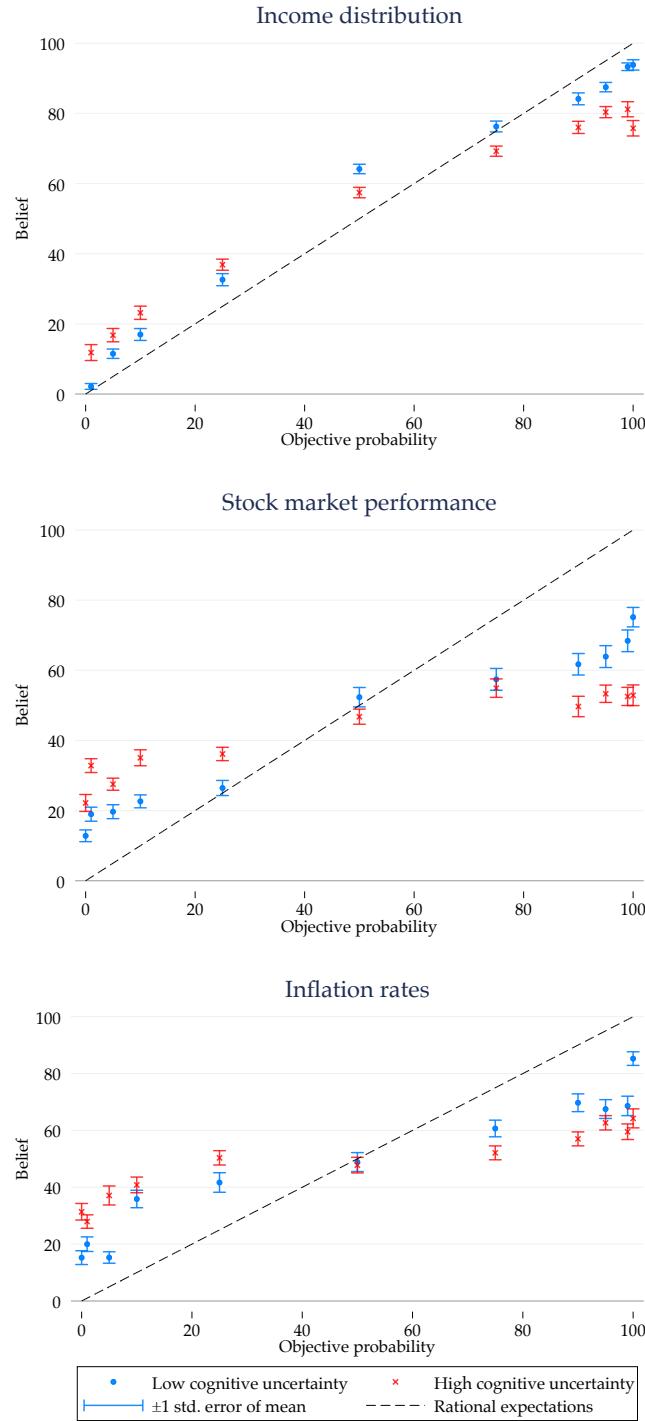


Figure 48: Survey beliefs as a function of objective probabilities, separately for subjects above / below average cognitive uncertainty (full sample). The partition is done separately for each probability bucket. In the top panel, the question asks for the probability that a randomly selected U.S. household earns less than \$x. In the middle panel, the question asks for the probability that in a randomly selected year the S&P500 increased by less than x%. In the bottom panel, the question asks for the probability that in a randomly selected year the inflation rate was less than x%. $N = 2,000$ observations each.

D.4 Results excluding Speeders

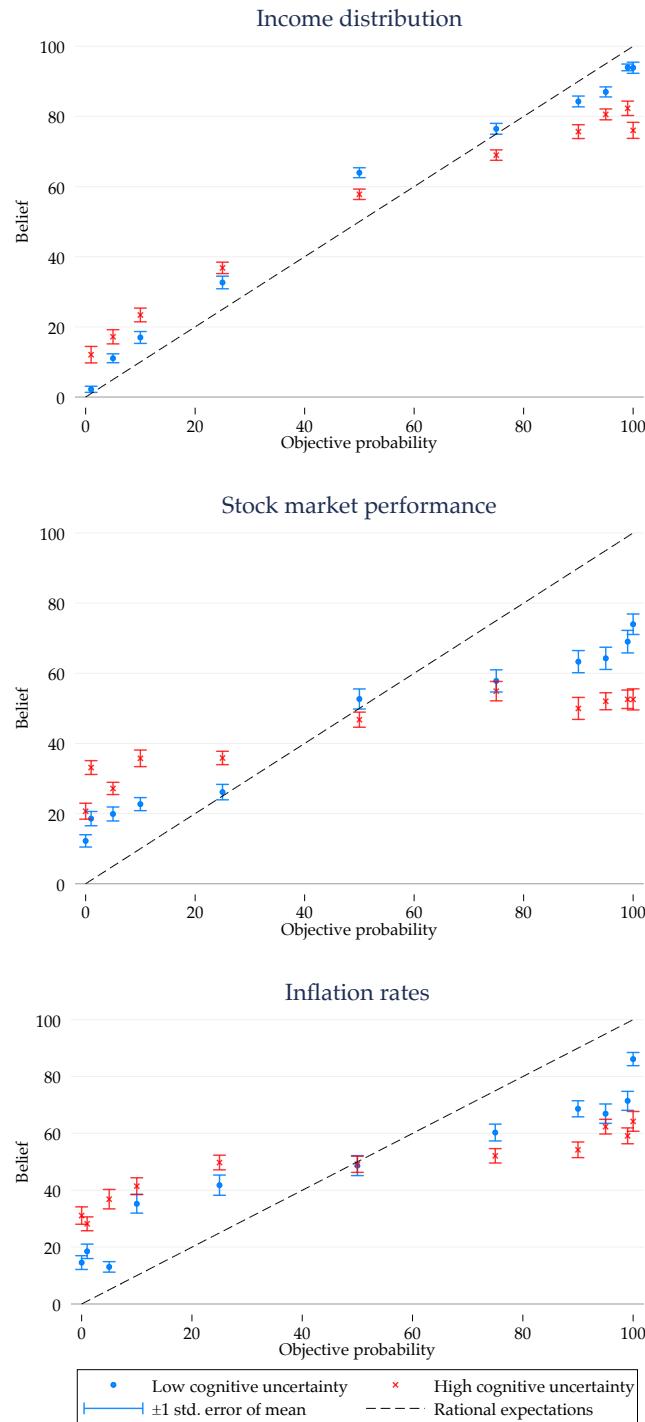


Figure 49: Survey beliefs as a function of objective probabilities, separately for subjects above / below average cognitive uncertainty (excl. speeders). The partition is done separately for each probability bucket. In the top panel, the question asks for the probability that a randomly selected U.S. household earns less than \$x. In the middle panel, the question asks for the probability that in a randomly selected year the S&P500 increased by less than x%. In the bottom panel, the question asks for the probability that in a randomly selected year the inflation rate was less than x%. $N = 1,896$ observations each.

E Forward-Looking Survey Expectations

In Section 5 in the main text, we elicited respondents' survey expectations about economic variables with respect to past values, which allowed us to easily incentivize participants. In a pre-registered robustness check, we implemented the same type of survey questions, but now regarding future values of these variables. These questions are hence theoretically more appropriate in that they elicit actual expectations, but they are not financially incentivized. The sample size is $N = 400$ for each of the three domains. We apply the same criteria regarding the exclusions of outliers as in Section 5.

The results are shown in Figure 50. Here, we define “objective probabilities” based on historical data, akin to Figure 10 in the main text. The results are almost identical to those reported in the main text.

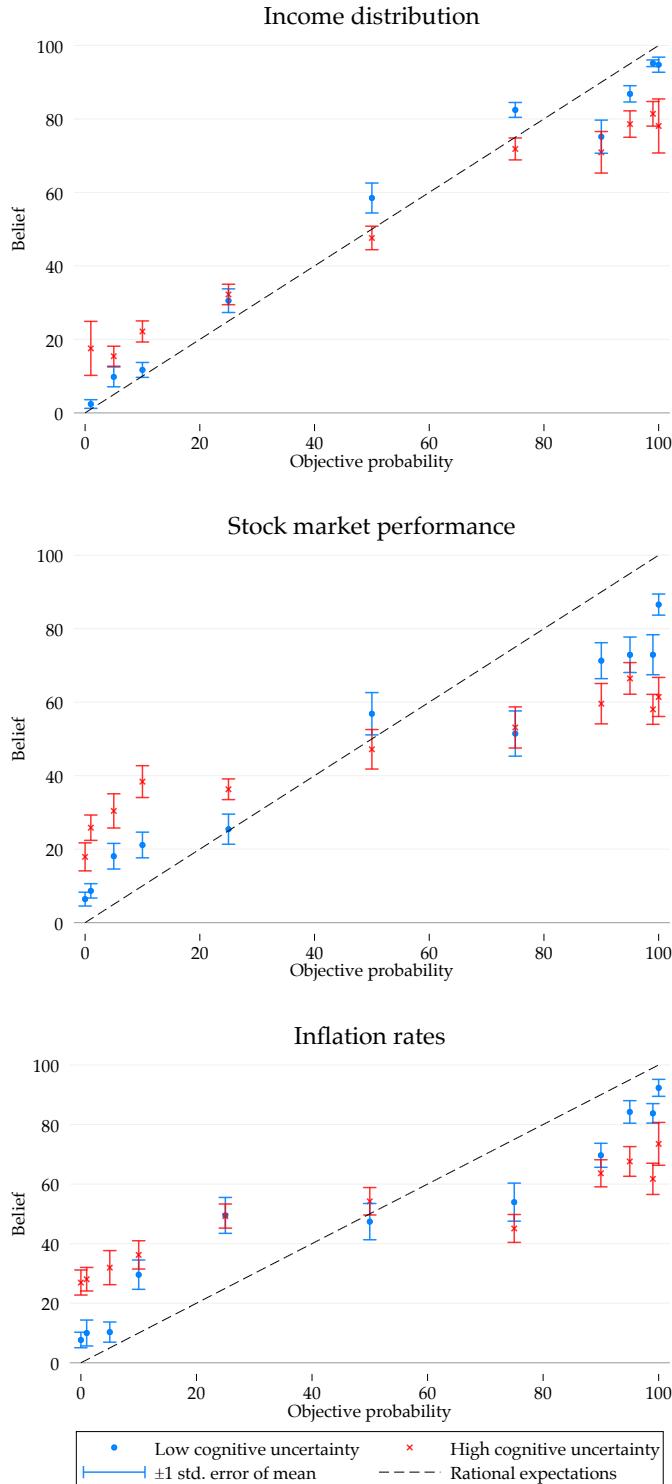


Figure 50: Survey beliefs about future variables as a function of “objective” probabilities, separately for subjects above / below average cognitive uncertainty. The partition is done separately for each probability bucket. “Objective” probabilities are defined using historical data, analogously to Figure 10. In the top panel, the question asks for the probability that a randomly selected U.S. household will earn less than \$x ($N = 491$). In the middle panel, the question asks for the probability that the S&P500 will increase by less than x% ($N = 463$) over the course of one year. In the bottom panel, the question asks for the probability that the inflation rate will be than x% ($N = 478$).

F Additional Ambiguity Experiment

In addition to the experiments reported in Section 3, we implemented an additional set of pre-registered ambiguity experiments. These experiments delivered statistically significant results in line with our pre-registered predictions. However, as explained below, we now believe that these experiments are conceptually less-than-ideal from the perspective of our framework, which is why we relegate them to an Appendix.

F.1 Experimental Design

The basic design builds on Dimmock et al. (2015) and aims at eliciting matching probabilities for ambiguous lotteries. In a given choice list, the left-hand side option A was constant and given by an ambiguous lottery. The ambiguous lottery was described as random draw from an urn that comprises 100 balls of ten different colors, where the precise composition of colors is unknown. A known number of these colors n were “winning colors” that resulted in the same payout $\$x$, while other colors resulted in a zero payout. Option B, on the right-hand side, varied across rows in the choice list and was also given by a lottery with upside $\$x$. Here, the number of “winning balls” was known and varied from 0% to 99% in 3% steps. Subjects were always given the option to pick their preferred winning colors.

A subject completed six choice lists, where the payout $x \in \{15, 20, 25\}$ and the number of winning colors $n \in \{1, 2, \dots, 9\}$ were randomly determined. Before each decision screen, subjects were always given the opportunity to pick their winning colors.

Cognitive uncertainty was measured analogously to choice under risk. After subjects had indicated their probability equivalent range for an ambiguous lottery, the subsequent screen asked them how certain they are that this range actually corresponds to how much the lottery is worth to them. Operationally, subjects used a slider to calibrate the statement “I am certain that to me the lottery is worth as much as playing a lottery over $\$x$ with a known number of between x and y winning balls.” 200 AMT workers participated in these experiments and earned an average of \$7.20.

F.2 Results

In the baseline analysis, we again exclude extreme outliers, defined as matching probability strictly larger than 75% for at most two winning colors, and matching probability strictly smaller than 25% for more than eight winning colors. This is the case for 1.6% of our data. We find that the response function of subjects with higher cognitive uncertainty is significantly less sensitive to variation in the number of winning colors (shallower), see the regressions in Table 23. This reduction in sensitivity corresponds to our

Table 23: Insensitivity to ambiguous “likelihood” and cognitive uncertainty

	<i>Dependent variable:</i> Matching probability	
	(1)	(2)
Number of winning colors * 10	0.68*** (0.04)	0.67*** (0.04)
Number of winning colors * 10 × Cognitive uncertainty	-0.26*** (0.10)	-0.25** (0.10)
Cognitive uncertainty	5.90 (5.14)	3.87 (5.14)
Session FE	No	Yes
Demographic controls	No	Yes
Observations	1181	1181
R^2	0.50	0.51

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. The dependent variable is a subject’s matching probability, computed as midpoint of the switching interval. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

main hypothesis, which is also what we re-registered. At the same time, we do not find that high cognitive uncertainty subjects are more ambiguity seeking than low cognitive uncertainty subjects for unlikely events.

F.3 Interpretive Problems

The analysis above focuses on whether reported matching probabilities of subjects with higher cognitive uncertainty are less sensitive to the variation in winning colors. However, our framework in Section 2 *only* makes this prediction if one assumes that the state space is binary (win-lose), so that subjects are hypothesized to “shrink” ambiguous probabilities towards 50:50. However, in the experiments, the state space was represented through ten different colors, some of which are winning and some of which are losing colors. As discussed in Section 3.4, a plausible alternative view is that in this situation there are actually ten states of the world, one for each color. In this case, our framework does not predict that subjects shrink their matching probabilities towards 50:50. To see this, take the example that there are three winning colors. In this case, the ignorance prior (for winning) would be given by 30%. In other words, subjects would be hypothesized to shrink an ambiguous probability of three winning colors towards a mental default of 30%, which does not produce any shrinking theoretically. For this reason, we view these experiments as imperfect.

G Results on Stake Size Increase

G.1 Stake Size and Choice Under Risk

To manipulate the size of financial incentives, we implement a within-subjects manipulation. We implemented the same procedures as described in Section 3, except that we only implemented gain lotteries. Subjects completed six choice lists, one of which determined a subject's payment in case the choice under risk part of the experiment got selected for payment (probability 1/3). Across the six choice lists, the probability of being payout-relevant varied in a transparent way. On top of the decision screen, we informed subjects about the probability that this choice list would determine their payout. For five tasks, this probability was given by 1% and for one task by 95%. As a measure of cognitive effort, we recorded subjects' (log) response times. 150 subjects participated in this treatment, which was also pre-registered.

The results are reported in Table 24.¹⁴ Exploiting variation within subjects across tasks, we find that response times increase significantly from 25 seconds on average to 36 seconds on average in the high stakes task. However, this increase in response times does not translate into a significant change in cognitive uncertainty.

Table 24: Effects of stake size increase in choice under risk

	Dependent variable:					
	Log [Response time]		Cognitive uncertainty		Normalized CE	
	(1)	(2)	(3)	(4)	(5)	(6)
1 if high stakes	0.26*** (0.06)	0.26*** (0.06)	-0.0041 (0.01)	-0.0037 (0.01)	0.41 (2.27)	-1.65 (2.16)
Probability of payout					0.69*** (0.03)	0.68*** (0.03)
Probability of payout × 1 if high stakes					0.022 (0.04)	0.065* (0.04)
Subject FE	No	Yes	No	Yes	No	Yes
Observations	893	893	893	893	893	893
R ²	0.02	0.50	0.00	0.53	0.60	0.79

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

¹⁴We again apply the same outlier exclusion criteria as in the main text.

G.2 Stake Size and Belief Updating

To manipulate the size of financial incentives, we again implement a within-subjects manipulation. We implemented the same procedures as described in Section 4, except that we did not elicit the WTP for the optimal guess. Subjects completed six updating tasks, one of which determined a subject's payment in case the belief updating part of the experiment got selected for payment (probability 1/3). Across the six tasks, the probability of being payout-relevant varied in a transparent way. On top of the decision screen, we informed subjects about the probability that this task would determine their payout. For five tasks, this probability was given by 1% and for one task by 95%. As a measure of cognitive effort, we recorded subjects' (log) response times. 150 subjects participated in this treatment, which was also pre-registered.

The results are reported in Table 25.¹⁵ Exploiting variation within subjects across tasks, we find that response times increase significantly. Cognitive uncertainty decreases, but only mildly so.

Table 25: Effects of stake size increase in belief updating

	Dependent variable:					
	Log [Response time]		Cognitive uncertainty		Posterior belief	
	(1)	(2)	(3)	(4)	(5)	(6)
1 if high stakes	0.19*** (0.06)	0.19*** (0.07)	-0.024* (0.01)	-0.025 (0.02)	-2.76 (2.53)	-3.56 (2.64)
Bayesian posterior					0.59*** (0.03)	0.58*** (0.03)
Bayesian posterior × 1 if high stakes					0.065 (0.04)	0.080* (0.04)
Subject FE	No	Yes	No	Yes	No	Yes
Observations	869	869	869	869	869	869
R ²	0.01	0.46	0.00	0.51	0.61	0.70

Notes. OLS estimates, robust standard errors (in parentheses) are clustered at the subject level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

¹⁵We again apply the same outlier exclusion criteria as in the main text.

H Calibrating the Cognitive Uncertainty Measurement

In all of our experiments, the elicitation of cognitive uncertainty did not specify which particular version of a subjective confidence interval we intend to elicit, such as a 90%, 95%, 99% or 100% confidence interval. We deliberately designed our experiments in this fashion because the hypothesis that underlines our research is that people have a hard time translating “99% confidence” into a statement about e.g. their certainty equivalent. In an attempt to trade off subject comprehension and quantitative interpretation, we hence refrained from inducing a particular version of a confidence interval.

To provide evidence for our conjecture that respondents cannot really tell the difference between different types of confidence intervals, we implemented an additional set of choice under risk experiments in which we elicited different versions of subjective confidence intervals. In these experiments, subjects were specifically instructed to state an interval such that they are “ $y\%$ certain” that to them the lottery is worth between a and b . Across experimental conditions, y varied from 75% to 90% to 95% to 99% to 100%. To analyze these data, we compare average cognitive uncertainty within a treatment with average cognitive uncertainty in our baseline treatments, in which we did not provide a specific quantitative version of a confidence interval. In total, we ran these experiments with $N = 293$ subjects.

Figure 51 summarizes the results. Here, we plot the coefficients of the different treatment dummies in a regression with stated cognitive uncertainty as dependent variable. In this regression, the omitted category is our (unspecific) baseline treatment. Each coefficient hence corresponds to the implied difference in cognitive uncertainty between a treatment and our baseline treatment. There are two main results. First, cognitive uncertainty does not vary in meaningful ways across conditions: subjects state statistically indistinguishable cognitive uncertainty intervals, regardless of whether we specify them as 75%, 90% etc. interval. Second, if anything, reported cognitive uncertainty is *higher* in the more precise quantitative versions relative to our baseline version, as can be inferred from the positive point estimates. This again suggests that subjects have a harder time thinking about specific quantitative versions of a confidence interval relative to our more intuitive question. We conclude from this exercise that a more precise quantitative implementation of our cognitive uncertainty interval is unlikely to deliver a more helpful quantitative interpretation of our measure.

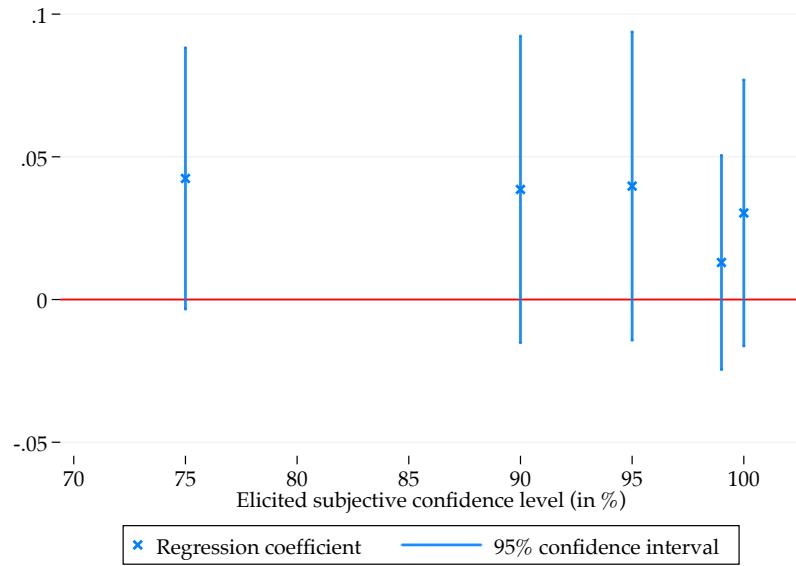


Figure 51: Comparison of average cognitive uncertainty across different elicitation modes in choice under risk. Each dot represents the coefficient of a treatment dummy in a regression with cognitive uncertainty as dependent variable. The explanatory variables are fixed effects for the different specifications of cognitive uncertainty, where the omitted category is our baseline wording. The plot controls for lottery amount fixed effects and probability of payout fixed effects.

I Censoring

In this section we replicate our weighting functions for subjects above and below average cognitive uncertainty after excluding observations that are affected by the boundaries of the response scales. Specifically, in the figures reported in section I.1 we exclude all observations in which the choices or beliefs exactly equaled one of boundaries of the response scale, whereas in section I.2 we exclude all observations in which the cognitive uncertainty range included one of the boundaries. The observed differences between high and low cognitive uncertainty choices or beliefs remain virtually unaffected by these exclusions.

I.1 Censored choices and beliefs

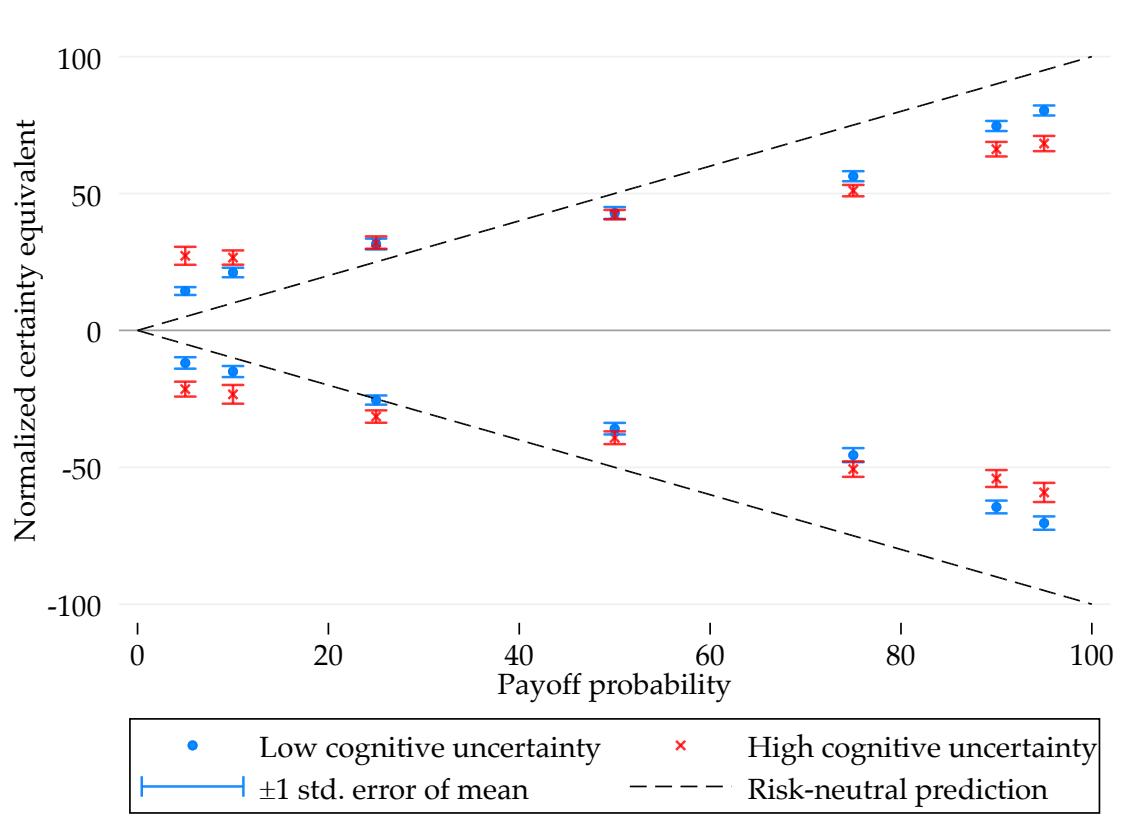


Figure 52: Probability weighting function excluding censored choices, separately for subjects above / below average cognitive uncertainty. The partition is done separately for each probability × gains / losses bucket. The plot shows averages and corresponding standard error bars. Normalized certainty equivalents (implied probability weights) are computed as certainty equivalent divided by payout probability. The figure excludes 4.28% of the original data that is based on 2,525 certainty equivalents of 700 subjects.

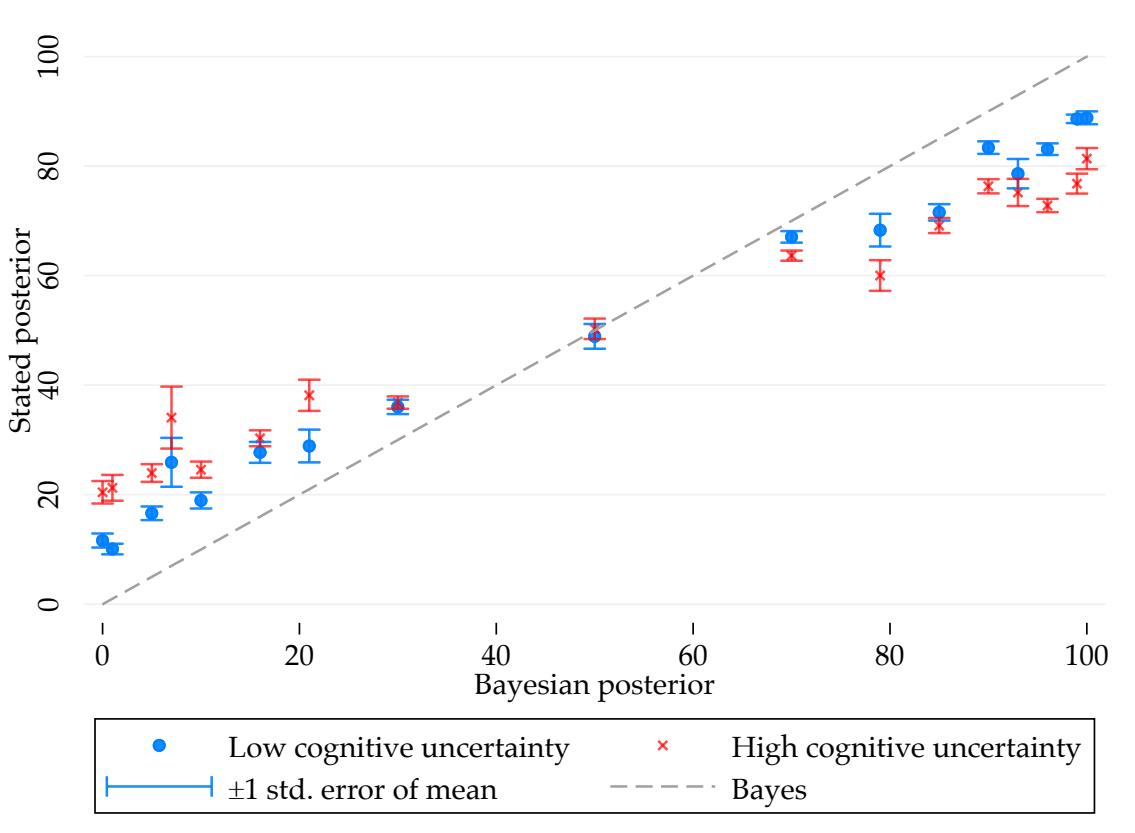


Figure 53: Relationship between average stated and Bayesian posteriors after excluding censored beliefs, separately for subjects above / below average cognitive uncertainty. The partition is done separately for each Bayesian posterior. Bayesian posteriors are rounded to the nearest integer. We only show buckets with more than ten observations. The figure excludes 2.6% of the original data that is based on 3,187 beliefs of 700 subjects.

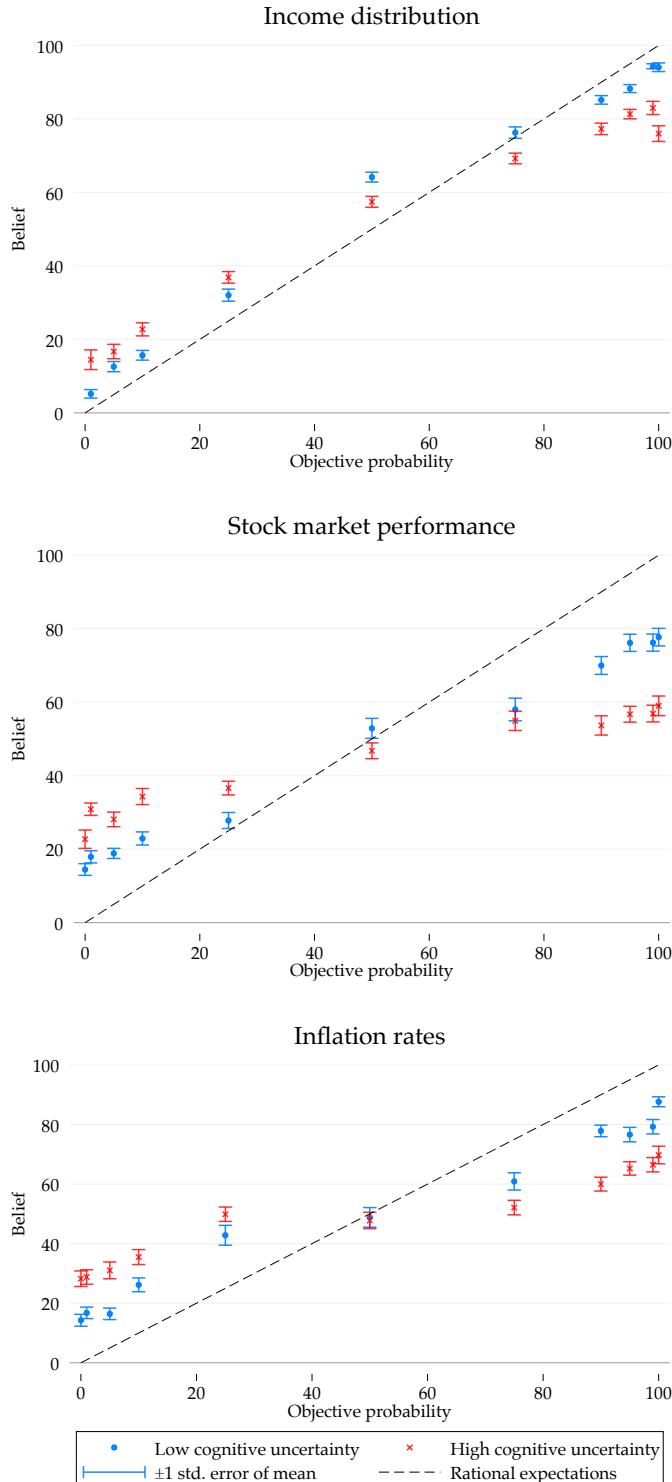


Figure 54: Survey beliefs as a function of objective probabilities after excluding censored beliefs, separately for subjects above / below average cognitive uncertainty. The partition is done separately for each probability bucket. In the top panel, the question asks for the probability that a randomly selected U.S. household earns less than \$x. In the middle panel, the question asks for the probability that in a randomly selected year the S&P500 increased by less than x%. In the bottom panel, the question asks for the probability that in a randomly selected year the inflation rate was less than x%. The figure excludes 6.61% of the original data that is based on 5,703 observations.

I.2 Censored cognitive uncertainty ranges

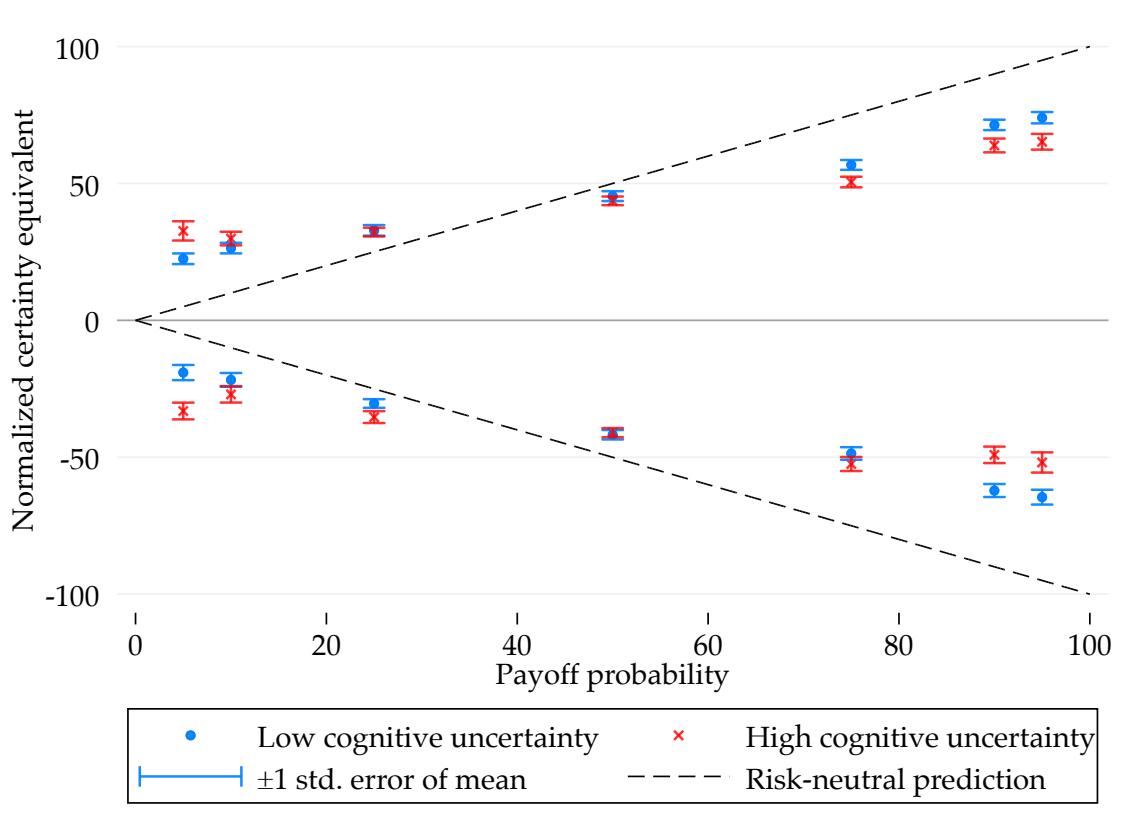


Figure 55: Probability weighting function excluding censored cognitive uncertainty ranges, separately for subjects above / below average cognitive uncertainty. The partition is done separately for each probability × gains / losses bucket. The plot shows averages and corresponding standard error bars. Normalized certainty equivalents (implied probability weights) are computed as certainty equivalent divided by payout probability. The figure excludes 23.25% of the original data that is based on 2,525 certainty equivalents from 700 subjects.

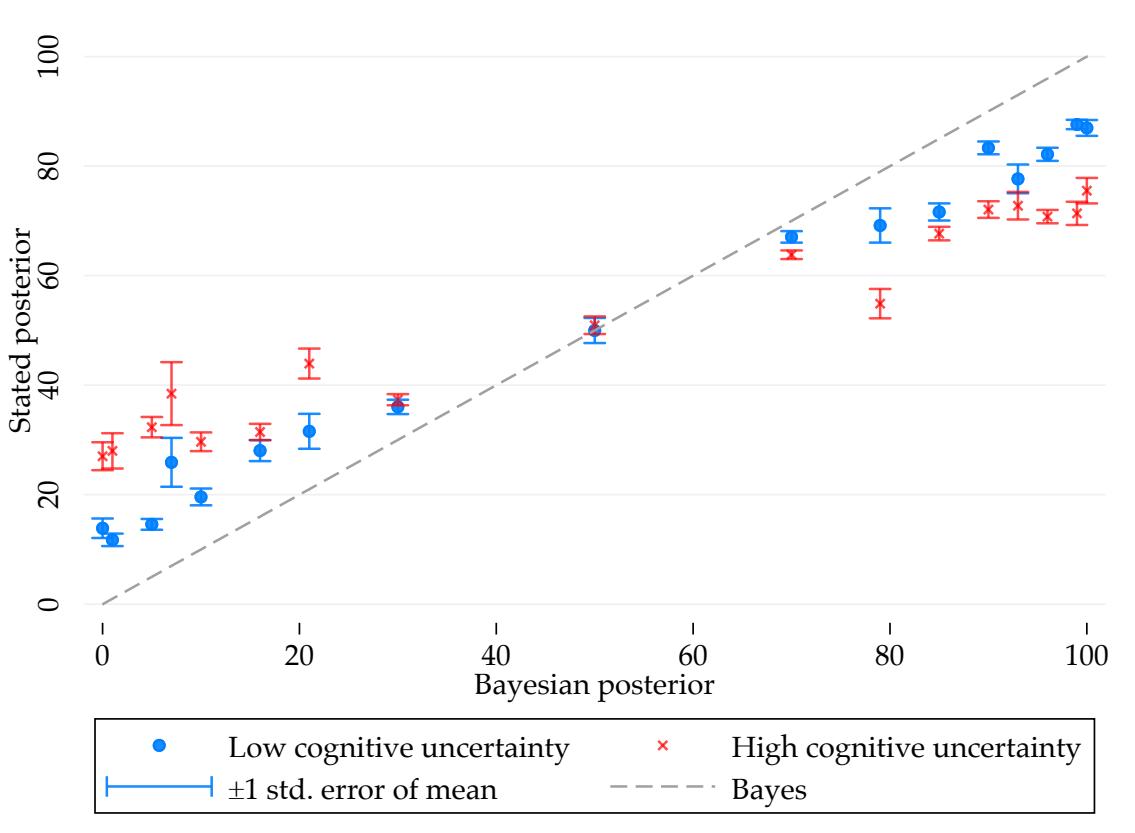


Figure 56: Relationship between average stated and Bayesian posteriors after excluding censored cognitive uncertainty ranges, separately for subjects above / below average cognitive uncertainty. The partition is done separately for each Bayesian posterior. Bayesian posteriors are rounded to the nearest integer. We only show buckets with more than ten observations. The figure excludes 25.93% of the original data that is based on 3,187 beliefs of 700 subjects.

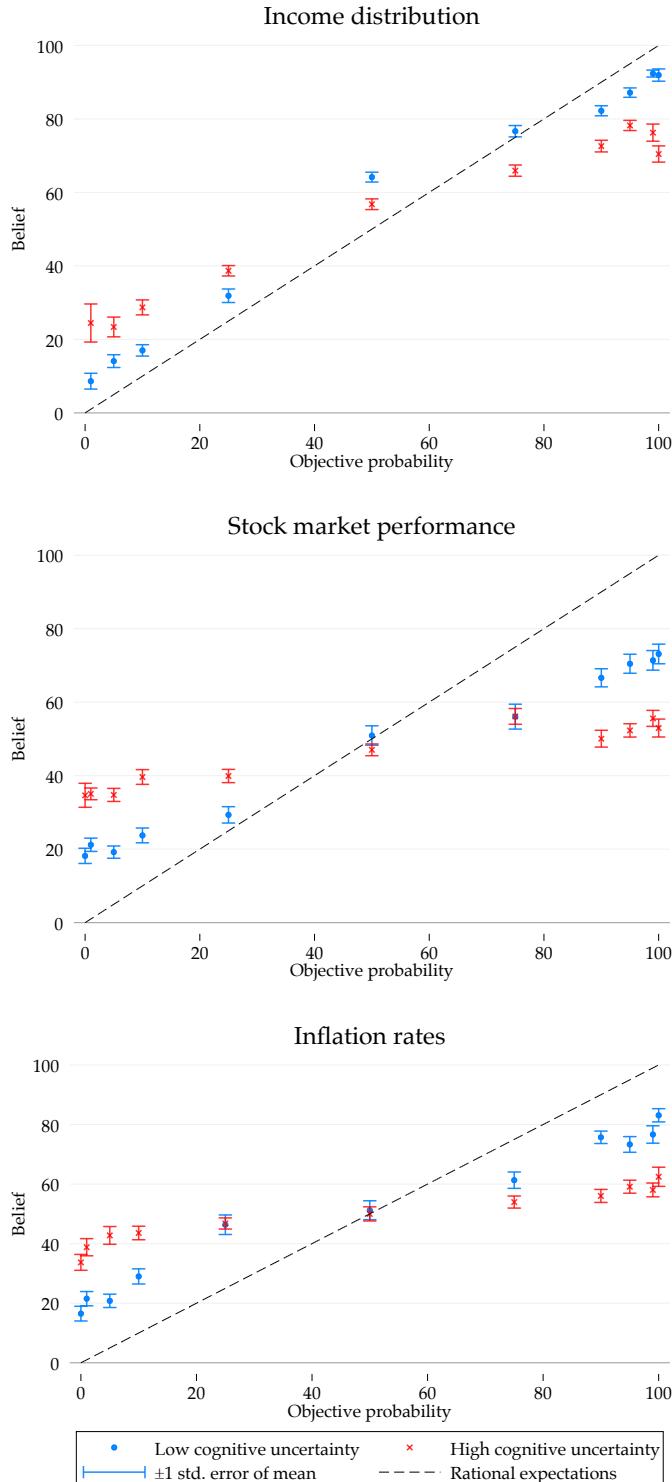


Figure 57: Survey beliefs as a function of objective probabilities after excluding censored cognitive uncertainty ranges, separately for subjects above / below average cognitive uncertainty. The partition is done separately for each probability bucket. In the top panel, the question asks for the probability that a randomly selected U.S. household earns less than x . In the middle panel, the question asks for the probability that in a randomly selected year the S&P500 increased by less than $x\%$. In the bottom panel, the question asks for the probability that in a randomly selected year the inflation rate was less than $x\%$. The figure excludes 16.66% of the original data that is based on 5,703 observations.

J Experimental Instructions and Control Questions

J.1 Treatment Baseline Risk

Welcome

Thank you for participating in this study.

This study will take approximately 25 minutes to complete.

You will earn a fixed reward of \$0.50. In addition, you can earn a bonus for completing the study. To complete the study, you will need to read all instructions carefully and answer the corresponding comprehension questions correctly. You receive a minimum bonus of \$1.20 when you complete the study.

You can earn an **additional bonus**. At the end of the study, one of the tasks will be randomly selected and your decision in this task determines your additional bonus. The chance that you get paid an additional bonus for the first part is 1 in 3.

Important information

- You should think about each task **independently** of all other tasks in this study. There is no point in strategizing across tasks.
- You will note that we sometimes ask you to work on similar-sounding tasks. These tasks might have similar answers, or very different ones. Please consider each individual task **carefully**.

Next

Part 1: Instructions

Please read these instructions carefully. There will be comprehension checks. If you fail these, you will be excluded from the study and only receive the reward of \$0.50.

In this study, there are various lotteries, all of which pay different amounts of money with different probabilities. An example lottery is:

With probability **70%**: Get \$ 20
With probability **30%**: Get \$ 5

This means that the lottery pays **either \$20 or \$5** (with different probabilities), but a lottery always only pays out one of the dollar amounts. The lotteries will actually be played out by the computer and determine your earnings in exactly the way we describe on the following pages.

In total, we have created 6 lotteries. For each lottery, you will make decisions on two consecutive decision screens.

- Decision screen 1: You will make various choices to indicate **how much** a lottery is worth to you.
- Decision screen 2: You will indicate **how certain** you are about how much **exactly** a lottery is worth to you.

Throughout the experiment, there are no right or wrong answers, because how much you like a lottery depends on your personal taste. We are only interested in learning about what you prefer.

Decision screen 1

On decision screen 1, you will be asked to choose which of two payment options you prefer. You will see choice lists such as the one below, where each row is a separate choice. In every list, the left-hand side option (Option A) is one of the lotteries that we generated, which is identical in all rows. The right-hand side option (Option B) is a safe payment that you would receive with certainty, so there is no risk attached. This safe payment in Option B increases as you go down the list.

You should consider the choice in each row **independently** of all other rows, because in the end one row will be randomly selected for payout and you will receive the option that you selected in that row. In some choice lists that you complete, you will receive a budget. If so, this budget will be listed at the top of your screen and paid out to you along with the payouts from the lottery.

Usually, people start by preferring Option A for small certain amounts (at the top of the list). At some point they switch to Option B as they proceed down the list, because the certain amount associated with Option B increases, so that Option B becomes more attractive. **Thus, an effective way to complete these choice lists is to determine in which row you would like to switch from Option A to Option B.**

Based on your decisions in this choice list, we assess **how much the lottery is worth to you by using those certain payments where you switch from Option A to Option B.** For example, in the example choice list below, the lottery would be worth between \$13 and \$14 to you, because this is where switching occurs.

Option A		Option B
	<input checked="" type="radio"/>	With certainty: Get \$ 0
	<input checked="" type="radio"/>	With certainty: Get \$ 1
	<input checked="" type="radio"/>	With certainty: Get \$ 2
	<input checked="" type="radio"/>	With certainty: Get \$ 3
	<input checked="" type="radio"/>	With certainty: Get \$ 4
	<input checked="" type="radio"/>	With certainty: Get \$ 5
	<input checked="" type="radio"/>	With certainty: Get \$ 6
	<input checked="" type="radio"/>	With certainty: Get \$ 7
	<input checked="" type="radio"/>	With certainty: Get \$ 8
	<input checked="" type="radio"/>	With certainty: Get \$ 9
With probability 70%: Get \$ 20	<input type="radio"/>	With certainty: Get \$ 10
With probability 30%: Get \$ 5	<input type="radio"/>	With certainty: Get \$ 11
	<input checked="" type="radio"/>	With certainty: Get \$ 12
	<input checked="" type="radio"/>	With certainty: Get \$ 13
	<input checked="" type="radio"/>	With certainty: Get \$ 14
	<input checked="" type="radio"/>	With certainty: Get \$ 15
	<input checked="" type="radio"/>	With certainty: Get \$ 16
	<input checked="" type="radio"/>	With certainty: Get \$ 17
	<input checked="" type="radio"/>	With certainty: Get \$ 18
	<input checked="" type="radio"/>	With certainty: Get \$ 19
	<input checked="" type="radio"/>	With certainty: Get \$ 20

Click "Next" to read about decision screen 2.

Next

Decision screen 2 (for the same lottery)

For any given lottery, you may actually be **uncertain about how much money it is really worth** to you. For some lotteries, you may know exactly how much they are worth to you. For other lotteries, you may feel uncertain about whether a lottery is worth, say, \$12, \$13, \$14, \$15, or \$16 to you.

On the second decision screen for a given lottery, we will hence ask you to use a slider to indicate **how certain** you are that your decisions on the first decision screen correspond **exactly** to how much the lottery is **worth to you**.

- The slider is linked to the statement below the slider. As you move around the slider, the values in this statement change. Your task is to position the slider in such a way that the statement corresponds to how you feel about the lottery.
- The **further you move the slider to the right, the more certain** you are about how much exactly the lottery is worth to you.
- For every step that you move the slider **further to the left (the less certain you are)**, the range of dollar values that you consider possible increases by \$0.50.

Example

Suppose that on the first decision screen you indicated that the lottery below is worth between \$13 and \$14 to you (because you switched from Option A to B between \$13 and \$14). Your screen would then look like this:

Example lottery:

With probability 70%: Get \$ 20
With probability 30%: Get \$ 5

How certain are you that to you this lottery is worth **exactly** the same as getting between \$13 and \$14 for sure?



I am certain that the lottery is worth **Use the slider!** to me.

Next

Your payment for part 1

If this part is randomly selected for payment in the end, your *additional* bonus will be determined as follows:

The computer will randomly select one choice list for payout, with equal probability. We then randomly select one of your decisions from this choice list (with equal probability), and this decision will be implemented for real payment.

Thus, you should take each decision independently of all other decisions as if it's the one that counts, because it may be.

Next

Comprehension questions

The questions below test your understanding of the instructions.

Important: If you fail to answer any one of these questions correctly, you will not be allowed to participate in the rest of the study, and you will not be able to earn a bonus.

1. Which one of the following statements is correct if the following lottery is played for you?

With probability **60%**: Get \$ 15
With probability **40%**: Get \$ 5

Please select one of the statements:

- It is possible that I get paid both \$15 and \$5, i.e., I may receive a total amount of \$20 from this lottery.
- I receive EITHER \$15 OR \$5 from this lottery.
- It is possible that I receive no money from this lottery.

2. Suppose a person made the decisions shown in the picture below. Which of the following statements is correct regarding these decisions?

Option A	Option B
<input checked="" type="radio"/>	With certainty: Get \$ 0
<input checked="" type="radio"/>	With certainty: Get \$ 1
<input checked="" type="radio"/>	With certainty: Get \$ 2
<input checked="" type="radio"/>	With certainty: Get \$ 3
<input checked="" type="radio"/>	With certainty: Get \$ 4
<input checked="" type="radio"/>	With certainty: Get \$ 5
<input checked="" type="radio"/>	With certainty: Get \$ 6
<input checked="" type="radio"/>	With certainty: Get \$ 7
<input checked="" type="radio"/>	With certainty: Get \$ 8
<input checked="" type="radio"/>	With certainty: Get \$ 9
<input checked="" type="radio"/>	With certainty: Get \$ 10
<input checked="" type="radio"/>	With certainty: Get \$ 11
<input checked="" type="radio"/>	With certainty: Get \$ 12
<input checked="" type="radio"/>	With certainty: Get \$ 13
<input checked="" type="radio"/>	With certainty: Get \$ 14
<input checked="" type="radio"/>	With certainty: Get \$ 15

Please select one of the statements:

- This person indicated that the lottery is worth more to them than \$9.
- This person indicated that the lottery is worth between \$3 and \$7 to them.
- This person indicated that the lottery is worth between \$8 and \$9 to them.

3. Now imagine that the person who made the decisions above is uncertain about how much exactly the lottery is worth to them.

This person is certain, however, that the lottery is worth between \$6 and \$11 to them. Please position the slider to accurately reflect this level of certainty:



This person is certain that the lottery is worth **Use the slider!** to them.

Next

J.2 Treatment Low Default Risk

Welcome

Thank you for participating in this study.

This study will take approximately 25 minutes to complete.

You will earn a fixed reward of \$0.50. In addition, you can earn a bonus for completing the study. To complete the study, you will need to read all instructions carefully and answer the corresponding comprehension questions correctly. You receive a **minimum bonus of \$1.20** when you complete the study.

You can earn an **additional bonus**. At the end of the study, one of the tasks will be randomly selected and your decision in this task determines your additional bonus. The chance that you get paid an additional bonus for the first part is 1 in 3.

Important information

- You should think about each task **independently** of all other tasks in this study. There is no point in strategizing across tasks.
- You will note that we sometimes ask you to work on similar-sounding tasks. These tasks might have similar answers, or very different ones. Please consider each individual task **carefully**.

Next

Part 1: Instructions

Please read these instructions carefully. There will be comprehension checks. If you fail these, you will be excluded from the study and only receive the reward of \$0.50.

In this study, there are various lotteries, all of which pay different amounts of money with different probabilities. To illustrate these probabilities, we will use the metaphor of colored balls in a bag. Imagine that there is a bag that contains 100 balls. Each of these balls has one of the following ten colors:

- red
- blue
- green
- orange
- brown
- pink
- black
- gold
- gray
- purple

The computer selects one of the 100 balls at random, where each ball is equally likely to get selected. Across lotteries, the number of balls of a given color might vary. Each color is associated with its own corresponding payout for you. An example lottery is:

70 out of 100 balls are red .	
If a red ball is drawn:	Get \$ 20
3 out of 100 balls are blue .	
If a blue ball is drawn:	Get \$ 5
3 out of 100 balls are green .	
If a green ball is drawn:	Get \$ 5
3 out of 100 balls are orange .	
If a orange ball is drawn:	Get \$ 5
3 out of 100 balls are brown .	
If a brown ball is drawn:	Get \$ 5
3 out of 100 balls are pink .	
If a pink ball is drawn:	Get \$ 5
3 out of 100 balls are black .	
If a black ball is drawn:	Get \$ 5
3 out of 100 balls are gold .	
If a gold ball is drawn:	Get \$ 5
3 out of 100 balls are gray .	
If a gray ball is drawn:	Get \$ 5
6 out of 100 balls are purple .	
If a purple ball is drawn:	Get \$ 5

You will always know how many balls of a given color are contained in the bag before making your decision. The lotteries will actually be played out by the computer and determine your earnings in exactly the way we describe on the following page.

In total, we have created 6 lotteries. For each lottery, you will make decisions on two consecutive decision screens.

- Decision screen 1: You will make various choices to indicate **how much** a lottery is worth to you.
- Decision screen 2: You will indicate **how certain** you are about how much **exactly** a lottery is worth to you.

Throughout the experiment, there are no right or wrong answers, because how much you like a lottery depends on your personal taste. We are only interested in learning about what you prefer.

Decision screen 1

On decision screen 1, you will be asked to choose which of two payment options you prefer. You will see choice lists such as the one below, where each row is a separate choice. In every list, the left-hand side option (Option A) is one of the lotteries that we generated, which is identical in all rows. Here, you can see how many balls of each color are contained in the bag that determines your payment. The right-hand side option (Option B) is a safe payment that you would receive *with certainty*, so there is no risk attached. This safe payment in Option B increases as you go down the list.

You should consider the choice in each row **independently** of all other rows, because in the end one row will be randomly selected for payout and you will receive the option that you selected in that row. In some choice lists that you complete, you will receive a budget. If so, this budget will be listed at the top of your screen and paid out to you along with the payouts from the lottery.

Usually, people start by preferring Option A for small certain amounts (at the top of the list). At some point they switch to Option B as they proceed down the list, because the certain amount associated with Option B increases, so that Option B becomes more attractive. **Thus, an effective way to complete these choice lists is to determine in which row you would like to switch from Option A to Option B.**

Based on your decisions in this choice list, we assess **how much the lottery is worth to you by using those certain payments where you switch from Option A to Option B**. For example, in the example choice list below, the lottery would be worth between \$13 and \$14 to you, because this is where switching occurs.

Option A		Option B
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 0
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 1
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 2
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 3
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 4
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 5
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 6
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 7
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 8
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 9
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 10
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 11
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 12
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$ 13
	<input type="radio"/> <input checked="" type="radio"/>	With certainty: Get \$ 14
	<input type="radio"/> <input checked="" type="radio"/>	With certainty: Get \$ 15
	<input type="radio"/> <input checked="" type="radio"/>	With certainty: Get \$ 16
	<input type="radio"/> <input checked="" type="radio"/>	With certainty: Get \$ 17
	<input type="radio"/> <input checked="" type="radio"/>	With certainty: Get \$ 18
	<input type="radio"/> <input checked="" type="radio"/>	With certainty: Get \$ 19
	<input type="radio"/> <input checked="" type="radio"/>	With certainty: Get \$ 20

Click "Next" to read about decision screen 2.

[Next](#)

Decision screen 2 (for the same lottery)

For any given lottery, you may actually be **uncertain about how much money it is really worth** to you. For some lotteries, you may know exactly how much they are worth to you. For other lotteries, you may feel uncertain about whether a lottery is worth, say, \$12, \$13, \$14, \$15, or \$16 to you.

On the second decision screen for a given lottery, we will hence ask you to use a slider to indicate **how certain** you are that your decisions on the first decision screen correspond **exactly** to how much the lottery is **worth to you**.

- The slider is linked to the statement below the slider. As you move around the slider, the values in this statement change. Your task is to position the slider in such a way that the statement corresponds to how you feel about the lottery.
- The **further you move the slider to the right, the more certain** you are about how much exactly the lottery is worth to you.
- For every step that you move the slider **further to the left (the less certain you are)**, the range of dollar values that you consider possible increases by \$0.50.

Example

Suppose that on the first decision screen you indicated that the lottery below is worth between \$13 and \$14 to you (because you switched from Option A to B between \$13 and \$14). Your screen would then look like this:

Example lottery:

70 out of 100 balls are red.
If a red ball is drawn: Get \$ 20

3 out of 100 balls are blue.
If a blue ball is drawn: Get \$ 5

3 out of 100 balls are green.
If a green ball is drawn: Get \$ 5

3 out of 100 balls are orange.
If a orange ball is drawn: Get \$ 5

3 out of 100 balls are brown.
If a brown ball is drawn: Get \$ 5

3 out of 100 balls are pink.
If a pink ball is drawn: Get \$ 5

3 out of 100 balls are black.
If a black ball is drawn: Get \$ 5

3 out of 100 balls are gold.
If a gold ball is drawn: Get \$ 5

3 out of 100 balls are gray.
If a gray ball is drawn: Get \$ 5

6 out of 100 balls are purple.
If a purple ball is drawn: Get \$ 5

How certain are you that to you this lottery is worth exactly the same as getting between \$13 and \$14 for sure?

very uncertain

completely certain

I am certain that the lottery is worth Use the slider! to me.

Next

Your payment for part 1

If this part is randomly selected for payment in the end, your *additional* bonus will be determined as follows:

The computer will randomly select one choice list for payout, with equal probability. We then randomly select one of your decisions from this choice list (with equal probability), and this decision will be implemented for real payment.

Thus, you should take each decision independently of all other decisions as if it's the one that counts, because it may be.

Next

Comprehension questions

The questions below test your understanding of the instructions.

Important: If you fail to answer any one of these questions correctly, you will not be allowed to participate in the rest of the study, and you will not be able to earn a bonus.

1. Which one of the following statements is correct if the following lottery is played for you?

60 out of 100 balls are **red**.
If a **red** ball is drawn: **Get \$ 15**

4 out of 100 balls are **blue**.
If a **blue** ball is drawn: **Get \$ 5**

4 out of 100 balls are **green**.
If a **green** ball is drawn: **Get \$ 5**

4 out of 100 balls are **orange**.
If a **orange** ball is drawn: **Get \$ 5**

4 out of 100 balls are **brown**.
If a **brown** ball is drawn: **Get \$ 5**

4 out of 100 balls are **pink**.
If a **pink** ball is drawn: **Get \$ 5**

4 out of 100 balls are **black**.
If a **black** ball is drawn: **Get \$ 5**

4 out of 100 balls are **gold**.
If a **gold** ball is drawn: **Get \$ 5**

4 out of 100 balls are **gray**.
If a **gray** ball is drawn: **Get \$ 5**

8 out of 100 balls are **purple**.
If a **purple** ball is drawn: **Get \$ 5**

Please select one of the statements:

- It is possible that I get paid both \$15 and \$5, i.e., I may receive a total amount of \$20 from this lottery.
- Only one of the colors will be drawn, hence I receive EITHER \$15 OR \$5 from this lottery.
- It is possible that I receive no money from this lottery.

2. Suppose a person made the decisions shown in the picture below. Which of the following statements is correct regarding these decisions?

Option A		Option B
60 out of 100 balls are red. If a red ball is drawn: Get \$15	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$0
4 out of 100 balls are blue. If a blue ball is drawn: Get \$5	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$1
4 out of 100 balls are green. If a green ball is drawn: Get \$5	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$2
4 out of 100 balls are orange. If a orange ball is drawn: Get \$5	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$3
4 out of 100 balls are brown. If a brown ball is drawn: Get \$5	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$4
4 out of 100 balls are pink. If a pink ball is drawn: Get \$5	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$5
4 out of 100 balls are black. If a black ball is drawn: Get \$5	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$6
4 out of 100 balls are gold. If a gold ball is drawn: Get \$5	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$7
4 out of 100 balls are gray. If a gray ball is drawn: Get \$5	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$8
8 out of 100 balls are purple. If a purple ball is drawn: Get \$5	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$9
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$10
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$11
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$12
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$13
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$14
	<input checked="" type="radio"/> <input type="radio"/>	With certainty: Get \$15

Please select one of the statements:

- This person indicated that the lottery is worth more to them than \$9.
- This person indicated that the lottery is worth between \$3 and \$7 to them.
- This person indicated that the lottery is worth between \$8 and \$9 to them.

3. Now imagine that the person who made the decisions above is uncertain about how much exactly the lottery is worth to them.

This person is certain, however, that the lottery is worth between \$6 and \$11 to them. Please position the slider to accurately reflect this level of certainty:



This person is certain that the lottery is worth **Use the slider!** to them.

Next

J.3 Treatment Baseline Beliefs

Welcome

Thank you for participating in this study.

This study will take approximately 25 minutes to complete.

You will earn a fixed reward of \$0.50. In addition, you can earn a bonus if you complete the study. To complete the study, you will need to read all instructions carefully and answer the corresponding comprehension questions correctly. You receive a minimum bonus of \$1.20 when you complete the study.

You can earn an **additional bonus**. At the end of the study, one of the tasks will be randomly selected and your decision in this task determines your additional bonus. The chance that you get paid an additional bonus for the first part is 1 in 3.

Important information

- You should think about each task **independently** of all other tasks in this study. There is no point in strategizing across tasks.
- You will note that we sometimes ask you to work on similar-sounding tasks. These tasks might have similar answers, or very different ones. Please consider each individual task **carefully**.
- Whenever a task involves a random draw, then this random draw will **actually be implemented for you by the computer** in exactly the way it is described to you in the task.

[Next](#)

Part 1: The guessing tasks

Please read these instructions carefully. We will test your understanding of them later.

In this study, you will be asked to complete **6 guessing tasks**.

In each guessing task, there are two bags, "bag A" and "bag B". Each bag contains 100 balls, some of which are red and some of which are blue. One of the bags is selected at random by the computer as described below. You will not observe which bag was selected. Instead, the computer will then randomly draw one or several balls from the secretly selected bag, and will show these balls to you. Your task is to **guess which bag was selected** based on the available information. The exact procedure is described below:

Task setup

- There is a deck of cards that consists of 100 cards. Each card in the deck either has "bag A" or "bag B" written on it. You will be informed about **how many** of these 100 cards have "bag A" or "bag B" written on them.
- There are two bags, "bag A" and "bag B". You will be informed about **how many red and blue balls** each bag contains.

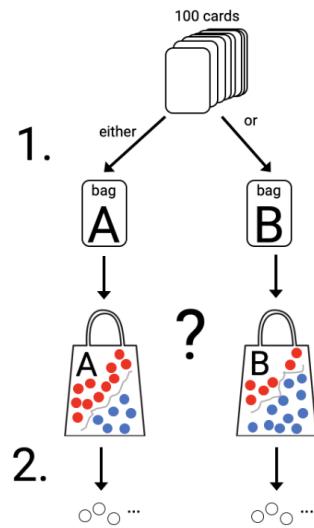
Sequence of events

1. The computer **randomly selects one** of the 100 cards, with equal probability.
If a "bag A" card was drawn, bag A is selected.
If a "bag B" card was drawn, bag B is selected.
2. Next, the computer **randomly draws one or more** balls from the secretly selected bag. Each ball is equally likely to get selected.
Importantly, if more than one ball is drawn, the computer draws these balls **with replacement**. This means that after a ball has been drawn and taken out of the bag, it gets replaced by a ball of the same color. Thus, the probability of a red or blue ball being drawn does not depend on whether previous draws were red or blue.
The computer will **inform you about the color** of the randomly drawn balls.

You then make your guess by **stating a probability between 0% and 100%** that bag A was drawn. The corresponding probability that bag B was drawn is 100 minus your stated probability that bag A was drawn.

Please note:

- The number of "bag A" and "bag B" cards **varies across the guessing tasks**.
- The number of red and blue balls in each bag **varies across the guessing tasks**.
- The computer **draws a new card in each task**, so you should **think about which bag was selected in a task independently** of all other tasks.



Next

Your payment for part 1

There is a **prize of \$10.00**. Whether or not you receive the \$10.00 depends on how much probability you assigned to the bag that was *actually drawn* in that problem.

This means: if bag A was selected, your chances of receiving \$10.00 are greater the higher the probability you assigned to bag A. If bag A was not selected, your chances of receiving \$10.00 are greater the lower the probability you assigned to bag A. In case you're interested, the specific method that determines whether you get the prize is explained below:

A number q between 0 and 2500 is randomly drawn by the computer.

If bag A was selected in that problem, you receive \$10.00 if the square of the probability (in percent) that you assigned to bag B is *lower than q* .

If bag A was not selected in that problem, you receive \$10.00 if the square of the probability (in percent) that you assigned to bag A is *lower than q* .

All this means that, in order to earn as much money as possible, you should try to give your best estimate of the probability that each bag was drawn. For example, if you are 80% sure that bag A was selected and 20% sure that bag B was selected, you should allocate probability 80% to bag A and 20% to bag B.

Your certainty about your guess

The optimal guess

Using the laws of probability, the computer computes a **statistically correct statement of the probability that bag A was drawn**, based on all the information available to you. This **optimal guess** does not rely on information that you do not have. It is just the best possible (this means: payoff-maximizing) estimate given the available information. In technical terms, this guess is based on a statistical rule called Bayes' Law.

Your certainty about your guess

In any given task, you may actually be **uncertain about whether your probability guess corresponds to the optimal guess**. On a separate decision screen, we will hence ask you to use a slider to indicate **how certain** you are that your guess equals the optimal guess.

- The slider is linked to the statement below the slider. As you move around the slider, the values in this statement change. Your task is to position the slider in such a way that the statement corresponds to how close you think your guess is to the optimal guess.
- The further you move the slider to the right, the more certain you are that your guess is close to the optimal guess.
- For every step that you move the slider further to the left (the less certain you are), the range of probability values that you consider possible increases by 2.

Example

Suppose that you stated a guess of 80%. Your screen would then look like this:

How **certain** are you that **the optimal guess is exactly 80%**?

Use the slider to complete the statement below.



I am certain that the optimal guess of the **probability that bag A was drawn** is **Use the slider!**.

Replacing your guess by the optimal guess

In each task, if you are **uncertain about what you should guess**, you may **increase your chances of winning** the prize of \$10.00 by paying money to **replace your guess with the optimal guess**.

For this purpose, in each task, you receive a budget of \$3.00. You then have to state the **highest amount (between \$0.00 and \$3.00) that you are willing to pay to replace your guess with the optimal guess**. In the end, a price p between \$0.00 and \$3.00 will be randomly determined by the computer. You will purchase the optimal guess at price p if p is below your stated amount, and you will not purchase the optimal guess and keep your endowment otherwise. If you buy the optimal guess, your own guess is replaced with the optimal guess with 10% probability. You only have to pay for the optimal guess if your guess actually gets replaced.

[Next](#)

Sequence of events in each task

You will be asked to complete 6 guessing tasks. For each task, there will be **3 decision screens**:

Decision screen 1

You will be asked to state the highest amount that you are **willing to pay to replace your own guess** (that you will make on decision screen 2) with the **optimal guess in this task**.

Decision screen 2

You have to guess which bag was selected by entering a probability (between 0 and 100) that expresses **how likely** you think it is that bag A as opposed to bag B have been selected.

Decision screen 3

You will be asked to indicate **how certain** you are that the guess you provided on decision screen 1 **equals the optimal guess in this task**.

[Next](#)

Comprehension questions

The following questions test your understanding of the instructions.

Important: If you fail to answer any one of these questions correctly, you will not be allowed to participate in the rest of the study, and you will not be able to earn a bonus.

Which statement about the number of cards corresponding to each bag is correct?

- The number of "bag A" cards is the same in all tasks.
- The exact number of cards corresponding to each bag may vary across tasks.

Which statement about the allocation of red and blue balls in the bags is correct?

- The exact fractions of red and blue balls in each bag may vary across tasks.
- The fraction of red balls in each bag is the same in all tasks, and bag A always contains the most red balls.

Which statement about your bonus payment is correct?

- Purchasing the optimal guess always means that I would make a loss.
- By purchasing the optimal guess I can potentially increase my earnings in the guessing tasks if my own guess is not sufficiently good.

In order to maximize your overall profit, how should you determine how much you are willing to pay for the optimal guess?

- I should be willing to pay a high amount if I think that my own guess will probably not be optimal.
- I should be willing to pay a high amount if I think that my own guess will probably be optimal.
- I should never pay money for the optimal guess because it costs money but has no benefits.

[Next](#)

J.4 Treatment Low Default Beliefs

Welcome

Thank you for participating in this study.

This study will take approximately 25 minutes to complete.

You will earn a fixed reward of \$0.50. In addition, you can earn a bonus if you complete the study. To complete the study, you will need to read all instructions carefully and answer the corresponding comprehension questions correctly. You receive a minimum bonus of \$1.20 when you complete the study.

You can earn an **additional bonus**. At the end of the study, one of the tasks will be randomly selected and your decision in this task determines your additional bonus. The chance that you get paid an additional bonus for the first part is 1 in 3.

Important information

- You should think about each task **independently** of all other tasks in this study. There is no point in strategizing across tasks.
- You will note that we sometimes ask you to work on similar-sounding tasks. These tasks might have similar answers, or very different ones. Please consider each individual task **carefully**.
- Whenever a task involves a random draw, then this random draw will **actually be implemented for you by the computer** in exactly the way it is described to you in the task.

Next

Part 1: The guessing tasks

Please read these instructions carefully. We will test your understanding of them later.

In this study, you will be asked to complete **6 guessing tasks**.

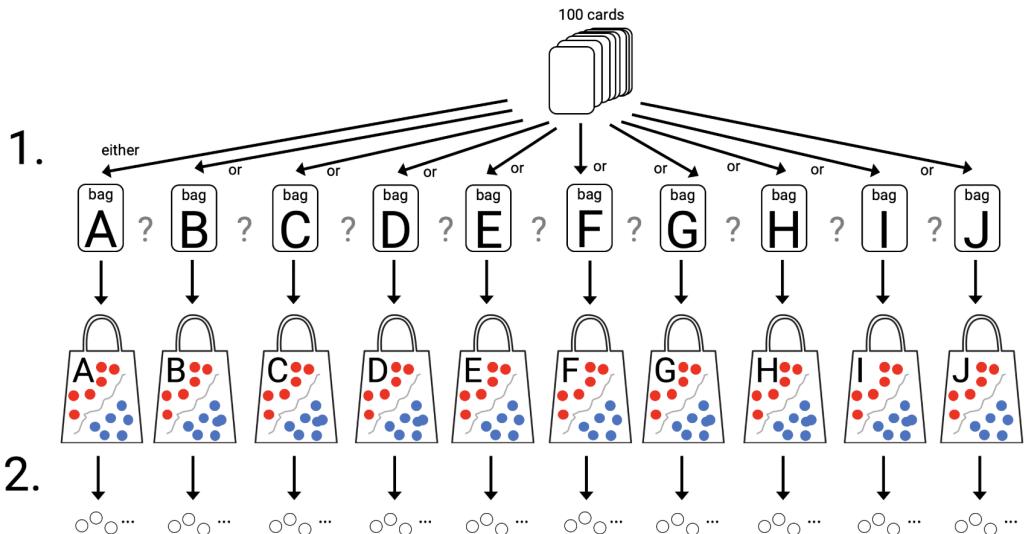
In each guessing task, there are ten bags, "bag A", "bag B", "bag C", "bag D", "bag E", "bag F", "bag G", "bag H", "bag I" and "bag J". Each bag contains 100 balls, some of which are red and some of which are blue. One of the bags is selected at random by the computer as described below. You will not observe which bag was selected. Instead, the computer will then randomly draw one or several balls from the secretly selected bag, and will show these balls to you. Your task is to **guess which bag was selected** based on the available information. The exact procedure is described below:

Task setup

- There is a deck of cards that consists of 100 cards. Each card in the deck either has "bag A", "bag B", "bag C", "bag D", "bag E", "bag F", "bag G", "bag H", "bag I" or "bag J" written on it.
- You will be informed about **how many** of these 100 cards have "bag A", "bag B", ..., or "bag J" written on them.
- There are ten bags, "bag A" through "bag J". You will be informed about **how many red and blue balls** each bag contains.

Sequence of events

1. The computer **randomly selects one** of the 100 cards, with equal probability.
If a "bag A" card was drawn, bag A is selected.
If a "bag B" card was drawn, bag B is selected.
... etc. ...
2. Next, the computer **randomly draws one or more** balls from the secretly selected bag. Each ball is equally likely to get selected. Importantly, if more than one ball is drawn, the computer draws these balls **with replacement**. This means that after a ball has been drawn and taken out of the bag, it gets replaced by a ball of the same color. Thus, the probability of a red or blue ball being drawn does not depend on whether previous draws were red or blue.
The computer will **inform you about the color** of the randomly drawn balls.



You then make your guess by **stating a probability between 0% and 100%** that bag A was drawn. The corresponding probability that one of bag B, C, D, E, F, G, H, I or J was drawn is 100 minus your stated probability that bag A was drawn.

Please note:

- The number of "bag A", "bag B", ..., and "bag J" cards **varies across the guessing tasks**.
- The number of red and blue balls in each bag **varies across the guessing tasks**.
- The computer **draws a new card in each task**, so you should **think about which bag was selected in a task independently** of all other tasks.

[Next](#)

Your payment for part 1

There is a **prize of \$10.00**. Whether or not you receive the \$10.00 depends on how much probability you assigned to the bag that was *actually* drawn in that problem.

This means: if bag A was selected, your chances of receiving \$10.00 are greater the higher the probability you assigned to bag A. If bag A was not selected, your chances of receiving \$10.00 are greater the lower the probability you assigned to bag A. In case you're interested, the specific method that determines whether you get the prize is explained below:

A number q between 0 and 2500 is randomly drawn by the computer.

If bag A was selected in that problem, you receive \$10.00 if the square of the probability (in percent) that you assigned to all other bags is *lower than* q .

If bag A was not selected in that problem, you receive \$10.00 if the square of the probability (in percent) that you assigned to bag A is *lower than* q .

All this means that, in order to earn as much money as possible, you should try to give your best estimate of the probability that each bag was drawn. For example, if you are 80% sure that bag A was selected and 20% sure that any bag other than bag A was selected, you should allocate probability 80% to bag A and 20% to bag B.

[Next](#)

Your certainty about your guess

The optimal guess

Using the laws of probability, the computer computes a **statistically correct statement of the probability that bag A was drawn**, based on all the information available to you. This **optimal guess** does not rely on information that you do not have. It is just the best possible (this means: payoff-maximizing) estimate given the available information. In technical terms, this guess is based on a statistical rule called Bayes' Law.

Your certainty about your guess

In any given task, you may actually be **uncertain about whether your probability guess corresponds to the optimal guess**. On a separate decision screen, we will hence ask you to use a slider to indicate **how certain** you are that your guess equals the optimal guess.

- The slider is linked to the statement below the slider. As you move around the slider, the values in this statement change. Your task is to position the slider in such a way that the statement corresponds to how close you think your guess is to the optimal guess.
- The further you move the slider to the right, the more certain you are that your guess is close to the optimal guess.
- For every step that you move the slider further to the left (the less certain you are), the range of probability values that you consider possible increases by 2.

Example

Suppose that you stated a guess of 80%. Your screen would then look like this:

How **certain** are you that **the optimal guess is exactly 80%**?

Use the slider to complete the statement below.



I am certain that the optimal guess of the **probability that bag A was drawn** is **Use the slider!**.

Replacing your guess by the optimal guess

In each task, if you are **uncertain about what you should guess**, you may **increase your chances of winning** the prize of \$10.00 by paying money to **replace your guess with the optimal guess**.

For this purpose, in each task, you receive a budget of \$3.00. You then have to state the **highest amount (between \$0.00 and \$3.00) that you are willing to pay to replace your guess with the optimal guess**. In the end, a price p between \$0.00 and \$3.00 will be randomly determined by the computer. You will purchase the optimal guess at price p if p is below your stated amount, and you will not purchase the optimal guess and keep your endowment otherwise. If you buy the optimal guess, your own guess is replaced with the optimal guess with 10% probability. You only have to pay for the optimal guess if your guess actually gets replaced.

[Next](#)

Sequence of events in each task

You will be asked to complete 6 guessing tasks. For each task, there will be **3 decision screens**:

Decision screen 1

You will be asked to state the highest amount that you are **willing to pay to replace your own guess** (that you will make on decision screen 2) with the **optimal guess in this task**.

Decision screen 2

You have to guess which bag was selected by entering a probability (between 0 and 100) that expresses **how likely** you think it is that bag A as opposed to any bag other than bag A have been selected.

Decision screen 3

You will be asked to indicate **how certain** you are that the guess you provided on decision screen 1 **equals the optimal guess in this task**.

Next

Comprehension questions

The following questions test your understanding of the instructions.

Important: If you fail to answer any one of these questions correctly, you will not be allowed to participate in the rest of the study, and you will not be able to earn a bonus.

Which statement about the number of cards corresponding to each bag is correct?

- The number of "bag A" cards is the same in all tasks.
- The exact number of cards corresponding to each bag may vary across tasks.

Which statement about the allocation of red and blue balls in the bags is correct?

- The exact fractions of red and blue balls in each bag may vary across tasks.
- The fraction of red balls in each bag is the same in all tasks, and bag A always contains the most red balls.

Which statement about your bonus payment is correct?

- Purchasing the optimal guess always means that I would make a loss.
- By purchasing the optimal guess I can potentially increase my earnings in the guessing tasks if my own guess is not sufficiently good.

In order to maximize your overall profit, how should you determine how much you are willing to pay for the optimal guess?

- I should be willing to pay a high amount if I think that my own guess will probably not be optimal.
- I should be willing to pay a high amount if I think that my own guess will probably be optimal.
- I should never pay money for the optimal guess because it costs money but has no benefits.

Next

J.5 Survey Expectations

Part 2 of this study

You have completed part 1. We will now continue with part 2 of this study.

Next

Your payment for part 2

In this part, there will be 3 tasks. At the end, one of the tasks will be randomly selected to count for your potential bonus. The chance that you get paid an additional bonus for this part is 1 in 3.

In each task, you will be asked to **state a guess in the form a probability estimate (between 0 and 100)**.

There is a **prize of \$2.00**. In each guessing task, there are two possible events, call them A and B. One of the two events actually occurred, the other did not. Whether or not you receive the \$2.00 depends on **how much probability** you assigned to the **event that actually occurred**.

If event A occurred, your chances of receiving \$2.00 are greater the higher the probability you assigned to event A. If event B occurred, your chances of receiving \$2.00 are greater the higher the probability you assigned to event B.

In case you're interested, the specific method that determines whether you get the prize is explained below:

A number **q** between 0 and 2,500 is randomly drawn by the computer.

- If event A occurred in that problem, you receive \$2.00 if the square of the probability (in percent) that you assigned to event B is *lower than q*.
- If event B occurred in that problem, you receive \$2.00 if the square of the probability (in percent) that you assigned to event A is *lower than q*.

All this means that, in order to earn as much money as possible, you should try to provide your best probability estimate in each task.

Next

The study begins on the next page

We will now start with part 2 of study.

Next