Moral Universalism and the Structure of Ideology*

Benjamin Enke Ricardo Rodríguez-Padilla
Florian Zimmermann

April 19, 2022

Abstract

Throughout the Western world, people's policy views are correlated across domains in a strikingly similar fashion. This paper proposes that what partly explains the structure of ideology is *moral universalism*: the extent to which people exhibit the same level of altruism and trust towards strangers as towards in-group members. In new large-scale multinational surveys, heterogeneity in universalism descriptively explains why some people support redistribution, health care, environmental protection, affirmative action and foreign aid, while others advocate for spending on the military, law enforcement and border protection. Universalism is a substantially stronger predictor of policy views and ideological constraint than variables such as income, wealth, education, religiosity, or beliefs about government efficiency. Consistent with the idea that universalism shapes policy views, we further document that the left-right divide on redistribution, environmental protection or foreign aid strongly attenuates or even reverses when people evaluate less universalist implementations of these policies.

JEL classification: D72, D01

Keywords: Moral universalism, ideological constraint, behavioral political economy

*For valuable comments and discussions we thank Nicola Gennaioli, four very constructive referees, Alberto Alesina, Roland Bénabou, Maxim Boycko, Matt Gentzkow, Ed Glaeser, Vincent Pons, Jesse Shapiro, Andrei Shleifer, Jim Snyder, Guido Tabellini, Marco Tabellini and many seminar audiences. Phoebe Cai, Aditi Chitkara, Robert Fluegge and Jan Knuf provided outstanding research assistance. We also thank Spencer Kwon, Joakim Semb and Julien Senn for help with the surveys. We acknowledge funding from the William F. Milton Fund, the Harvard Asia Center, and a Harvard Inequality in America Grant. Enke: Harvard University and NBER, enke@fas.harvard.edu; Rodríguez-Padilla: Harvard University, rodriguezpadilla@g.harvard.edu; Zimmermann: University of Bonn and Institute on Behavior and Inequality, florian.zimmermann@briq-institute.org.

1 Introduction

A key stylized fact in the study of political economy is the existence of *ideological constraint*: people's policy views are systematically correlated across economic and social domains. As we confirm using new large-scale survey data, people in one cluster generally desire government expenditure on foreign aid, affirmative action, environmental protection, welfare, and universal health care, while people in another cluster support government spending on the military, police and law enforcement, and border control.

While these ideological clusters are sufficiently obvious that they have almost come to define how the public thinks about political disagreements, their structure is also somewhat puzzling. First, as we review in detail below, the prevailing ideological clusters are very similar across many rich Western countries, even though these often exhibit considerable differences in electoral systems, party structures, and ethnic composition. Second, despite the widespread occurrence of these clusters, there is no obvious "law of nature" that explains why these *particular* bundles of policy views should prevail in the first place. For example, it is not obvious why support for redistribution is always correlated with support for environmental protection rather than support for a strong military. Still, the striking similarity of the correlations in issue positions across countries suggests that these bundles reflect more than coincidence.

The central proposition of this paper is that what imposes this particular structure on the space of policy views is heterogeneity in *moral universalism*: the extent to which people exhibit the same level of altruism and trust towards strangers as towards in-group members. Universalism is not about a person's overall *level* of altruism or trust, but instead about its *slope* as a function of social distance, such that group- or place-based identities are less important to universalists than to communitarians.

Our analysis of the role of universalism for policy views is motivated by a growing body of work that argues that people's political ideology and voting behavior to a large extent reflect heterogeneity in moral values related to universalism, rather than mostly material concerns (e.g., Graham et al., 2009; Haidt, 2012; Waytz et al., 2019; Enke, 2020). Building on this body of work, our investigation is based on the simple intuition that one of the two aforementioned ideological clusters consists of policies that require high altruism towards, and trust in, socially distant individuals, while the other cluster comprises policies that are primarily aimed at protecting, and caring for, certain in-groups. For instance, foreign aid, affirmative action, environmental protection and federal redistribution are all highly universalist policies in that their beneficiaries will often be socially or geographically distant strangers. On the other hand, public spending on the military or border control is often designed to erect boundaries between "us" and "them," which may at least partly reflect the communitarian moral stance that com-

patriots (in-group members) deserve higher priority, as well as the belief that outsiders cannot be trusted.

Based on these hypotheses, this paper presents a new set of descriptive stylized facts to make four contributions. First, we collect rich new survey data to show that people's desired expenditure shares are strongly partitioned into two distinct clusters, and that these clusters are extremely similar across Australia, France, Germany, Sweden and the United States. Second, by supplementing our data-collection exercise with rich and experimentally-validated measures of universalism in altruism and universalism in trust, we show that universalism is strongly predictive of policy views in a way that rationalizes the observed clusters, in an almost identical fashion across countries. Third, we provide an extensive set of benchmarking analyses that suggest that universalism is quantitatively substantially more important for explaining policy views and ideological constraint than traditional political economy variables such as income, wealth, equity-efficiency preferences or beliefs about government efficiency. Fourth, in an attempt to disentangle supply- and demand-side explanations, we document that people's views on broad policy domains such as redistribution or foreign aid strongly depend on how universalist vs. communitarian the specific implementation of a policy is. For example, while left- and right-wingers differ in their views on national redistribution in familiar ways, this relationship strongly attenuates or even reverses when people are polled about local redistribution.

Our internet surveys are pre-registered and comprise nationally representative samples in each of five countries, for a total of about 11,000 respondents. To measure universalism in altruism, we implement structured decision tasks. In each task, a respondent is endowed with the hypothetical sum of $100 and is asked to split the money between two equally rich individuals: (i) a randomly selected member of a specific social (in-) group who lives in their own country of residence and (ii) a randomly-selected person who lives in their own country of residence. For example, in one question, a U.S. participant is asked to split hypothetical $100 between a member of their extended family and a randomly-selected person from the United States. In addition to these questions that measure "domestic universalism", we also measure "foreign universalism" and "global universalism" through money allocation tasks that involve different types of foreigners. While all of our survey questions are hypothetical in nature, they underwent an extensive selection and experimental validation procedure, and have been shown to be correlated with real donation decisions (see Enke et al., 2021).

Using an analogous procedure, we estimate respondents' universalism in trust by asking them to indicate who of the individuals described above they trust more. In our data, universalism in altruism and universalism in trust are highly correlated, which suggests that they capture the same underlying psychology, which we refer to as "moral

universalism." Respondents exhibit large variation in universalism: some participants always split their money or trust equally, while others consistently share more money with, and trust more, members of their in-groups.

We supplement these measures of universalism with detailed questions on respondents' policy views. To balance richness and quantitative interpretability against the cognitive burden placed on participants, we elicit both quantitative information on desired government expenditure levels for specific policies, and simple Likert scale measures.

Looking at the link between policy views and universalism, we find that desired expenditures are strongly correlated with universalism in the ways we hypothesized and pre-registered. Universalism is *positively* correlated with desired expenditure on welfare payments, environment, affirmative action, foreign aid, and – to a lesser extent – universal health care. Moreover, universalism is *negatively* correlated with desired expenditure on border control, military, and law enforcement and police. In this sense, universalism reproduces the structure of policy views that we attempt to explain in this paper. In terms of quantitative magnitudes, moving from universalism of zero (allocating all money or trust to the in-group) to equal splits is associated with changes in policy views of between 10% and 80% of a standard deviation.

These correlations are robust and general in the following three ways. (i) The results are qualitatively very similar when we consider either universalism in altruism or universalism in trust, though quantitatively universalism in altruism is typically about twice as important. (ii) The relationship between universalism and policy preferences is robust against controlling for rich measures of income, wealth, religiosity, education, urbanicity and beliefs about government efficiency, among others. (iii) The results are strikingly similar across the United States, Australia, France, Germany, and Sweden.

To put our results in perspective, we implement a series of exercises that benchmark universalism against a rich set of individual characteristics. In a first step, we assess the degree to which universalism and other characteristics contribute to the magnitude of ideological constraint, defined as the average intracorrelation of people's policy views. In conservative estimations, we find that ideological constraint would decrease by 25% if there was no heterogeneity in universalism. To put this magnitude into perspective, our estimations suggest that ideological constraint would not change at all if there was no heterogeneity in, for example, income, wealth, education, equity-efficiency preferences or beliefs about government efficiency.

In a second step of benchmarking our results on universalism, we use simple machine learning techniques to identify a candidate set of "predictors" of policy views, which partly consists of complex interactions of individual characteristics. We then study whether any of these candidate predictors meaningfully predict the structure of policy views. We find that the candidate predictors are often correlated with desired expendi-

tures in important and known ways. At the same time, universalism is the only variable in our data that meaningfully organizes the key pattern we are trying to explain: simultaneous support for government spending in the domains of welfare, universal health care, environmental protection, affirmative action, and foreign aid, but opposition to large government spending in the domains of military, police, and border control.

A potential alternative view of the formation of ideological clusters that contrasts with our emphasis on a demand-side mechanism are supply-side accounts, according to which people identify with a party and then "learn" from elites which bundles of policy positions they are supposed to hold. While such an account does not explain why we observe the *specific* ideological clusters that we do, we acknowledge and discuss in detail the potential role of supply-side mechanisms.

To provide direct evidence that a demand-side mechanism is at least a part of the story, we leverage the observation that many broad policies can be implemented in a more or less universalist vs. communitarian fashion. For instance, redistribution can be organized at an impersonal, federal level or it can be implemented locally. Similarly, environmental protection policies can aim at preventing global climate change or saving local forests and rivers. We hypothesize that the traditional left-right divide on these and other topics attenuates or even reverses when canonical left-wing policies are implemented in a more communitarian fashion and when canonical right-wing policies are implemented in more universalist ways.

To test these hypotheses, we elicit respondents' desired spending levels for specific policy implementations, where some policies are more universalist than others. For example, within the broad domain of welfare payments, we separately elicit desired spending on "Redistributing local tax revenues as welfare payments across all communities nationwide" and "Redistributing local tax revenues as welfare payments only within the local communities they were raised." We design these survey questions holding fixed the efficiency of the redistributive system. In these exercises, the relationship between respondents' political orientation and policy views predictably attenuates or even reverses, depending on whether the specific policy proposal is more or less universalist. For example, respondents who identify as right-wing are equally likely to support redistribution or environmental protection as left-wingers once it takes place locally.

We view these results as suggesting that the link between universalism and policy views at least partly reflects a demand-side mechanism. To further corroborate this, we also show that all of our results on the link between universalism and policy views hold controlling for respondents' political ideology on a left-right spectrum.

Linking our work to the literature, much research in economics, political science, and moral and political psychology has highlighted the role of morality, identity and social preferences for political attitudes. Most closely related to our approach is recent

work on moral universalism, in both psychology (e.g., Graham et al., 2009; Haidt, 2012) and economics (Enke, 2020). We contribute to this line of research (i) by providing a quantitative analysis examining the internal structure of policy views rather than voting; (ii) not just in the U.S. but in the Western world more generally; and (iii) by working with direct economic measures of preferences and beliefs rather than psychological questionnaires. Much of our approach is inspired by the model in Tabellini (2008).

The idea that social groups and identity play an important role in understanding contemporary policy views runs through various recent contributions and reviews (Shayo, 2009; Grossman and Helpman, 2018; Bonomi et al., 2021; Besley and Persson, 2019; Guriev and Papaioannou, 2020). Relatedly, large literatures explain variation in demand for redistribution through ethnic divisions, religion and citizenship (e.g., Alesina et al., 1999; Luttmer, 2001; Alesina and Glaeser, 2004; Gilens, 2009; Scheve and Stasavage, 2006; Chen and Lind, 2019).[1] Our central contribution to these literatures is to highlight the importance of studying *heterogeneity in how much people care about group- and place-based identities*, as this shapes an entire vector of policy views in a strikingly similar fashion across Western democracies.

In political science, a large amount of work has been devoted to studying the internal structure of elite opinion (Poole and Rosenthal, 2000), but there is no extant theory that convincingly explains the internal structure of mass opinion. A popular view in the sub-field of political behavior was (and is) that citizens usually do not hold internally coherent policy views, and that, by and large, genuine ideological constraint does not exist. Rather, people are believed to follow the cues of party leaders (e.g. Lenz, 2013; Erikson and Tedin, 2015; Broockman and Butler, 2017). However, Ansolabehere et al. (2008) documented that such inferences about lack of preference stability and ideological constraint are largely driven by measurement error in surveys. We build on this insight by illuminating how universalism structures an entire vector of (richly measured) policy views, and by showing that demand-side explanations contribute to this phenomenon.

The remainder of the paper proceeds as follows. Section 2 summarizes the internal structure of ideology. Section 3 states our hypotheses. Sections 4–6 describe the design and results of our surveys and Section 7 concludes.

[1]The broader concept of social capital has received substantial attention in the political economy literature (Putnam, 2000). For example, Dal Bó et al. (2018) and Algan et al. (2018) document that far-right voters exhibit lower trust. Somewhat relatedly, a number of social theorists have argued that what fundamentally distinguishes the left from the right is that people on the left believe that human nature is fundamentally "good," while people on the right believe that people are "flawed" and need control (e.g., Sowell, 2007; Lakoff, 2010). We differ from these contributions in that we emphasize the relevance of the *slope* of social capital, rather than its level.

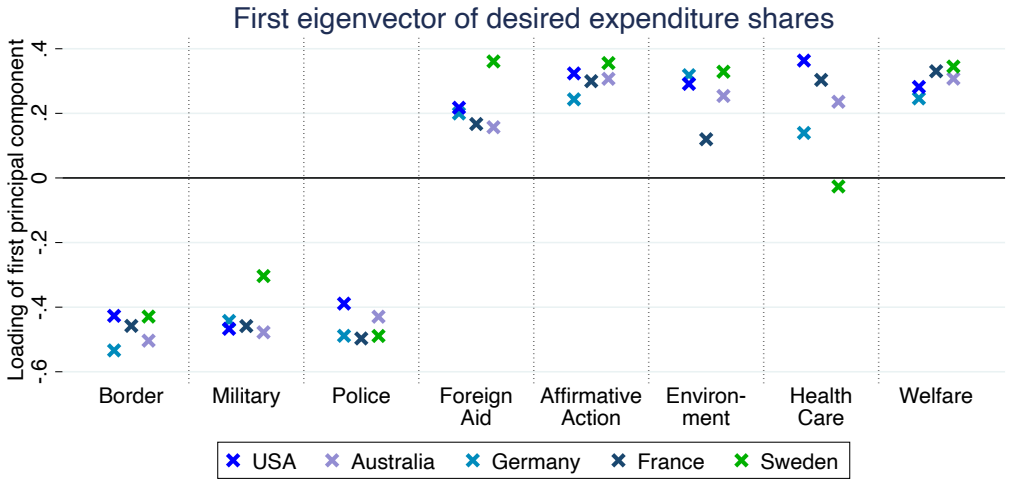

Figure 1: Factor loadings of the first principal component of desired expenditure shares. Sign convention: the loading on "Border" is always non-positive, and the other signs are determined accordingly.

2 The Structure of Western Political Ideology

We illustrate the structure of political ideology in rich Western societies using our own survey data, as they are substantially richer than standard cross-national political science datasets such as the CSES. The data cover the United States, Australia, France, Germany, and Sweden, for a total of approximately 11,000 respondents. We elicited respondents' desired per capita expenditure for eight domains: welfare payments; universal health care; affirmative action; environmental protection; foreign aid; military; police and law enforcement; and border control.

To probe the correlation structure of policy views, we implement principal component analyses (PCA), separately in each country. We find that, in each country, the first principal component of (arsinh) desired expenditures[2] across domains exhibits an unsurprising and almost identical structure: it loads positively and with essentially equal weights on desired expenditure levels in the eight categories. This first component captures "big vs. small government" views. The second principal component closely corresponds to our object of interest: in each country, it loads negatively on desired expenditure levels for military, police and law enforcement, and border control, and almost always positively on welfare, universal health care, affirmative action, environmental protection, and foreign aid. This second component, by virtue of being orthogonal to the first one, intuitively captures desired expenditure *shares*.

To make this point more explicit, we perform a principal component analysis directly on desired shares of overall spending, see Figure 1. Border control, military, and police and law enforcement all receive negative weights in each country, while foreign aid,

[2]Throughout the paper, we transform desired expenditures as $arsinh(x) = ln(x + \sqrt{x^2 + 1})$.

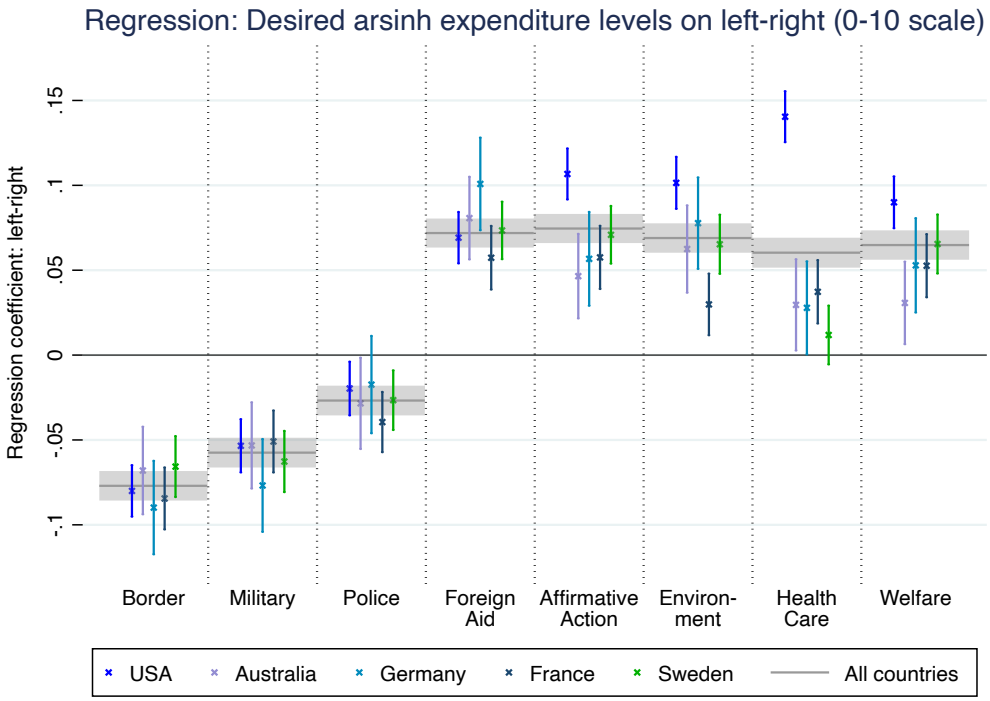

Figure 2: The figure plots the OLS regression coefficients of univariate regressions of desired arsinh expenditure levels for each policy domain on self-positioning on a left-right scale (0–10). The dependent variables are standardized into z-scores. Error bars indicate 95% confidence intervals using robust standard errors. The "All countries" specification includes country fixed effects.

affirmative action, environmental protection, welfare payments, and universal health care almost always receive positive weights.

This structure is reminiscent of intuitive notions of "left" and "right." To confirm this intuition, we elicited from our respondents how they would position themselves on a left-vs.-right Likert scale. Figure 2 summarizes the relationship between respondents' ideology and their desired expenditure levels. In all countries, more pronounced left-wing identification is correlated with *higher* desired expenditure levels for canonical left-wing policies and *lower* desired expenditure levels for canonical conservative policies.

Indeed, Figure 2 informally suggests that when respondents tell us that they are "left" or "right," they appear to refer more to *how* they would like to use a given government budget rather than the *overall size* of government. To make this argument more formal, we compute the pairwise correlations between people's left-vs.-right self-positioning, the first principal component of desired expenditure levels (the "big-vs.-small-government" component), and the first principal component of desired expenditure shares. We find that the correlation between the left-right-scale and the big-vs.-small-government component ranges between $\rho = 0.14$ in the U.S. and $\rho = 0.02$ in France. In contrast, the correlation between the left-right-scale and the expenditure-shares-component ranges between $\rho = 0.49$ in the U.S. and $\rho = 0.30$ in Australia. This suggests that a large part

of people's self-identification as "left" and "right" relates to *how* a given budget is spent, rather than how big the budget is in the first place.

While we make use of our own survey data to exposit the structure of ideological clusters, the existence of ideological constraint is well-documented. While early scholars believed that ideological constraint is relatively weak, Ansolabehere et al. (2008) showed that such inferences are largely driven by measurement error. Once political views are elicited using rich and multiple measures, they tend to be highly intracorrelated. For example, in the 2016 American National Election Survey, standard composite measures of respondents' economic conservatism and their social conservatism (Ansolabehere et al., 2006) exhibit a correlation of $\rho = 0.48$, suggesting strong ideological constraint. Draca and Schwarz (2020) and Wu (2020) also provide evidence for the existence of ideological constraint across multiple Western democracies. While some recent research suggests that the magnitude of the intra-correlations between people's policy views has increased over the last 40 years, the basic qualitative structure of ideology has been relatively stable over time (e.g., Kozlowski and Murphy, 2019; Wu, 2020).

3 Hypotheses

Conceptualizing universalism. Figure 3 illustrates how we think about heterogeneity in universalism, which is a slight modification of the setup in Tabellini (2008). A person's degree of universalism concerns the slope of their altruism (or trust) as social distance increases, holding fixed the overall level of altruism (or trust). According to this conceptualization, which will be reflected in our empirical measurement, universalists are not "more or less moral," they just allocate a given level of altruism more uniformly, and trust people at varying social distances more uniformly. Intuitively, a universalist might argue that it is appealing to treat everyone equally, while communitarians might point out that the universalist does not treat her friends very well. Indeed, in Enke et al. (2021) we show that universalists have fewer friends and spend less time with them.

In this conceptualization, "social distance" is a stand-in for different types of social identities and group memberships, including family, friendship, ethnicity, religious beliefs, values, hobbies, nationality, etc. Since our interest is in both domestic and foreign polices, it will be useful to broadly distinguish between domestic in-groups, domestic strangers as well as global in-groups and global strangers.

Universalism and policy views. To articulate why universalism should affect the formation of an entire vector of policy views, we rely on two complementary strategies. First, we here briefly discuss simple intuitions. Second, in Appendix A we provide a formal model based on Tabellini (2008) that also generates our hypotheses.

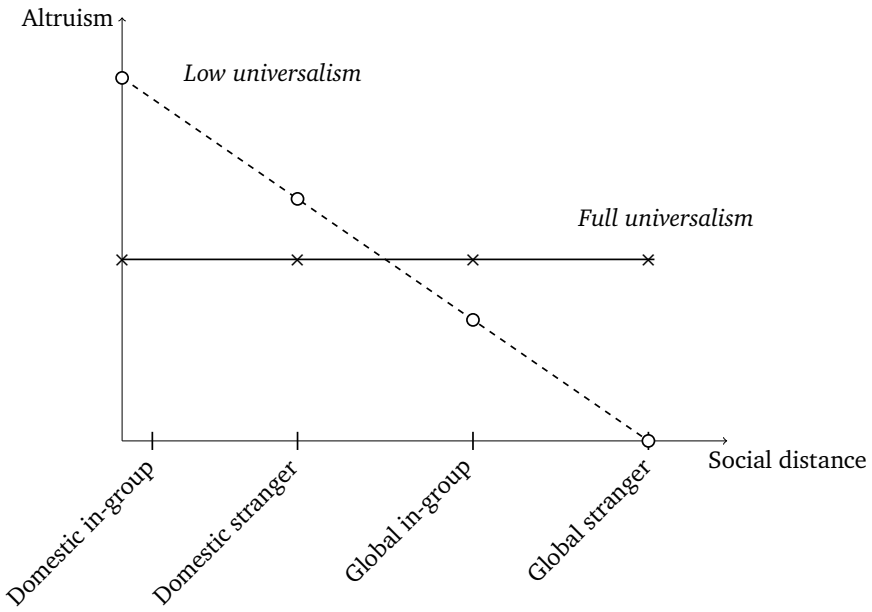

Figure 3: Illustration of heterogeneity in universalism. In the empirical analysis, we do not impose that the domestic stranger is socially closer than the global in-group.

Our starting point is the observation that many policy domains involve the welfare and anticipated behavior of people who differ in how socially close they are from the decision-maker. This suggests that the decision-maker's universalism in altruism (who he cares about) and his universalism in trust (who he believes will misbehave) will matter for how he thinks about policy proposals.

Consider first domestic policies that have a redistributive flavor, such as *welfare payments* or *universal health care*. In Western democracies, these are implemented in a highly impersonal fashion, whereby people's tax money is used to fund payments to others that they have potentially no connection to. This is in stark contrast to the more community-based redistributive systems that have dominated human history for centuries. We argue that the decision-maker's support for such highly impersonal redistributive systems will partly depend on (i) whether he actually cares about people that are socially distant from him and (ii) whether he believes that these socially distant people will "cheat on the system" (e.g., by claiming benefits they are not entitled to, or reducing their supply of labor as a form of moral hazard).

Very similar arguments apply to *affirmative action* and *foreign aid*: these are – almost by definition – policies that increase the welfare of people who are socially distant from most decision-makers. Thus, universalist individuals should be more likely to support them, both because they internalize the welfare of socially distant people to a larger degree, and because they are less inclined to believe that these strangers will "cheat" by reducing effort after receiving favorable treatment. Again relatedly, from the perspective of Western citizens, *environmental protection* arguably also largely benefits socially dis-

tant strangers, including future generations and people in developing countries whose geographic location or lack of funds makes them especially vulnerable to climate change.

All of the preceding hypotheses concern policy domains in which universalists are more supportive of government spending than communitarians. Yet, we hypothesize that the opposite holds true for expenditure categories that carry signatures of "us vs. them." First, because universalists internalize the welfare of foreigners more, and believe that they are as trustworthy as compatriots, universalists will be less supportive of strong *border protection* than decision-makers who place little weight on the welfare of foreigners (and potentially believe that they are more likely to commit crimes or to misbehave in other ways). Similarly, universalists who place high weight on the welfare of foreign nations, and who believe that these nations are trustworthy partners rather than potential aggressors, will be less inclined to spend money on a strong *military*.

Finally, regarding *police and law enforcement*, communitarian decision-makers believe that socially distant groups are more likely to cheat on society and commit crimes, which comes at the expense of the communitarian's in-groups. Yet, the communitarian fundamentally dislikes this idea because he primarily cares about his in-group members and hence supports a strong police force.

Hypothesis 1. *Universalism is positively linked to support for welfare, universal health care, affirmative action, environmental protection and foreign aid, but negatively linked to support for border protection, a strong military and police and law enforcement. This holds for universalism in both altruism and trust.*

Discussion: Default implementation of policies. In practice, many of the policies discussed above contain both universalist and communitarian elements, or could be implemented in more or less universalist ways. For example, the U.S. military not only defends American security (a communitarian cause), it also frequently engages in humanitarian or peacekeeping missions (a universalist cause). Redistribution could be implemented not just at the national level (a universalist system) but also at the level of small, local communities (a communitarian system). Environmental protection programs can be targeted at preventing global climate change (a universalist objective) but also at conserving local forests and rivers (a more communitarian objective). Foreign aid could be given to those in most need (a universalist implementation) but also to those who are international allies (a more communitarian implementation).

As a result, the development of our hypotheses above implicitly relies on the idea that policies are usually implemented – or framed – in a particular fashion, which we refer to as a policy's *default implementation*. For instance, when people think of welfare, they arguably have in mind the type of redistributive system that is currently in

place (a national one). Indeed, our analysis not only rests on the assumption of there being default implementations – we also assume that these are similar across the five countries that we study. We deem this assumption plausible, also given the cultural and economic similarities between the countries in our sample. Indeed, for some domains, this assumption is demonstrably true: in all countries in our sample, redistribution and health care do take place at the national level rather than locally.

Manipulating policy implementations. The fact that policies often contain both universalist and communal elements is useful because it allows us to elicit support for specific implementations of policies that are more or less universalist in nature.

Hypothesis 2. *The positive correlation between left- versus right-wing political ideology and support for canonical "left-wing" policies (such as support for welfare, environmental protection and foreign aid) attenuates or even reverses once these policies are implemented in less universalist ways. Similarly, the negative correlation between left-wing ideology and support for canonical "right-wing" policies (such as a strong military) attenuates or reverses once these policies are implemented in more universalist ways.*

To sum up, our empirical analyses will leverage variation in both (i) moral preferences and beliefs and (ii) the policies that respondents are asked to assess. This is useful because approach (i) is necessarily correlational in nature (we cannot change people's universalism), while approach (ii) allows for experimental manipulations.

4 Survey Design

4.1 Logistics

In the summer of 2019, we implemented internet surveys in Australia, France, Germany, Sweden and the United States through the infrastructure of the market research panel of *Dynata*. The original survey was developed in English, translated into other languages by *Dynata*, and then checked by us using native speakers.

The survey consisted of four components: (i) an introductory screen that elicited demographics and routed respondents into or out of the survey; (ii) decision screens to measure universalism and other social preferences; (iii) screens to measure policy views; and (iv) a questionnaire to elicit additional information and covariates. The order of parts (ii) and (iii) was randomized across respondents.[3] We also randomized the order

[3]In Appendix B.7 we provide histograms that show that the distribution of policy views is unaffected by whether policy views are elicited before or after universalism.

in which universalism in altruism and universalism in trust were elicited.[4]

We took two measures to ensure quality control. First, every respondent who completed the survey in less than 400 seconds was dropped and replaced by *Dynata*. Second, the survey contained two attention check questions, interspersed throughout the survey. Whenever a respondent answered an attention check incorrectly, they were immediately routed out of the survey and replaced by *Dynata*.

We contracted with *Dynata* for nationally representative samples of $N = 1,700$ citizens aged at least 18 in each country (see details on the pre-registration below). However, because constructing a sample that is nationally representative along the lines of age, gender, ethnicity, income, employment status, and education is logistically difficult, *Dynata* eventually supplied a larger sample to us (total $N = 11,063$), a subset of which ($N = 8,500$) makes up the more representative samples that we pre-registered. Since we view throwing away data as scientifically questionable, all analyses reported in the main text make use of the full sample. In the Appendix we replicate all analyses using the pre-registered (smaller) representative samples. The results are always very similar. Sample characteristics are summarized in Appendix B.1.

Finally, we implemented identical surveys also in Brazil and South Korea, see Section 6 and Appendix C.9 for an analysis of these "non-Western" countries.

4.2 Measurement of Universalism

4.2.1 Decision Tasks

Universalism in altruism. We rely on a set of structured experimentally-validated decision tasks to measure universalism (Enke et al., 2021). Respondents completed a total of 16 hypothetical money allocation tasks that allow us to construct a summary statistic of universalism in altruism. Analogous to Figure 3 and our formal framework in Appendix A, the construction of the decision tasks is organized along four different types of groups: domestic in-groups, domestic strangers, global in-groups, and global strangers.

First, to estimate *domestic* universalism, respondents made ten decisions. In each of them, they were asked to split hypothetical \$100 between (i) a randomly-selected person from their country of residence and (ii) a randomly-selected member of one of their social groups, who also resides in the respondent's country of residence. Across the ten questions, the social groups included extended family, friends of family, neighbors, colleagues at work or school, same organization (e.g., club), same age, same ethnic background or race, same political views, same hobbies, and same religious beliefs. For example, in one question, respondents in the U.S. were asked to split \$100 between a

[4] A permanent link for the U.S. version of our survey is: `https://harvard.az1.qualtrics.com/jfe/form/SV_aftuqgHsyIAShkp`.

randomly-selected person who lives in the U.S. and a member of their extended family, such as a cousin. The average allocation to the randomly-selected person across the ten questions then makes up the domestic universalism measure.

Second, to estimate *foreign* universalism, respondents were asked to split $100 between (i) a randomly-selected person from their country of residence and (ii) a randomly-selected person who lives anywhere in the world. Foreign universalism then corresponds to the monetary amount sent to the global stranger.

Third, to estimate *global* universalism, respondents made five decisions, in each of which they were asked to split hypothetical $100 between (i) a randomly-selected person who lives anywhere in the world and (ii) a randomly-selected person who lives anywhere in the world and is a member of the respondent's social groups. Across the five questions, the social groups included same language, same religious beliefs, same ethnic background, same values, and same occupation. The average amount of money sent to the randomly-selected world citizen makes up the global universalism measure.

For the purpose of these tasks, respondents were always asked to assume (i) that both individuals are equally rich (addressing income effects) and (ii) that neither of these individuals would find out who sent them the money (ruling out reciprocity considerations). The order of questions was randomized across respondents. Figures 12 and 13 in Appendix B.2 show example decision screens.

The money allocation decisions, and in particular the domestic, foreign, and global universalism summary components are all highly positively correlated with each other. To reduce the dimensionality of the data and minimize measurement error, we average the three components into a pre-registered summary statistic of universalism in altruism.

Universalism in trust. Respondents again completed a total of 16 tasks. The procedure was identical to the one described for altruism above, except that in a given task respondents were asked to allocate 100 points (rather than $100) between two individuals, to express whom of the two they trust more. This was explained as indicating "how much you trust that different people will not cheat on you or take advantage of you." This procedure again yields domestic, foreign, and global universalism components, which we average into a summary statistic of universalism in trust. Figure 14 in Appendix B.2 shows an example decision screen.

Composite measure of universalism. Universalism in altruism and trust exhibit a correlation of $\rho = 0.70$ after accounting for measurement error using the obviously-related instrumental variables technique of Gillen et al. (2019). To reduce the dimensionality of the analysis, we often work with a composite measure, which consists of the unweighted

average of universalism in trust and universalism in altruism. We always reference robustness checks that use the altruism and trust measures separately, see Section 5.4.

Relationship to concepts in moral psychology. Prior work on moral universalism in politics usually relies on psychological questionnaires such as the Moral Foundations Questionnaire (e.g., Graham et al., 2009; Haidt, 2012; Enke, 2020). While these questionnaires have the advantage of capturing the full richness of the human moral mind, the underlying concepts are only vaguely related to economic concepts and utilitarian models of preferences and beliefs. Indeed, many concepts in psychological questionnaires such as the MFQ capture deontological motivations. At the same time, arguably the main idea behind, for example, the moral foundations of Haidt (2012) is the distinction between a relationship- or group-specific morality and a universalist morality.[5] Thus, our main objectives in designing our measures were (i) to separately capture preferences and beliefs; and (ii) to capture the broad main distinction between a universalist and a relationship-specific morality.

Construct Validity. We validate the universalism measures along two dimensions. See Enke et al. (2021) for details. (i) We implemented an ex-ante experimental validation procedure. First, we show that, over a one-week horizon, our hypothetical measure of universalism in altruism is highly correlated with a financially-incentivized measure of universalism, which consists of the same questions with real incentives. Second, we document that behavior in our trust task is highly correlated with trust beliefs in a structured cheating task that is standard in the experimental economics literature. (ii) We also show that our survey measure of moral universalism predicts real donation decisions: while universalists donate less to local community organizations, they donate more to nationwide and international charities.

4.2.2 Descriptives

Figure 4 shows a histogram of the composite universalism measure, pooled across countries. Numbers around 50 imply on average equal allocations of money and trust points to in-groups and strangers. Numbers below 50 indicate a tendency to allocate more money and trust points towards in-groups. Numbers above 50 correspond to the (largely counterfactual) case that someone allocates more money and trust points to socially more distant individuals. Appendix B.3 shows histograms for each country separately.

[5]For example, Enke (2020) collapsed Haidt's first four moral foundations into a one-dimensional summary measure because much psychological research suggests that such a "principal component" captures a large fraction of the variation in human morality (as far as ideas related to universalism are concerned).

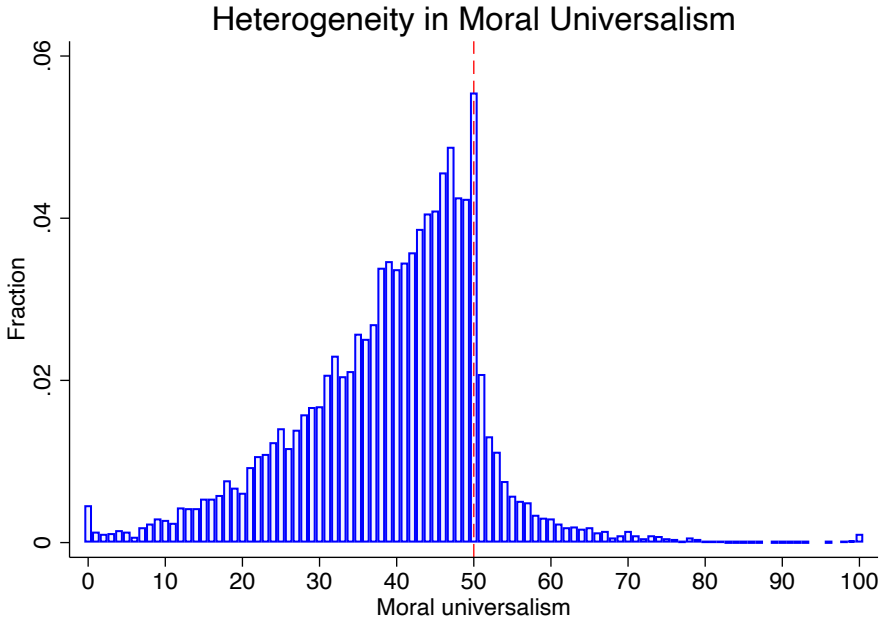

Figure 4: Distribution of the composite measure of moral universalism, pooled across all Western countries. The amounts reflect allocations to random strangers, so that the measure is decreasing in in-group favoritism. 50 corresponds to an equal split of money or trust points.

Table 1 reports correlations with demographics. To avoid focusing on variables that we select based solely on our intuitions, we instead select demographics (or their interactions) through a LASSO regression, see Appendix Table 12. The strongest correlations are with age and wealth, both of which correlate negatively with moral universalism. Similarly, men, higher-income individuals, and the religious exhibit lower universalism. Of potential interest is also the significant conditional correlation between education and universalism: fixing people's income and wealth, higher education is associated with higher universalism. This is potentially relevant for understanding recent changes in political cleavages (Gethin et al., 2022). Appendix B.5 reports correlations between universalism and an expanded set of survey variables.

Finally, while our study was designed to study within- rather than across-country variation, we note that average composite universalism in each country is as follows: Australia 0.38, France 0.41, Germany 0.40, Sweden 0.38 and United States 0.39.

4.3 Measurement of Political Attitudes

We measure policy views in two complementary ways: (i) by eliciting measures of desired government expenditure amounts and (ii) through standard Likert scale questions in which people express their support for certain policies. Both of these techniques have well-understood advantages and disadvantages, which is why we leverage both. The

main upside of (i) is that the measures are quantitative in nature and do not rely on the interpretational ambiguity associated with responses on Likert scales. The main downside of these measures, however, is that they are cognitively more complex.

Desired expenditures. Respondents were instructed to imagine they could decide the average amount of money that their federal or national government collects per year from each citizen to spend on each of eight policy categories. We asked respondents to assume that all dollar amounts collected for a category would be spent only on this particular category, without any waste. In addition, we provided respondents with a reference value: annual per capita spending on education in their country of residence.

Respondents were asked to enter eight monetary amounts to indicate their desired per capita spending levels for each of welfare, universal health care, foreign aid, environmental protection, affirmative action,[6] military and counterintelligence, police and law enforcement, and border control.[7] The order of these categories on the computer screen was randomized. Figure 15 in Appendix B.2 shows an example decision screen, and histograms of desired expenditure amounts for each policy category and country are presented in Appendix B.6.[8]

Frequently, we will be interested in linking universalism to desired expenditure *shares* rather than simple levels. For this purpose, we simply compute the fraction of money a respondent desires to spend on a given policy domain.

Likert scale questions. Respondents were asked to indicate their support for each of the eight policies described above on a 0–10 scale. We again frequently work with normalized measures, referred to as *relative policy support*, that consists of support for a given policy divided by the sum of support across all policies.

4.4 Covariates

Even though this paper is descriptive in nature, we seek to assess the extent to which a potential relationship between universalism and policy views is driven by omitted

[6]The term "affirmative action" was never used in the survey. Instead, this was described as "Measures to ensure no individual is disadvantaged in access to education, the labor force, and marriage."

[7]This selection of policies was motivated by two considerations. First, to include some of the most important policies in terms of government spending. Second, to restrict attention to those that are politically contentious in all countries that we study (which for example excludes gun control and abortion).

[8]We implement the following procedure to all desired expenditure amounts referenced in the paper. First, we perform a PPP conversion to USD. Because of the free-entry format, responses to these questions are subject to large outliers. As such, we then winsorize the desired PPP spending levels at +/− 3 SD of the within-country mean, as specified in our pre-registration. This affects 2.2% of all responses. Third, we compute the inverse hyperbolic sine (arsinh) of these winsorized, PPP desired expenditure amounts, and finally standardize into z-scores within each country.

Table 1: Individual-level correlates of universalism

	Correlation between composite measure of universalism and:						
	Age	Religiosity (z-score)	Wealth Index (z-score)	Income Index (z-score)	Male (0-1)	Age × Male × Income (z-score)	College (0-1)
Raw correlation	-0.16^{***}	-0.10^{***}	-0.12^{***}	-0.07^{***}	-0.07^{***}	-0.01	0.01
OLS coeff. (w/ Country FEs)	-0.12^{***}	-1.16^{***}	-1.51^{***}	-0.84^{***}	-1.78^{***}	-0.06	0.21
OLS coeff. (multivariate) (w/ Country FEs)	-0.09^{***}	-1.04^{***}	-0.82^{***}	-0.50^{***}	-1.11^{***}	0.33^{***}	1.23^{***}

Notes. The first row reports the Pearson raw correlation between individual characteristics and the composite measure of universalism ($N = 11,063$). The second row reports OLS coefficients from individual regressions of the composite measure of universalism on the given characteristic, including country fixed effects; this row thus presents by how many dollars / trust points universalism increases for a one unit change in the demographic variable. The third row reports OLS coefficients from a multivariate regression of the composite measure of universalism on all characteristics at once, including country fixed effects. See Appendix E for details on the construction of the demographic variables. All z-scores are computed separately within each country. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

variables. Our survey hence elicits rich measures of covariates, including: age, gender, ethnicity / race, educational attainment, income (two measures), wealth and asset ownership (three measures), religiosity (three measures), urbanicity, employment status, marital status, migration background, belief about whether the government is efficient or wasteful (on a scale 0–10), beliefs about whether the respondent is likely to personally benefit from government expenditure in a given category, and measures of altruism, generalized trust, and equity-efficiency preferences. See Appendix E.

4.5 Pre-Registration

The survey was pre-registered on EGAP, see `http://egap.org/registration/5792`. The pre-registration contained (i) the desired sample size; (ii) the precise construction of the summary statistics of universalism in altruism and trust; (iii) predictions about how we expected universalism to be correlated with support for each of the eight policy domains, as stated in Section 3; (iv) the construction of a summary statistic of policy views discussed below; and (v) the analysis of the specific policy proposals in Section 6.1.

5 Results

5.1 Universalism and Self-Reported Ideology

Our ultimate interest in this paper is to understand how universalism potentially shapes an entire vector of policy views. However, because today clusters of policy views are commonly associated with the labels “left” and “right,” we begin by documenting the link between universalism and people’s self-assessment on a 0–10 left-vs.-right scale.

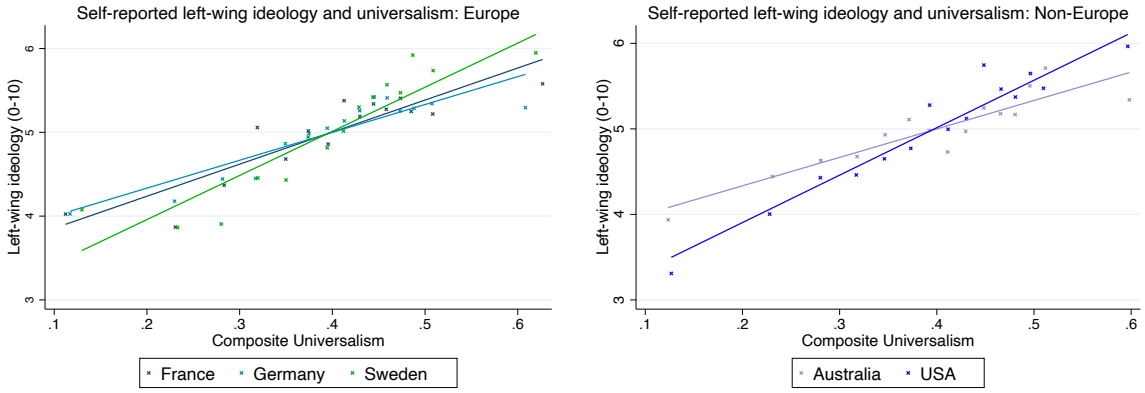

Figure 5: Binscatter plot of individuals' left-right self-assessment against composite universalism. Both universalism and left-wing ideology are residualized of country fixed-effects.

Figure 5 summarizes the results separately for each country by showing binned scatter plots. We see that, in each country, the composite universalism measure is strongly correlated with left-wing identification. The correlation in the full sample is $\rho = 0.25$, conditional on country fixed effects. This correlation is statistically highly significant in each country and ranges between $\rho = 0.19$ in France and $\rho = 0.29$ in the United States. While we provide more sophisticated benchmarking analyses later, it is perhaps informative that the corresponding correlations between left-vs.-right-wing ideology and other commonly analyzed variables are substantially smaller in magnitude: $\rho = -0.01$ for the belief that government is efficient vs. wasteful, $\rho = 0.02$ for college degree, $\rho = -0.06$ for age, $\rho = -0.09$ for an income index, and $\rho = -0.13$ for a wealth index (all correlations computed conditional on country fixed effects).

An obvious question is whether particular components of the universalism measure drive this result. Instead, as Appendix C.1 shows, *every single one* of our money and trust tasks is significantly correlated with people's self-reported ideology in the same direction. This suggests that a universalist mindset is a general characteristic of people who view themselves as "left."

5.2 Universalism and Policy Views

Figure 6 shows the link between composite universalism and policy views. Panels A and B report the results of regressions of desired expenditure shares on composite universalism, the only difference being that the regressions in Panel B control for age, gender, income, wealth, college, urbanicity, religiosity, equity-efficiency preferences, altruism, trust, beliefs about the efficiency of government, and beliefs about whether one will personally benefit from government expenditure in each domain. Panels C and D follow an analogous logic, except that the dependent variables are given by the relative policy sup-

port measures as computed from the Likert scale questions. All dependent variables are standardized into z-scores. Each panel shows the results of 48 regressions: eight policy domains in five countries each, plus a full sample specification.

As hypothesized, in all countries, we observe a strong negative relationship between universalism and desired expenditure shares or policy support for the three "right-wing" policy domains, while the relationship is generally positive and statistically significant for the five "left-wing" domains.[9] In Panel A of Figure 6, out of the 40 regression coefficients for the individual countries reported, 37 have the expected (pre-registered) sign. Of these, 33 are statistically significant at least at the 10% level. Once our battery of controls is added in Panel B, 38 of these coefficients have the expected sign, out of which 31 are statistically significant at least at the 10% level.

An immediate question is whether these results are only driven by universalism in altruism or universalism in trust alone. Indeed, while we construct Figure 6 based on the composite universalism measure, we highlight that our data allow us to separately consider preferences and trust beliefs. Table 2 summarizes the results of OLS regressions, in which we link a respondent's policy views (normalized into z-scores) to their universalism in altruism and trust, controlling for our full set of covariates. In panel A, we investigate desired expenditure shares, which directly corresponds to Figure 6. In panel B, the dependent variables are the analogous Likert scale questions, again normalized into measures of relative support across policy domains. All regressions include the full set of controls discussed above.

We find that the coefficients of universalism in altruism and universalism in trust are almost always statistically significant and sizable in magnitude. This suggests that even though universalism in altruism and universalism in trust are positively correlated, they each capture distinct variation that is relevant for understanding policy views.

The coefficient magnitudes are consistently larger for universalism in altruism than for universalism in trust, which suggests that group-specific altruism weights are more important for policy views than group-specific trust levels. We view it as one strength of our approach that we can separately quantify the importance of beliefs and preferences, while for example the psychological approach of Haidt (2012) does not allow such conceptual distinctions.

Comparing magnitudes across policy domains, for both the desired expenditure share data and for the Likert scale questions, we find that universalism appears to be

[9]A notable exception occurs in the domain of universal health care as elicited using expenditure shares, where the relationship is strongly positive in the U.S. but either not statistically significant or even negative in the other countries. This pattern might arise because, in contrast to the United States, all of these countries have had versions of universal health care for decades, which may generate less heterogeneity in views on universal health care across the political spectrum. Indeed, Figure 2 shows that, in all countries except the U.S., the link between self-reported left-wing ideology and desired health care expenditure levels is quantitatively very small, and about 1/4 the magnitude as in the U.S.

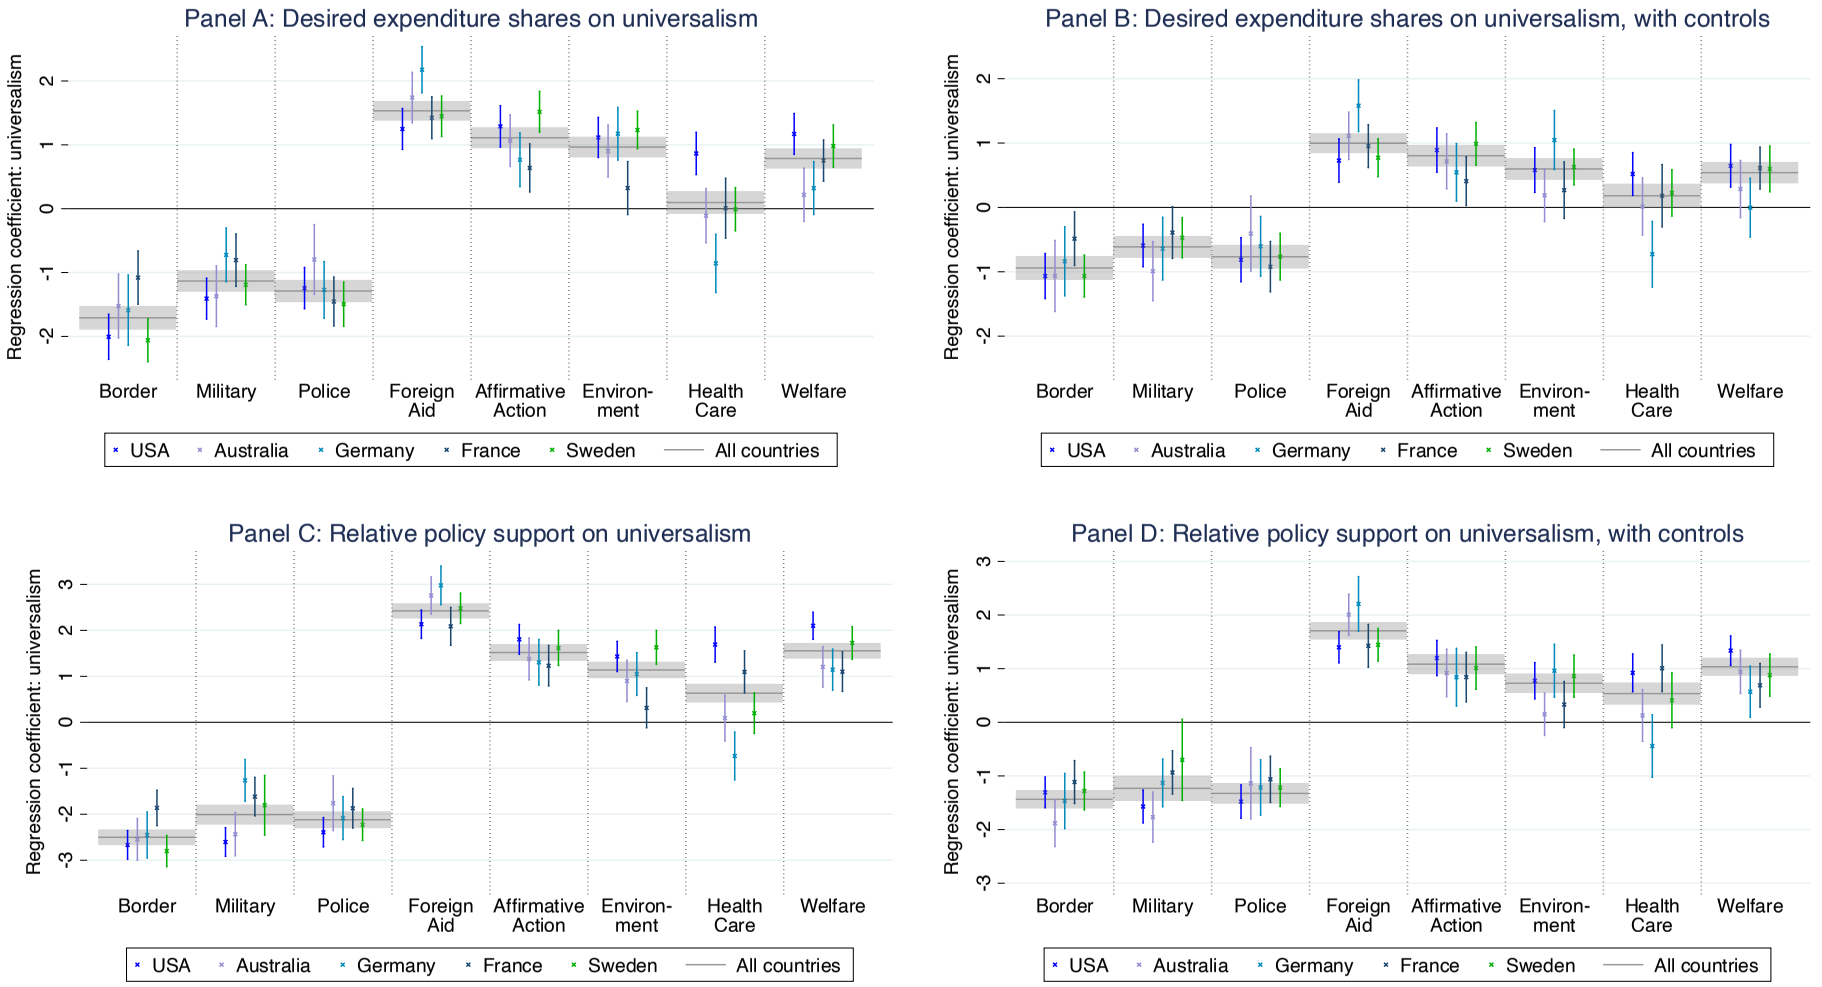

Figure 6: OLS regression coefficients of desired expenditure shares (Panels A and B) or relative policy support (Panels C and D) on composite universalism. The left panels only control for country FE, while the right panels also control for age, gender, income, wealth, college, urbanicity, religiosity, equity-efficiency preferences, altruism, trust, beliefs about the efficiency of government, and beliefs about whether one will personally benefit from government expenditure in each domain. In the top panels, the dependent variables are desired expenditure shares for each policy domain, while in the bottom panels, the dependent variables are relative policy support as computed from the Likert scale questions. Universalism is in [0,1] and the dependent variables are standardized into z-scores within each country. Error bars indicate 95% confidence intervals using robust standard errors.

Table 2: Universalism and policy views

PANEL A

Dependent variable: Desired expenditure shares (Z-scores)

	Border control	Military	Police	Foreign aid	Aff. action	Environment	Health care	Welfare
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Universalism in altruism	-0.84***	-0.55***	-0.57***	0.69***	0.51***	0.46***	0.16*	0.25***
	(0.08)	(0.08)	(0.08)	(0.08)	(0.08)	(0.08)	(0.08)	(0.09)
Universalism in trust	-0.59***	-0.28***	-0.36***	0.39***	0.30***	0.25***	0.030	0.35***
	(0.10)	(0.10)	(0.10)	(0.09)	(0.09)	(0.09)	(0.10)	(0.11)
Country FE	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Controls	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Observations	10881	10881	10881	10881	10881	10881	10881	10881
R^2	0.07	0.05	0.05	0.08	0.03	0.03	0.01	0.04

PANEL B

Dependent variable: Relative policy support (Z-scores)

	Border control	Military	Police	Foreign aid	Aff. action	Environment	Health care	Welfare
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Universalism in altruism	-1.33***	-1.09***	-0.97***	1.40***	0.84***	0.60***	0.32***	0.54***
	(0.10)	(0.09)	(0.09)	(0.08)	(0.09)	(0.12)	(0.11)	(0.11)
Universalism in trust	-0.62***	-0.41***	-0.56***	0.36***	0.24**	0.30*	0.27**	0.59***
	(0.14)	(0.12)	(0.11)	(0.09)	(0.11)	(0.16)	(0.12)	(0.15)
Country FE	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Controls	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Observations	10971	10971	10971	10971	10971	10971	10971	10971
R^2	0.13	0.10	0.11	0.18	0.05	0.04	0.03	0.09

Notes. OLS estimates, robust standard errors in parentheses. In panel A, the dependent variables are desired expenditure shares for each of the eight policy categories, normalized into z-scores within each country. In panel B, the dependent variables are given by responses to the policy Likert scale questions, normalized by the sum of responses across Likert scale questions. These measures are also normalized into z-scores within each country. Universalism in altruism and universalism in trust are both in [0,1]. Controls include age, gender, income, wealth, college, urbanicity, religiosity, equity-efficiency preferences, altruism, trust, beliefs about the efficiency of government, and beliefs about whether one will personally benefit from government expenditure in each domain. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

most important for foreign aid, affirmative action, border control, military and police and law enforcement. We did not pre-register predictions about differential magnitudes across policy domains, and we can only speculate about what drives them. One potential reason is that some policies (like foreign aid) are primarily about in-group vs. stranger tradeoffs, while others (such as welfare) also involve many other considerations and preferences, such as one's own relative income or views on self-sufficiency.

We consistently find that (i) the magnitude of the OLS coefficient of universalism is 50%–100% larger for the Likert scale questions and (ii) the R-squared is also usually twice as large. A plausible interpretation of this is that measurement error in the Likert scale questions is lower than in the more complex quantitative desired expenditure amounts questions. Another contributing factor is that the Likert scale questions have fewer outliers. While we winsorized the desired expenditure data at +/- 3 SD of the mean, this still generates large outliers that reduce the variance explained.

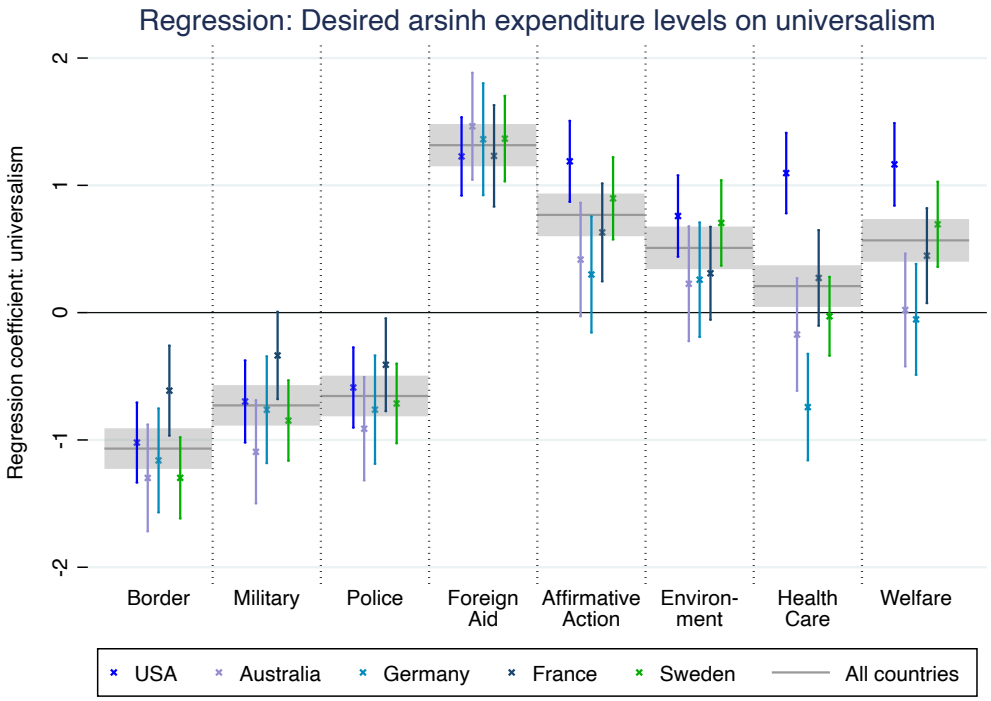

Figure 7: The figure plots the OLS regression coefficients of univariate regressions of desired arsinh expenditure levels for each policy domain on composite universalism. Universalism is in [0,1] and the dependent variables are standardized into z-scores within each country. Error bars indicate 95% confidence intervals using robust standard errors. The "All countries" specification includes country fixed effects. Figure 31 in Appendix C.2 reproduces this figure controlling for arsinh total desired expenditure both linearly and using quantile fixed effects.

Desired expenditure levels. Figure 7 reproduces the left panel of Figure 6, except that now the dependent variables are desired (arsinh) expenditure levels rather than implied shares. Again, we standardize desired spending into z-scores within each country to keep the magnitudes comparable across countries. The results show that universalists desire higher government spending in the canonical left-wing policy domains, yet lower government spending in the canonical conservative domains.[10] Thus, universalists do not *always* desire higher government spending, but rather only in those policy domains that we argue have a universalist "default implementation" (see Section 3).

5.3 Benchmarking Exercises

In the following, we benchmark the link between universalism and policy views against the link between universalism and other individual characteristics, showing that universalism is quantitatively substantially more important for explaining policy views and ideological constraint.

[10]Appendix C.2 shows that similar results hold when we directly control for arsinh total desired expenditure.

Benchmarking of "effect sizes." We first compare our estimated magnitudes of the link between universalism and policy views with the link between policy views and people's left-vs-right ideology (0–10). We do not do this because we believe that all variation in left-vs.-right ideology should itself be considered primitive – as discussed above, we suspect that when people say they are "left" they partly mean that they are universalist. Instead, we view people's self-reported ideology as close to an upper bound for how much universalism can possibly explain, given how strongly people's policy views are clustered around their political ideology. We additionally consider income as a benchmarking variable, given its salience in political economy analyses.

Figure 8 summarizes the results. Each coefficient stems from a separate regression of desired expenditure shares on either composite universalism, people's left-right self-assessment, or their income. To aid comparability, all variables are expressed as z-scores, and the signs on the coefficients on income are all negated (that is, they show the expected change in desired expenditure shares given a one standard deviation *decrease* in our income index). We see that, on average across policy domains, the point estimate of universalism is about two-thirds as large as that of people's self-reported political ideology, while that of our income index is only about one-fourth as large. We view these results as supportive of the idea that universalism quantitatively matters for policy views.[11]

Benchmarking of contribution to ideological constraint. In a second step, rather than asking how much universalism matters for policy views considered in isolation from each other, we now study how much it matters for the *intracorrelations* among policy views observed in the data. To study this, we ask by how much the correlations among policy views (ideological constraint) would decrease if there was no heterogeneity in universalism. That is, we aim to construct a counterfactural measure of ideological constraint for a world in which everyone has the same degree of universalism.

To this effect, we proceed in two steps. First, we construct a measure of predicted ideological constraint. We compute the fitted values of regressions of desired expenditure shares on a set of predictors. The average intracorrelation of these predicted policy views tells us how strongly the predictable component of policy views are correlated in our data, on average across all policy domains.[12] Second, to derive a quantitative indication of the importance of universalism, we repeat the entire procedure, except that

[11]Appendix Figure 32 shows the variance explained of each of the three predictors across policy domains.

[12]To compute this measure of average predicted intracorrelation, we recode the correlations such that they are positive when the prevailing structure of ideology predicts it, and as negative otherwise. For example, if predicted views on welfare and military were negatively correlated, we would recode the correlation as positive. If predicted views on welfare and foreign aid were negatively correlated, we code it as negative in computing average correlations.

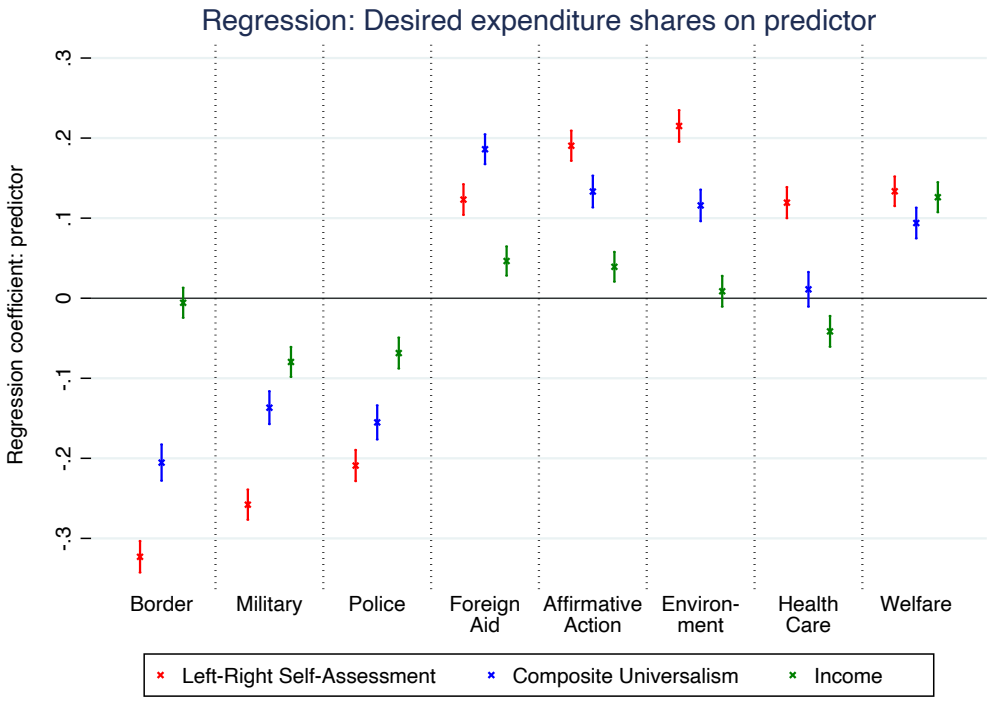

Figure 8: This figure plots coefficients from separate regressions of desired expenditure shares on respondents' left-right self-assessment, composite universalism, and income, controlling for country fixed effects. All variables are standardized into z-scores within each country. Error bars indicate 95% confidence intervals using robust standard errors. Sign convention: the coefficients on income are all negated for easier comparability of magnitudes across predictors. As such, they should be interpreted as the expected change in desired expenditure share for each policy domain following a one standard deviation *decrease* in income.

the predicted policy views are computed after collapsing variation in universalism to the sample mean. Thus, this measure tells us how strongly the predictable component of policy views would be intracorrelated if everyone had the same degree of universalism. We perform the same procedure by collapsing other demographc variables to the sample mean, as natural benchmarks. An obvious limitation of such an analysis is that most variables are measured with error – as a result, even if universalism was the only determinant of ideological constraint (which we don't claim it is), we would still see intracorrelations among policy views when all heterogeneity in (partly mismeasured) universalism is collapsed.

Figure 9 summarizes the results. The left panel shows the results from predicting policy views based on universalism and the set of individual characteristics discussed above. The leftmost bar shows how strongly predicted policy views are correlated in our data when none of the variables is collapsed to its sample mean. The second bar shows the average intracorrelation when age is collapsed to its sample mean, and equivalently for the other variables. We see that heterogeneity in universalism is the only variable that visibly reduces predicted ideological constraint, from around 0.57 to 0.43 (which corre-

sponds to 25%). All other variables, including measures of income, wealth, religiosity and beliefs in government efficiency, do not appreciably contribute to predicted ideological constraint, and sometimes collapsing variation in these variables even *increases* the observed intracorrelations.[13]

The right panel provides another comparison, by repeating the same analysis except that now predicted policy views are computed based on people's self-reported left-right assessment, in combination with the same individual characteristics as in the left panel. Again, we here view people's left-right self-assessment as a useful comparison because it is a plausible upper bound for how much can be explained with (noisy) survey measures of individual characteristics. There are two main takeaways. First, predicted ideological constraint increases by about 25 percent when self-reported ideology rather than universalism is taken into account. Second, predicted ideological constraint decreases by about 48 percent when variation in ideology is collapsed to its sample mean.

Our takeaway from this analysis is that universalism contributes substantially more to the phenomenon of ideological constraint than all other "primitive" individual characteristics in our data. Universalism does not explain as much as self-reported ideology, which is plausibly an upper bound for how much could in principle be explained based on individual characteristics.

Do other individual characteristics produce the observed policy clusters? In a final step of our benchmarking analyses, we study whether other individual-level variables also produce the characteristic pattern shown in Figure 6: positive correlations with the five "left-wing" policies and negative correlations with the three "right-wing" policies.

Given that our dataset includes a large set of individual characteristics, all of which could matter either in isolation or in the form of interaction effects, we proceed by selecting a set of candidate variables using simple machine learning techniques. Specifically, we implement a LASSO regression of respondents' left-right self-assessment on all individual characteristics except universalism, in a fully interacted model. We then select those ten regressors that exhibit the largest regression coefficients (all variables were standardized into z-scores). The selected regressors are: religiosity, wealth, residual trust, college degree, efficiency-equity preferences times beliefs about government efficiency, age times beliefs about government efficiency, income times beliefs about government efficiency, income, male times beliefs about government efficiency, and male.

In Figure 10, the leftmost panel serves as a reminder about the patterns for universalism. We find that none of the other ten individual characteristics that got picked by the

[13]This happens, for example, when a variable is positively correlated with one policy view and negatively with another policy view, even though the two policy views "should" be positively correlated from the perspective of today's political clusters.

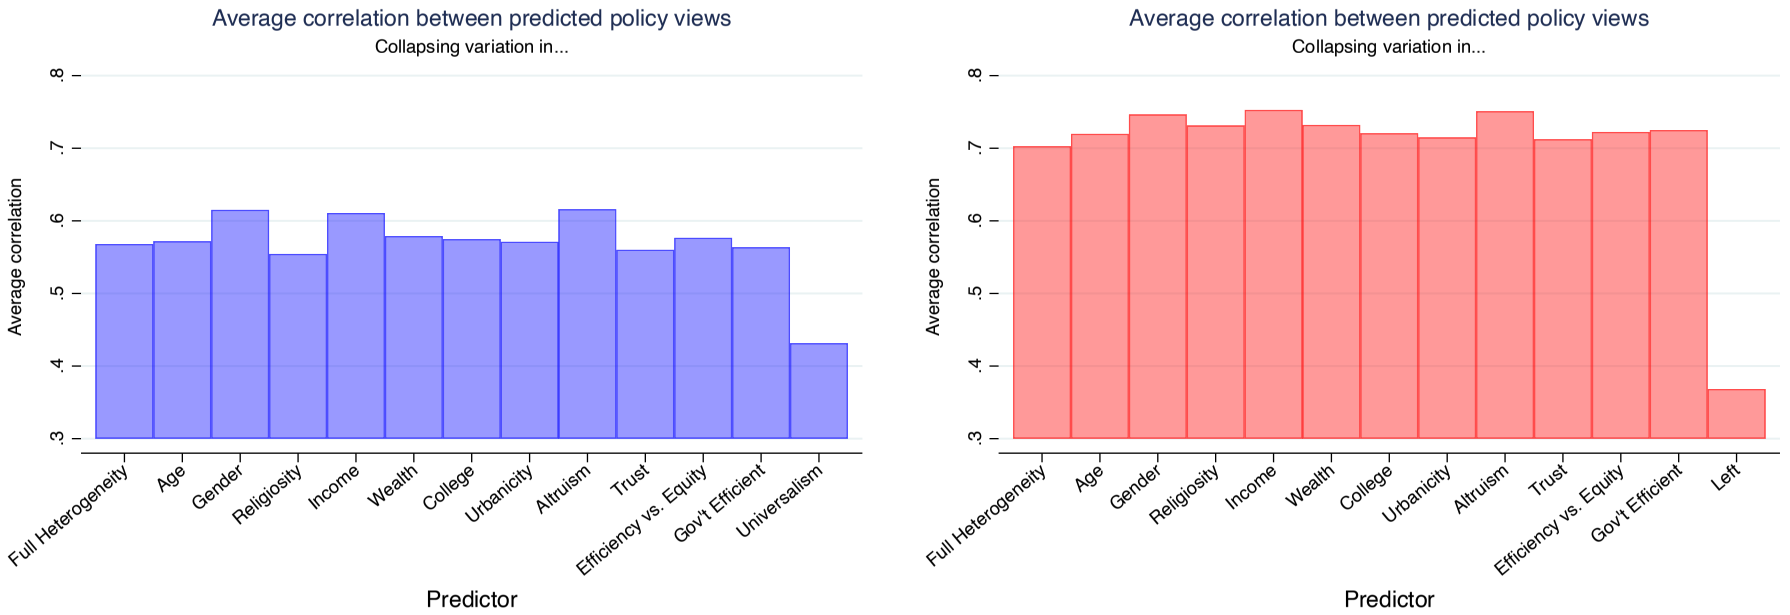

Figure 9: This figure plots counterfactual, average correlations between desired expenditure shares across policy domains when we collapse variation in demographics, beliefs, and preferences one at a time. The figure is constructed as follows. First, we regress the z-scores of desired expenditure shares for each of our eight policy domains on all of the variables shown in the left panel (or right panel), controlling for country FE. We then obtain the fitted values from these regressions when we collapse one of the variables in the regression at a time. Finally, we compute the correlation between these predicted desired expenditure shares for each policy domain (including country fixed effects). We average these correlations across policy domains (recoding correlations so they are positive if in the hypothesized direction and negative otherwise), so that we obtain *counterfactual* clusters of ideology when the effect of heterogeneity in a given demographic, preference, or belief is removed. In the left panel, we document how predicted ideological constraint changes when we collapse variation in these variables in a regression that includes composite universalism. In the right panel, the regression instead includes the left-right self-assessment.

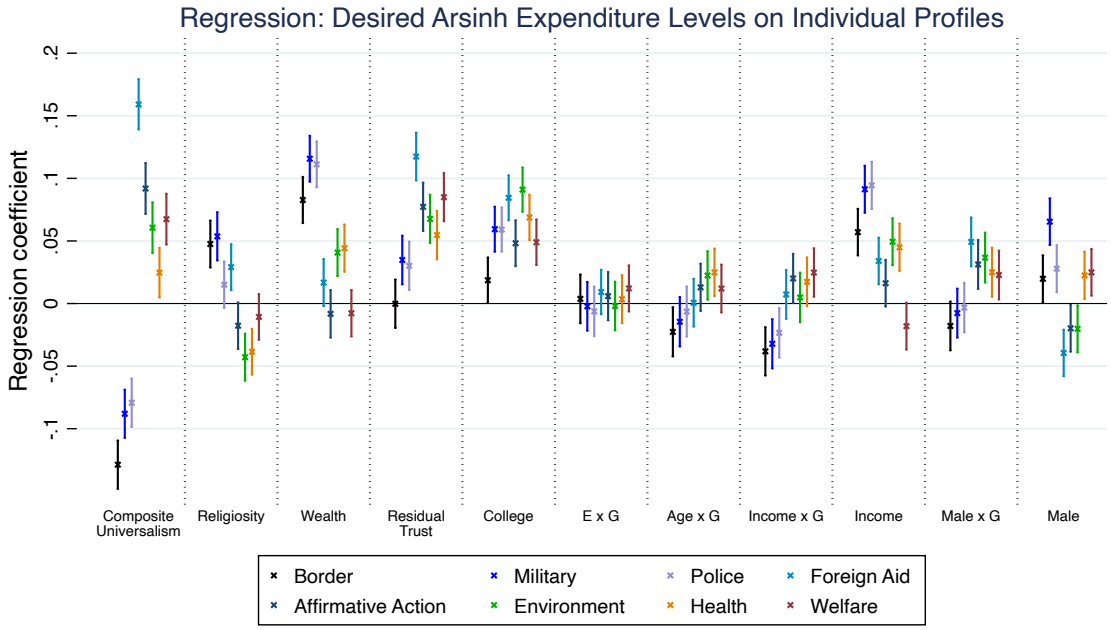

Figure 10: Benchmarking analyses. We report the standardized beta coefficients and confidence intervals for regressions of arsinh desired expenditure level for a policy domain on different individual-level characteristics, conditional on country fixed effects. These characteristics were selected through a LASSO regression of people's left-vs.-right self-assessment on all possible interactions of standardized demographics, beliefs, and preferences. Out of those variables selected by the regression, we plot the top 10. Each estimate corresponds to a separate regression. To obtain residual altruism and trust, we respectively computed the residuals of dictator game allocations and generalized trust with respect to universalism.

LASSO produces the characteristic pattern that universalism successfully reproduces, at least in a statistically- or even economically-significant manner. In other words, other variables are often significantly correlated with policy views in meaningful and known ways – we are not trying to argue that they are unimportant for understanding policy views. However, our results show that they do not generate the characteristic internal structure of ideology that we are interested in here.

5.4 Robustness Checks

The Appendix contains five sets of robustness checks. First, some of our main analyses employed the composite measure of universalism. As specified in our pre-registration, Appendix C.4 shows that very similar results hold if we work with universalism in altruism or universalism in trust separately.

Second, to tentatively probe the time stability of our results, Appendix C.5 documents that the link between universalism and policy views is very similar among young and old survey participants.

Third, as we pre-registered, we employ instrumentation strategies from Gillen et al.

(2019) to address the effects of measurement error in our elicitations of policy views and universalism. Results using multiple elicitations for both outcome and explanatory variables are very similar, see Appendix D.

Fourth, we contracted with *Dynata* for $N = 1,700$ respondents in each country, stratified to match the population on a number of dimensions. In Appendix C.6, we replicate the analysis using these more representative samples, with very similar results.

Fifth, we address potential concerns over multiple hypothesis testing that may arise from our procedure of linking universalism to a vector of eight policy views (though we note that each of these analyses was pre-registered). To reduce such concerns, we first collapse the eight policy variables into a summary statistic, such that only one set of regressions needs to be run. This summary statistic was also pre-registered and is given by the average desired expenditure share for the five left-wing domains minus the average desired expenditure share for the three right-wing domains. Appendix Table 14 shows that this summary statistic is strongly correlated with both universalism in altruism and universalism in trust.

Finally, as an additional remedy against multiple testing concerns, Appendix C.8 implements the false discovery rate (FDR) procedure proposed by Benjamini et al. (2006); Anderson (2012). The adjusted p-values that result from this procedure are very similar to the unadjusted ones reported above.

6 Demand- and Supply-Side Explanations

In trying to better understand the structure of policy views in the Western world, this paper focuses on a demand-side explanation: that universalism systematically shapes (dis)approval of policies. We recognize, however, that supply-side mechanisms could also be driving some of the observed patterns. As posited in the political behavior literature, voters might identify with a political party and then simply follow the party line on different policy domains. While it is not entirely clear why such a supply-side story should generate the particular clusters that we observe (in an identical fashion across countries), supply and demand-side mechanisms could also interact in important ways. For instance, it is conceivable that heterogeneity in universalism determines people's policy views on "core" issues and therefore drives the sorting of voters into parties, but that voters simply follow the party line in other domains. Or, political views may exert a feedback effect on basic moral values (Hatemi et al., 2019), which suggests complex interactions between the supply- and demand-side of politics. We seek to provide additional evidence that underscores the relevance of the demand-side mechanism we

propose, rather than to rule out that the supply side also matters.[14]

6.1 Specific Policy Implementations

As we discussed in Section 3, many of the policy domains we consider entail both universalist and communitarian elements, or could be implemented in more or less universalist ways. Hence, our hypotheses implicitly rely on the assumption that there exist *default implementations* for these policies. In the following, we leverage that policies often contain both universalist and communal elements and propose to respondents specific implementations of policies that are more or less universalist in nature. One way of thinking about these empirical exercises is that they are meant to "overwrite" people's mental default implementations of the different policies.

To do so, our survey additionally elicited respondents' desired government spending level for specific policy proposals (two for each broad policy domain). The general objective of this exercise was to manipulate the implementation of a policy, such that one was unambiguously more universalist than the other. We then investigate whether self-reported left-wingers become more supportive of canonical conservative policies once they are implemented in a more universalist fashion, and whether right-wingers become more supportive of canonical left-wing policies once they are implemented in a communitarian fashion.

To systematically manipulate the implementation of a policy, we manipulated the social distance between the respondent and the beneficiaries of a policy. For welfare, environmental protection and health care, we implemented this by manipulating *geographic scope*. To illustrate, consider the example of welfare payments. We elicited desired spending levels separately for (i) "Redistributing local tax revenues as welfare payments across all communities nationwide" and (ii) "Redistributing local tax revenues as welfare payments only within the local communities they were raised."[15] We emphasized to respondents that every other aspect of the policy, e.g., collection of money and

[14]One piece of suggestive evidence that supply-side mechanism may play a role stems from our own surveys. In addition to the Western countries for which we reported results above, we also implemented the same survey in Brazil and South Korea. As is well-known in comparative political science, intracorrelations among policy views outside the cultural West do not follow the same patterns as in the West and cannot be neatly organized according to a left-vs.-right axis. We therefore pre-registered that the link between universalism and policy views would be different in Brazil and South Korea. Appendix C.9 replicates our main analysis for these two countries. The relationships between universalism and policy preferences are all weaker in magnitude and sometimes opposite in sign relative to those observed in Western countries. While these results do not allow us to disentangle between supply- and demand-side mechanisms, they are at least consistent with the view that a difference between Western and non-Western political competition is that parties in the West more aggressively frame political issues around themes related to universalism. This may of course itself be partly driven by parties' recognition of demand-side forces related to universalism.

[15]Figure 16 in Appendix B.2 provides a screenshot.

efficiency of redistribution, was identical between the two implementations. Thus, the two policies plausibly only differ in universalism-related aspects: who the money goes to (which concerns universalism in altruism) and whose tax cheating matters (which concerns universalism in trust). We used similar language to elicit desired spending levels for health care that is funded through local vs. national tax revenues, and for local environmental protection vs. preventing global climate change, see Appendix Table 20.

For foreign aid and affirmative action, we likewise manipulated the social distance between respondent and beneficiaries of a policy, except that this was done through other means than geography. For foreign aid, we separately elicit desired spending for (i) "Sending foreign aid to countries that are in most need of help" and (ii) "Sending foreign aid to foreign countries that are our international allies," which manipulates social distance. Similarly, for affirmative action, we separately elicit support for policies that ensure that (i) "no individual is disadvantaged" vs. that "no one of your same background (e.g., gender, ethnicity or ancestry) is disadvantaged." Again, this manipulates the universalist appeal of the policy domain because the latter proposal is directed at improving the welfare of in-group members.[16] See Appendix Table 20. We acknowledge that not all of these specific policy implementations are as tight as the ones for welfare, health care and environmental protection, in the sense that they sometimes change more than one aspect of a policy at the same time. However, we nonetheless view these exercises as helpful because they suggest that making policies more universalist (in one or multiple ways) changes *who* supports a policy.

Figure 11 illustrates the results by showing binned scatter plots that link arsinh desired spending levels for specific policy proposals to respondents' self-reported left-vs.-right ideology (0–10), where higher values mean that the respondent considers himself more left-wing. To conserve space, we show the results for welfare, environmental protection and foreign aid; the results for healthcare and affirmative action look very similar, see Appendix C.10. In each row, the left panel shows the link between left-vs.-right ideology and the more universalist policy implementation, while the right panel shows results for the more communitarian implementations.

Unsurprisingly, we see that a left-wing political orientation is strongly correlated with desired spending for national redistribution, global climate change prevention and foreign aid to the most needy. These are essentially the standard correlations that we see in the political discourse, as the more universalist policies arguably closely correspond to

[16]We process the data as follows. First, we again first perform a PPP conversion to USD. Second, because there are some huge outliers in the free text entry format, we delete all responses that are larger than 500 times the within-country median response. This affects 0.2% of responses. Third, as pre-registered, we winsorize the data at +/− 3 sd. of the within-country mean, which affects 1.3% of all responses. Fourth, we take the inverse hyperbolic sine of these winsorized PPP amounts, and finally standardize these variables into z-scores, separately within each country.

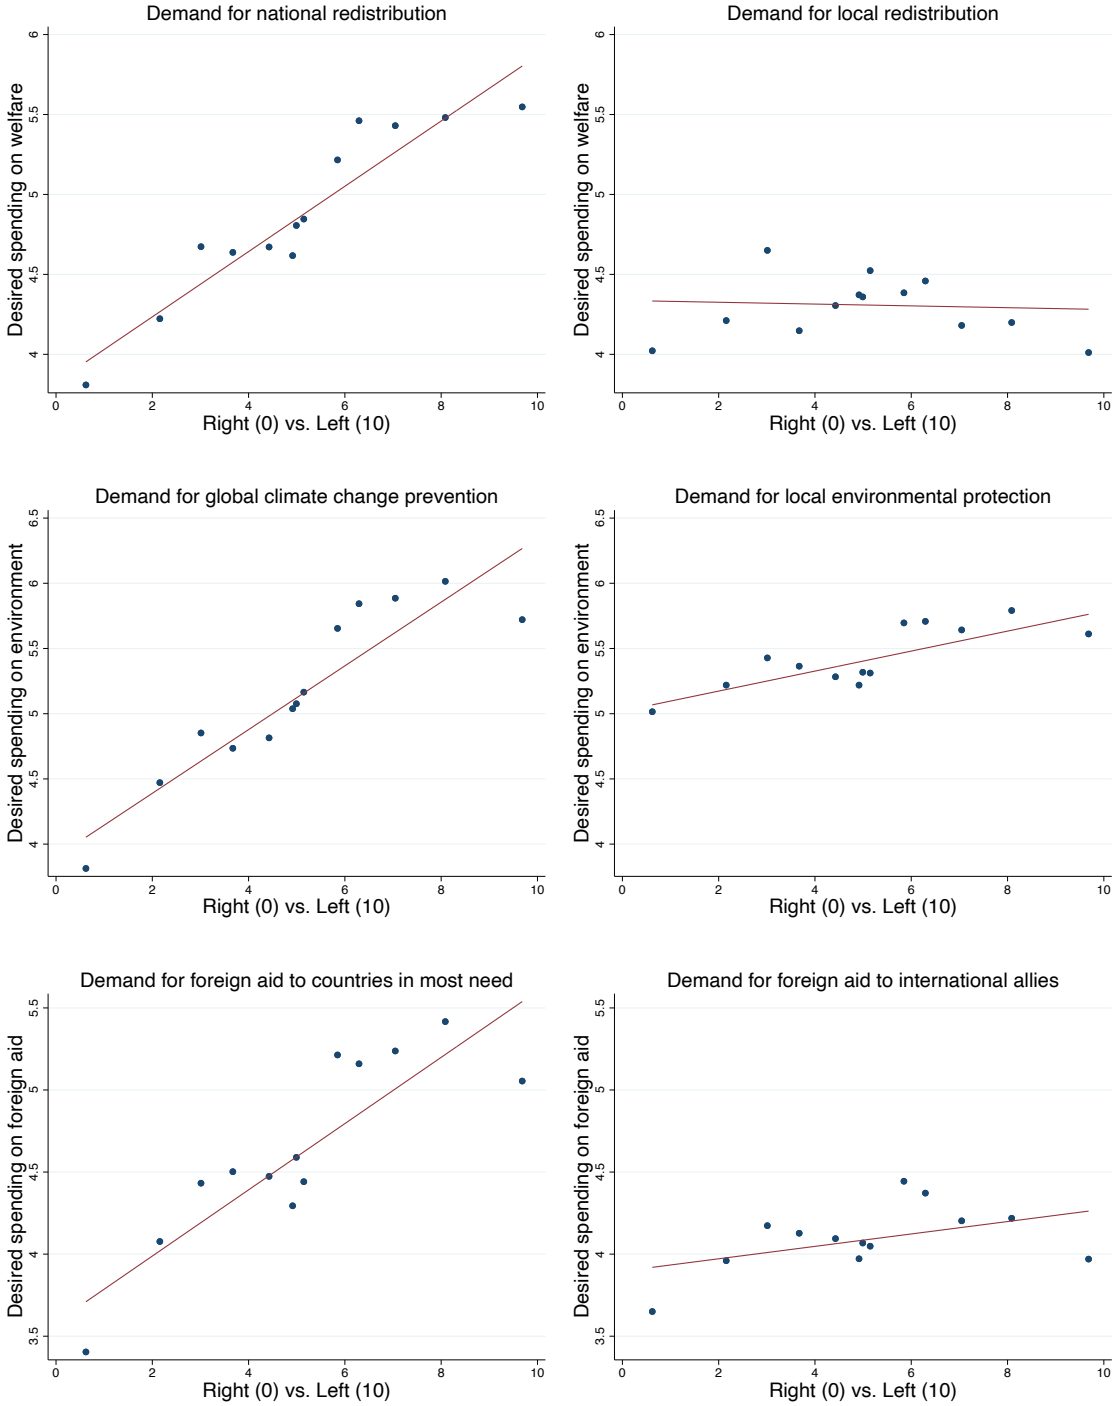

Figure 11: This figure shows binned scatter plots between arsinh desired spending levels for specific policy proposals and respondents' political ideology. In each row, the left panel shows the more universalist and the right panel the more communitarian implementation. The data are pooled across countries, and all panels are constructed controlling for country fixed effects.

what we referred to as "default manipulation" (or default framing) in Section 3. However, much more interestingly, and as we hypothesized, we also see that self-identified conservatives become almost as supportive of redistribution, environmental protection and

foreign aid once these policies are implemented in more communitarian ways. The partial correlations (conditional on country fixed effects) between the self-reported left-vs.-right scale and desired spending levels drops from $\rho = 0.16$ for national redistribution to $\rho = 0.00$ for local redistribution, from $\rho = 0.20$ for global climate change prevention to $\rho = 0.07$ for local environmental protection, and from $\rho = 0.17$ for foreign aid to the most needy to $\rho = 0.03$ for foreign aid to international allies.

In Appendix C.10, we implement similar exercises for the traditional conservative spending categories of border control, military and police and law enforcement. In line with the results summarized above, we find that self-identified left-wingers become substantially more likely to support these expenditures once they focus on universalist objectives, such as peacekeeping by the military.

Overall, these patterns are suggestive that people's policy views systematically depends on whether the policies are implemented in a universalist or communitarian fashion. This seems at odds with a pure supply-side account according to which people support certain government expenditure categories because they "learned" from their party that this is the right thing to do.

6.2 Universalist Minorities and Support for Affirmative Action

As a second strategy to document at least a partial role for demand-side explanations, we leverage that the role of moral universalism in some policy domains should be mediated by demographics. This is most evident for affirmative action, where, arguably, the link between universalism and support for affirmative action should be weaker among the beneficiaries of affirmative action (e.g, racial minorities, women).

To formally investigate this, in columns (1) and (3) of Table 3, we regress support for affirmative action on moral universalism, interacted with a gender dummy. Columns (2) and (4) implement a similar analysis for racial minorities. Here, we restrict attention to the United States, for which we arguably have a clear hypothesis about who the racial minorities are that stereotypically benefit from affirmative action (Hispanics, African Americans and Native Americans).[17] According to our demand-side account, the coefficient on universalism should be smaller for the minority group, whereas a supply-side story would predict similar effects for minorities and non-minorities.

Overall, we see that the "effect" of universalism on affirmative action is weaker among racial minorities and women. We interpret these results as providing additional evidence for the partial relevance of a demand-side story.

[17]In contrast, in the other countries in our sample, the race / ethnicity variable does not allow us to cleanly identify minorities / beneficiaries of affirmative action. For example, when a respondent in Germany indicates that s/he is not German but "Other European," then this person may be a minority (because s/he may be a foreigner) but also a majority (because s/he is likely white).

Table 3: Universalism and support for affirmative action

	Dependent variable: Z-scores of support for affirmative action:			
	Desired arsinh expenditure		Likert scale	
	(1)	(2)	(3)	(4)
Composite universalism	0.54***	0.90***	0.74***	1.06***
	(0.12)	(0.18)	(0.12)	(0.17)
Female	0.044		0.21***	
	(0.07)		(0.07)	
Female x Universalism	-0.054		-0.37**	
	(0.17)		(0.16)	
Racial minority		0.68***		0.59***
		(0.19)		(0.16)
Racial minority x Universalism		-1.36***		-1.09***
		(0.46)		(0.37)
Country FE	Yes	No	Yes	No
Controls	Yes	Yes	Yes	Yes
Observations	11063	2934	11063	2934
R^2	0.05	0.09	0.19	0.28

Notes. OLS estimates, robust standard errors in parentheses. In columns (2) and (4), the sample is restricted to U.S. respondents. Racial minorities include African-Americans, Hispanics and Native Americans. Universalism is in [0,1]. Controls include age, gender, income, wealth, college education, urbanicity, religiosity, equity-efficiency preferences, altruism, trust, beliefs about the efficiency of government, and beliefs about whether one will personally benefit from government expenditure in affirmative action. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

6.3 Replicating Results Controlling for Political Ideology

Finally, we attempt to shut down the supply-side channel by looking at the relationship between universalism and policy preferences *conditional* on political affiliation. To this effect, Appendix C.11 replicates all results reported above, controlling for fixed effects for respondents' self-positioning on a 0–10 left-right scale. There are three results. First, universalism remains a statistically significant predictor of policy views in the same manner as identified in Section 5, see Appendix Figure 43.

Second, in a difference-in-difference analysis, universalism is strongly correlated with the degree to which respondents prefer the more universalist implementation of a policy relative to the more communitarian one, conditional on political ideology fixed effects, see Appendix Figure 44.

Third, and relatedly, Appendix Figures 45 through 47 show the link between universalism and support for the specific universalist or communitarian policy implementations discussed in Section 6.1. We find that the relationship between policy views and universalism attenuates or reverses as we switch from a universalist to a more communitarian implementation of a policy, even when we include left-right fixed effects.

7 Conclusion

For decades, philosophers have argued about the scope of our moral obligations. Should we always be impartial? Do we have some special connection to those that are socially close to us, in a way that should make us trust them more and feel more morally responsible for them? These deep normative questions have direct empirical counterparts because many policies and government expenditure categories that are of central concern to economists implicitly or explicitly concern tradeoffs between strangers and those that are close to us. Yet, empirically, people exhibit large disagreement over what the morally appropriate way to treat different social groups is. This paper has proposed that this individual-level heterogeneity in universalism accounts for the particular structure of policy views that Western democracies have seen over the last few decades.

We believe that this empirical insight is valuable for at least three reasons. First, it helps us to understand where political differences come from. Second, it sheds light on how policy-makers can appeal to voters by proposing policies with a particularly universalist or communitarian flavor. Third, our paper arguably does some heavy lifting in terms of a dimensionality reduction of political views: rather than merely explain attitudes on one topic (e.g., redistribution) with people's underlying moral preferences and beliefs, we endogenize a vector of eight (and perhaps more) policy views as a function of morality.

As discussed in Section 2, our analysis is conditional on two restrictions. First, we only analyze the structure of ideology as it has prevailed over the last 40-50 years. We do not have much to say about whether or how universalism mattered for policy in more distant history. This being said, there is some evidence that suggests that the relevance of universalism for politics has increased over time. The Democratic "loss of the South" and subsequent polarization were largely tied to ideas related to (non-) universalism (Kuziemko and Washington, 2018). Furthermore, Enke (2020) documents using text analyses that Republicans and Democrats used universalist vs. communal moral language in roughly equal frequencies until the mid-60's but steadily diverged thereafter, which could be understood as suggesting that heterogeneity in universalism is more relevant politically today than in the past.

Second, our analysis deliberately focused on the Western world. The connection between the structure of ideology and morality might be different outside the West for various reasons. Future research is needed to understand why the political systems of non-"WEIRD" societies (Henrich, 2020) don't give rise to the ideological clusters that dominate the West, and which role morality plays in these contexts.

References

Alesina, Alberto and Edward Glaeser, *Fighting poverty in the US and Europe: A world of difference*, Oxford University Press, 2004.

_ , Reza Baqir, and William Easterly, "Public goods and ethnic divisions," *The Quarterly Journal of Economics*, 1999, *114* (4), 1243–1284.

Algan, Yann, Elizabeth Beasley, Daniel Cohen, and Martial Foucault, "The rise of populism and the collapse of the left-right paradigm: lessons from the 2017 French presidential election," 2018.

Anderson, Michael L., "Multiple Inference and Gender Differences in the Effects of Early Intervention: A Reevaluation of the Abecedarian, Perry Preschool, and Early Training Projects," *Journal of the American Statistical Association*, 2012, *103*.

Ansolabehere, Stephen, Jonathan Rodden, and James M Snyder, "The strength of issues: Using multiple measures to gauge preference stability, ideological constraint, and issue voting," *American Political Science Review*, 2008, *102* (2), 215–232.

_ , _ , and James M Snyder Jr, "Purple america," *Journal of Economic Perspectives*, 2006, *20* (2), 97–118.

Benjamini, Yoav, Abba M Krieger, and Daniel Yekutieli, "Adaptive Linear Step-Up Procedures that Control the False Discovery Rate," *Biometrika*, 2006, *93* (3), 491–507.

Besley, Timothy and Torsten Persson, "The Rise of Identity Politics," 2019.

Bó, Ernesto Dal, Frederico Finan, Olle Folke, Torsten Persson, and Johanna Rickne, "Economic Losers and political winners: Sweden's radical right," *Working Paper*, 2018.

Bonomi, Giampaolo, Nicola Gennaioli, and Guido Tabellini, "Identity, beliefs, and political conflict," *The Quarterly Journal of Economics*, 2021, *136* (4), 2371–2411.

Broockman, David E and Daniel M Butler, "The Causal Effects of Elite Position-Taking on Voter Attitudes: Field Experiments with Elite Communication," *American Journal of Political Science*, 2017, *61* (1), 208–221.

Chen, Daniel L and Jo Thori Lind, "The Political Economy of Beliefs: Why Fiscal and Social Conservatives/Liberals (Sometimes) Come Hand-in-Hand," *Working Paper*, 2019.

Draca, Mirko and Carlo Schwarz, "How polarised are citizens? Measuring ideology from the ground up," *Working Paper*, 2020.

Enke, Benjamin, "Kinship, cooperation, and the evolution of moral systems," *The Quarterly Journal of Economics*, 2019, *134* (2), 953–1019.

_ , "Moral values and voting," *Journal of Political Economy*, 2020, *128* (10), 3679–3729.

_ **, Ricardo Rodríguez-Padilla, and Florian Zimmermann**, "Moral Universalism: Measurement and Economic Relevance," *Management Science*, 2021.

Erikson, Robert S and Kent L Tedin, *American public opinion: Its origins, content and impact*, Routledge, 2015.

Gethin, Amory, Clara Martínez-Toledano, and Thomas Piketty, "Brahmin left versus merchant right: Changing political cleavages in 21 Western Democracies, 1948–2020," *The Quarterly Journal of Economics*, 2022, *137* (1), 1–48.

Gilens, Martin, *Why Americans hate welfare: Race, media, and the politics of antipoverty policy*, University of Chicago Press, 2009.

Gillen, Ben, Erik Snowberg, and Leeat Yariv, "Experimenting with Measurement Error: Techniques with Applications to the Caltech Cohort Study," *Journal of Political Economy*, 2019, *127* (4), 1826–1863.

Graham, Jesse, Jonathan Haidt, and Brian A. Nosek, "Liberals and Conservatives Rely on Different Sets of Moral Foundations," *Journal of Personality and Social Psychology*, 2009, *96* (5), 1029.

Grossman, Gene and Elhanan Helpman, "Identity Politics and Trade Policy," *Working Paper*, 2018.

Guriev, Sergei and Elias Papaioannou, "The political economy of populism," 2020.

Haidt, Jonathan, *The Righteous Mind: Why Good People are Divided by Politics and Religion*, Vintage, 2012.

Hatemi, Peter K, Charles Crabtree, and Kevin B Smith, "Ideology justifies morality: Political beliefs predict moral foundations," *American Journal of Political Science*, 2019, *63* (4), 788–806.

Henrich, Joseph, "The WEIRDest People in the World: How the West Became Psychologically Peculiar and Particularly Propserous," *New York: Farrar, Strauss, and Giroux*, 2020.

Kozlowski, Austin and James P Murphy, "Issue Alignment and Partisanship in the American Public: Revisiting the'Partisans without Constraint'Thesis," 2019.

Kuziemko, Ilyana and Ebonya Washington, "Why Did the Democrats Lose the South? Bringing New Data to an Old Debate," *American Economic Review*, 2018.

Lakoff, George, *Moral politics: How liberals and conservatives think*, University of Chicago Press, 2010.

Lenz, Gabriel S, *Follow the leader?: how voters respond to politicians' policies and performance*, University of Chicago Press, 2013.

Levine, David K, "Modeling altruism and spitefulness in experiments," *Review of economic dynamics*, 1998, *1* (3), 593–622.

Luttmer, Erzo FP, "Group loyalty and the taste for redistribution," *Journal of political Economy*, 2001, *109* (3), 500–528.

Poole, Keith T and Howard Rosenthal, *Congress: A political-economic history of roll call voting*, Oxford University Press on Demand, 2000.

Putnam, Robert D, *Bowling alone: The collapse and revival of American community*, Simon and schuster, 2000.

Scheve, Kenneth and David Stasavage, "Religion and Preferences for Social Insurance," *Quarterly Journal of Political Science*, 2006, *1*, 255–286.

Shayo, Moses, "A model of social identity with an application to political economy: Nation, class, and redistribution," *American Political science review*, 2009, *103* (2), 147–174.

Sowell, Thomas, *A conflict of visions: Ideological origins of political struggles*, Basic Books, 2007.

Tabellini, Guido, "The Scope of Cooperation: Values and Incentives," *Quarterly Journal of Economics*, 2008, *123* (3), 905–950.

Waytz, Adam, Ravi Iyer, Liane Young, Jonathan Haidt, and Jesse Graham, "Ideological differences in the expanse of the moral circle," *Nature Communications*, 2019, *10* (1), 1–12.

Wu, Alex, "Political polarization," *Undergraduate senior thesis, Harvard University*, 2020.

ONLINE APPENDIX

A Formal Framework

Here we provide a simple model based on Tabellini (2008) that generates the hypotheses stated in Section 3.

The main ingredients of the model are: (i) individuals live on a rectangle and hence at different (social) distance to different members of humanity, where social distance could capture distance along the lines of family, ethnicity, religion, language, values, geography etc.; (ii) the two policy options differ in the extent to which they enable or rule out free-riding or cheating; (iii) in terms of timeline, individuals first vote on a policy and then decide whether to cheat on society; and (iv) individuals differ in the extent to which their altruism and trust are universalist. Our object of interest is how an individual's universalism affects their choice between the two policies.

A.1 Social Distance and Preferences

Let I be a set consisting of N individuals from two separate countries, where for analytical convenience we assume that $N \to \infty$ and that N is a multiple of four. We formalize countries and social distances by allocating individuals in equal proportion to the vertices of a rectangle of length d_l and width d_w where $d_w < d_l$ and $d_w + d_l = 1$. The social distance between individuals i and j is $d_{i,j}$, where distance is measured along the edges of the rectangle. People in the same country are connected by the short end of the rectangle.

We assume that each vertex of the rectangle corresponds to a social group. People who populate the same vertex are said to belong to the same domestic in-group (say, the same neighborhood or the same set of religious beliefs). Individuals at distance d_w can be thought of as domestic out-group. Likewise, we think of individuals at distance d_l as global in-group (say, people who live in a different country but adhere to the same values) and at distance $d_l + d_w = 1$ as global out-group.

Individuals care about their own consumption and the consumption of others, though to potentially heterogeneous degrees. Our formalization of universalism is similar to Tabellini (2008); also see Enke (2019) for a recent cultural economics application. Define $J_i = I \setminus \{i\}$ to be the set of $N-1$ people in the population other than i and by D_i the set of $N/2-1$ domestic people other than i. Let x_i denote the consumption of individual

i. The utility function of individual i is given by

$$u_i(x_i, x_{-i}) = x_i + \beta_i \sum_{j \in J_i} x_j a_{i,j}(d_{i,j}, \theta_i) \tag{1}$$

$$a_{i,j}(d_{i,j}, \theta_i) = \frac{1 + \theta_i}{2} - \theta_i d_{i,j} \tag{2}$$

The parameter $\beta_i \in (0, 1]$ scales individual i's *level* of altruism, while $\theta_i \in (0, 1]$ governs the *slope* of altruism as a function of social distance. Figure 3, introduced also in Enke et al. (2021), illustrates. We construct $a_{i,j}$ such that (i) altruism declines linearly as a function of distance; and (ii) the function integrates to a constant (1/2). This clarifies that the universalism parameter θ_i does not scale who is "more or less moral," but only how uniformly an individual distributes a given altruism budget.[18] Intuitively, a full universalist might argue that it is appealing to treat everyone equally, while others might point out that the universalist's moral compass is distorted in that she treats her friends not very well. Indeed, in Enke et al. (2021) we show that universalists have fewer friends and spend less time with them, compared to less universalist people.

A.1.1 Domestic Policy Options

Individuals first vote for one of two policies in a simple majority system, where voting is assumed to be sincere. After each individual casts a vote $v_i \in \{A, B\}$, depending on which policy was selected, individuals potentially take an action $q_i \in \{0, 1\}$ that we will think of as free-riding or cheating. The "safe" policy option A enforces that nobody can cheat on society. However, the enforcement of this policy is costly, and that cost is shared equally among all domestic individuals for a per capita cost c that is deducted from the baseline consumption level that is normalized to zero. The "risky" option B does not impose a per capita cost on each person, yet each individual can free-ride or cheat on society. Cheating by person i delivers an extra rent $s > 0$ for person i but imposes an overall externality of $e > 2s$, which is shared equally among all domestic individuals for a per capita externality of $2e/N$.

Let us emphasize that we only use the terminology "safe" and "risky" policy to point out the scope for cheating that is implied by the policies – it will sometimes be the case that what we call the "safe policy" is riskier in respects other than cheating, but this is immaterial for our purposes.

As explained in Table 4, we argue that these abstract features of the two policies

[18] A potential micro-foundation for such type-dependent altruism is that individuals exhibit greater altruism towards those people that they believe to be "good" types, as in the model of Levine (1998). Then, our utility function corresponds to a reduced-form version of a model in which beliefs about the types of others vary as a function of social distance, as in Section A.1.2 below.

map into some of the structural features of each of the eight policy domains discussed above. For example, in the case of welfare payments, the safe option A corresponds to a system with few welfare payments, so that individuals cannot cheat on society by claiming benefits they are not entitled to. On the other hand, this causes a societal loss because random income shocks cannot be equalized. The risky option B, on the other hand, corresponds to a more expansive welfare state, which opens up the possibility of cheating.

In this model, option A and option B are not defined by the implied level of government spending. Indeed, as can be seen in Table 4, in the domain of welfare, the safe option A corresponds to lower spending, while in the domain of police and law enforcement, the safe policy A corresponds to higher spending.

A.1.2 Beliefs and Equilibrium

In order to calculate valuations of each policy, a decision maker must form beliefs about who would cheat under policy option B. We model decision makers that have rational expectations about the *overall* fraction of individuals who will cheat under option B. However, to allow for an analysis of the role of universalism in trust beliefs, we assume that decision makers may not form correct beliefs about *which* individuals will cheat. For example, casual introspection suggests that people differ dramatically in whether they believe that immigrants are more likely to be criminals than natives. Formally, the subjective probability that individual i assigns to individual j not cheating under policy B is

$$b_{i,j}(d_{i,j}, \delta_i) = \gamma^* + \frac{d_w}{2}\delta_i - \delta_i d_{i,j} \tag{3}$$

where $\delta_i \in (0, 1]$ controls the rate at which the belief that an individual will not cheat falls as a function of social distance. We think of δ_i as the inverse of universalism in trust. Note that beliefs are defined analogously to altruism above, and can be graphically represented analogously to Figure 3. As in the case of altruism, universalists and non-universalists do not differ in their overall level of trust: the belief function in equation (3) integrates to the constant γ^*, which is endogenous and corresponds to the fraction of individuals who do not cheat in equilibrium. Thus, in equilibrium, individuals have rational expectations about the overall rate of cheating in society but not necessarily about how cheating is correlated with social distance from them. We assume that $(\beta_i, \theta_i, \delta_i)$ are positive independent joint uniform.

Table 4: Mapping of policy domains to abstract model policies

Policy domain	**Safe Option A**	**Risky Option B**
Abstract framework	Cheating impossible, but individuals pay fixed per capita cost of c	No fixed cost, but individuals can cheat and hence earn rent s by imposing per capita externality of e
Domestic policies		
Welfare	No welfare state: Individuals cannot cheat by claiming benefits they are not entitled to, yet this imposes a per capita cost c because in the absence of welfare payments, random income shocks cannot be equalized ex post	Expansive welfare state: Individuals can cheat on society by claiming benefits they are not entitled to (and hence reap rent s) or through moral hazard (reducing labor supply); this causes per capita externality e
Universal health care	Same logic as for welfare	
Affirmative Action	No AA: Individuals cannot reduce effort while still getting promoted; yet absence of AA also entails a social cost because disadvantaged groups in society cannot live up to their potential	Extensive AA: Individuals who benefit from AA can cheat by reducing effort because they know that they will get promoted either way; this imposes a cost on other individuals
Police and law enforcement	Strong police: Stealing is impossible; but entails a cost because police needs to be paid for	Weak police: Stealing and fraud possible
Foreign policies		
Effective border control	Strong border control: less immigration, but this entails a per capita cost because border control is expensive, and because some immigrants are truly in need	Weak border control: Increase in number of people who could come into country and free ride on others' efforts
Military	Strong military: Foreigners cannot exploit domestic people; entails per-capita cost because military needs to get paid for	Weak military: Other countries can cheat or exploit
Environmental protection	Weak regulation: Foreign countries cannot exploit domestic regulation; yet this entails cost because environmental degradation might have economic or health impacts on domestic individuals	Strong regulation: Other countries can cheat by de-regulating and hence growing their economy at expense of domestic individuals
Foreign aid	No aid: Foreigners cannot cheat; yet this entails a cost because lack of aid could cause increased migration or wars	Extensive aid: Foreigners can cheat by claiming aid they are not entitled to or by misusing funds

A.1.3 Domestic Policy Views

We solve the game by backward induction. Denote by $E_i[\cdot]$ the subjective "expectations operator" that applies the belief function in equation (3). Further denote by L_{-i} the hypothetical losses that individuals incur due to the cheating of individuals other than i. In the second stage of the game, if the risky policy is implemented, individual i cheats iff

$$E_i[u_i(q_i=1)] = \left(s - \frac{2e}{N} - L_{-i}\right) + \beta_i \sum_{j \in D_i} \left\{[1 - b_{i,j}(\delta_i)] \cdot s - \frac{2e}{N} - L_{-i}\right\} \cdot a_{i,j}(\theta_i) \quad (4)$$

$$> E_i[u_i(q_i=0)] = -L_{-i} + \beta_i \sum_{j \in D_i} \{[1 - b_{i,j}(\delta_i)] \cdot s - L_{-i}\} \cdot a_{i,j}(\theta_i) \quad (5)$$

which delivers the vector q^* of individual cheating decisions $q_i^*(\theta_i)$. The resulting losses (externalities) that each individual incurs are denoted by $L^*(q^*(\theta)) \equiv (1-\gamma^*(\theta))e$.

In the first stage of the game, an individual votes for the safe policy A iff

$$u_i(A) = -c + \beta_i \sum_{j \in D_i} \{-c\} \cdot a_{i,j}(\theta_i) \quad (6)$$

$$> E_i[u_i(B)]$$

$$= [s \cdot q_i^*(\theta_i) - L^*(q^*(\theta))] + \beta_i \sum_{j \in D_i} \{[1 - b_{i,j}(\delta_i)] \cdot s - L^*(q^*(\theta))\} \cdot a_{i,j}(\theta_i) \quad (7)$$

which delivers the vote $v_i^*(\theta_i, \delta_i)$ as a function of universalism and other parameters.

We impose two sets of assumptions that have straightforward economic intuitions. First, we impose restrictions on model parameters that ensure that the (equilibrium) efficiency losses under the risky policy are larger than the efficiency losses under the safe policy. Formally, define $z \equiv \frac{2\frac{s}{e}\log(2-d_w)}{1-d_w}$. As will be shown in Section A.2.2, z captures the share of cheaters in the domestic policy equilibrium. We state the following assumptions:

$$z(e-s) > c \quad (8)$$

$$ed_w/2 \geq c \quad (9)$$

In addition, we impose two restrictions on parameters to ensure that individuals' beliefs about others' cheating are in $[0,1]$:

$$\min\{2z, 2(1-z)\} \geq d_w \quad (10)$$

$$\gamma_f \in [d_w/2, 1 - d_w/2] \quad (11)$$

We obtain the following predictions:

Predictions. *Individuals with higher universalism exhibit a stronger preference for the risky domestic policy B: welfare, universal health care, affirmative action, and weak police and law enforcement. These predictions hold for universalism in both altruism and trust.*

A.2 Proof

Altruism and trust decline linearly in social distance at rate θ_i and δ_i, respectively. Since distances are discrete, we can define altruism and beliefs about other individuals in the same country as

$$a_0 = \tfrac{1+\theta_i}{2} \qquad a_1 = \tfrac{1+\theta_i}{2} - \theta_i d_w$$

$$b_0 = \gamma + \tfrac{d_w}{2}\delta_i \qquad b_1 = \gamma - \tfrac{d_w}{2}\delta_i$$

where throughout these derivations, the 0 subscript refers to domestic in-group members, and the 1 subscript to domestic out-group members.

A.2.1 Cheating Decision

We proceed by backward induction. Suppose the risky policy got implemented. We now ask under which conditions an individual cheats.

Suppose that some arbitrary fraction φ_i of the other $\frac{N}{2}-1$ domestic individuals is perceived by person i to be cheating on the policy. For every individual i, per-capita loss caused by all cheating individuals *other* than i totals $\varphi_i\left(\frac{N}{2}-1\right)\frac{e}{\frac{N}{2}} = \varphi_i\left(\frac{N-2}{N}\right)e$. Thus:

$$E_i[u_i(q_i=0)] =$$

$$\underbrace{-\varphi_i\left(\frac{N-2}{N}\right)e}_{\substack{\text{Person } i\text{'s}\\ \text{consumption utility}}} + \beta_i\Bigg[\underbrace{\left(\frac{N}{4}-1\right)a_0\tilde{b}_0\left(-\varphi_i\left(\frac{N-2}{N}\right)e\right)}_{\substack{\text{Utility from consumption of domestic}\\ \text{in-group members (other than individual } i)\\ \text{who do not cheat}}} + \underbrace{\left(\frac{N}{4}\right)a_1\tilde{b}_1\left(-\varphi_i\left(\frac{N-2}{N}\right)e\right)}_{\substack{\text{Utility from consumption of domestic}\\ \text{out-group members who do not cheat}}} \tag{12}$$

$$+\underbrace{\left(\frac{N}{4}-1\right)a_0(1-\tilde{b}_0)\left(-\varphi_i\left(\frac{N-2}{N}\right)e+s\right)}_{\substack{\text{Utility from consumption of domestic in-group members}\\ \text{(other than individual } i\text{) who cheat}}} + \underbrace{\left(\frac{N}{4}\right)a_1(1-\tilde{b}_1)\left(-\varphi_i\left(\frac{N-2}{N}\right)e+s\right)}_{\substack{\text{Utility from consumption of domestic out-group}\\ \text{members who cheat}}}\Bigg]$$

$$E_i[u_i(q_i=1)] =$$

$$\underbrace{-\varphi_i\left(\frac{N-2}{N}\right)e-\frac{2e}{N}+s}_{\substack{\text{Person } i\text{'s}\\ \text{consumption utility}}} + \beta_i\Bigg[\underbrace{\left(\frac{N}{4}-1\right)a_0\tilde{b}_0\left(-\varphi_i\left(\frac{N-2}{N}\right)e-\frac{2e}{N}\right)}_{\substack{\text{Utility from consumption of domestic}\\ \text{in-group members (other than individual } i)\\ \text{who do not cheat}}}$$

$$+\underbrace{\left(\frac{N}{4}\right)a_1\tilde{b}_1\left(-\varphi_i\left(\frac{N-2}{N}\right)e-\frac{2e}{N}\right)}_{\substack{\text{Utility from consumption of domestic out-group}\\ \text{members who do not cheat}}}+\underbrace{\left(\frac{N}{4}-1\right)a_0(1-\tilde{b}_0)\left(-\varphi_i\left(\frac{N-2}{N}\right)e-\frac{2e}{N}+s\right)}_{\substack{\text{Utility from consumption of domestic in-group}\\ \text{members (other than individual } i) \text{ who cheat}}} \quad (13)$$

$$+\underbrace{\left(\frac{N}{4}\right)a_1(1-\tilde{b}_1)\left(-\varphi_i\left(\frac{N-2}{N}\right)e-\frac{2e}{N}+s\right)}_{\substack{\text{Utility from consumption of domestic}\\ \text{out-group members who cheat}}}\Bigg]$$

An individual doesn't cheat iff $E_i[u_i(q_i=1)] \leq E_i[u_i(q_i=0)]$, i.e. a person doesn't cheat iff their draw of β and θ satisfies the following no-cheat condition:

$$\frac{Ns}{2e} \leq 1+\beta_i\left[\frac{N}{4}(1+\theta_i(1-d_w))-\frac{1+\theta_i}{2}\right] \quad (14)$$

Under our maintained assumption that $N \to \infty$, this delivers

$$\frac{2s}{e} \leq \underbrace{\beta(1+\theta_i(1-d_w))}_{\equiv\Psi} \quad (15)$$

A.2.2 Equilibrium Fraction of Cheaters

The previous condition defines the equilibrium fraction of individuals that don't cheat γ^*:

$$\gamma^* = \mathbb{P}\left(\Psi \geq \frac{2s}{e}\right) = 1-\mathbb{P}\left(\Psi < \frac{2s}{e}\right) = 1-\iint_{Area} f_{\beta\theta}(\beta,\theta)d\beta d\theta$$

Where $Area = \left\{(\beta,\theta)|\Psi < \frac{2s}{e}\right\}$ and $f_{\beta\theta}(\beta,\theta)$ is the joint probability density function of β and θ. Note that since β and θ are i.i.d. $\mathcal{U}(0,1]$, $f_{\beta\theta}(\beta,\theta) = f_\beta \cdot f_\theta = 1\cdot 1 = 1$ over $\beta \in (0,1], \theta \in (0,1]$.

Under the assumptions stated above (in particular that $2s < e$), we get that:

$Area = A + B$, where:

$$A \equiv \beta \in \left(0, \frac{2\frac{s}{e}}{2-d_w}\right] \times \theta \in (0,1] \tag{16}$$

$$B \equiv \beta \in \left[\frac{2\frac{s}{e}}{2-d_w}, 2\frac{s}{e}\right] \times \theta \in \left(0, \frac{2\frac{s}{e}-\beta}{\beta(1-d_w)}\right] \tag{17}$$

Integrating over $Area$ gives:

$$\gamma^*(e,s,d_w) = 1 - \frac{2\frac{s}{e}\cdot\log(2-d_w)}{1-d_w} \in [0,1] \tag{18}$$

A.2.3 Policy Views

Having derived the cheating decision and the equilibrium fraction of cheating, we now determine voting behavior (relative support for policies A and B). The utility from policy options A and B is given by:

$$\pi_A = \underbrace{-c}_{\substack{\text{Person } i\text{'s} \\ \text{consumption utility}}} + \beta_i\Big[\underbrace{\left(\frac{N}{4}-1\right)a_0(-c)}_{\substack{\text{Utility from consumption of} \\ \text{domestic in-group members} \\ \text{(other than individual } i\text{)}}} + \underbrace{\left(\frac{N}{4}\right)a_1(-c)}_{\substack{\text{Utility from consumption of} \\ \text{domestic out-group members}}}\Big] \tag{19}$$

$$\pi_B(q_i^*(\beta,\theta)) =$$

$$\underbrace{-(1-\gamma^*)e + s(q_i^*(\theta_i))}_{\substack{\text{Person } i\text{'s} \\ \text{consumption utility}}} + \beta_i\Big[\underbrace{\left(\frac{N}{4}-1\right)a_0 b_0(-(1-\gamma^*)e)}_{\substack{\text{Utility from consumption of domestic} \\ \text{in-group members (other than individual i)} \\ \text{who do not cheat}}} + \underbrace{\left(\frac{N}{4}\right)a_1 b_1(-(1-\gamma^*)e)}_{\substack{\text{Utility from consumption of domestic} \\ \text{out-group members who do not cheat}}}$$

$$+\underbrace{\left(\frac{N}{4}-1\right)a_0(1-b_0)(-(1-\gamma^*)e+s)}_{\substack{\text{Utility from consumption of domestic} \\ \text{in-group members (other than individual i)} \\ \text{who cheat}}} + \underbrace{\left(\frac{N}{4}\right)a_1(1-b_1)(-(1-\gamma^*)e+s)}_{\substack{\text{Utility from consumption of domestic} \\ \text{out-group members who cheat}}}\Big] \tag{20}$$

Given this, the relative support for policy A over policy B is given by:

$$\Pi_{A,B} \equiv E[u_i(A)] - E[u_i(B)] = (1-\gamma^*)e - s(q_i^*(\theta_i)) - c \tag{21}$$

$$+\beta_i\left[\left(\frac{N}{4}-1\right)a_0\{(1-\gamma^*)e-(1-b_0)s-c\}\right] \tag{22}$$

$$+\beta_i\left[\left(\frac{N}{4}\right)a_1\{(1-\gamma^*)e-(1-b_1)s-c\}\right] \tag{23}$$

We now investigate how the relative support for policies A and B depends on universalism in trust (δ_i) and universalism in altruism (θ_i).

A.2.4 Policy Views and Universalism in Trust

$$\frac{\partial \Pi_{A,B}}{\partial \delta_i} = \beta_i s \frac{d_w}{4}\left[\theta_i\left(\frac{N d_w}{2}-1\right)-1\right] \tag{24}$$

which is greater than zero as $N \to \infty$.

A.2.5 Policy Views and Universalism in Altruism

Case 1: Individuals that are not on the margin. Consider those individuals that are not on the margin of cheating vs. not cheating ($\Psi \neq \frac{2s}{e}$). In other words, $q_i^*(\theta_i)$ is constant for marginal changes of θ_i, $\frac{\partial q^*}{\partial \theta_i} = 0$. Thus,

$$\begin{aligned}\frac{\partial \Pi_{A,B}}{\partial \theta_i}|_{\Psi\neq\frac{2s}{e}} =& \beta_i\left[\left(\frac{N}{4}-1\right)\frac{1}{2}\left\{(1-\gamma^*)e-(1-b_0)s-c\right\}\right] &(25)\\ &+\beta_i\left[\left(\frac{N}{4}\right)\left(\frac{1}{2}-d_w\right)\left\{(1-\gamma^*)e-(1-b_1)s-c\right\}\right] &(26)\end{aligned}$$

This simplifies to:

$$\begin{aligned}\frac{\partial \Pi_{A,B}}{\partial \theta_i}|_{\Psi\neq\frac{2s}{e}} =& \beta_i\left[\left(\frac{N}{4}-1\right)\frac{1}{2}\left\{(1-\gamma^*)(e-s)+\delta_i\frac{d_w}{2}s-c\right\}\right] &(27)\\ &+\beta_i\left[\left(\frac{N}{4}\right)\left(\frac{1}{2}-d_w\right)\left\{(1-\gamma^*)(e-s)-\delta_i\frac{d_w}{2}s-c\right\}\right] &(28)\end{aligned}$$

The coefficient on δ_i in $\frac{\partial \Pi_{A,B}}{\partial \theta_i}|_{\Psi\neq\frac{2s}{e}}$ is:

$$\beta_i\left[\left(\frac{N}{4}-1\right)\cdot\frac{1}{2}\cdot\frac{d_w s}{2} - \frac{N}{4}\left(\frac{1}{2}-d_w\right)\cdot\frac{d_w s}{2}\right] = \beta_i\frac{d_w s}{4}\left(\frac{N d_w}{2}-1\right)\geq 0$$

This expression is non-negative as $N \to \infty$ and therefore $\frac{\partial \Pi_{A,B}}{\partial \theta_i}|_{\Psi\neq\frac{2s}{e}} \geq \frac{\partial \Pi_{A,B}}{\partial \theta_i}|_{\Psi\neq\frac{2s}{e},\delta=0}$. In other words, for any value of β_i, $\frac{\partial \Pi_{A,B}}{\partial \theta_i}$ takes its smallest value when $\delta_i = 0$ since the coefficient on δ_i is non-negative.

We can therefore focus on evaluating:

$$\begin{aligned}\frac{\partial \Pi_{A,B}}{\partial \theta_i}|_{\Psi\neq\frac{2s}{e},\delta_i=0} =& \beta_i\left[\left(\frac{N}{4}-1\right)\frac{1}{2}\left\{(1-\gamma^*)(e-s)-c\right\}\right] &(29)\\ &+\beta_i\left[\left(\frac{N}{4}\right)\left(\frac{1}{2}-d_w\right)\left\{(1-\gamma^*)(e-s)-c\right\}\right] &(30)\\ =& \beta_i\left(\frac{N}{4}(1-d_w)-\frac{1}{2}\right)\left((1-\gamma^*)(e-s)-c\right) &(31)\end{aligned}$$

Notice that $(\frac{N}{4}(1-d_w)-\frac{1}{2}) > 0$, and so $\frac{\partial \Pi_{A,B}}{\partial \theta_i}|_{\Psi \neq \frac{2s}{e}, \delta_i = 0} > 0$ (and by extension $\frac{\partial \Pi_{A,B}}{\partial \theta_i}|_{\Psi \neq \frac{2s}{e}} > 0$) if $(1-\gamma^*)(e-s)-c > 0$. This is ensured by the assumption in equation (35). We hence have that $\frac{\partial \Pi_{A,B}}{\partial \theta_i}|_{\Psi \neq \frac{2s}{e}} \geq 0$, meaning that support for the safe policy A increases in θ_i everywhere except at $\Psi = \frac{2s}{e}$, to which we now turn attention.

Case 2: Individuals on the margin of cheating. Consider those individuals for whom $\Psi = \frac{2s}{e}$, so that a marginal change in θ_i induces the individual to switch from cheating to not cheating (recall from equation (15) that cheating decreases in θ_i).

For a marginal change of θ_i, the utility derived from the safe policy option A does not change. For the risky policy option B, observed actions change from $q_i^* = 1$ to $q_i^* = 0$, yet utility varies smoothly. To see this, define $\Lambda_i(\beta_i, \theta_i) = u_B(q=0) - u_B(q=1)$, the difference between the utility from cheating and not cheating under policy option B. From equation (15):

$$\Lambda_i(\beta_i, \theta_i) = \beta_i \frac{e}{2}(1 + \theta_i(1-d_w)) - s \tag{32}$$

It is obvious that $\Lambda(\beta, \theta)$ is continuous at 0, i.e., at $\Psi = \frac{2s}{e}$. Intuitively, when the individual switches from $q_i^* = 1$ to $q_i^* = 0$, utility changes smoothly because – by assumption – the individual is indifferent at $\Psi = \frac{2s}{e}$, where $u_B(q=0) = u_B(q=1)$.

Therefore, even though $\Pi_{A,B}$ is composed of a piecewise function (q^*), $\Pi_{A,B}$ is continuous because Λ is continuous at 0:

$$\lim_{\Psi \to \frac{2s}{e}^-} u_A - u_B(q^* = 1) = \lim_{\Psi \to \frac{2s}{e}^+} u_A - u_B(q^* = 0) \tag{33}$$

$$\implies \lim_{\Psi \to \frac{2s}{e}^-} \Pi_{A,B} = \lim_{\Psi \to \frac{2s}{e}^+} \Pi_{A,B} = \lim_{\Psi = \frac{2s}{e}} \Pi_{A,B} \tag{34}$$

Since Π is continuous, and strictly increasing in θ_i on either side of the indifference boundary, Π increases in θ_i over all values of θ_i.

A.3 Foreign Policy

Decision makers are again presented with a choice between two policy options. Under the safe policy option A, domestic and foreign individuals receive their baseline consumption x. Domestic individuals additionally pay a per capita cost c.[19] Under the risky policy B, domestic individuals do not have to pay c. However, in this regime foreign individuals can cheat and get s by imposing an overall cost of e on all domestic people,

[19] In some of the foreign policy domains we consider, c is likely to be paid by both domestic and foreign individuals. Our main predictions remain unchanged if we assume that c is paid by both domestic and foreign people.

which is again equally shared. Table 4 explains how this abstract structure maps into the domains of military, border control, foreign aid, and environmental protection. As with the domestic policies above, note that the risky policy B sometimes corresponds to big and sometimes to small government. Again, the key defining characteristic that matters for our analysis is whether a policy introduces or prevents cheating opportunities.

The mechanics of the foreign policy analysis are very similar to the domestic case. We again assume that individuals' beliefs about the overall fraction of cheaters are correct, but that they have heterogeneous beliefs about how cheating is correlated with distance from them. As for the domestic case, we impose restrictions on model parameters that ensure that the (equilibrium) efficiency losses under the risky policy are larger than the efficiency losses under the safe policy. Specifically, we state the following assumptions:

$$z(e-s) > c \tag{35}$$

$$ed_w/2 \geq c \tag{36}$$

In addition, we impose that:

$$\min\{2z, 2(1-z)\} \geq d_w \tag{37}$$

$$\gamma_f \in [d_w/2, 1-d_w/2] \tag{38}$$

We get:

Predictions. *Individuals with higher universalism exhibit a stronger preference for the risky policy B: weak border control, weak military, stringent environmental protection, and expansive foreign aid. These predictions hold for universalism in both altruism and trust.*

A.4 Proof

In the foreign policy context it is much simpler to derive the relationship between universalism and policy views. This is because in the context of foreign policies, only *foreign* people are allowed to cheat on one's own country.

In order to evaluate each policy option, decision-makers must form beliefs about who cheats. As in the domestic case, we will take the stance that the subjective probability of not cheating declines linearly in social distance at a rate δ_i:

$$b_{i,j}(d_{i,j}, \delta_i) = \gamma_f + \frac{1+d_l}{2}\delta_i - \delta_i d_{i,j} \tag{39}$$

where decision-makers are again correct about the overall fraction of cheaters in the foreign country (γ_f), but may be incorrect in their beliefs about which foreigners cheat.

Levels of altruism and beliefs by distance are given by the below:

$$\begin{array}{cccc} a_0 = \frac{1+\theta_i}{2} & a_1 = \frac{1+\theta_i}{2} - \theta_i d_w & a_2 = \frac{1+\theta_i}{2} - \theta_i d_l & a_3 = \frac{1+\theta_i}{2} - \theta_i \\ b_0 = 1 & b_1 = 1 & b_2 = \gamma_f + \frac{1+d_l}{2}\delta_i - d_l\delta_i & b_3 = \gamma_f + \frac{1+d_l}{2}\delta_i - \delta_i \end{array}$$

where, as above, the 0 subscript refers to domestic in-group members, the 1 subscript to domestic out-group members, the 2 subscript to foreign in-group members, and the 3 subscript to foreign out-group members.

We define L_f to be the total per capita cost of cheating imposed by foreigners. That is, $L_f = (1-\gamma_f)e$. For a domestic individual i, the relative value of Option A to Option B in the foreign policy domain is therefore:

$$\begin{aligned} \Pi \equiv & \underbrace{(-c)}_{\substack{\text{Person } i\text{'s} \\ \text{consumption utility} \\ \text{under Policy A}}} + \beta_i \Big[\underbrace{\left(\frac{N}{4} - 1\right) a_0(-c)}_{\substack{\text{Utility from consumption of} \\ \text{domestic in-group members (other} \\ \text{than individual } i\text{) under Policy A}}} + \underbrace{\left(\frac{N}{4}\right) a_1(-c)}_{\substack{\text{Utility from consumption of} \\ \text{domestic out-group members} \\ \text{under Policy A}}} \Big] \\ & - \underbrace{(-L_f)}_{\substack{\text{Person } i\text{'s} \\ \text{consumption utility} \\ \text{under Policy B}}} - \beta_i \Big[\underbrace{\left(\frac{N}{4} - 1\right) a_0(-L_f)}_{\substack{\text{Utility from consumption of} \\ \text{domestic in-group members (other} \\ \text{than individual } i\text{) under Policy B}}} + \underbrace{\left(\frac{N}{4}\right) a_1(-L_f)}_{\substack{\text{Utility from consumption of} \\ \text{domestic out-group members} \\ \text{under Policy B}}} \\ & + \underbrace{\left(\frac{N}{4}\right) a_2(1-b_2)s}_{\substack{\text{Utility from consumption of} \\ \text{foreign in-group members} \\ \text{who cheat under Policy B}}} + \underbrace{\left(\frac{N}{4}\right) a_3(1-b_3)s}_{\substack{\text{Utility from consumption of} \\ \text{foreign out-group members} \\ \text{who cheat under Policy B}}} \Big] \end{aligned}$$

$$= (L_f - c) + \beta_i \left[\left(\frac{N}{4}(a_0 + a_1) - a_0 \right)(L_f - c) - \left(\frac{N}{4} \right)(a_2(1-b_2) + a_3(1-b_3))s \right]$$

$$= (L_f - c) + \beta_i \Bigg[\left(\frac{N}{4}(1 + \theta_i(1-d_w)) - \frac{1+\theta_i}{2} \right)(L_f - c)$$

$$- \left(\frac{N}{4} \right) \left\{ \left(\frac{1+\theta_i}{2} - \theta_i d_l \right) \left(1 - \left(\gamma_f + \frac{1+d_l}{2}\delta_i - d_l\delta_i \right) \right) + \left(\frac{1+\theta_i}{2} - \theta_i \right) \left(1 - \left(\gamma_f + \frac{1+d_l}{2}\delta_i - \delta_i \right) \right) \right\} s \Bigg] \tag{40}$$

We first take the comparative static of the relative valuation of Option A compared to Option B with respect to universalism in trust δ_i:

$$\frac{\partial \Pi}{\partial \delta_i} = \beta_i \left(\frac{N}{4} \right) s \left(\frac{1-d_l}{2} \right) (\theta_i(1-d_l)) > 0 \tag{41}$$

Next, we take the comparative static of the relative valuation of Option A relative to Option B with respect to universalism in altrusim θ_i:

$$\frac{\partial \Pi}{\partial \theta_i} = \beta_i \left[(L_f - c)\left(\frac{N}{4}(1-d_w) - \frac{1}{2}\right) + \left(\frac{N}{4}\right)\left(d_l - \frac{1}{2}\right)(1-b_2)s + \left(\frac{N}{4}\right)\left(\frac{1}{2}\right)(1-b_3)s \right] \tag{42}$$

Note that $(L_f - c) = ((1-\gamma_f)e - c) \geq 0$ by the assumption in equation (36), so that the entire expression is positive. We hence see that the relative support for the "safe" policy option A increases as universalism in trust and universalism in altruism decrease.

B Additional Details and Analyses for Survey

B.1 Sample Characteristics

This section presents sample characteristics for all seven countries included in our survey. For reasons beyond our control and related to Dynata's reach in data collection, several of our samples are relatively too educated in comparison with a representative sample. This is specifically the case for Australia, Brazil, Germany, South Korea, and Sweden.

B.1.1 Australia

		Study Sample (%)	
Category	Population (%)	*Full*	*Representative*
Gender			
Male	49	47.1	48.6
Female	51	52.9	51.4
Age			
18–29	24	14.5	14.9
30–39	17	19.2	19.2
40–49	17	19.3	19.1
50–59	16	19.4	18.3
60–69	13	14.1	14.6
≥ 70	13	13.5	13.9
Income (annual; AUD)			
Below 20,000	7	5.0	5.1
20,000–34,999	13	12.4	12.8
35,000–49,999	12	12.8	13.0
50,000–64,999	12	12.9	12.8
65,000–79,999	10	10.2	10.5
80,000–99,999	10	11.1	11.1
100,000–124,999	10	11.2	11.0
125,000–149,999	8	10.0	8.9
150,000–199,999	9	8.3	8.6
200,000 or more	8	6.0	6.2
Ancestry			
English	26	11.5	11.8
Australian	25	68.5	67.5

Other	49	20.1	20.7
Education			
No high school	28	7.5	7.7
High school	18	26.4	25.4
Vocational training	29	33.2	34.2
Bachelor's degree or higher	25	33.0	32.6
Employment Status (for those at most 65)			
Employed full-time	55	53.2	55.0
Not employed full-time	45	46.8	45.0

Note: We were advised by Dynata that it is not common practice to ask respondents in Australia about their race or ethnicity. Accordingly, we found data from the Australian census corresponding to ancestry, which we condensed into "Australian", "English", or "Other". Our final sample characteristics correspond closely to guidance from Dynata that 74% of Australian citizens are Australian-born, which leads us to believe respondents interpreted our ancestry question as eliciting their country of birth or nationality, as is more common practice in Australia.

B.1.2 Brazil

		Study Sample (%)	
Category	Population (%)	*Full*	*Representative*
Gender			
Male	48	52.1	50.2
Female	52	47.9	49.8
Age			
18–29	30	36.7	36.9
30–39	22	30.3	28.8
40–49	19	18.6	19.3
50–59	14	11.0	11.4
≥60	15	3.5	3.6
Income (annual; Brazilian reals)			
Below 3,000	5	15.8	16.1
3,000–5,999	12	13.0	13.5
6,000–11,999	22	9.7	10.1
12,000–17,999	17	8.6	8.9
18,000–29,999	20	9.9	10.3
30,000–59,999	16	19.2	20.0
≥60,000	8	23.9	21.1
Ancestry			
White	49	61.3	59.8
Multi-racial	41	27.4	28.5
Other	10	11.3	11.8
Education			
No formal education	45	0.3	0.3
Elementary school	17	3.3	3.5
High school	28	48.1	50.0
Bachelor's degree or higher	10	48.3	46.2
Employment Status (for those at most 65)			
Employed full-time	41	67.8	66.5
Not employed full-time	59	32.2	33.5

Note: Our samples in Brazil are relatively educated, young, wealthy, and employed. We have reason to believe that some subsamples of the Brazilian population are inaccessible to *Dynata*. For example, the Brazilian census likely includes indigenous populations that likely make up a sizable portion of the "No formal education" bucket.

B.1.3 France

Category	Population (%)	Study Sample (%)	
		Full	*Representative*
Gender			
Male	48	47.9	47.7
Female	52	52.1	52.3
Age			
18–29	18	12.9	17.4
30–39	16	20.7	18.0
40–49	16	23.0	18.0
50–59	17	23.3	19.6
≥60	33	20.0	27.0
Income (annual, EUR)			
Below 10,000	7	9.3	9.5
10,000–14,999	6	7.5	7.3
15,000–19,999	13	11.0	12.4
20,000–24,999	12	13.5	13.1
25,000–29,999	11	11.1	10.6
30,000–34,999	10	10.9	10.2
35,000–39,999	8	8.1	7.2
40,000–49,999	13	13.1	12.1
50,000–64,999	10	8.7	8.6
65,000 or more	10	6.9	9.0
Ancestry			
French or other European	85	96.8	95.6
Other	15	3.2	4.4
Education			
No high school	22	17.9	21.4
High school	43	30.9	41.5
Some college	14	20.8	14.8
Bachelor's degree or higher	21	30.4	22.3
Employment Status (for those at most 65)			
Employed full-time	56	65.6	56.5
Not employed full-time	44	34.4	43.5

Note: "High school" corresponded to "Baccalauréat", "Some college" to "Enseignement supérieur, niveau Bac+2 max", and "Bachelor's degree or higher" to "Enseignement supérieur, niveau Bac+3 et plus".

B.1.4 Germany

		Study Sample (%)	
Category	Population (%)	*Full*	*Representative*
Gender			
Male	49	51.0	50.3
Female	51	49.0	49.7
Age			
18–29	21	13.5	13.7
30–39	14	19.4	19.7
40–49	19	21.0	21.3
50–59	17	26.0	25.0
60–69	13	16.0	16.2
≥70	17	4.0	4.1
Income (monthly; EUR)			
Below 1,300	19	15.1	15.3
1,300–2,599	33	33.6	34.0
2,600–3,599	19	22.1	22.4
3,600–5,000	15	21.8	20.7
More than 5,000	14	7.5	7.6
Ancestry			
German	79	96.6	96.6
European (not German)	15	2.3	2.3
Other	6	1.1	1.1
Education			
No vocational training	27	5.6	5.7
Vocational training	57	58.3	59.1
University degree	16	36.1	35.2
Employment Status (for those at most 65)			
Employed full-time	59	64.3	63.8
Not employed full-time	41	35.7	36.2

Note: The option included in the survey equivalent to vocational training was “Lehre oder Berufsausbildung im dualen System”. For “University degree”, the option provided was “Hochschulabschluss”.

B.1.5 South Korea

		Study Sample (%)	
Category	Population (%)	*Full*	*Representative*
Gender			
Male	50	49.3	48.3
Female	50	50.7	51.7
Age			
18–29	19	19.6	23.0
30–39	18	30.8	26.2
40–49	20	27.3	24.7
50–59	20	16.3	19.1
60–69	13	5.5	6.4
≥70	10	0.5	0.6
Income (annual; ten-thousand Won)			
Below 200 ten-thousand Won	19	9.7	11.4
200–350 ten-thousand Won	23	27.4	27.7
350–500 ten-thousand Won	21	26.3	23.4
500–750 ten-thousand Won	17	22.0	20.5
More than 750 ten-thousand Won	20	14.5	17.1
Ancestry			
Korean	96	99.8	99.8
Other	4	0.2	0.2
Education			
No high school	13	1.0	1.2
High school	40	26.9	31.6
Some college	13	7.4	8.6
Bachelor's degree or higher	34	64.8	58.7
Employment Status (for those at most 65)			
Employed full-time	59	85.3	82.7
Not employed full-time	41	14.7	17.3

Note: Our samples in Korea are relatively too educated, too young, and too employed.

B.1.6 Sweden

		Study Sample (%)	
Category	Population (%)	*Full*	*Representative*
Gender			
Male	50	58.1	50.3
Female	50	41.9	49.7
Age			
18–29	24	12.5	17.8
30–39	15	10.5	14.8
40–49	15	13.8	16.7
50–59	15	19.2	16.2
60–69	13	20.0	14.8
≥70	18	24.0	19.8
Income (annual; Swedish kronor)			
Below 100,000 kr	14	6.9	9.8
100,000–200,000 kr	13	16.0	13.9
200,000–299,999 kr	18	20.8	18.3
300,000–399,999 kr	25	22.7	25.7
400,000–499,999 kr	16	16.8	16.1
500,000–599,999 kr	7	7.8	7.2
600,000–749,999 kr	4	4.6	4.8
750,000–999,999 kr	2	2.6	2.7
1,000,000 kr or more	1	1.8	1.6
Ancestry			
Swedish	82	92.8	89.7
Other	18	7.3	10.3
Education			
No high school	40	8.1	11.5
High school	22	32.3	31.8
Some college	15	30.1	24.7
Bachelor's degree or higher	23	29.6	32.0
Employment Status (for those at most 65)			
Employed full-time	67	63.3	66.9
Not employed full-time	33	36.7	33.1

Note: "High school" corresponded to "Gymnasieexamen", while "Some college" to "Viss universitets-/högskoleutbildning". The option equivalent to a university degree or higher was "Kandidatexamen."

B.1.7 United States

Category	Population (%)	Study Sample (%)	
		Full	*Representative*
Gender			
Male	49	36.4	48.9
Female	51	63.6	51.1
Age			
18–29	21	12.6	19.1
30–39	16	11.3	14.4
40–49	16	14.1	15.2
50–59	17	24.5	19.4
60–69	14	25.8	15.5
≥70	16	11.6	16.5
Income (annual; USD)			
Below 15,000	11	14.3	13.9
15,000–24,999	9	14.3	9.2
25,000–34,999	9	14.7	9.8
35,000–49,999	12	14.6	11.7
50,000–74,999	17	14.8	17.3
75,000–99,999	13	10.4	13.0
100,000–149,999	15	9.9	14.4
150,000–199,999	7	3.9	6.0
200,000 or more	7	3.1	4.7
Ancestry			
White	63	81.9	69.1
African-American	17	8.1	13.6
Hispanic	12	4.7	7.9
Asian	5	3.3	5.9
Other	3	2.0	3.5
Education			
No high school	11	3.9	6.4
High school	29	41.8	30.2
Some college	29	29.7	30.6
Bachelor's degree or higher	31	24.6	32.8
Employment Status (for those at most 65)			
Employed full-time	67	37.3	63.4
Not employed full-time	33	62.7	36.6

B.2 Screenshots

B.2.1 Universalism tasks

Domestic universalism in altruism.

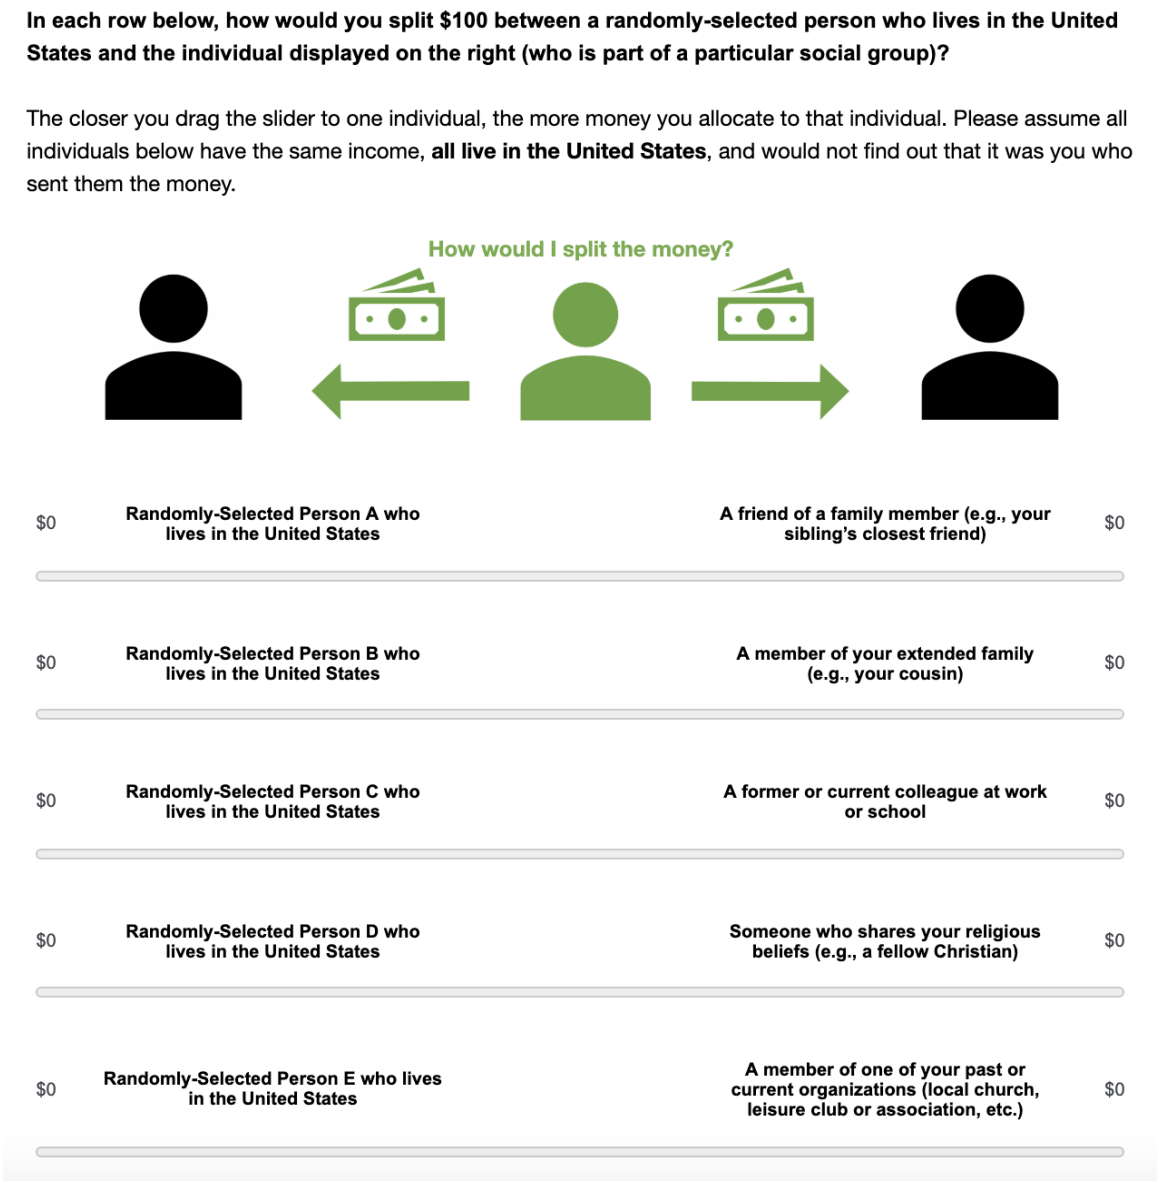

Figure 12: Screenshot of decision screen for money allocation tasks meant to elicit domestic universalism in altruism. Respondents would see two of these screens consecutively, where five of the ten groups would be presented on each screen. Note that across all respondents, the order of the ten social groups was randomized, and whether all social groups appeared on the left or all appeared on the right was also randomized for any given choice domain. The layout for tasks eliciting global universalism in altruism is identical to that of domestic groups.

Foreign universalism in altruism.

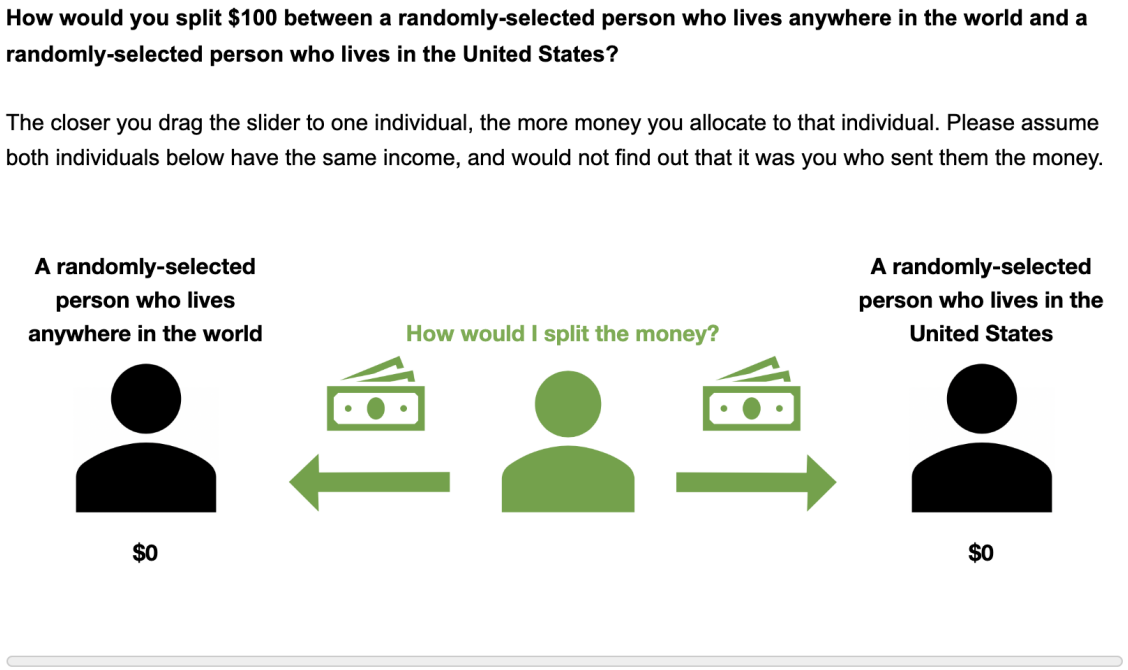

Figure 13: Screenshot of decision screen for money allocation task meant to elicit foreign universalism in altruism. Across respondents, it was randomized whether the domestic social group appeared on the left or on the right. The layout for the task eliciting foreign universalism in *trust* is identical to this layout, with the exception of necessary changes to the instructions and to graphics, as consistent with the layout for trust tasks presented in Figure 14.

Domestic universalism in trust.

Figure 14: Screenshot of decision screen for tasks meant to elicit domestic universalism in trust. Respondents would see two of these screens consecutively, where five of the ten groups would be presented on each screen. Note that across all respondents, the order of the ten social groups was randomized, and whether all social groups appeared on the left or all appeared on the right was also randomized for any given choice domain. The layout for tasks eliciting global universalism in trust is identical to that of domestic groups.

B.2.2 Policy preferences

Desired government spending.

Suppose that you could determine how the United States government spends money on various different categories of the federal budget, such as the military or redistribution.

Specifically, imagine you could decide the **average amount of money that the federal government collects per year from each American to spend on each of the eight categories below**. For the purposes of this question, you should assume that all dollar amounts collected for a category are spent only on this particular category, without any waste.

How much money would you have the federal government collect on average from each American, in order to spend on each of the following eight categories of expenditure in the federal budget?

To provide a reference, it is estimated that altogether, all levels of government in the United States spend a combined average amount of $2,750 per American every year for the purposes of education.

	Amount of money ($) collected on average from each American to spend on category, per year
Police and law enforcement	$
Foreign aid	$
Universal healthcare	$
Environmental protection	$
Measures to ensure no individual is disadvantaged in access to education, the labor force, and marriage	$
Military and counterintelligence	$
Welfare payments	$
Effective border control	$

Figure 15: Screenshot of decision screen eliciting respondents' policy preferences through the means of desired per-capita spending on categories of the federal/national government's budget. Across respondents, the order of categories was randomized.

Support for Specific Policy Implementations.

You just indicated your preferences over several broad categories of expenditure. We will now ask you to indicate how much money you would like to collect and spend on **specific projects or policy proposals**. After all, even within broad categories such as military or redistribution, **you may like some policies and projects more than others. We are now interested in which specific projects or policies you favor.**

Please consider the policy proposals presented in the table below. Note that the table will expand with additional rows as you fill in your answers and until you see **eight** rows in total. Again assume that all money collected for the purposes of a policy are spent only on implementing that particular policy, without waste.

How much money would you have the federal government collect on average from each American, in order to spend on each of the **specific projects or policies** presented?

Policy	Amount of money ($) collected on average from each American to spend on policy, per year	Policy	Amount of money ($) collected on average from each American to spend on policy, per year
Sensitivity training for the police to **ensure justice and equal treatment of all**	$	Increasing the capabilities of the police to **prevent and prosecute criminal or suspicious behavior**	$

Figure 16: Screenshot of decision screen eliciting respondents' preferences towards particular policy implementations of national government expenditure. Across respondents, the order of categories was randomized, and it was randomized whether all more universalist policies appeared on the left or on the right. Additional policies continued to fill the screen as the respondent filled in desired spending levels for each category of policies.

B.3 Histograms of Composite Universalism

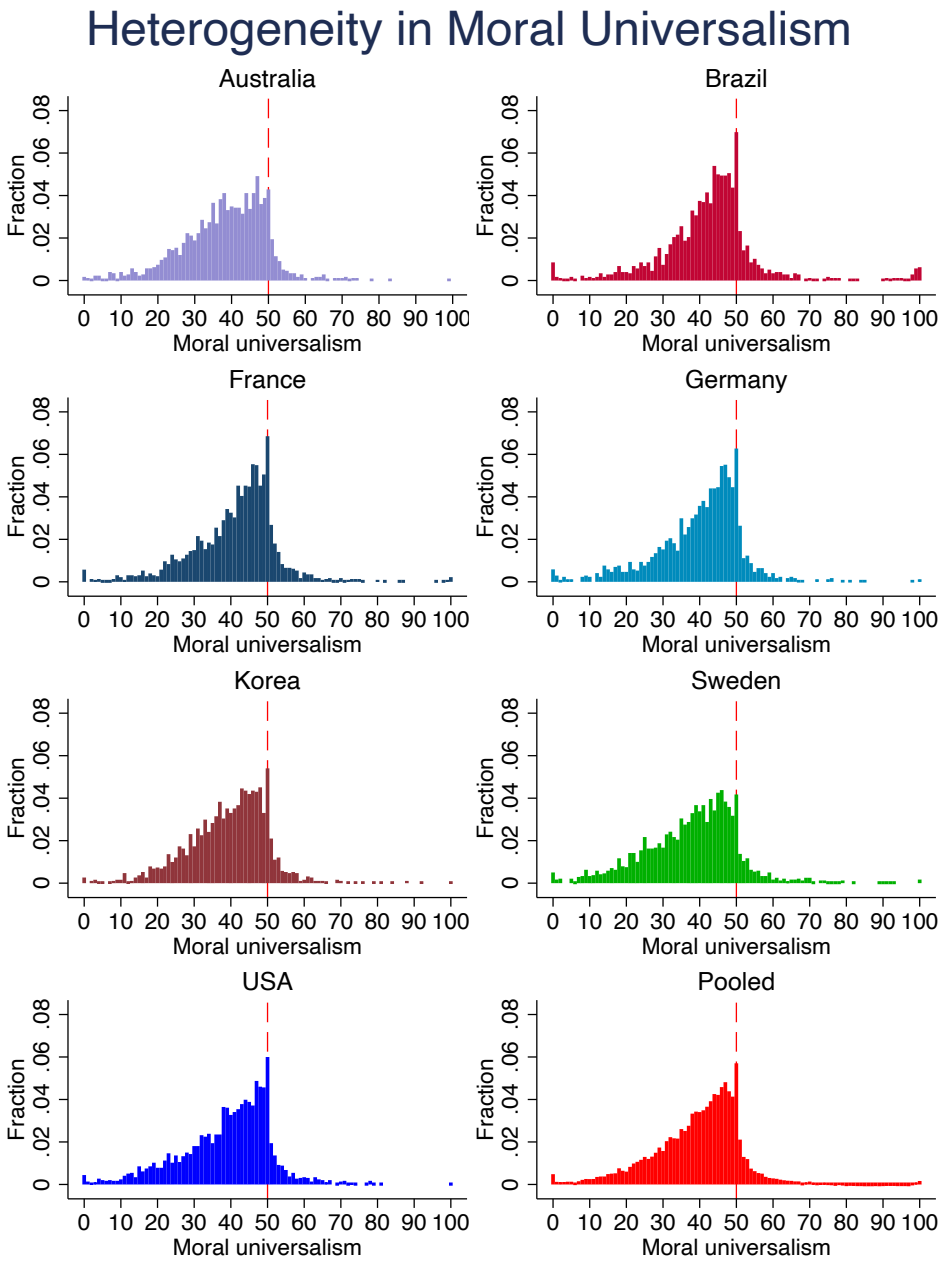

Figure 17: This figure plots a comparative set of distributions of our composite measure of moral universalism. All individual plots are scaled to the same x-axis and y-axis.

B.4 Selection of demographics through LASSO

Table 12: LASSO: Universalism on demographics, Western countries

	Dependent variable:
	Composite Universalism (Z-scores)
	(1)
Age (z-scores, within country)	-0.121***
	(0.010)
Male (z-scores, within country)	-0.039***
	(0.011)
Religion Index (z-scores, within country)	-0.080***
	(0.011)
Wealth Index (z-scores, within country)	-0.066***
	(0.012)
Income Index (z-scores, within country)	-0.042***
	(0.011)
Age (z-scores) × Male (z-scores) × Income Index (z-scores)	0.032**
	(0.010)
College (z-scores, within country)	0.030**
	(0.010)
Observations	11063

Notes. Coefficients presented are post-selection, OLS coefficients. Independent variables were selected by running a LASSO regression of standardized composite universalism on all possible interactions of: age, male, religiosity, wealth, income, college, urbanicity, first-generation immigrant, second-generation immigrant, and married. Only the top 7 of 56 selected features (in descending order of absolute value of coefficient in the post-selection, OLS regression) are displayed. See Appendix E for details on the construction of the demographic variables. All z-scores are computed separately within each country.

B.5 Correlations between Universalism and Other Variables

Table 13: Individual-level correlates of universalism: Western countries

Covariate	Universalism w.r.t...	Correlation coefficient
Age	Altruism	-0.158
Age	Composite	-0.161
Age	Trust	-0.116
Arsinh[Total Desired Gov't Expenditure] (Z-score)	Altruism	-0.00869
Arsinh[Total Desired Gov't Expenditure] (Z-score)	Composite	-0.0250
Arsinh[Total Desired Gov't Expenditure] (Z-score)	Trust	-0.0373
Belief in efficiency of gov't (0-10)	Altruism	0.0910
Belief in efficiency of gov't (0-10)	Composite	0.0633
Belief in efficiency of gov't (0-10)	Trust	0.0106
College	Altruism	0.0266
College	Composite	0.00980
College	Trust	-0.0136
Dictator game: allocation to stranger (0-100)	Altruism	0.287
Dictator game: allocation to stranger (0-100)	Composite	0.282
Dictator game: allocation to stranger (0-100)	Trust	0.191
Generalized trust in stranger (0-100)	Altruism	0.111
Generalized trust in stranger (0-100)	Composite	0.109
Generalized trust in stranger (0-100)	Trust	0.0729
Income Index (Z-score)	Altruism	-0.0610
Income Index (Z-score)	Composite	-0.0688
Income Index (Z-score)	Trust	-0.0574
Male	Altruism	-0.0719
Male	Composite	-0.0735
Male	Trust	-0.0532
Preference for efficiency over inequity (0-50)	Altruism	-0.0309
Preference for efficiency over inequity (0-50)	Composite	-0.0301
Preference for efficiency over inequity (0-50)	Trust	-0.0200
Religion Index (Z-score)	Altruism	-0.0605
Religion Index (Z-score)	Composite	-0.0957
Religion Index (Z-score)	Trust	-0.109
Right-Left Scale (Right: 0; Left:10)	Altruism	0.212
Right-Left Scale (Right: 0; Left:10)	Composite	0.248
Right-Left Scale (Right: 0; Left:10)	Trust	0.216
Summary statistic of policy views (Z-score)	Altruism	0.241
Summary statistic of policy views (Z-score)	Composite	0.250
Summary statistic of policy views (Z-score)	Trust	0.187
Urbanicity (0-9)	Altruism	0.0328
Urbanicity (0-9)	Composite	0.0164
Urbanicity (0-9)	Trust	-0.00849
Wealth Index (Z-score)	Altruism	-0.106
Wealth Index (Z-score)	Composite	-0.124
Wealth Index (Z-score)	Trust	-0.109

Notes. See Appendix E for details on the construction of the demographic variables. All z-scores are computed separately within each country.

B.6 Histograms of Desired Expenditure Amounts

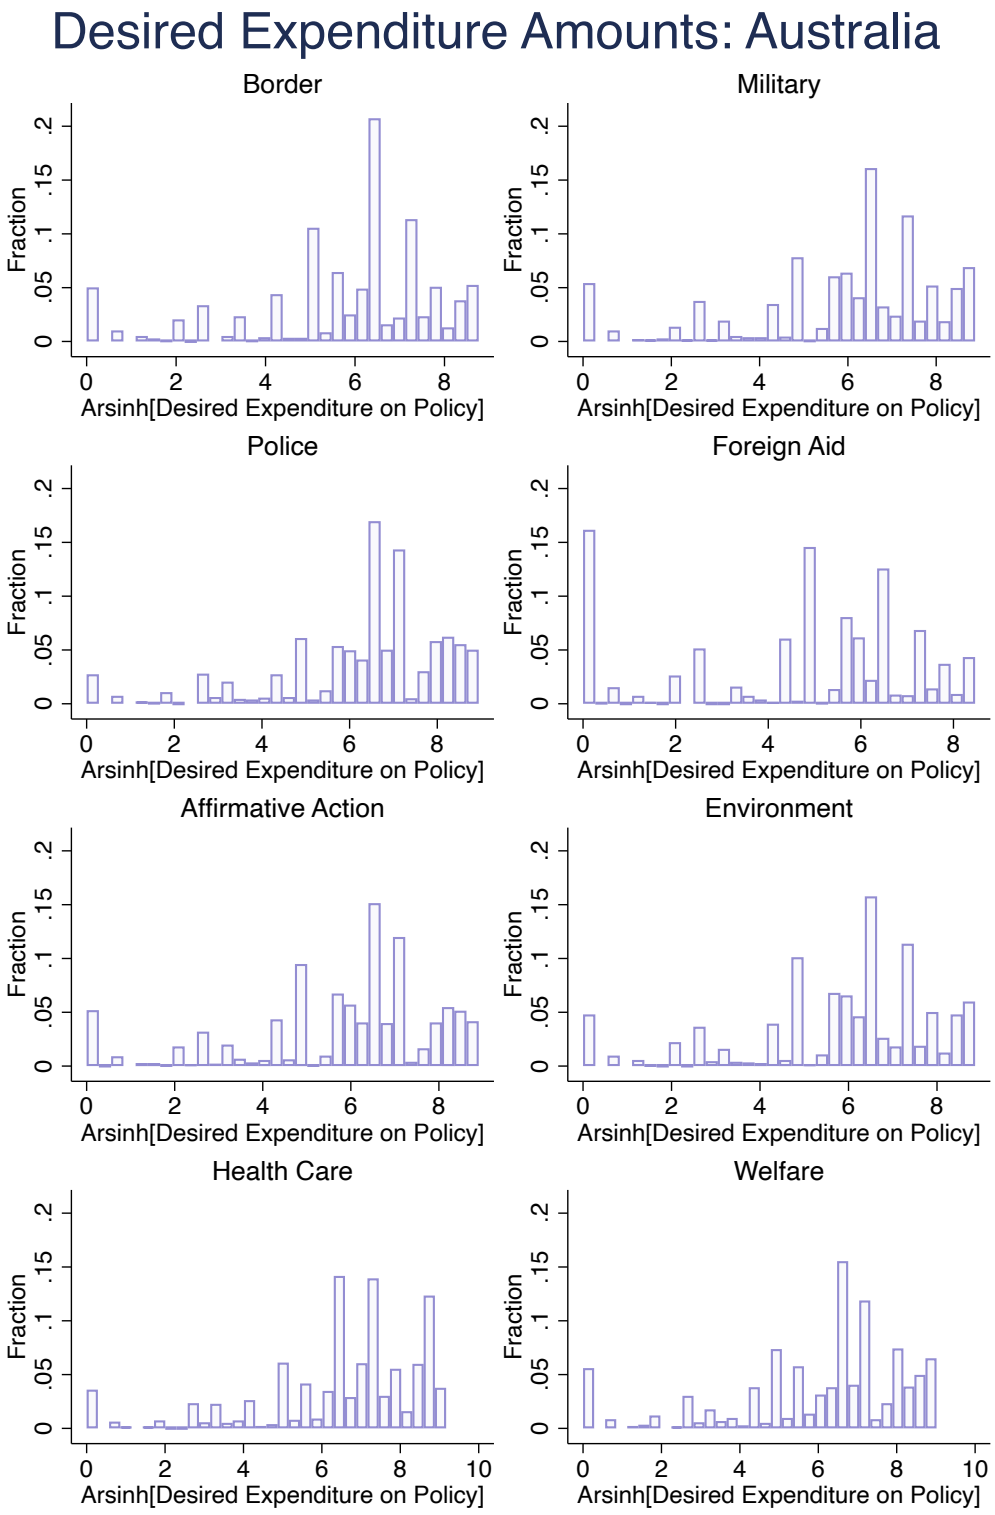

Figure 18: This figure plots histograms of desired expenditure amounts for each policy category by respondents from Australia.

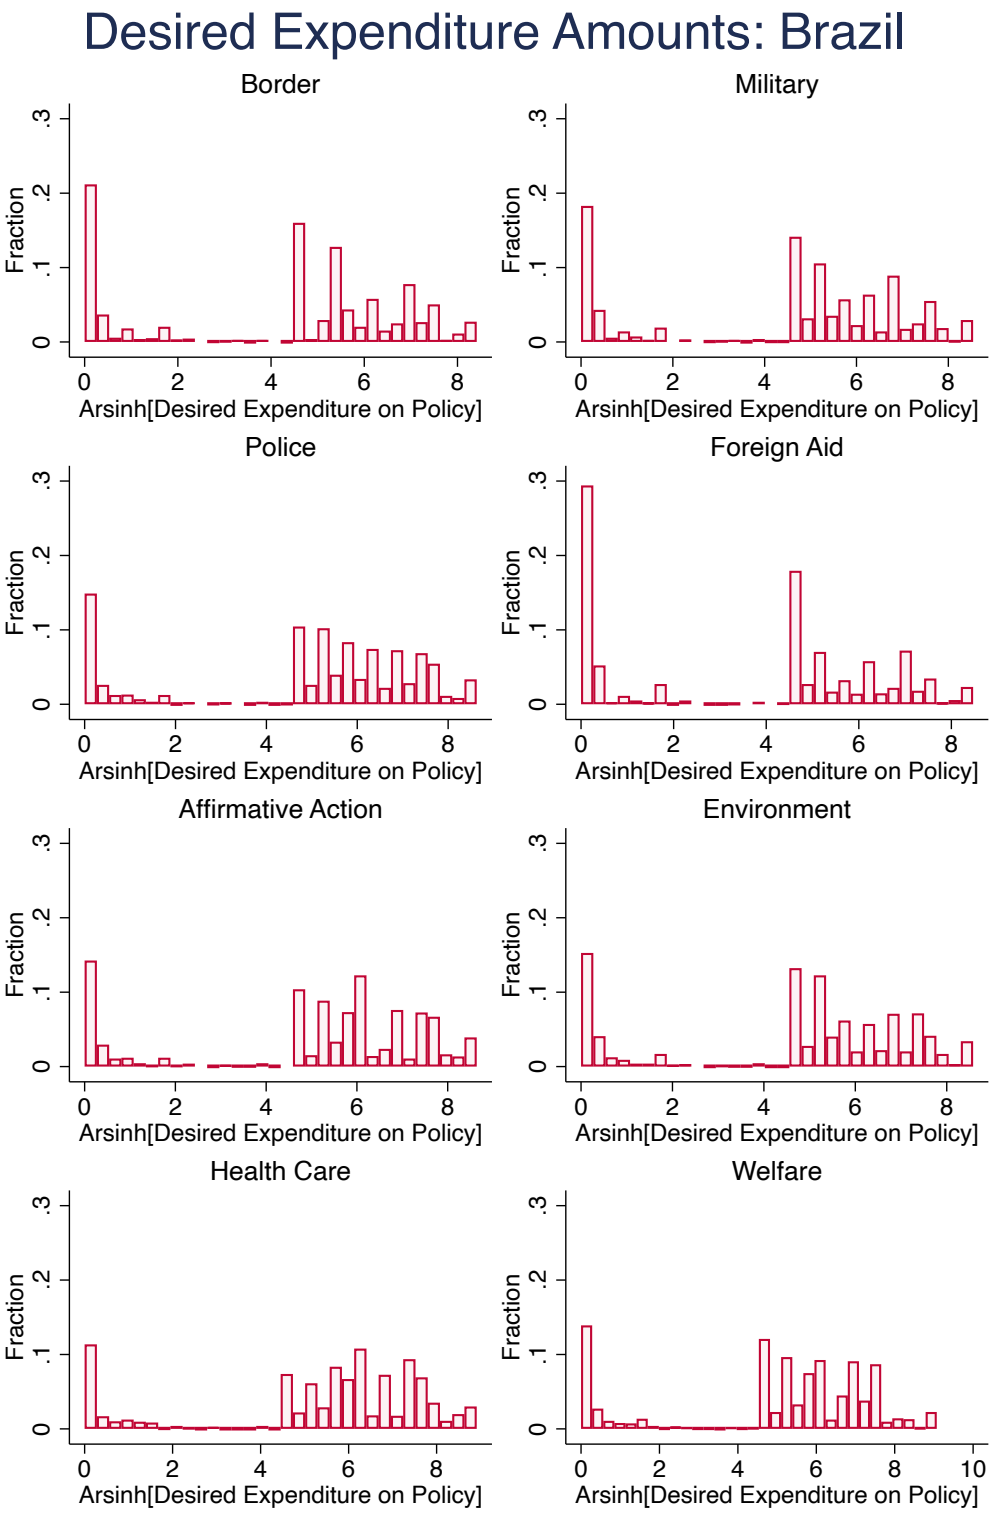

Figure 19: This figure plots histograms of desired expenditure amounts for each policy category by respondents from Brazil.

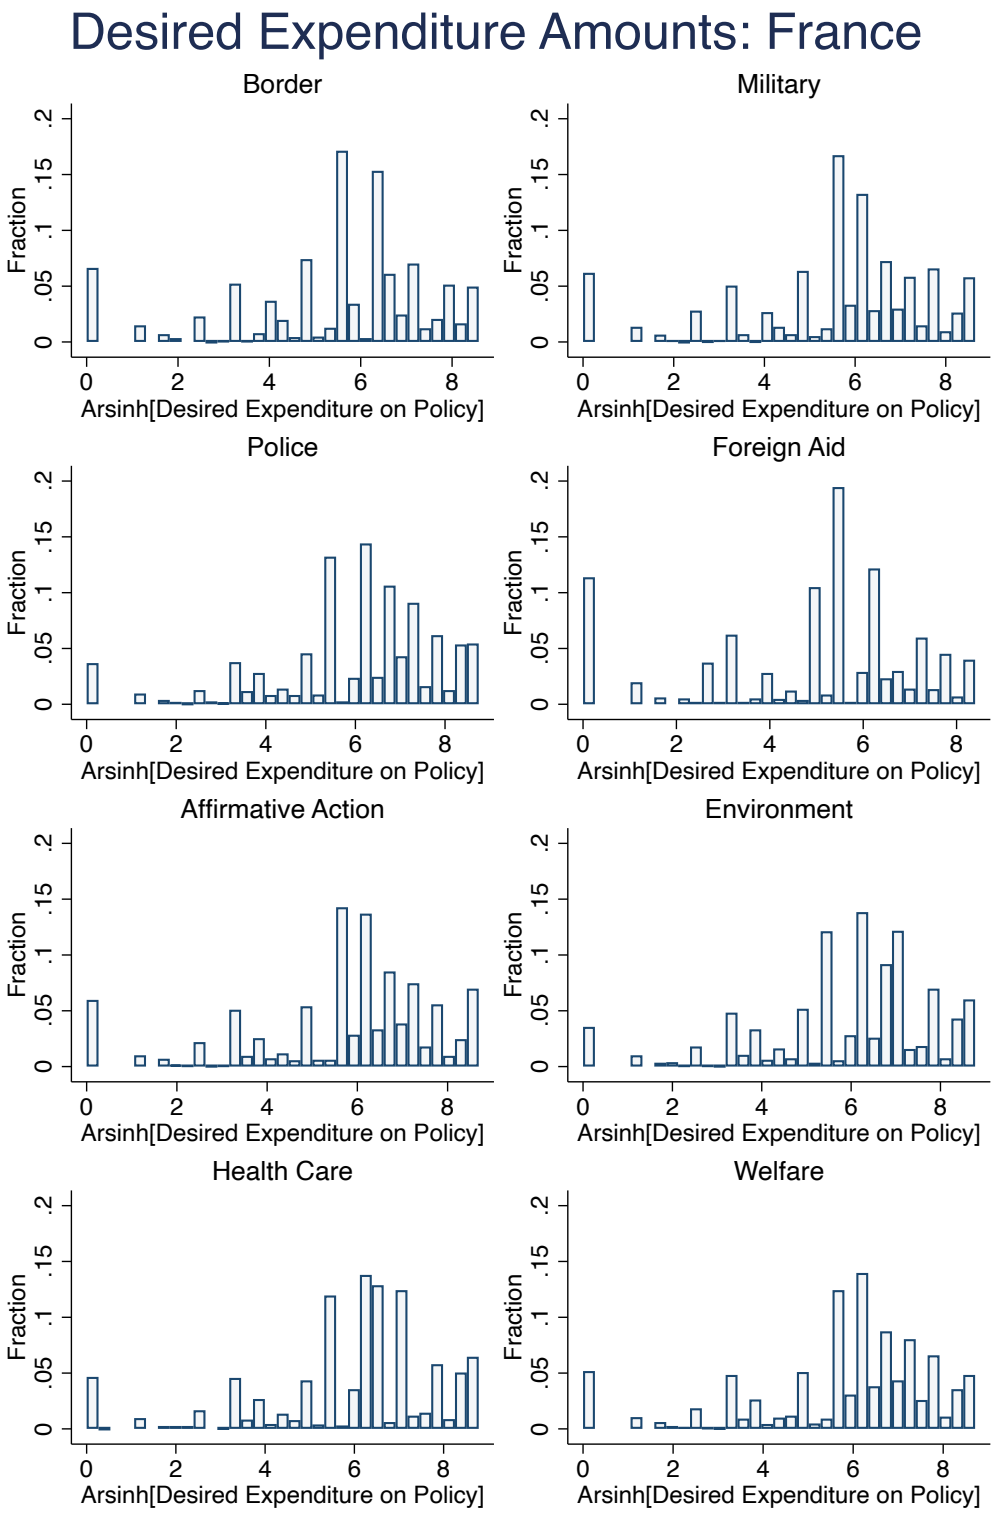

Figure 20: This figure plots histograms of desired expenditure amounts for each policy category by respondents from France.

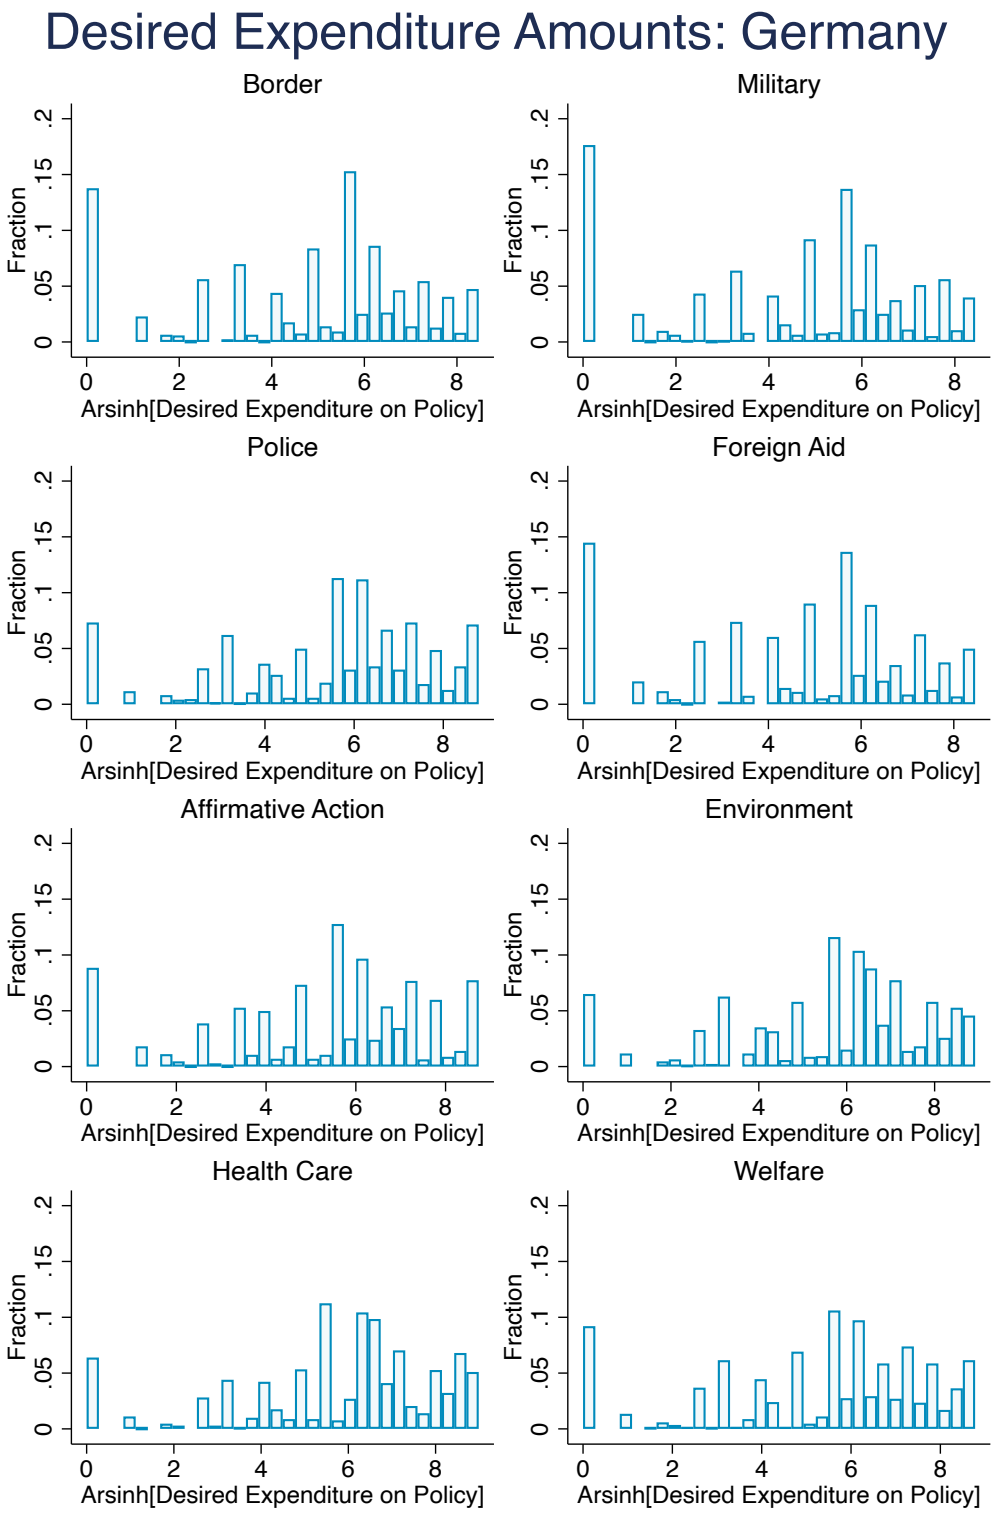

Figure 21: This figure plots histograms of desired expenditure amounts for each policy category by respondents from Germany.

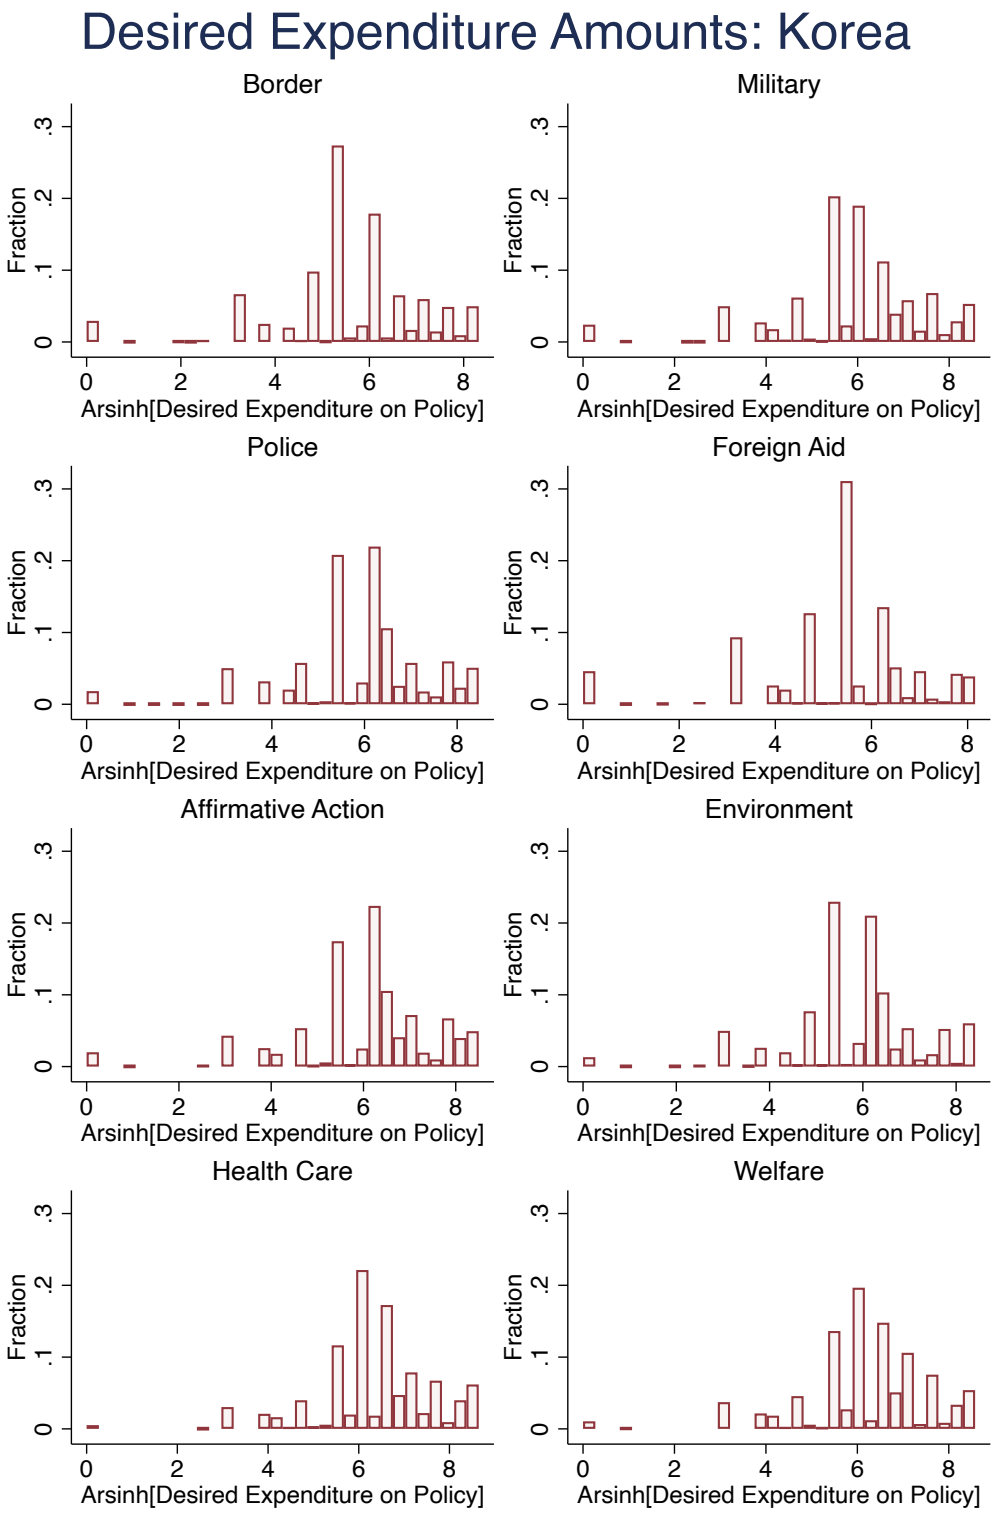

Figure 22: This figure plots histograms of desired expenditure amounts for each policy category by respondents from Korea.

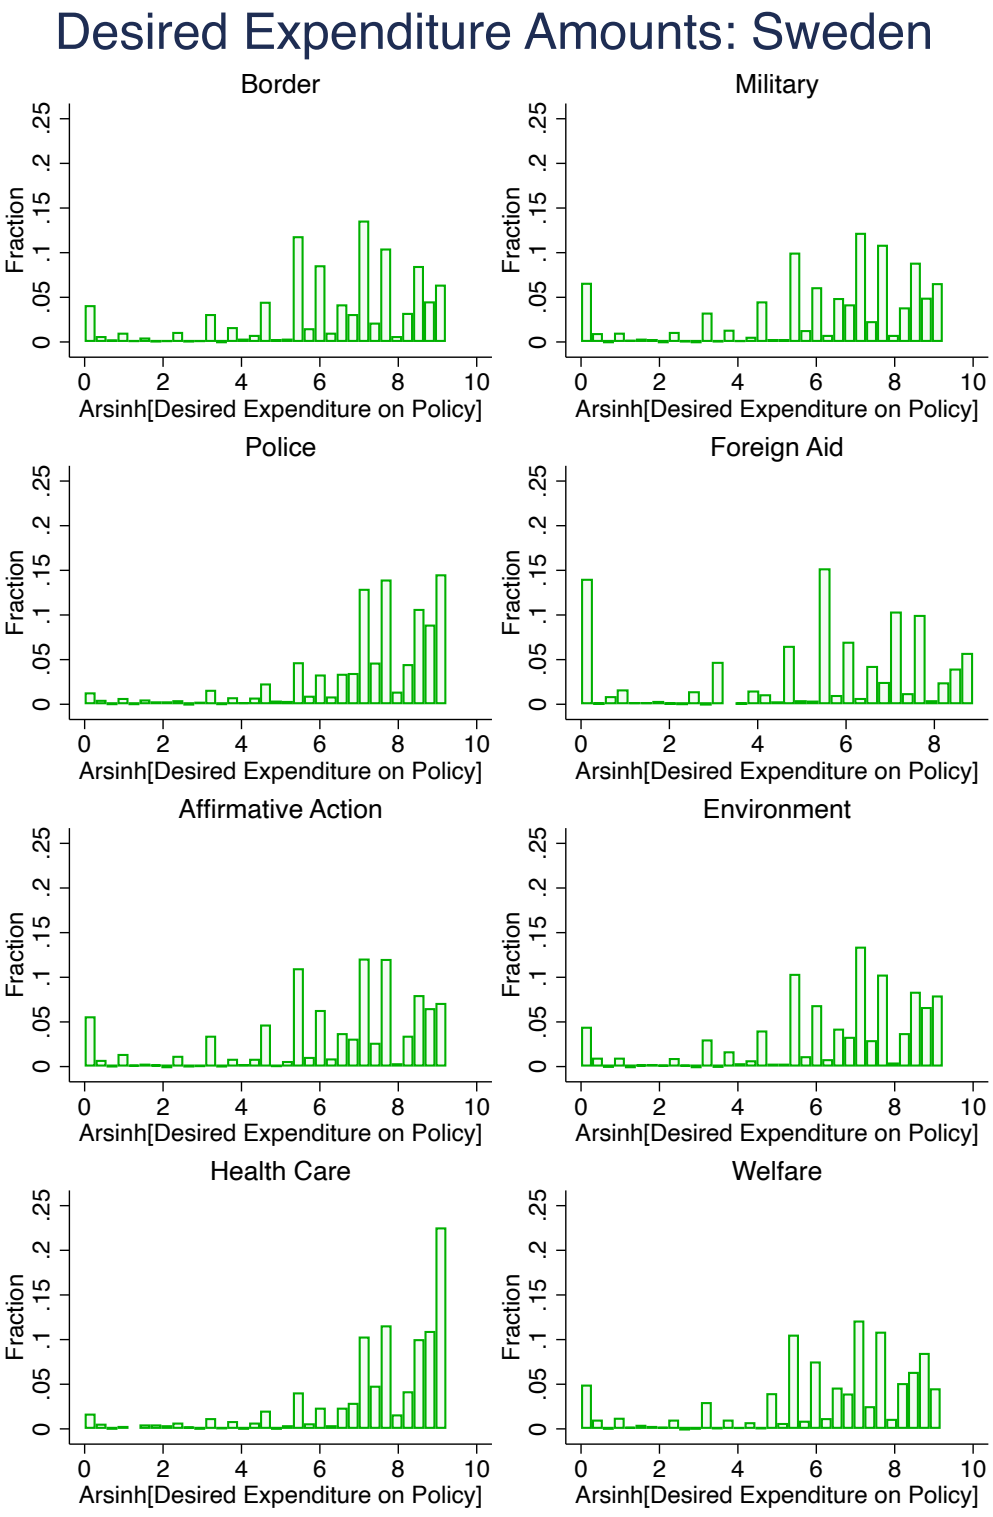

Figure 23: This figure plots histograms of desired expenditure amounts for each policy category by respondents from Sweden.

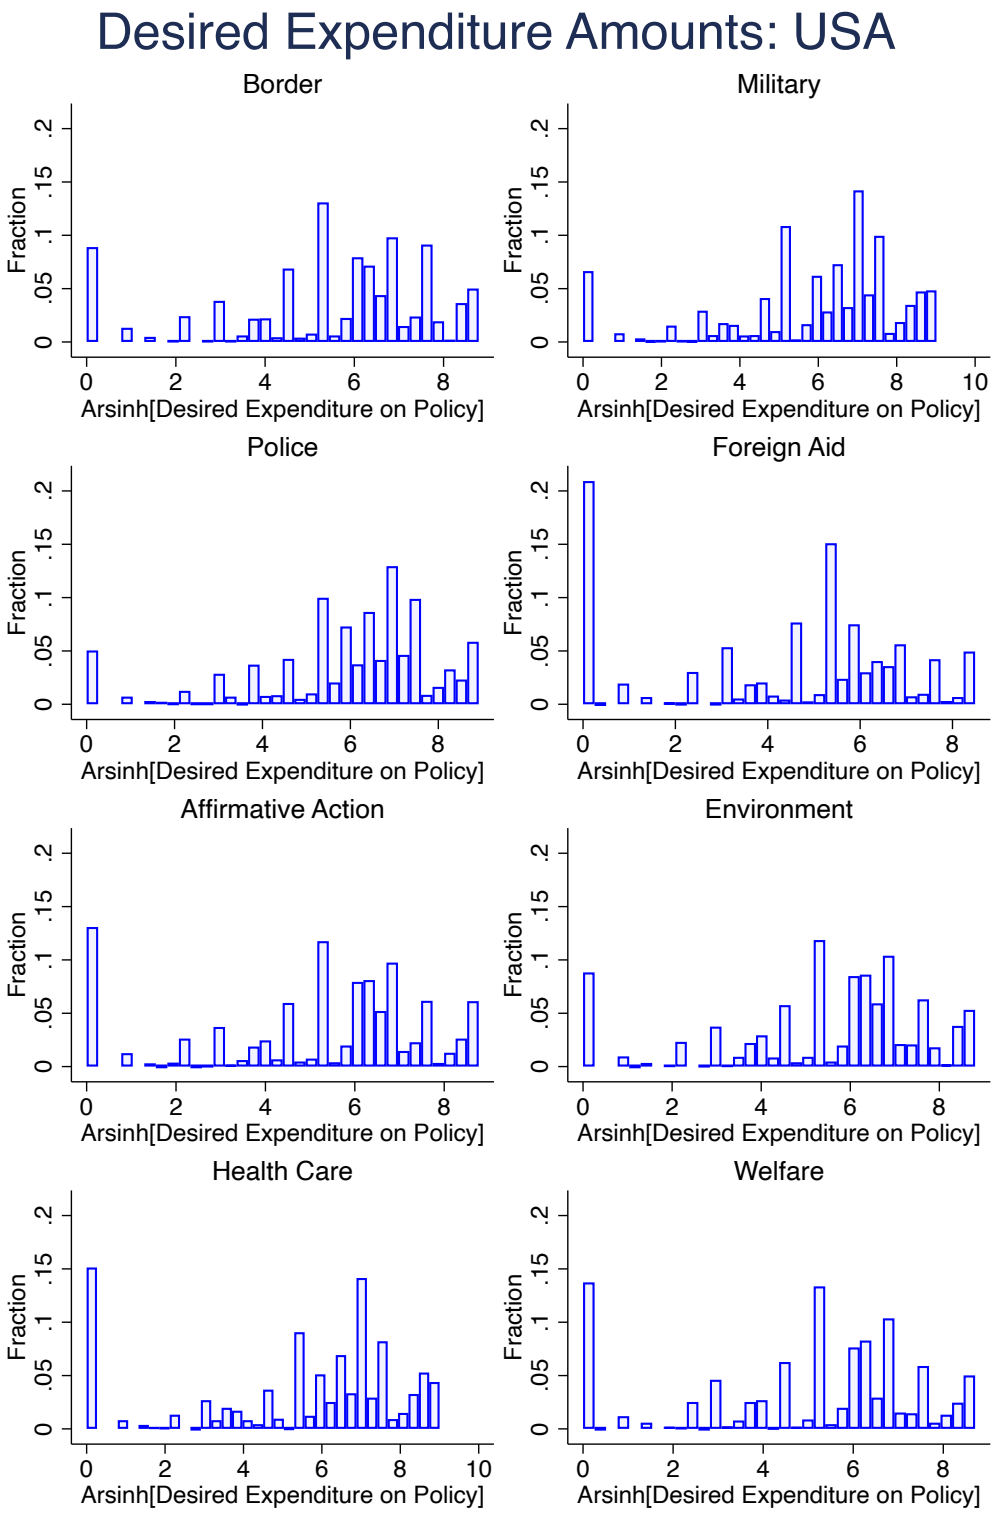

Figure 24: This figure plots histograms of desired expenditure amounts for each policy category by respondents from USA.

B.7 Order Effects

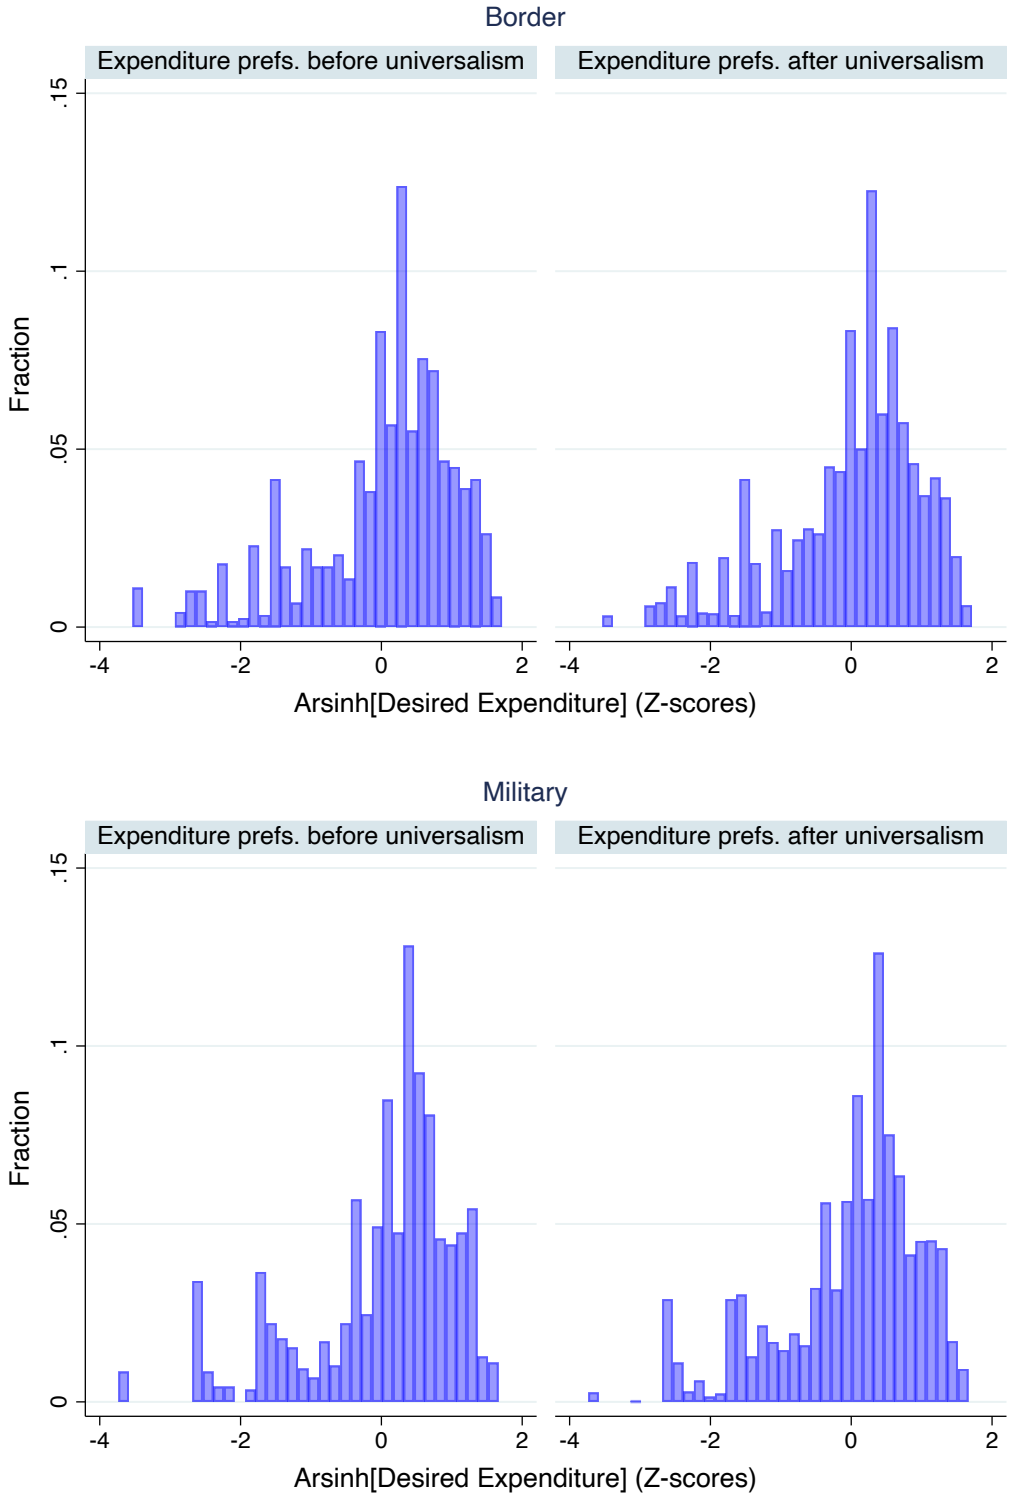

Figure 25: This figure plots histograms of standardized arsinh desired expenditure amounts for border control and the military pooled across all countries, conditional on whether the respondent was asked about their policy preferences before elicitations of univeralism or not.

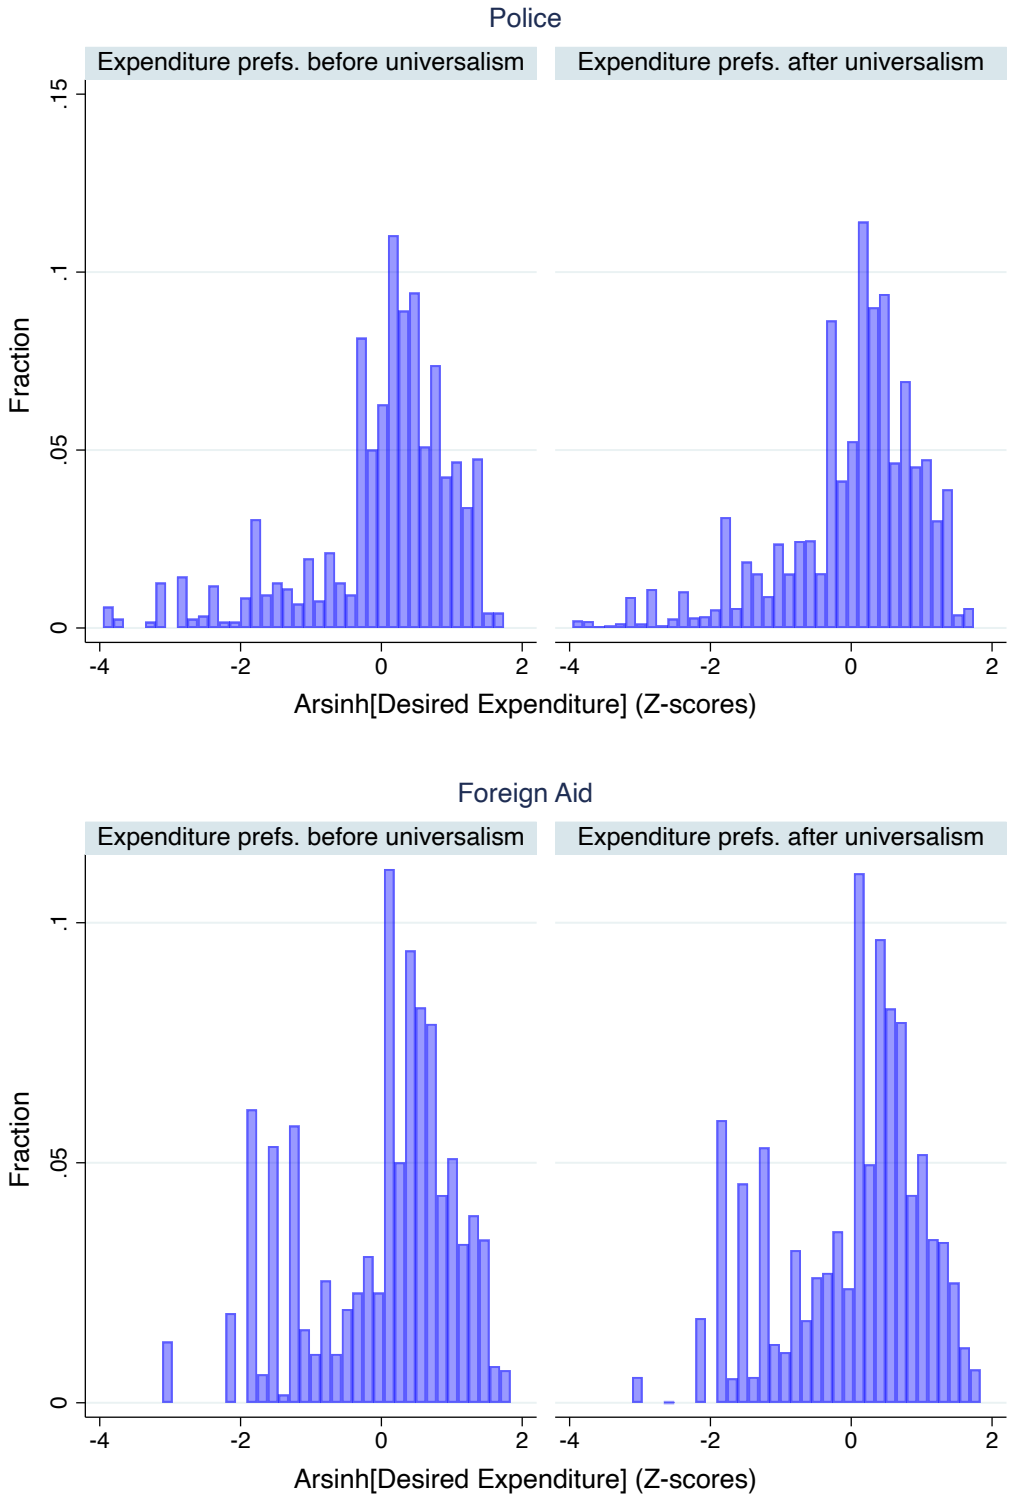

Figure 26: This figure plots histograms of standardized arsinh desired expenditure amounts for police and foreign aid pooled across all countries, conditional on whether the respondent was asked about their policy preferences before elicitations of univeralism or not.

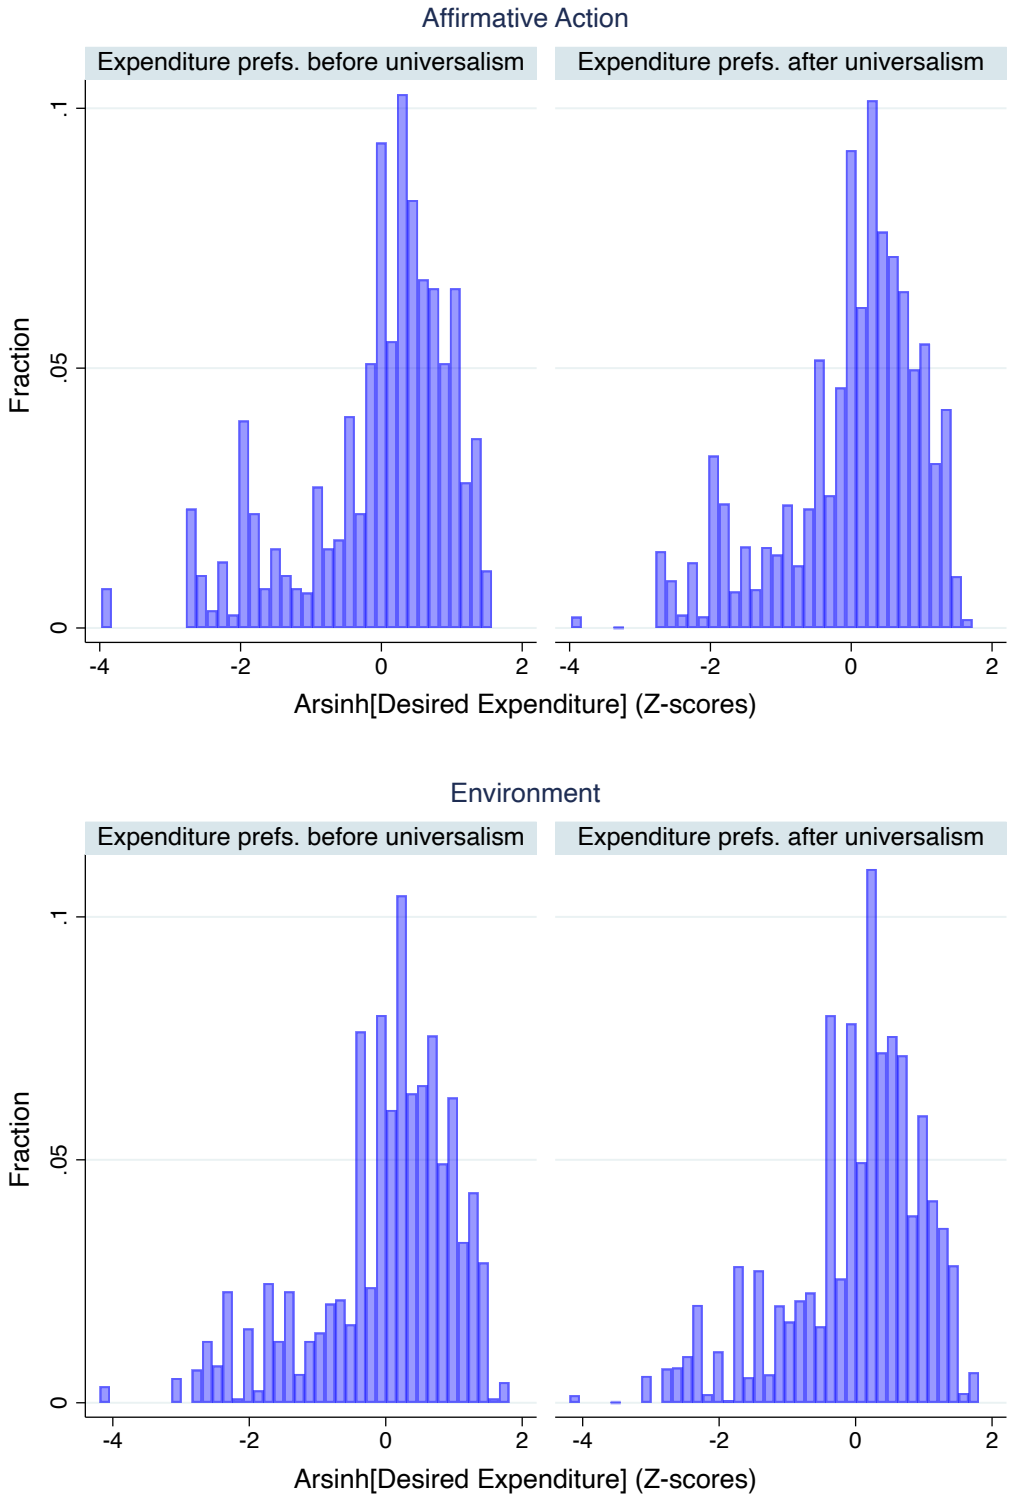

Figure 27: This figure plots histograms of standardized arsinh desired expenditure amounts for affirmative action and the environment pooled across all countries, conditional on whether the respondent was asked about their policy preferences before elicitations of univeralism or not.

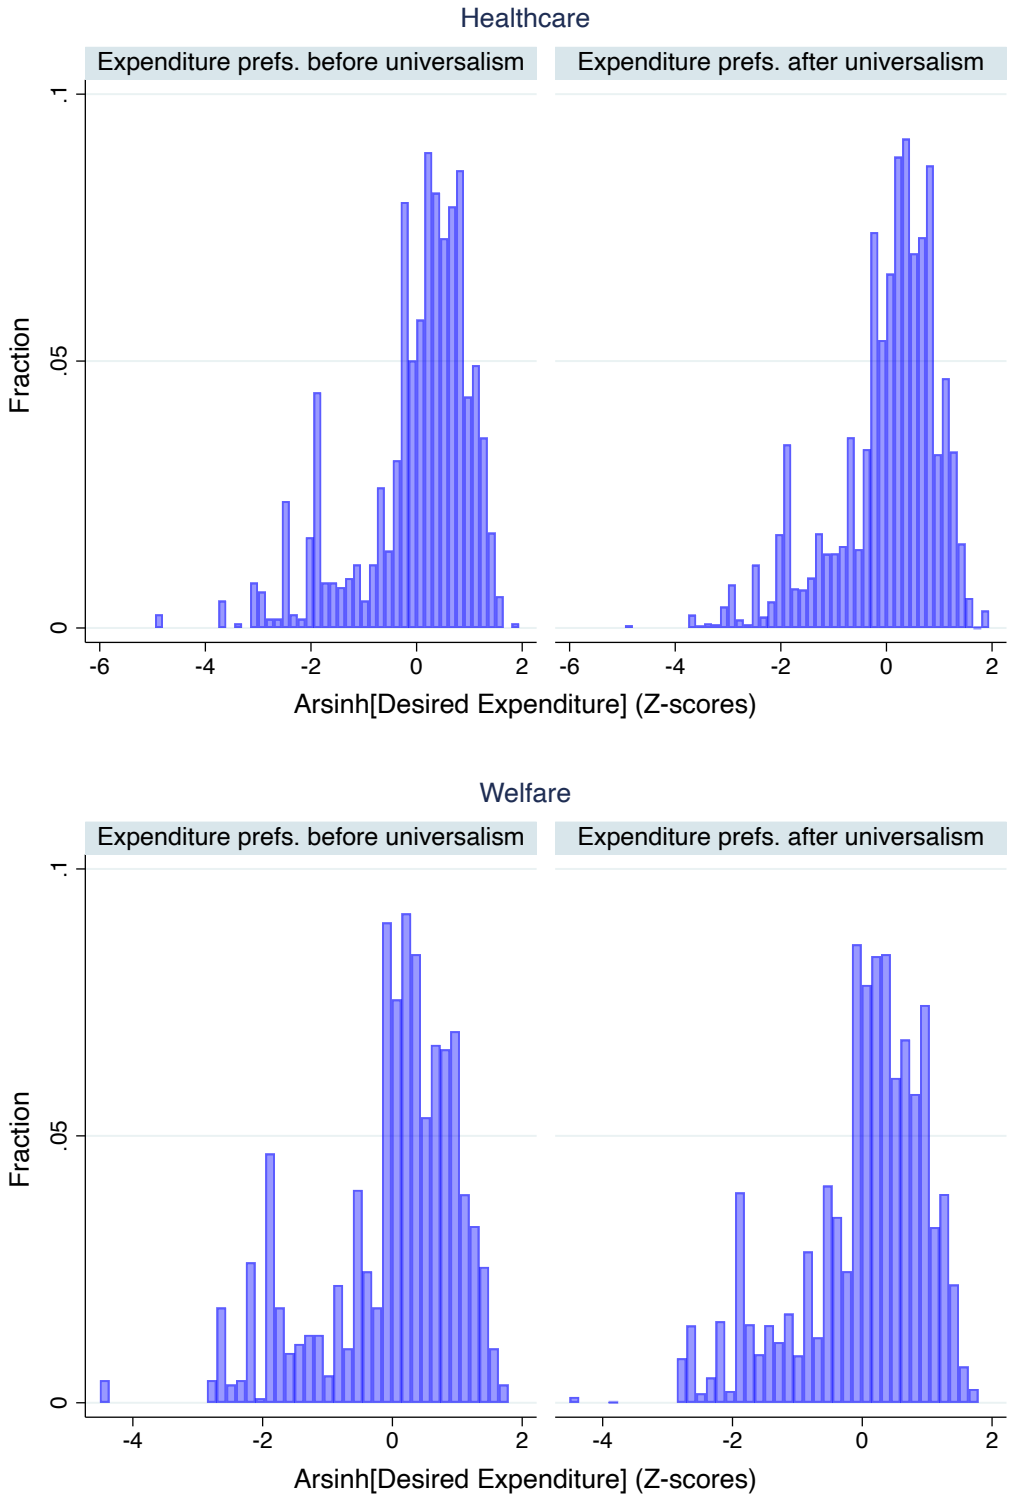

Figure 28: This figure plots histograms of standardized arsinh desired expenditure amounts for healthcare and welfare pooled across all countries, conditional on whether the respondent was asked about their policy preferences before elicitations of univeralism or not.

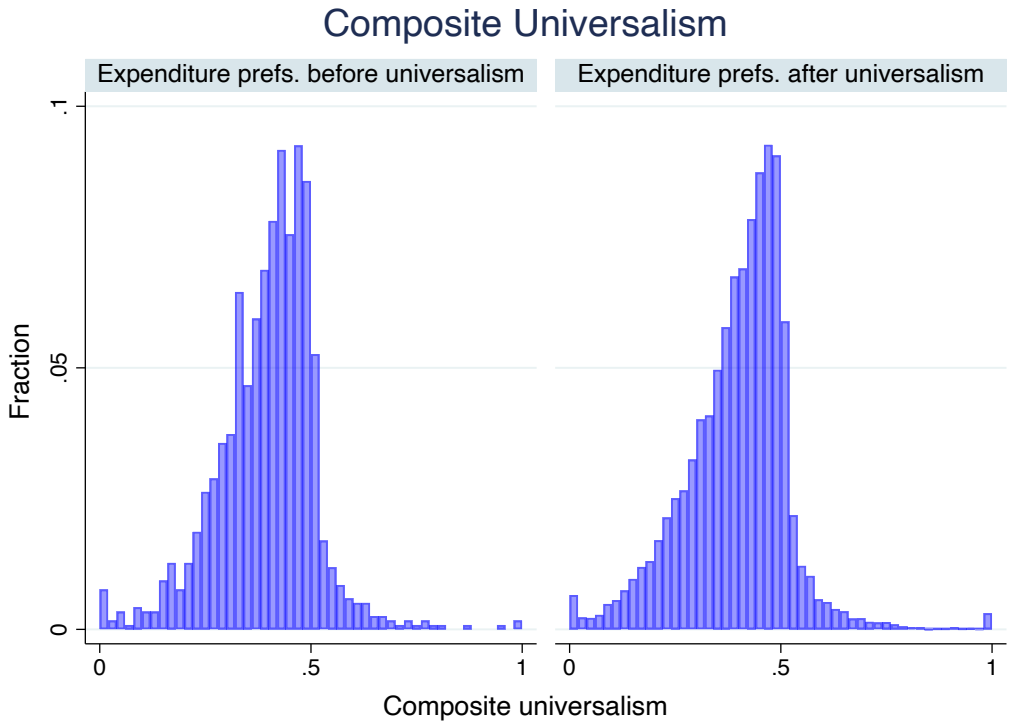

Figure 29: This figure plots histograms of composite universalism pooled across all countries, conditional on whether the respondent was asked about their policy preferences before elicitations of univeralism or not.

C Analysis of Ideological Clusters

C.1 Separate Allocation Decisions

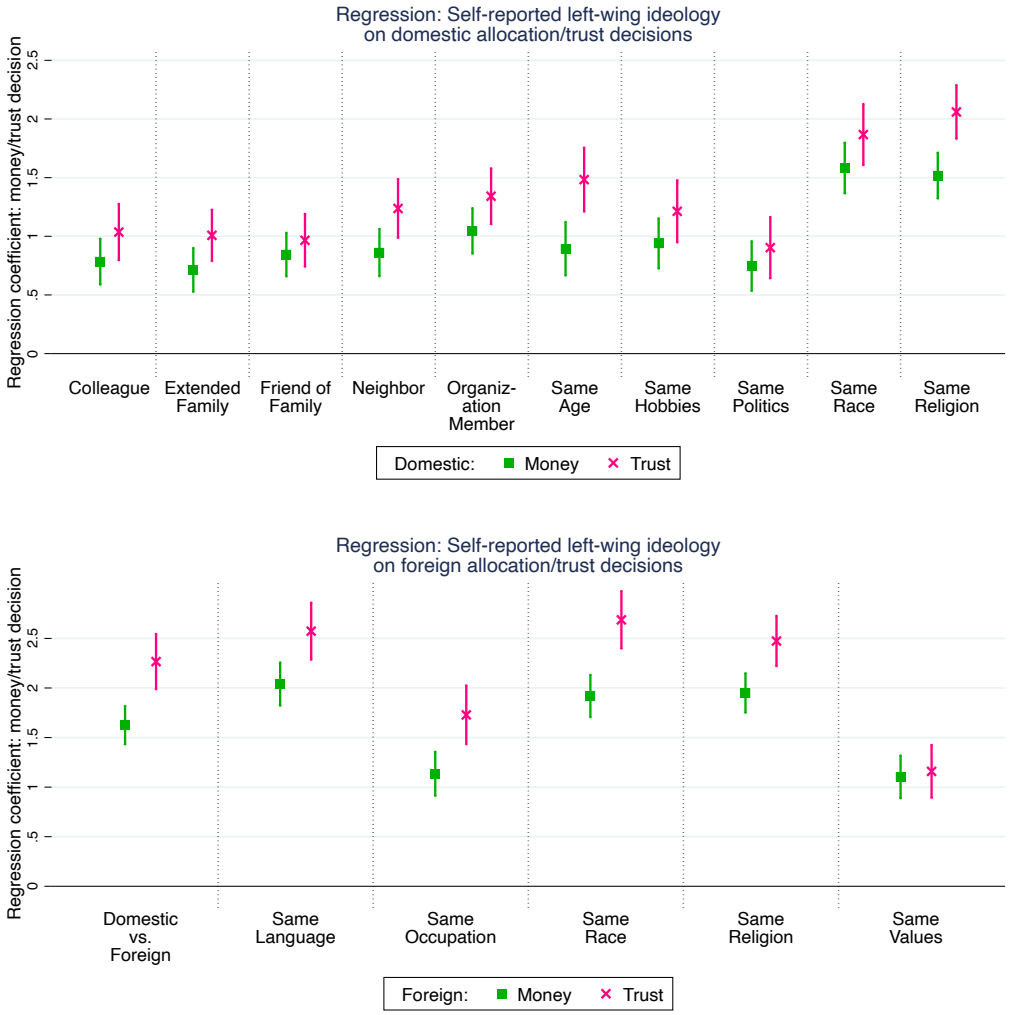

Figure 30: This figure plots OLS coefficients from regressions of respondents' left-right self-assessment on each separate allocation decision at a time, for all Western countries and including country-fixed effects. A positive regression coefficient indicates that a higher allocation to the more distant individual (i.e. a more "universalist" allocation) is positively correlated with "left-leaning" policy preferences. Error bars indicate 95% confidence intervals using robust standard errors.

C.2 Expenditure Levels

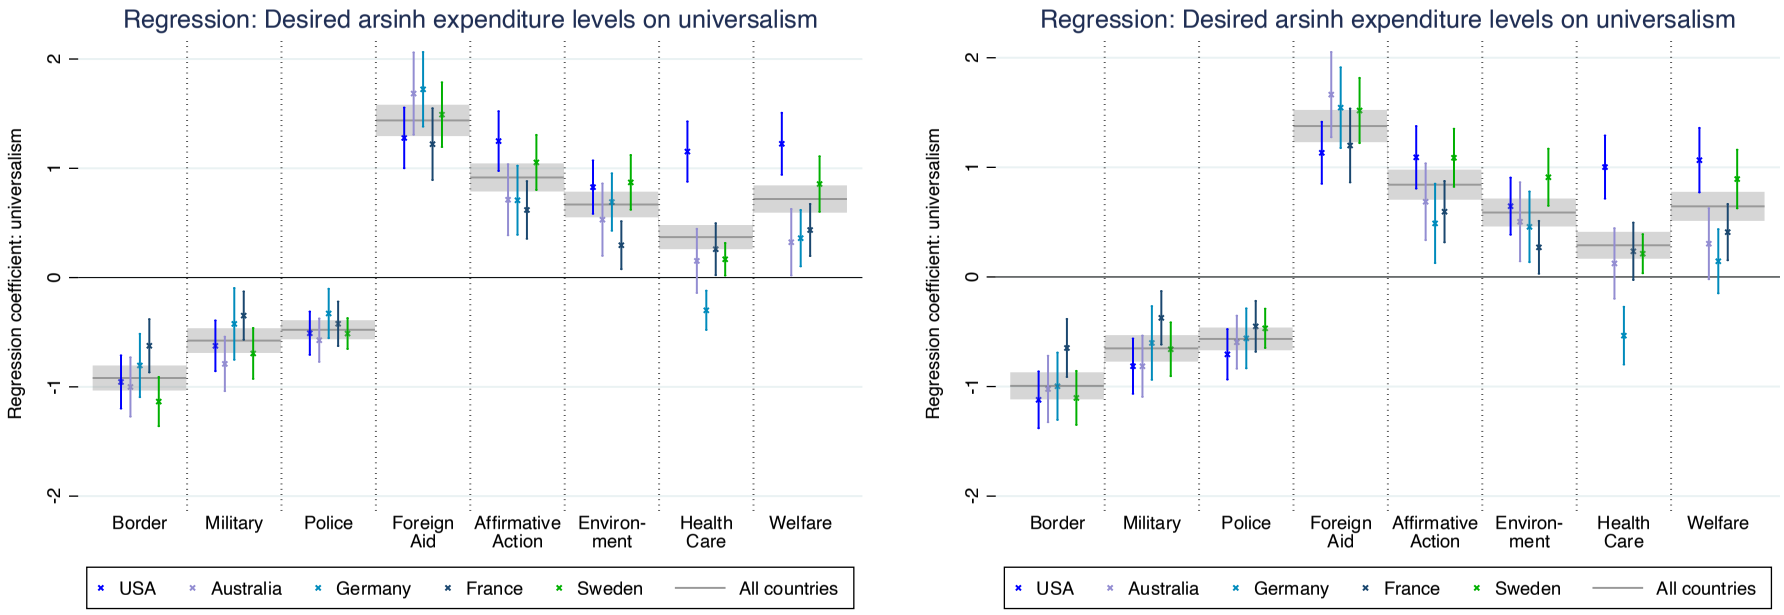

Figure 31: The figure plots the OLS coefficients of regressions of desired arsinh expenditure levels for each policy domain on composite universalism, controlling linearly for arsinh *total* desired expenditure (in the left panel), and controlling for *quantiles* of arsinh total desired expenditure (in the right panel). Universalism is in [0,1] and the dependent variables are standardized into z-scores within each country. Error bars indicate 95% confidence intervals using robust standard errors. The "All countries" specification includes country fixed effects.

C.3 Benchmarking Explanatory Power

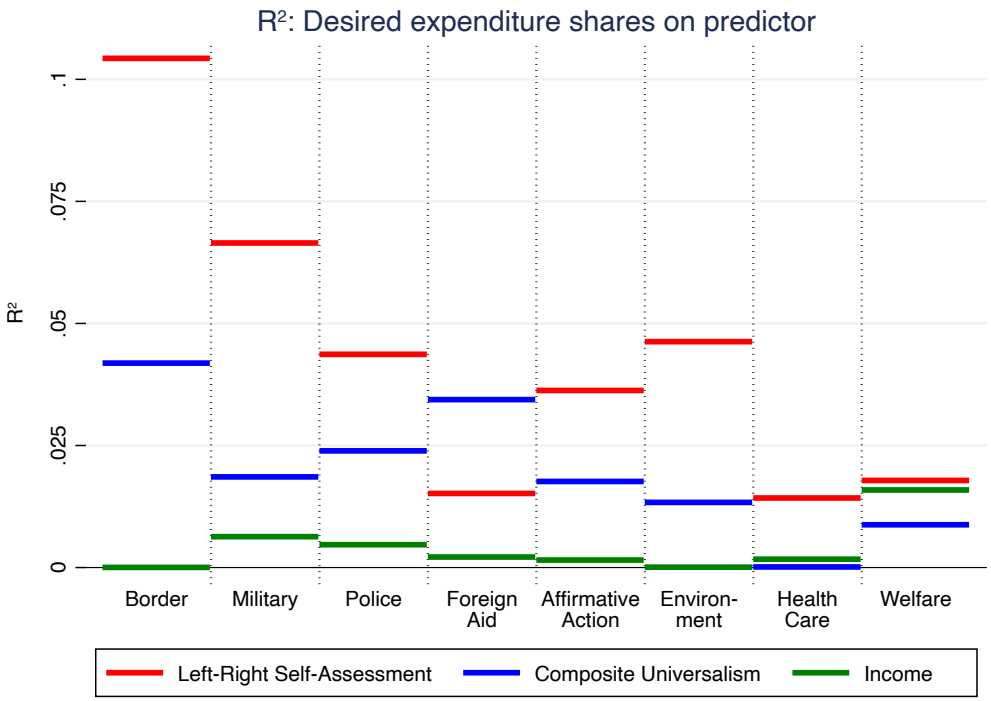

Figure 32: This figure plots the R^2 from separate regressions of desired expenditure shares for each policy domain on respondents' left-right self-assessment, composite universalism, and income, controlling for country fixed effects. All variables are standardized into z-scores within each country.

C.4 Universalism in Altruism and Universalism in Trust

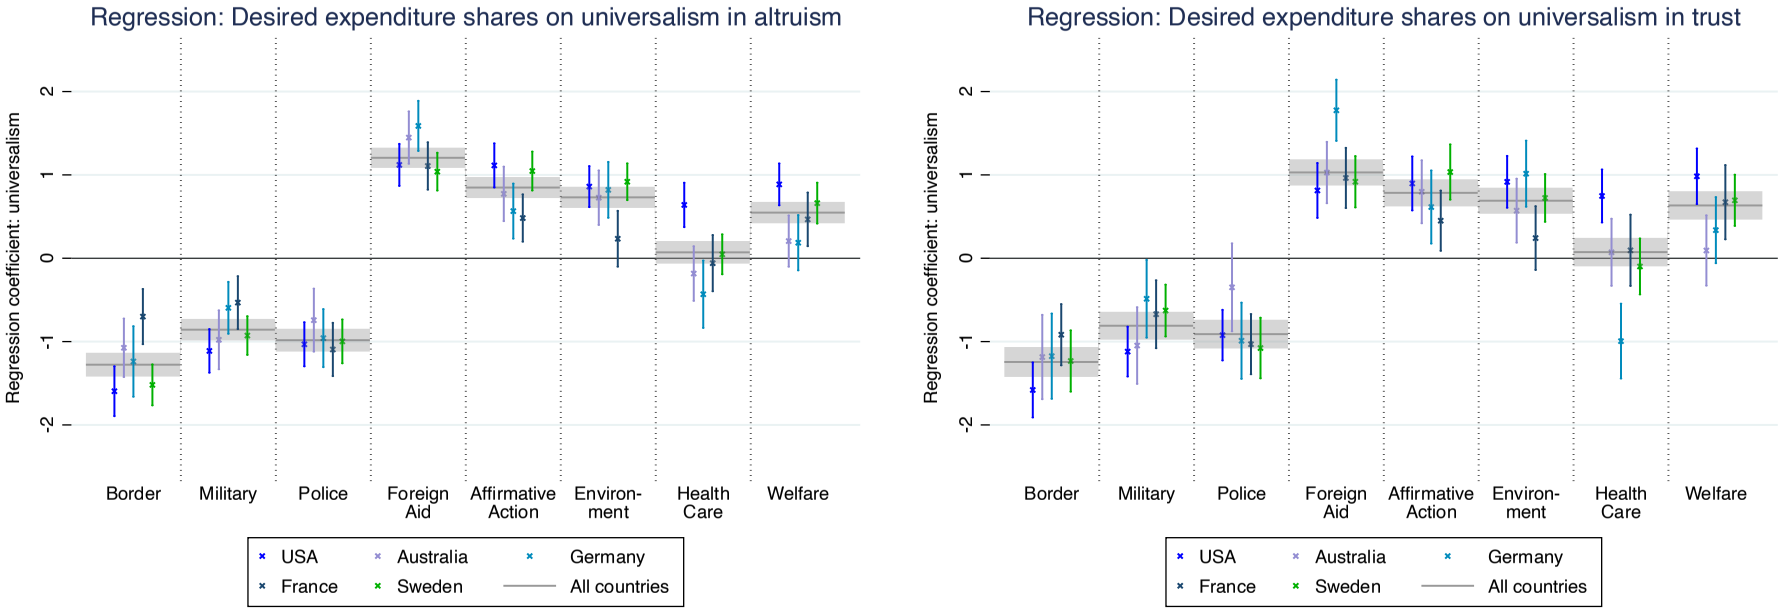

Figure 33: This figure plots the OLS regression coefficient of univariate regressions of desired expenditure shares on the separate universalism measures (universalism in altruism and universalism in trust). Universalism is in [0,1] and the dependent variable is standardized into z-scores. Error bars indicate 95% confidence intervals using robust standard errors. The "All countries" specification includes country fixed effects.

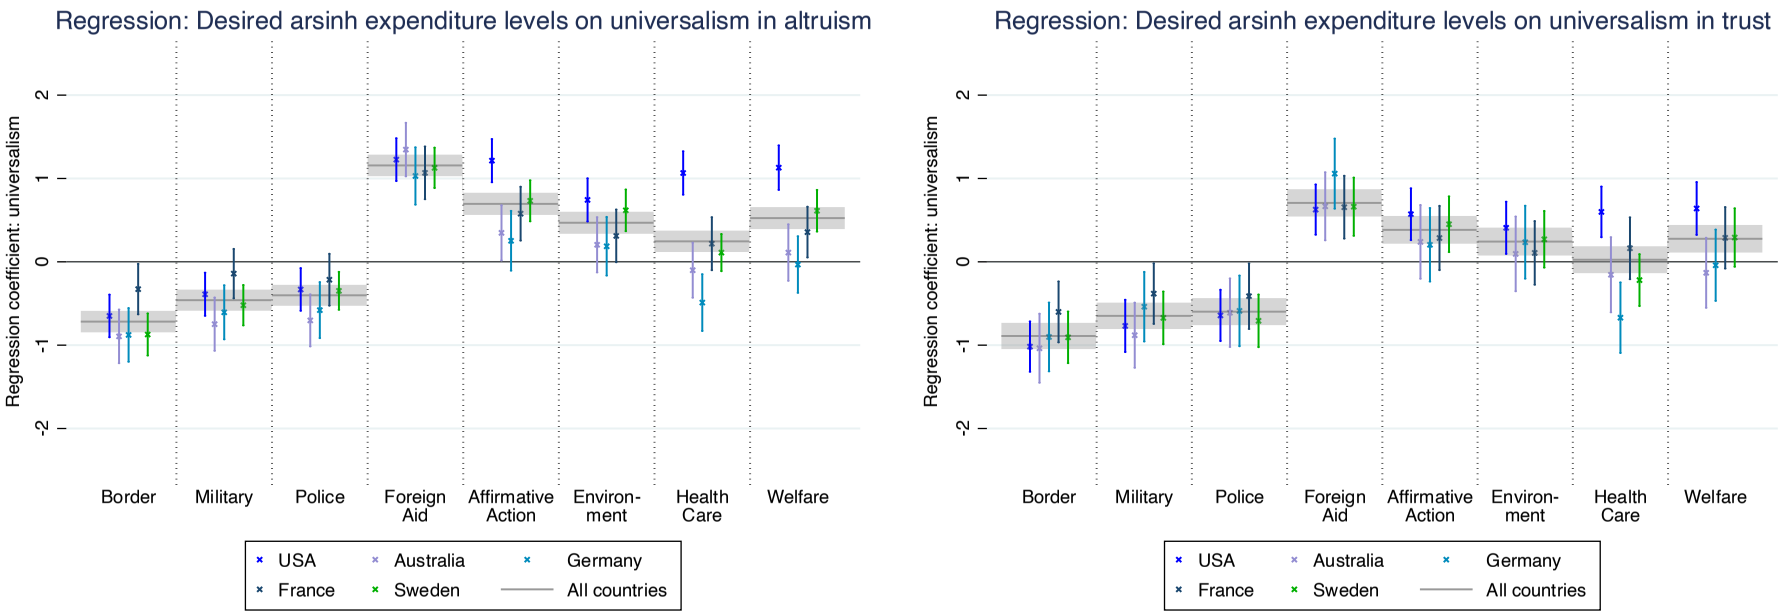

Figure 34: This figure plots the OLS regression coefficient of univariate regressions of desired arsinh expenditure levels on the separate universalism measures (universalism in altruism and universalism in trust). Universalism is in [0,1] and the dependent variable is standardized into z-scores. Error bars indicate 95% confidence intervals using robust standard errors. The "All countries" specification includes country fixed effects.

C.5 Results Across Respondent Age

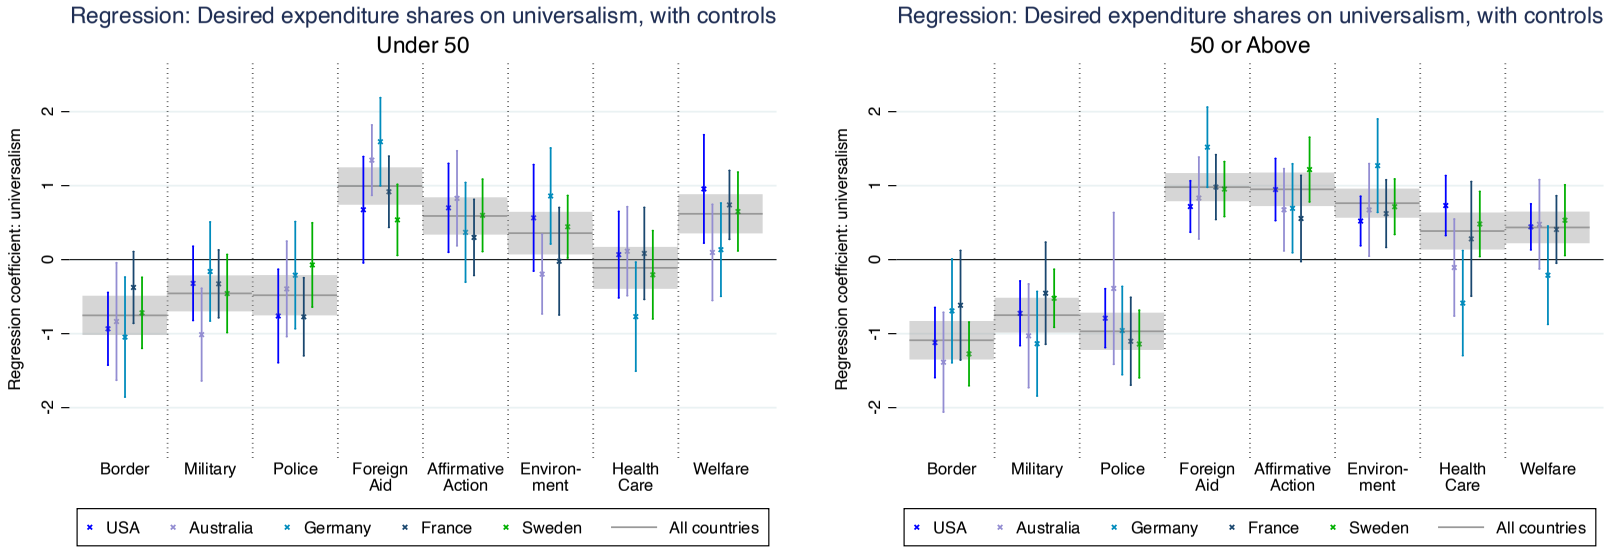

Figure 35: This figure plots the OLS coefficients from regressions of desired expenditure shares for each policy category on universalism, separately for each country, and broken down into respondents that are below (left panel) or at least (right panel) the age of 50. Universalism is in [0,1] and the dependent variables are standardized into z-scores. Error bars indicate 95% confidence intervals using robust standard errors. The "All countries" specification includes country fixed effects.

C.6 Representative Sample

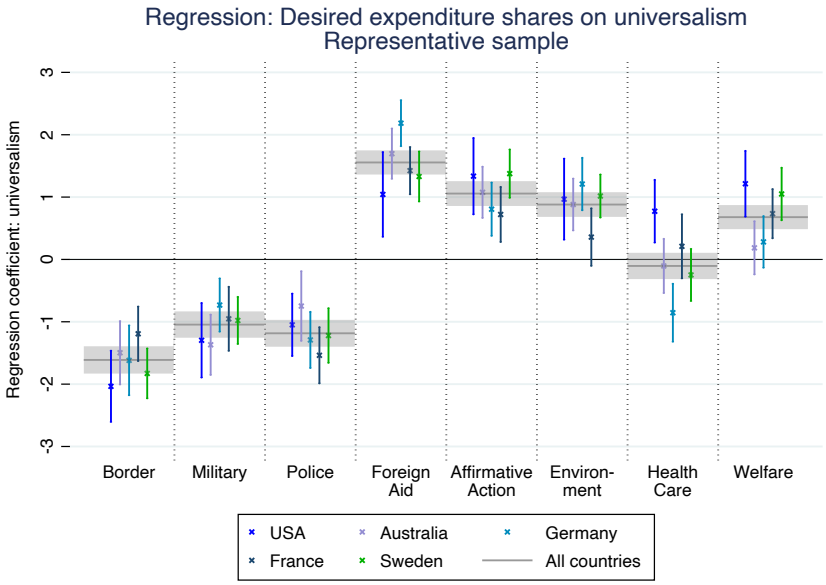

Figure 36: This figure plots the OLS regression coefficients of regressions of desired expenditure shares for each policy domain on universalism, using only the representative sample described in Section 4.1. Universalism is in [0,1] and the dependent variables are standardized into z-scores. Error bars indicate 95% confidence intervals using robust standard errors. The "All countries" specification includes country fixed effects.

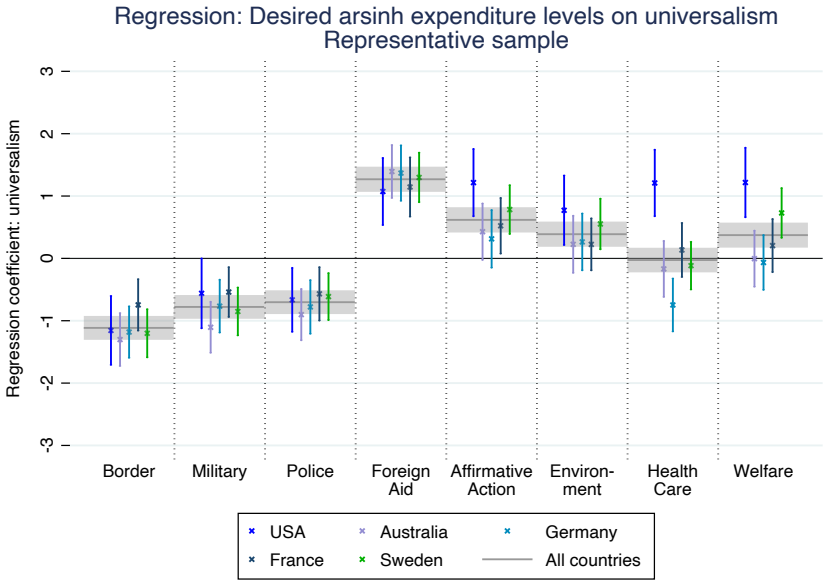

Figure 37: This figure plots the OLS regression coefficients of regressions of desired arsinh expenditure levels for each policy domain on universalism, using only the representative sample described in Section 4.1. Universalism is in [0,1] and the dependent variables are standardized into z-scores. Error bars indicate 95% confidence intervals using robust standard errors. The "All countries" specification includes country fixed effects.

C.7 Summary Statistic of Policy Views and Universalism Measures

Table 14: Summary statistic of policy views and different universalism measures, pooled across countries

	Dependent variable:					
	Summary statistic of policy views					
	(1)	(2)	(3)	(4)	(5)	(6)
Composite universalism in altruism	1.57***	1.22***			1.28***	0.98***
	(0.07)	(0.07)			(0.08)	(0.08)
Composite universalism in trust			1.50***	1.13***	0.74***	0.62***
			(0.09)	(0.09)	(0.10)	(0.09)
Country FE	Yes	Yes	Yes	Yes	Yes	Yes
Demographic Controls	No	Yes	No	Yes	No	Yes
Observations	10881	10881	10881	10881	10881	10881
R^2	0.06	0.11	0.04	0.10	0.06	0.11

Notes. OLS estimates, robust standard errors in parentheses. Data are pooled across all five Western countries. The dependent variable is the summary statistic of policy views, constructed as described in Section 4.3 and standardized into a z-score within each country. The construction of each universalism measure is outlined in Section 4.2. Demographic controls include age, gender, income, wealth, college, urbanicity, religiosity, equity-efficiency preferences, altruism, trust, and beliefs about the efficiency of government. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

C.8 Adjustment for Multiple Hypothesis Testing

This section presents sharpened two-stage q-values to control for the false discovery rate, following Benjamini et al. (2006); Anderson (2012).

Table 15: Two-Stage q-values: Universalism and policy views (expenditure shares)

Policy Domain	Universalism w.r.t...	Original p-value	q-value
Border	Altruism	0.0000	0.001
Border	Trust	0.0000	0.001
Military	Altruism	0.0000	0.001
Military	Trust	0.0038	0.002
Police	Altruism	0.0000	0.001
Police	Trust	0.0003	0.001
Foreign Aid	Altruism	0.0000	0.001
Foreign Aid	Trust	0.0000	0.001
Affirmative Action	Altruism	0.0000	0.001
Affirmative Action	Trust	0.0013	0.001
Environment	Altruism	0.0000	0.001
Environment	Trust	0.0058	0.003
Healthcare	Altruism	0.0519	0.007
Healthcare	Trust	0.7584	0.059
Welfare	Altruism	0.0048	0.002
Welfare	Trust	0.0010	0.001

Table 16: Two-Stage q-values: Universalism and policy views (Likert scale)

Policy Domain	Universalism w.r.t...	Original p-value	q-value
Border	Altruism	0.0000	0.001
Border	Trust	0.0000	0.001
Military	Altruism	0.0000	0.001
Military	Trust	0.0004	0.001
Police	Altruism	0.0000	0.001
Police	Trust	0.0000	0.001
Foreign Aid	Altruism	0.0000	0.001
Foreign Aid	Trust	0.0002	0.001
Affirmative Action	Altruism	0.0000	0.001
Affirmative Action	Trust	0.0274	0.006
Environment	Altruism	0.0000	0.001
Environment	Trust	0.0714	0.014
Healthcare	Altruism	0.0026	0.001
Healthcare	Trust	0.0254	0.006
Welfare	Altruism	0.0000	0.001
Welfare	Trust	0.0001	0.001

Table 17: Two-Stage q-values: Desired expenditure shares on universalism

Policy Domain	Country	Original p-value	q-value
Border	USA	0.0000	0.001
Border	Australia	0.0000	0.001
Border	Germany	0.0000	0.001
Border	France	0.0000	0.001
Border	Sweden	0.0000	0.001
Border	All countries	0.0000	0.001
Military	USA	0.0000	0.001
Military	Australia	0.0000	0.001
Military	Germany	0.0008	0.001
Military	France	0.0001	0.001
Military	Sweden	0.0000	0.001
Military	All countries	0.0000	0.001
Police	USA	0.0000	0.001
Police	Australia	0.0045	0.002
Police	Germany	0.0000	0.001
Police	France	0.0000	0.001
Police	Sweden	0.0000	0.001
Police	All countries	0.0000	0.001
Foreign Aid	USA	0.0000	0.001
Foreign Aid	Australia	0.0000	0.001
Foreign Aid	Germany	0.0000	0.001
Foreign Aid	France	0.0000	0.001
Foreign Aid	Sweden	0.0000	0.001
Foreign Aid	All countries	0.0000	0.001
Affirmative Action	USA	0.0000	0.001
Affirmative Action	Australia	0.0000	0.001
Affirmative Action	Germany	0.0004	0.001
Affirmative Action	France	0.0009	0.001
Affirmative Action	Sweden	0.0000	0.001
Affirmative Action	All countries	0.0000	0.001
Environment	USA	0.0000	0.001
Environment	Australia	0.0000	0.001
Environment	Germany	0.0000	0.001
Environment	France	0.1283	0.022
Environment	Sweden	0.0000	0.001
Environment	All countries	0.0000	0.001
Healthcare	USA	0.0000	0.001
Healthcare	Australia	0.6137	0.104
Healthcare	Germany	0.0003	0.001
Healthcare	France	0.9712	0.166
Healthcare	Sweden	0.9586	0.166
Healthcare	All countries	0.2868	0.048
Welfare	USA	0.0000	0.001
Welfare	Australia	0.3063	0.051
Welfare	Germany	0.1284	0.022
Welfare	France	0.0000	0.001
Welfare	Sweden	0.0000	0.001
Welfare	All countries	0.0000	0.001

Table 18: Two-Stage q-values: Desired expenditure shares on universalism, with controls

Policy Domain	Country	Original p-value	q-value
Border	USA	0.0000	0.001
Border	Australia	0.0002	0.001
Border	Germany	0.0023	0.002
Border	France	0.0226	0.01
Border	Sweden	0.0000	0.001
Border	All countries	0.0000	0.001
Military	USA	0.0005	0.001
Military	Australia	0.0000	0.001
Military	Germany	0.0103	0.006
Military	France	0.0575	0.02
Military	Sweden	0.0034	0.003
Military	All countries	0.0000	0.001
Police	USA	0.0000	0.001
Police	Australia	0.1734	0.047
Police	Germany	0.0110	0.006
Police	France	0.0000	0.001
Police	Sweden	0.0000	0.001
Police	All countries	0.0000	0.001
Foreign Aid	USA	0.0000	0.001
Foreign Aid	Australia	0.0000	0.001
Foreign Aid	Germany	0.0000	0.001
Foreign Aid	France	0.0000	0.001
Foreign Aid	Sweden	0.0000	0.001
Foreign Aid	All countries	0.0000	0.001
Affirmative Action	USA	0.0000	0.001
Affirmative Action	Australia	0.0011	0.001
Affirmative Action	Germany	0.0171	0.008
Affirmative Action	France	0.0353	0.015
Affirmative Action	Sweden	0.0000	0.001
Affirmative Action	All countries	0.0000	0.001
Environment	USA	0.0011	0.001
Environment	Australia	0.3670	0.076
Environment	Germany	0.0000	0.001
Environment	France	0.2314	0.056
Environment	Sweden	0.0000	0.001
Environment	All countries	0.0000	0.001
Healthcare	USA	0.0023	0.002
Healthcare	Australia	0.9500	0.193
Healthcare	Germany	0.0054	0.004
Healthcare	France	0.4625	0.088
Healthcare	Sweden	0.2214	0.055
Healthcare	All countries	0.0585	0.02
Welfare	USA	0.0001	0.001
Welfare	Australia	0.2084	0.053
Welfare	Germany	0.9909	0.198
Welfare	France	0.0003	0.001
Welfare	Sweden	0.0010	0.001
Welfare	All countries	0.0000	0.001

Table 19: Two-Stage q-values: Desired arsinh expenditure levels on universalism

Policy Domain	Country	Original p-value	q-value
Border	USA	0.0000	0.001
Border	Australia	0.0000	0.001
Border	Germany	0.0000	0.001
Border	France	0.0009	0.001
Border	Sweden	0.0000	0.001
Border	All countries	0.0000	0.001
Military	USA	0.0000	0.001
Military	Australia	0.0000	0.001
Military	Germany	0.0005	0.001
Military	France	0.0604	0.022
Military	Sweden	0.0000	0.001
Military	All countries	0.0000	0.001
Police	USA	0.0004	0.001
Police	Australia	0.0000	0.001
Police	Germany	0.0003	0.001
Police	France	0.0304	0.012
Police	Sweden	0.0000	0.001
Police	All countries	0.0000	0.001
Foreign Aid	USA	0.0000	0.001
Foreign Aid	Australia	0.0000	0.001
Foreign Aid	Germany	0.0000	0.001
Foreign Aid	France	0.0000	0.001
Foreign Aid	Sweden	0.0000	0.001
Foreign Aid	All countries	0.0000	0.001
Affirmative Action	USA	0.0000	0.001
Affirmative Action	Australia	0.0678	0.024
Affirmative Action	Germany	0.1508	0.043
Affirmative Action	France	0.0023	0.002
Affirmative Action	Sweden	0.0000	0.001
Affirmative Action	All countries	0.0000	0.001
Environment	USA	0.0000	0.001
Environment	Australia	0.3202	0.083
Environment	Germany	0.1704	0.047
Environment	France	0.1002	0.031
Environment	Sweden	0.0001	0.001
Environment	All countries	0.0000	0.001
Healthcare	USA	0.0000	0.001
Healthcare	Australia	0.4163	0.092
Healthcare	Germany	0.0007	0.001
Healthcare	France	0.1888	0.051
Healthcare	Sweden	0.8860	0.18
Healthcare	All countries	0.0132	0.006
Welfare	USA	0.0000	0.001
Welfare	Australia	0.9152	0.18
Welfare	Germany	0.9019	0.18
Welfare	France	0.0171	0.007
Welfare	Sweden	0.0000	0.001
Welfare	All countries	0.0000	0.001

C.9 Results for Non-Western Countries

As discussed in Section 6, the intracorrelated cluster of policy views that we seek to explain in this paper is not neatly documented in the two non-Western countries included in our survey, Brazil and South Korea. This section documents that: (i) in Brazil and South Korea, the structure of the first eigenvector of desired expenditure shares does not resemble the same intuitive notion of "left" and "right" that we see in the West; (ii) accordingly, the link between left-vs.-right wing ideology and desired arsinh expenditure levels on traditionally "right-wing" or "left-wing" domains is much attenuated in these two countries, and (iii) the relationship between universalism and policy views is significantly weaker and often opposite in sign relative to what we observe in the Western countries in our survey.

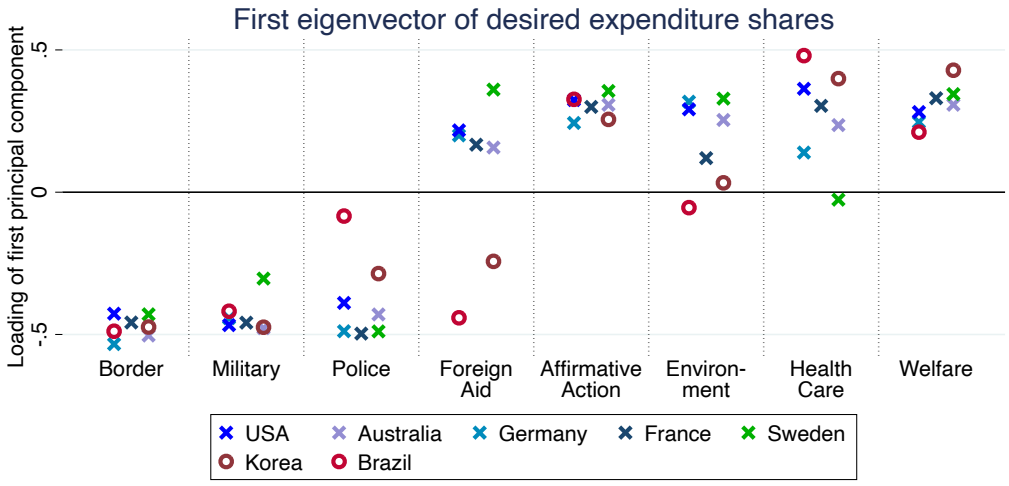

Figure 38: Factor loadings of the first principal component of desired expenditure shares. Sign convention: the loading on "Border" is always non-positive, and the other signs are determined accordingly.

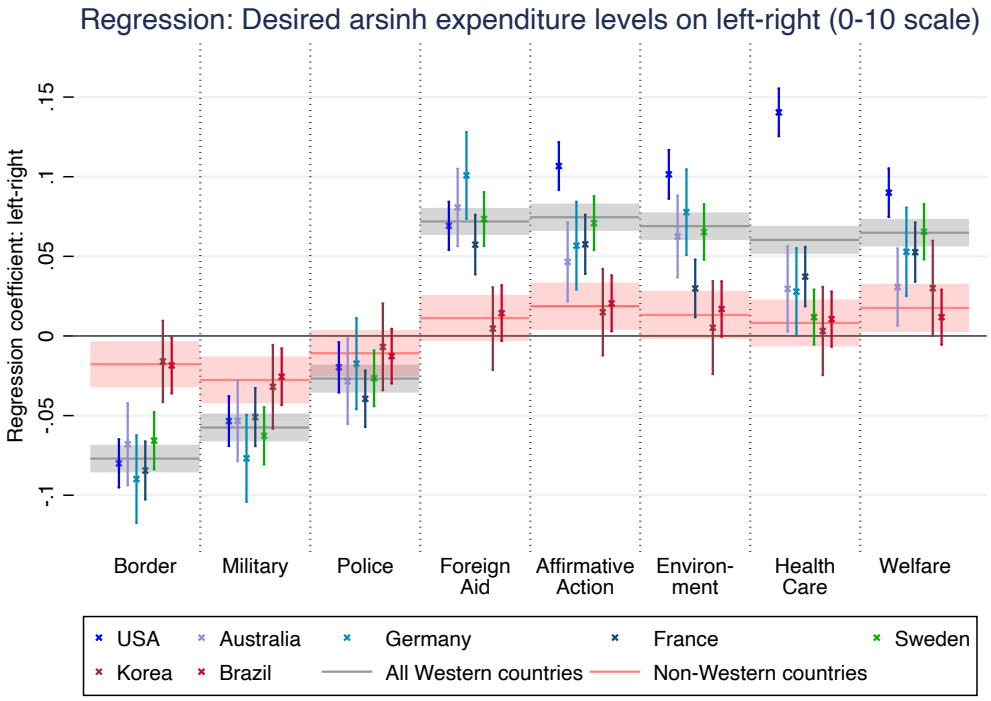

Figure 39: The figure plots the OLS regression coefficients of univariate regressions of desired arsinh expenditure levels for each policy domain on self-positioning on a left-right scale (0–10). The dependent variables are standardized into z-scores. Error bars indicate 95% confidence intervals using robust standard errors. The "All Western countries" and "Non-Western countries" specifications include country fixed effects.

Baseline results. Figure 40 replicates Figure 6 but additionally shows the results for Brazil and South Korea. The results are visibly different from those in the Western countries, with both the magnitude and the sign of the relationship between universalism and policy views usually different. Moreover, we note that while in the Western countries the composite universalism measure is strongly correlated with left-wing identification ($\rho = 0.25$, conditional on country fixed effects), the correlation between composite universalism and left-wing identification is only $\rho = 0.09$ in the two non-Western countries, conditional on country fixed effects.

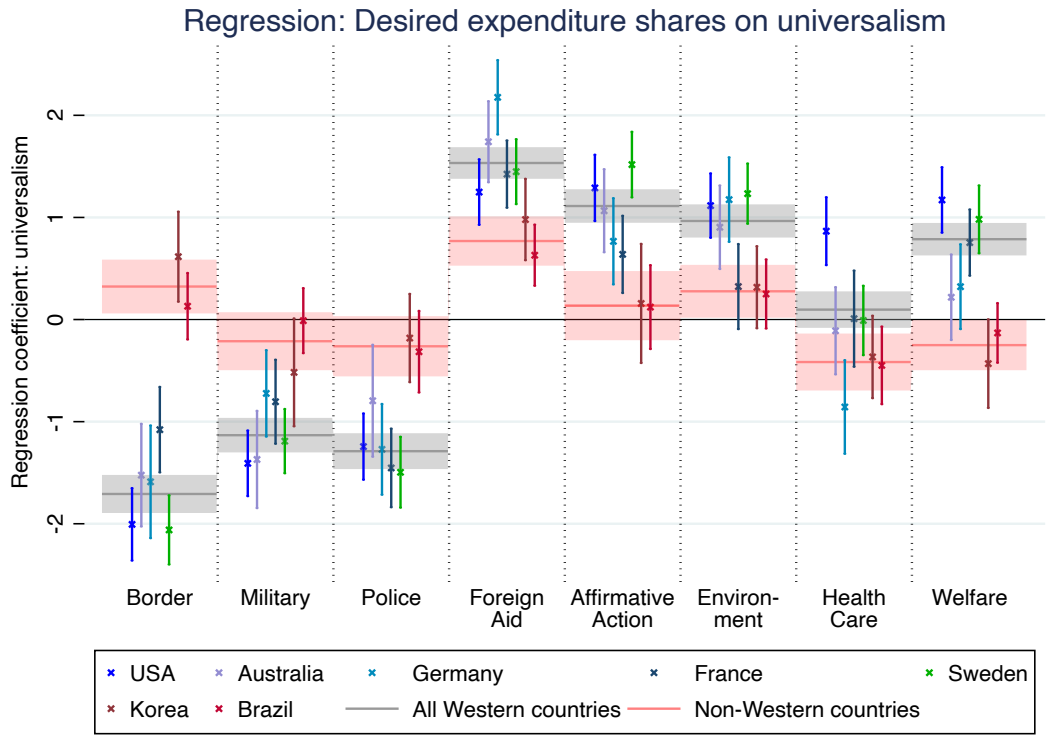

Figure 40: This figure plots OLS coefficients from regressions of desired expenditure shares on universalism. Universalism is in [0,1] and the dependent variables are standardized into z-scores. Error bars indicate 95% confidence intervals using robust standard errors. The "All Western countries" and "Non-Western countries" specifications include country fixed effects.

C.10 Analysis of Specific Policy Proposals

C.10.1 Tables

Table 20: Specific policy proposals

Policy domain	**More universalist**	**Less universalist**
Welfare payments	Redistributing local tax revenues as welfare payments across all communities nationwide	Redistributing local tax revenues as welfare payments only within the local communities they were raised
Universal healthcare	Using local tax revenues to fund health insurance across all communities nationwide	Using local tax revenues to fund health insurance only within the local communities they were raised
Environmental protection	Preventing global climate change	Cleaning and conserving forests and rivers in local communities in [the U.S., France, etc.]
Foreign aid	Sending foreign aid to countries that are in most need of help	Sending foreign aid to countries that are our international allies
Affirmative action	Measures to ensure no individual is disadvantaged in access to education, the labor force, and marriage	Measures to ensure no one of your same background (e.g., gender, ethnic background or ancestry) is disadvantaged in access to education, the labor force, and marriage
Military and counterintelligence	Peacekeeping and humanitarian missions by the military abroad	Ensuring [American, French, etc.] defense and security
Effective border control	Identifying and admitting into the country only those immigrants with the highest need for help	Identifying and admitting into the country only those immigrants who would be good citizens (e.g., be likely to pay taxes and refrain from engaging in criminal activities)
Police and law enforcement	Sensitivity training for the police to ensure justice and equal treatment of all	Increasing the capabilities of the police to prevent and prosecute criminal or suspicious behavior

C.10.2 Binned scatter plots: Specific Policy Proposals

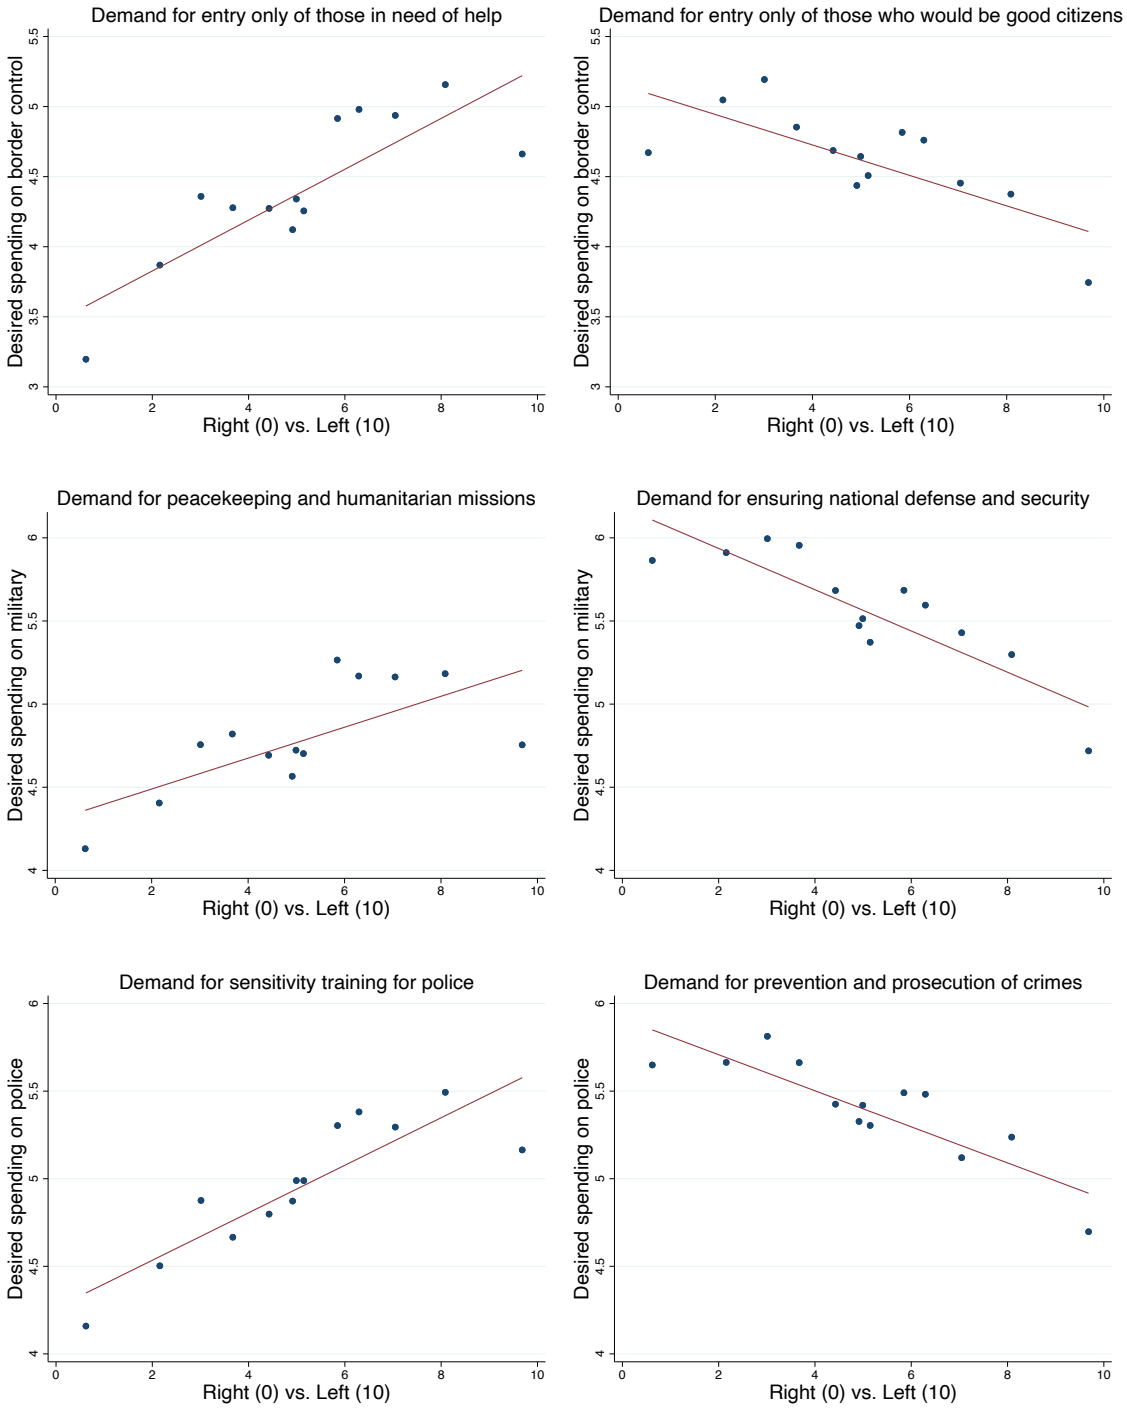

Figure 41: This figure shows binned scatter plots between arsinh desired spending levels for specific policy proposals and respondents' universalism. In each row, the left panel shows the more universalist and the right panel the more communitarian implementation. The first row shows results for border control, the second one for military and the third one for police. The data are pooled across Western countries, and all panels are constructed controlling for country fixed effects.

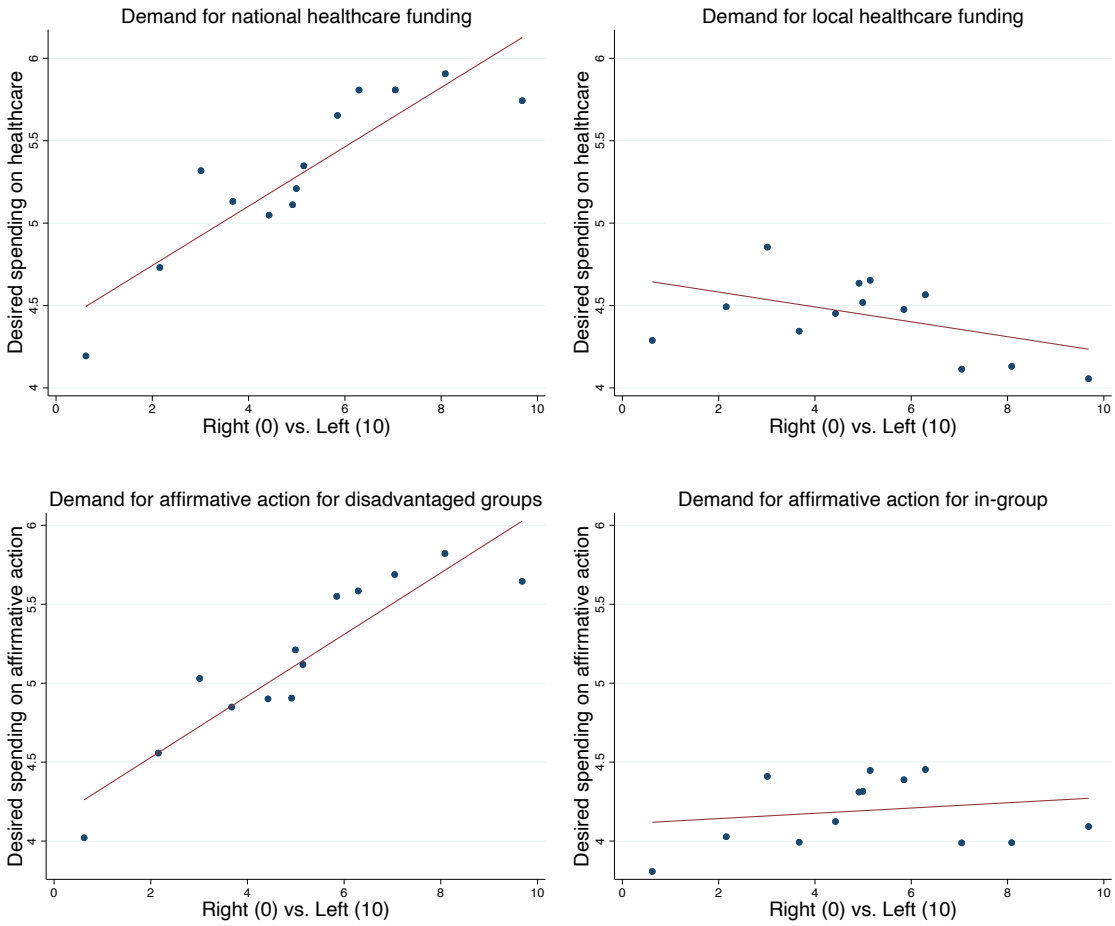

Figure 42: This figure shows binned scatter plots between arsinh desired spending levels for specific policy proposals and respondents' universalism. In each row, the left panel shows the more universalist and the right panel the more communitarian implementation. The first row shows results for healthcare, and the second one for affirmative action. The data are pooled across Western countries, and all panels are constructed controlling for country fixed effects.

C.11 Controlling for Left-Right Fixed Effects

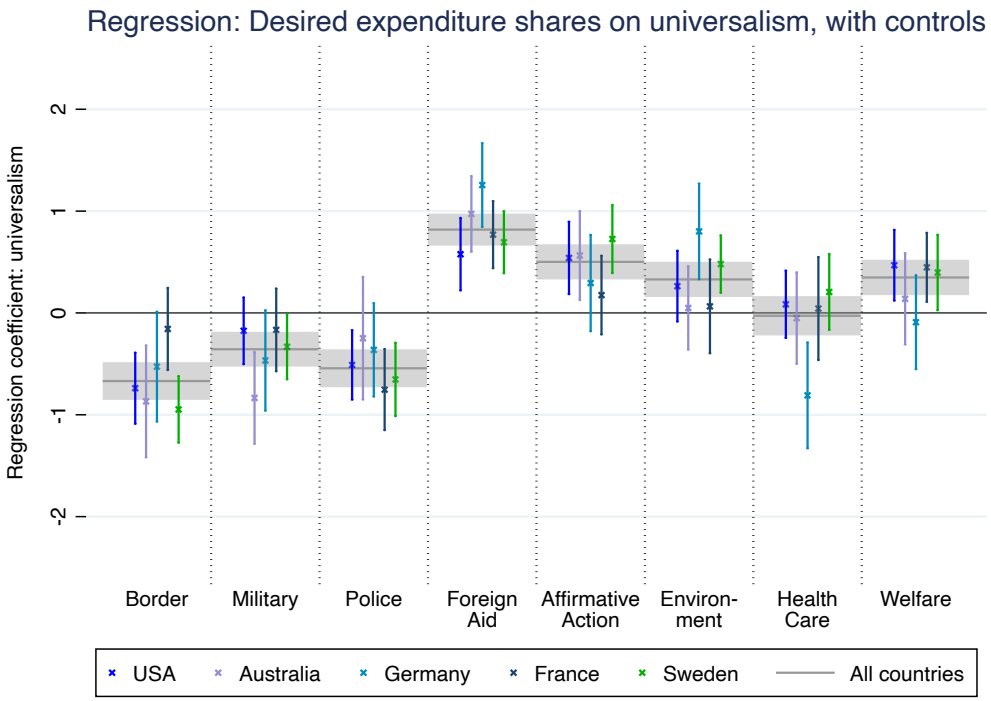

Figure 43: This figure plots the OLS coefficients of regressions of desired expenditure shares for each policy domain (as a fraction of overall desired government spending for the eight policy domains) on composite universalism, controlling for age, gender, income, wealth, college, urbanicity, religiosity, equity-efficiency preferences, altruism, trust, beliefs about the efficiency of government, beliefs about whether one will personally benefit from government expenditure in each domain, and left-right fixed effects. See Appendix E for details on the construction of these variables. Universalism is in [0,1] and the dependent variables are standardized into z-scores within each country. Error bars indicate 95% confidence intervals using robust standard errors. The "All countries" specifications include country fixed effects.

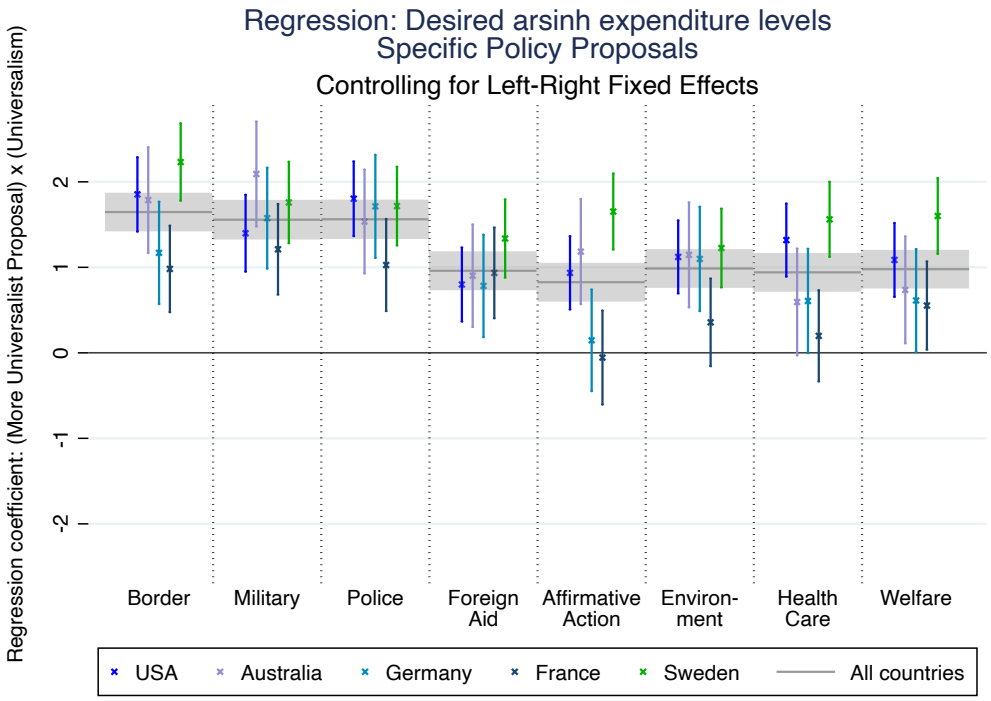

Figure 44: This figure presents OLS coefficients from regressions of arsinh desired expenditure for each specific policy on: universalism, an indicator for the more universalist implementation of the policy, and an interaction between these two variables. We control for left-right self-assessment in all of these regressions. We plot the coefficients on the interaction variable; that is, higher universalism is associated with a larger increase in support for the universalist implementation of a policy, even when collapsing variation to individuals with the same left-right self-assessment. Error bars indicate 95% confidence intervals using robust standard errors. The "All countries" specification includes country fixed effects.

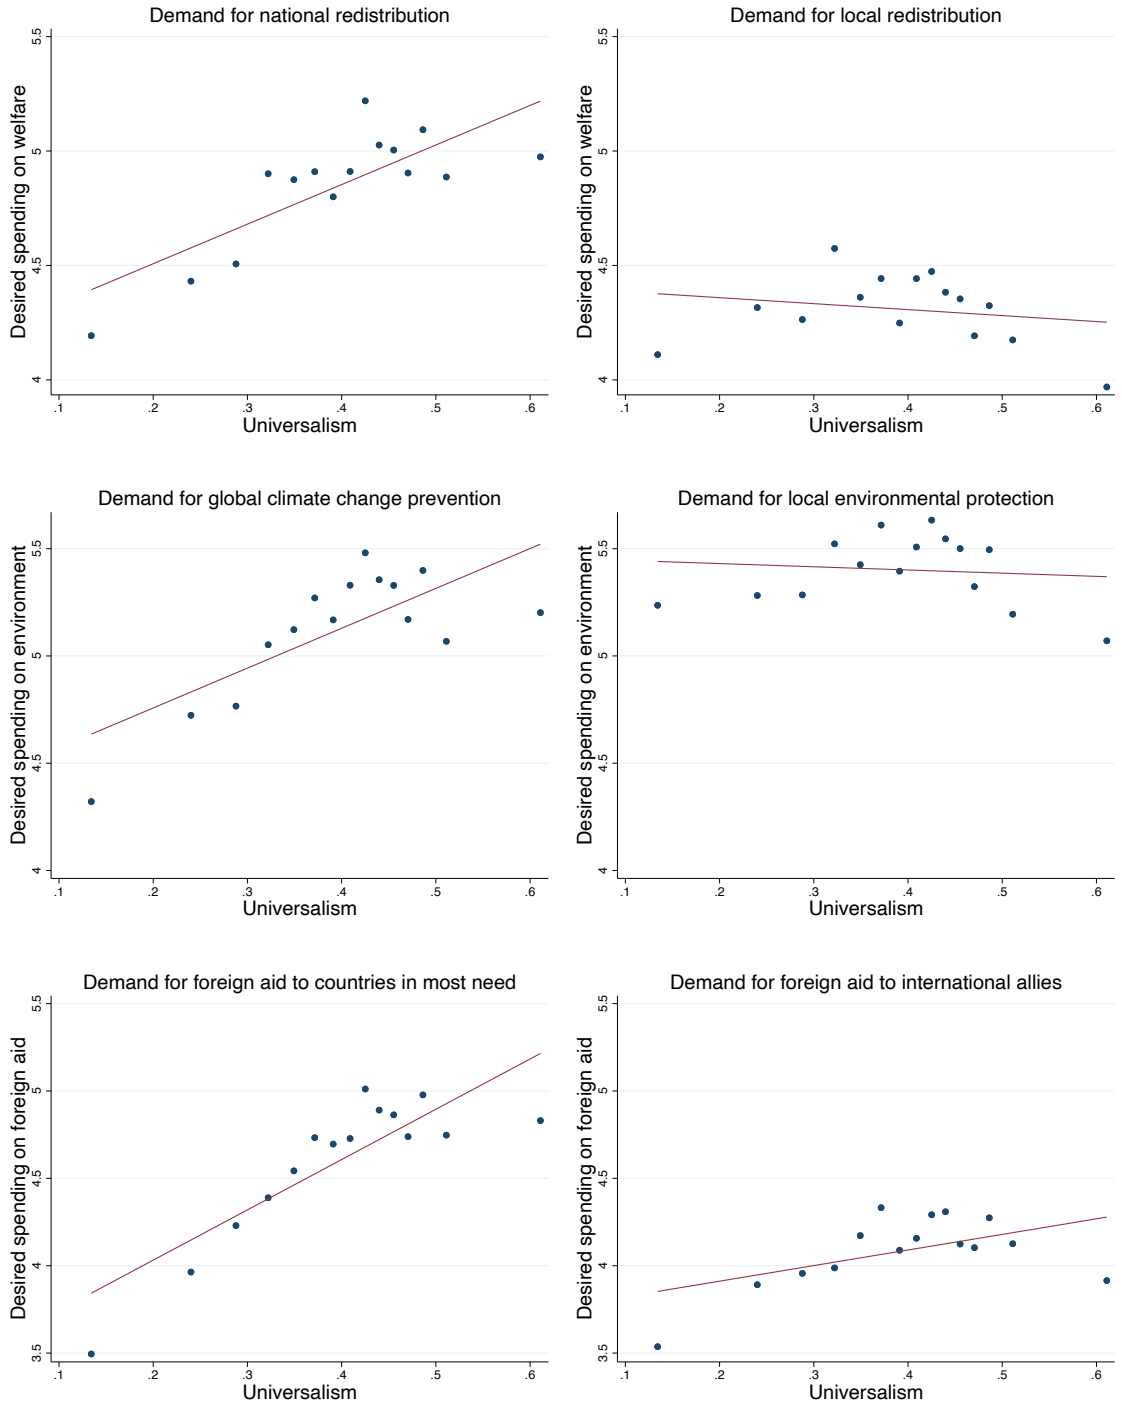

Figure 45: This figure shows binned scatter plots between arsinh desired spending levels for specific policy proposals and respondents' universalism, controlling for respondents' left-right self-assessment. In each row, the left panel shows the more universalist and the right panel the more communitarian implementation. The first row shows results for welfare, the second one for environmental protection and the third one for foreign aid. The data are pooled across Western countries, and all panels are constructed controlling for country fixed effects.

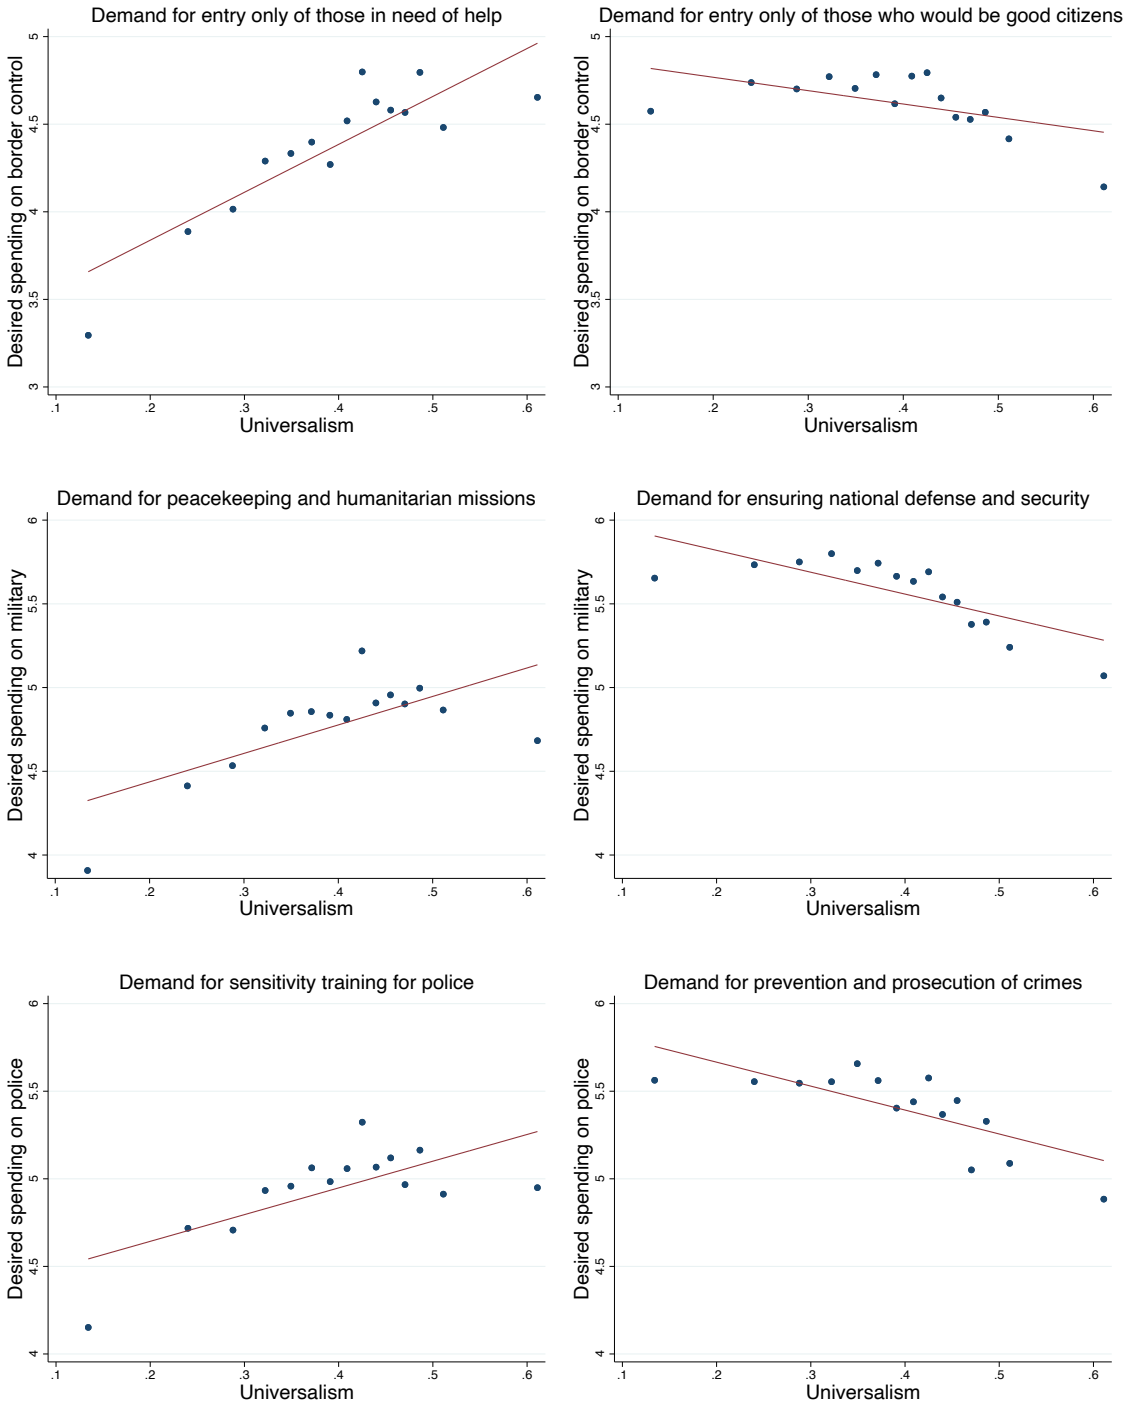

Figure 46: This figure shows binned scatter plots between arsinh desired spending levels for specific policy proposals and respondents' universalism, controlling for respondents' left-right self-assessment. In each row, the left panel shows the more universalist and the right panel the more communitarian implementation. The first row shows results for border control, the second one for military and the third one for police. The data are pooled across Western countries, and all panels are constructed controlling for country fixed effects.

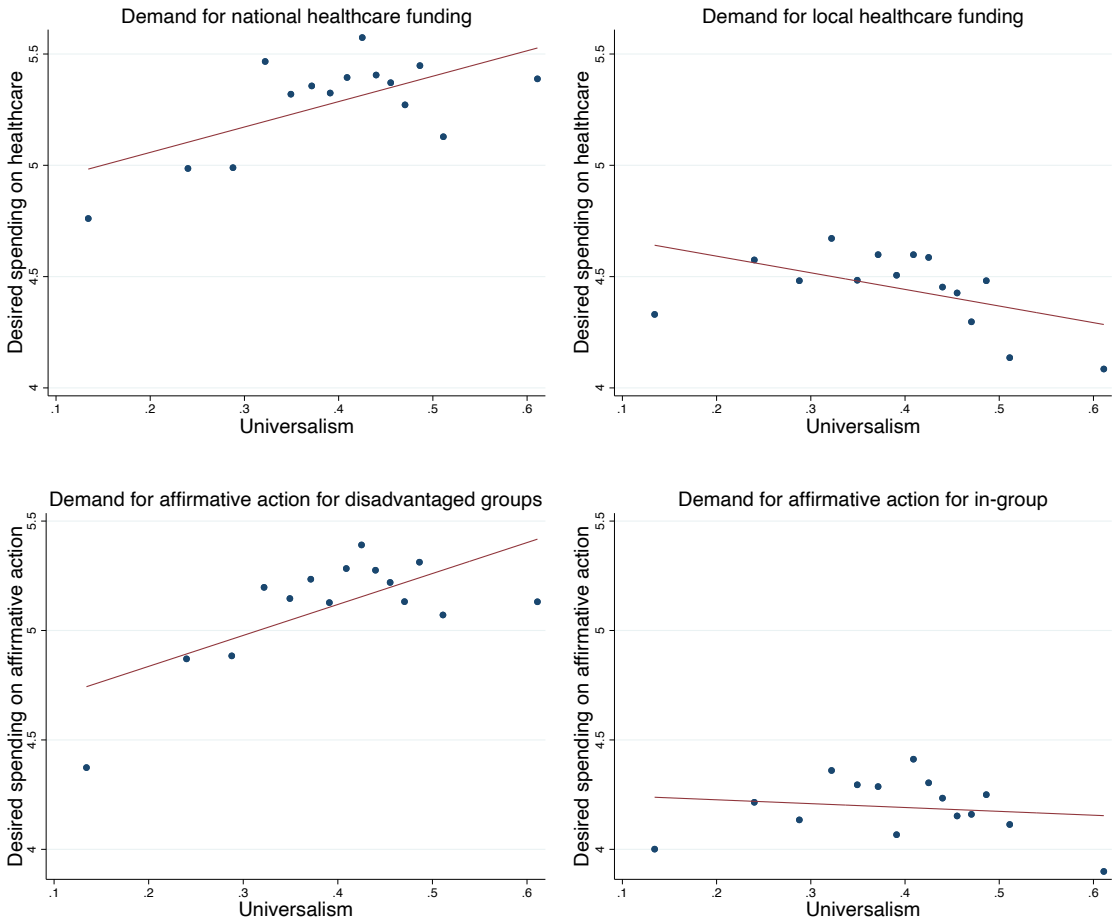

Figure 47: This figure shows binned scatter plots between arsinh desired spending levels for specific policy proposals and respondents' universalism, controlling for respondents' left-right self-assessment. In each row, the left panel shows the more universalist and the right panel the more communitarian implementation. The first row shows results for healthcare, and the second one for affirmative action. The data are pooled across Western countries, and all panels are constructed controlling for country fixed effects.

D Measurement Error and ORIV Analyses

Measurement error is ubiquitous in lab and survey settings. To ensure that the estimates presented in this paper are neither artefacts of nor attenuated by the presence of measurement error in our elicitations of outcome and explanatory variables, we make use of the instrumentation strategies laid out in Gillen et al. (2019).

That is, we employ the obviously-related instrumental variables (ORIV) technique by eliciting quasi-duplicate measurements of the set of specific policy views and of our measures of universalism. This analysis was pre-registered and detailed the formulation of instruments for our variables of interest as follows.

Choice of instruments. In Section 4.3, we document how we elicited support for our eight broad policy categories (affirmative action, border control, environment, foreign aid, health, military, police, and welfare) with two complementary strategies. The first elicited respondents' desired, per capita annual spending by their national government on each of these categories. The second strategy elicited respondents' support for government spending in each of these categories on an 11-point Likert scale. These two elicitations were separated by a series of tasks, including elicitations of support for specific, framed policies, and a sociodemographic questionnaire. This leaves us with duplicate measurements of support for the eight individual policy domains.

For the set of predictors (our measures of universalism), we leverage the fact that the order of social groups presented in our survey is randomized within the domestic and global categories. As such, the first measure of universalism is constructed just like the main measure described in Section 4.2, except that it only uses the five domestic groups that (randomly) appear first and the three global groups that (randomly) appear first in the survey for each respondent. We do not include the foreign decision as there was only *one* of these elicitations.

Analogously, the second measure of universalism is constructed just like the main measure described in Section 4.2, except that it only uses the five domestic groups that (randomly) appear last and the two global groups that (randomly) appear last in the survey. We construct these two proxies for both universalism in altruism and universalism in trust separately, and for the composite measure of universalism that averages the two.

Results. With this set of instruments in hand, we replicated our analysis of the relationships between universalism and the structure of ideology with the stacked ORIV regressions described in Gillen et al. (2019). We were interested in ensuring that measurement error neither attenuates nor artifically produces the relationship between views regard-

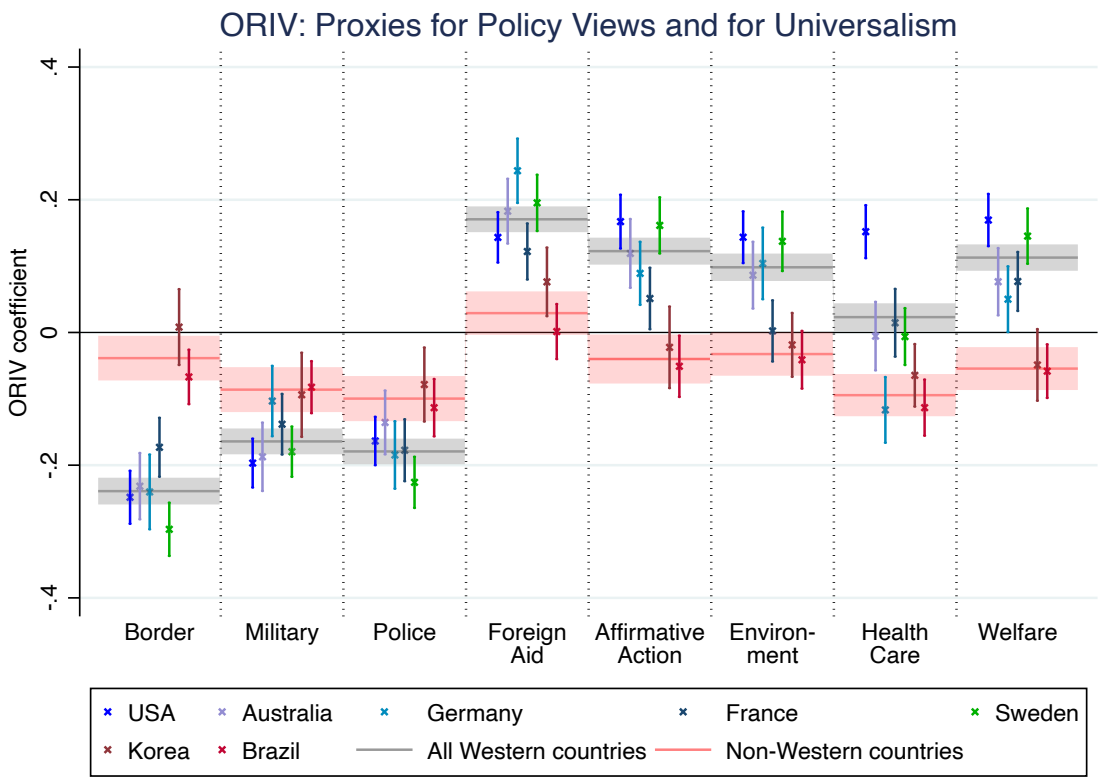

Figure 48: This figure presents ORIV coefficients for the regression of duplicate elicitations of policy views on duplicate elicitations of our summary measure of moral universalism. As recommended by Gillen et al. (2019), both the universalism measures and outcome variables are standardized into z-scores so they have the same scale. Standard errors are clustered at the respondent level. The "All western countries" and "Non-Western countries" specifications include country fixed effects.

ing each of our eight individual policies and universalism.

We follow the recommendations in Gillen et al. (2019) and use standardized versions of both our universalism measures and policy views. Moreover, since each respondent appears twice when implementing ORIV, standard errors are clustered at the respondent level. In all cases the results with the ORIV estimator are very similar to those presented in the main text with OLS.

E Definition of Main Survey Variables

Left vs. right. Respondent's self-positioning on the left-right political spectrum in response to the following prompt: "Oftentimes, people speak of relatively left-wing and relatively right-wing political views. On a scale from 0 (very left-wing) to 10 (very right-wing), where would you place yourself on this scale?"

Summary statistic of policy views. Summary statistic of policy views, given by

$$\text{Left vs. right summary statistic} = \frac{\text{Foreign aid} + \text{Environment} + \text{Aff. action} + \text{Welfare} + \text{Health care}}{5} - \frac{\text{Military} + \text{Police} + \text{Border control}}{3} \tag{43}$$

Each policy denotes the desire expenditure share for the given policy, defined also in this section. The summary statistic generally increases with attitudes towards left-wing views, and correlates with self-positioning on a 0 to 10 scale. This summary statistic is standardized into z-scores within countries.

Domestic universalism in altruism. Universalism with respect to altruism (preferences), measured through bystander dictator games over the local currency analogue of hypothetical $100, between a domestic member of one's in-groups relative to a domestic stranger. The measure averages the ten corresponding money allocation decisions.

Foreign universalism in altruism. Universalism with respect to altruism (preferences), measured through a bystander dictator games over the local currency analogue of hypothetical $100 between a domestic stranger and a global stranger.

Global universalism in altruism. Universalism with respect to altruism (preferences), measured through bystander dictator games over the local currency analogue of hypothetical $100, between a global member of one's in-groups relative to a global stranger. The measure averages the five corresponding money allocation decisions.

Summary measure of universalism in altruism. Unweighted average of domestic universalism in altruism, foreign universalism in altruism, and global universalism in altruism. Because these three individual components correlate highly with each other, the summary measure reduces the dimensionality of the data and describes a respondent's broad universalism in altruism as a general type.

Domestic universalism in trust. Trust analogue of domestic universalism in altruism, where the bystander dictator game is instead over 100 trust points.

Foreign universalism in trust. Trust analogue of foreign universalism in altruism, where the bystander dictator game is instead over 100 trust points.

Global universalism in trust. Trust analogue of global universalism in altruism, where the bystander dictator game is instead over 100 trust points.

Summary measure of universalism in trust. Trust analogue of the summary measure of universalism in altruism. That is, unweighted average of domestic universalism in trust, foreign universalism in trust, and global universalism in trust.

Composite measure of universalism. Unweighted average of (i) summary measure of universalism in altruism and (ii) summary measure of universalism in trust. Reduces the dimensionality of the data.

Revealed altruism. Altruism as elicited through a standard dictator game over $100 between the self and a domestic stranger.

Residual altruism. Residuals from within-country regressiosn of dictator game behavior (revealed altruism) on the summary statistic of universalism in altruism. Because the dictator game is framed vis-à-vis a randomly-selected stranger, the raw measure of altruism partly includes universalism; residualizing of universalism measures that portion of revealed altruism that cannot be explained by behavior in our universalism decisions.

Revealed generalized trust. Generalized trust in others as elicited through an allocation of trust points on a scale from 0 to 100. Respondents were prompted to consider their trust in a domestic stranger, where 0 meant that they believe they "cannot trust a randomly-selected person very much", and 100 meant they believe "a randomly-selected person can in general be trusted a great deal."

Residual trust. Residuals from within-country regressions of revealed generalized trust on the summary statistic of universalism in trust. Because generalized trust is framed vis-à-vis a randomly-selected stranger, the raw measure of trust partly includes universalism; residualizing of universalism measures that portion of generalized trust that cannot be explained by behavior in our universalism decisions.

Equity-efficiency preferences. Elicitation of preferences for efficiency over equity, as given by a bystander dictator game between two "randomly-selected people" who live in the respondent's country, in which the most unequal split of money maximizes total payoffs. The measure captures how much a respondent deviates from an equal, 50:50 split of the money.

Desired government spending on policy categories. Measure of support for eight distinct policy domains: (i) affirmative action, (ii) border control, (iii) environment, (iv) foreign aid, (v) healthcare, (vi) military, (vii) police, and (viii) welfare payments.

Respondents were prompted to respond in free-form text entry with their desired level of annual, per-capita spending (in local currency) by their corresponding national level of government on each of the eight domains. They were provided a reference value of the annual per capita spending amount on education by their national level of government.

These dollar amounts were then translated into desired shares, out of a total amount of per capita spending by the national government. Both desired expenditure shares and expenditure levels were standardized into z-scores within-country.

Support for policy categories. Measure of support for eight distinct policy domains: (i) affirmative action, (ii) border control, (iii) environment, (iv) foreign aid, (v) healthcare, (vi) military, (vii) police, and (viii) welfare payments.

Respondents were prompted to respond on a 0 (strongly oppose) to 10 (strongly support) Likert scale with their level of support for national government spending on each of the eight domains.

Relative support for policy categories. Measure of relative support for each of the eight distinct policy domains, given by support for a given policy divided by the sum of support across all policies.

Desired government spending on individual policies. Measure of support for sixteen distinct policies, two per each of the eight broad policy domains. Per each of these policy domains, one specific policy had a less universalist implementation, while the other a more universalist one. See Table 20.

Respondents were prompted to respond with their desired level of annual, per-capita spending (in local currency) by their corresponding national level of government on each of these sixteen policies. Both ensuing desired expenditure shares and expenditure levels were standardized into z-scores within-country.

Religiosity. Composite measure from a principal component analysis of: (i) self-described religiosity on a scale from 0 (not at all religious) to 10 (very religious); (ii) church attendance on a scale from 0 to 5; and an indicator for atheism, agnosticism, or no religion.

Income Index. Composite measure from a principal component analysis of: (i) percentile rank of income (from free-form text entry), and (ii) income on a scale from 0 to 4 (roughly corresponding to income quintiles in each country). Because there are some large outliers in the free text entry format for income, we first delete all responses that are larger than 500 times the within-country median response. This affects 0.02% of responses. We then winsorize this data at 3 standard deviations from the within-country mean, which affects 0.6% of responses.

Wealth Index. Composite measure from a principal component analysis of: (i) an indicator for stock ownership, (ii) an indicator for home ownership, and (iii) percentile rank of net worth (from free-form text entry). Because there are some large outliers in the free text entry format for net worth, we first delete all responses that are larger than 500 times the within-country median response. This affects 0.3% of responses. We then winsorize this data at 3 standard deviations from the within-country mean, which affects 1.6% of responses.

Urbanicity. Respondent's neighborhood size on a 10-step variable: > 1 million, 200k-1m, 50k-200k, 20k-50k and close to metro, 20k-50k and not close to metro, 3k-20k and close to metro, 3k-20k and not close to metro, 500-300k and close to metro, 500-3k and not close to metro, <500.

Educational attainment. Respondent's educational attainment. Across all countries but Brazil and Germany, the four educational categories were the local equivalents of: (i) no high school, (ii) high school, (iii) some college or vocational training, (iv) bachelor's degree or higher. In Brazil, the four educational categories were: (i) no formal education, (ii) elementary school, (iii) high school, and (iv) bachelor's degree or higher. In Germany, the three educational categories were: (i) no vocational training, (ii) vocational training, and (iii) university degree.

Beliefs in the efficiency of government. Respondent's rating on a scale from 0 (the government is wasteful) to 10 (the government is generally efficient) on the efficiency of the government in implementing policies and providing for public services.

Belief in personal benefit from government expenditure on policy categories. Respondent's report on the probability (0%-100%) that they would personally benefit over the twelve months following the survey from the corresponding services of the eight broad policy domains: (i) affirmative action, (ii) border control, (iii) environment, (iv) foreign aid, (v) healthcare, (vi) military, (vii) police, and (viii) welfare payments.